

Hortonworks DataFlow

Getting Started with Streaming Analytics

(November 9, 2017)

Hortonworks DataFlow: Getting Started with Streaming Analytics

Copyright © 2012-2017 Hortonworks, Inc. Some rights reserved.



Except where otherwise noted, this document is licensed under
Creative Commons Attribution ShareAlike 4.0 License.
<http://creativecommons.org/licenses/by-sa/4.0/legalcode>

Table of Contents

1. Building an End-to-End Stream Application	1
1.1. Understanding the Use Case	1
1.2. Reference Architecture	2
2. Prepare Your Environment	3
2.1. Deploying Your HDF Clusters	3
2.2. Registering Schemas in Schema Registry	3
2.2.1. Create the Kafka Topics	3
2.2.2. Register Schemas for the Kafka Topics	4
3. Creating a Dataflow Application	7
3.1. Data Producer Application Generates Events	7
3.2. NiFi: Create a Dataflow Application	8
3.2.1. NiFi Controller Services	8
3.2.2. NiFi Ingests the Raw Sensor Events	9
3.2.3. Publish Enriched Events to Kafka for Consumption by Analytics Applications	10
3.2.4. Start the NiFi Flow	11
4. Creating a Stream Analytics Application	12
4.1. Create a Service Pool and Environment	12
4.2. Create Your First Application	12
4.3. Creating and Configuring the Kafka Source Stream	13
4.4. Connecting Components	15
4.5. Joining Multiple Streams	16
4.6. Filtering Events in a Stream using Rules	17
4.7. Using Aggregate Functions over Windows	18
4.8. Implementing Business Rules on the Stream	19
4.9. Transforming Data using a Projection Processor	21
4.10. Creating Alerts with Notifications Sink	23
4.11. Streaming Alerts to an Analytics Engine for Dashboarding	24
4.12. Streaming Violation Events to an Analytics Engine for Descriptive Analytics	25
4.13. Streaming Violation Events into a Data Lake and Operational Data Store	27
5. Deploy an Application	31
5.1. Configure Deployment Settings	31
5.2. Deploy the App	31
5.3. Running the Stream Simulator	33
6. Stream Operations	34
6.1. My Applications View	34
6.2. Application Performance Monitoring	34
6.3. Troubleshooting and Debugging a Stream application	35
6.3.1. Streaming Engine Infrastructure Metrics	36
6.3.2. Changing Log Levels Dynamically and with Expiration Policies	37
6.3.3. Distributed Log Search	37
6.4. Exporting and Importing Stream applications	38
7. Advanced: Doing Predictive Analytics on the Stream	40
7.1. Logistical Regression Model	41
7.2. Export the Model into SAM's Model Registry	42
7.3. Enrichment and Normalization of Model Features	43

7.4. Setting up your Enrichment Store and Building Custom UDFs and Processors	43
7.5. Upload Custom Processors and UDFs for Enrichment and Normalization	44
7.5.1. Upload Custom UDFs	44
7.5.2. Upload Custom Processors	46
7.6. Scoring the Model in the Stream using a Streaming Split Join Pattern	51
7.7. Streaming Split Join Pattern	52
7.8. Score the Model using the PMML Processor and Alert	59
8. Creating Visualizations Using Superset	63
8.1. Creating Insight Slices	63
8.2. Adding Insight Slices to a Dashboard	65
8.2.1. Dashboards for the Trucking IOT App	65

1. Building an End-to-End Stream Application

A good way to get started using Hortonworks DataFlow (HDF) with Streaming Analytics Manager and Schema Registry is to imagine a real life use case, and to learn about the common HDF stream processing tasks and concepts through this use case. This guide sets up a fictional use case, and walks you through the deployment and common tasks you would perform while engaging in many of HDF's stream processing use cases.

Use this guide as a tutorial to get you started with SAM and Schema Registry. All the resources required to complete the tasks are provided in line.

1.1. Understanding the Use Case

To build a complex streaming analytics application from scratch, we will work with a fictional use case. A trucking company has a large fleet of trucks, and wants to perform real-time analytics on the sensor data from the trucks, and to monitor them in real time. Their analytics application has the following requirements:

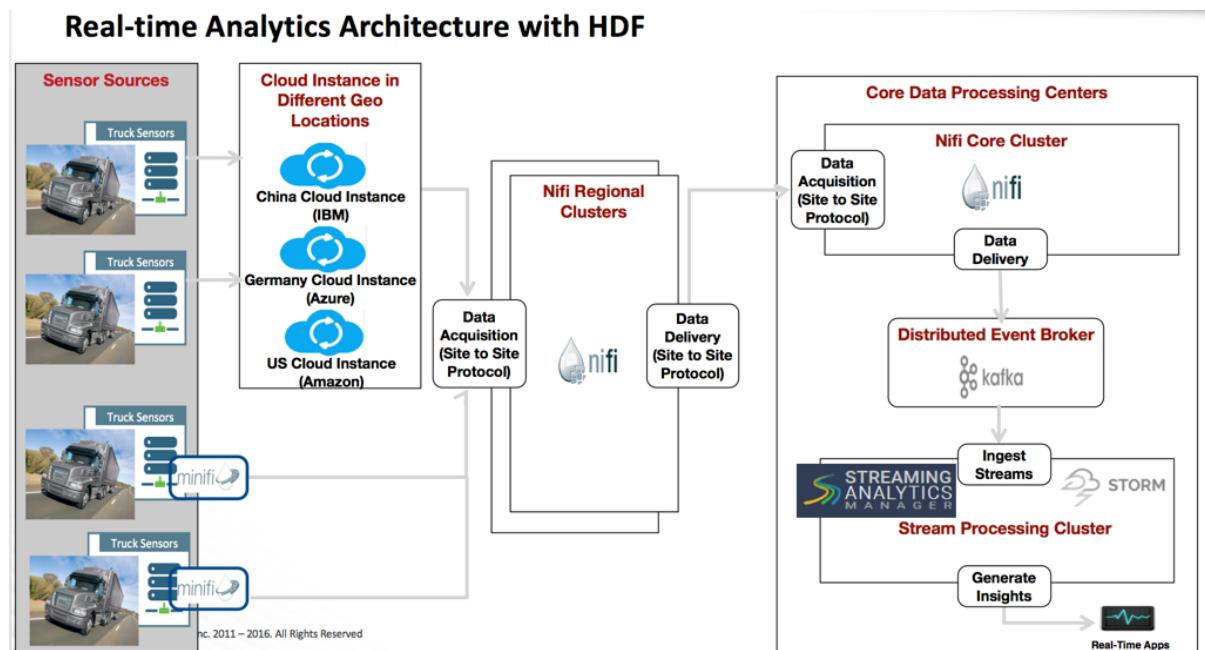
1. Outfit each truck with two sensors that emit event data such as timestamp, driver ID, truck ID, route, geographic location, and event type.
 - The geo event sensor emits geographic information (latitude and longitude coordinates) and events such as excessive braking or speeding.
 - The speed sensor emits the speed of the vehicle.
2. Stream the sensor events to an IoT gateway, which serializes the events as Avro objects and streams them into separate Kafka topics, one for each Kafka sensor.
3. Use NiFi to consume the serialized Avro events from the Kafka topics, and then route, transform, enrich, and deliver the data to a downstream Kafka instance.
4. Connect to the two streams of data to do analytics on the stream.
5. Join the two sensor streams using attributes in real-time. For example, join the geo-location stream of a truck with the speed stream of a driver.
6. Filter the stream on only events that are infractions or violations.
7. All infraction events need to be available for descriptive analytics (dash-boarding, visualizations, or similar) by a business analyst. The analyst needs the ability to do analysis on the streaming data.
8. Detect complex patterns in real-time. For example, over a three-minute period, detect if the average speed of a driver is more than 80 miles per hour on routes known to be dangerous.
9. When each of the preceding rules fires, create alerts and make them instantly accessible.

10. Execute a logistical regression Spark ML model on the events in the stream to predict if a driver is going to commit a violation. If violation is predicted, then alert on it.

The below sections walks you through how to implement all ten requirements. Requirements 1-3 are done using NiFi and Schema Registry. Requirements 4 through 10, are implemented using the new Streaming Analytics Manager.

1.2. Reference Architecture

This reference architecture diagram gives you a general idea of how to build an HDF cluster for your trucking use case. Review this suggested architecture before you begin deployment.



2. Prepare Your Environment

2.1. Deploying Your HDF Clusters

About This Task

Now that you have reviewed the reference architecture and planned the deployment of your trucking application, you can begin installing HDF according to your use case specifications. To fully build the trucking application as described in this *Getting Started with Stream Analytics* document, use the following steps.

Steps

1. Install Ambari 2.5.1.
2. Install HDP 2.6.1 Cluster via Ambari.
3. Install HDF 3.0 Management Pack.
4. Update HDF 3.0 Base URL.
5. Add HDF 3.0 Services to HDP 2.6.1 cluster.

Find instructions for these installation steps in [Installing HDF Services on a New HDP Cluster](#).

More Information

[Planning Your Deployment](#)

2.2. Registering Schemas in Schema Registry

The trucking application streams raw events that are serialized into Avro from the two sensors to its respective Kafka topics. NiFi consumes from these topics, and then routes, enriches, and delivers them to another set of Kafka topics for consumption by the streaming analytics applications. To do this, you must perform the following tasks:

- Creating the 4 Kafka topics
- Registering Schemas for each of the Kafka topics in the Schema Registry

2.2.1. Create the Kafka Topics

About This Task

Kafka topics are categories or feed names to which records are published.

Steps

1. Log into the node where Kafka broker is running.
2. Create the Kafka topics using the following commands:

```
cd /usr/[hdf/]current/kafka-broker/bin/  
  
./kafka-topics.sh \  
--create \  
--zookeeper <zookeeper-host>:2181 \  
--replication-factor 2 \  
--partition 3 \  
--topic raw-truck_events_avro  
  
./kafka-topics.sh \  
--create \  
--zookeeper <zookeeper-host>:2181 \  
--replication-factor 2 \  
--partition 3 \  
--topic raw-truck_speed_events_avro  
  
./kafka-topics.sh \  
--create \  
--zookeeper <zookeeper-host>:2181 \  
--replication-factor 2 \  
--partition 3 \  
--topic truck_events_avro  
  
./kafka-topics.sh \  
--create \  
--zookeeper <zookeeper-host>:2181 \  
--replication-factor 2 \  
--partition 3 \  
--topic truck_speed_events_avro
```

More Information

[Apache Kafka Component Guide](#)

2.2.2. Register Schemas for the Kafka Topics

About This Task

Register the schemas for the 2 Kafka topics that NiFi will consume from and the two other Kafka topics that NiFi will publish the enriched events to. Registering the Kafka topic schemas is beneficial in several ways. Schema Registry provides a centralized schema location, allowing you to stream records into topics without having to attach the schema to each record.

Steps

1. Go to the Schema Registry UI by selecting the Registry service in Ambari and under 'Quick Links' selecting 'Registry UI'
2. Click the "+" button to add a schema, schema group and schema metadata for the Raw Geo Event Sensor Kafka topic:
 - Name = raw-truck_events_avro
 - Description = Raw Geo events from trucks in Kafka Topic

- Type = Avro schema provider
 - Schema Group = truck-sensors-kafka
 - Compatibility: BACKWARD
 - Check the evolve check box
 - Copy the schema from [here](#) and paste it into the Schema Text area.
 - Click Save
3. Click the "+" button to add a schema, schema group (exists from previous step), and schema metadata for the Raw Speed Event Sensor Kafka topic:
- Name = raw-truck_speed_events_avro
 - Description = Raw Speed Events from trucks in Kafka Topic
 - Type = Avro schema provider
 - Schema Group = truck-sensors-kafka
 - Compatibility: BACKWARD
 - Check the evolve check box
 - Copy the schema from [here](#) and paste it into the Schema Text area.
 - Click Save
4. Click the "+" button to add a schema, schema group and schema metadata for the Geo Event Sensor Kafka topic:
- Name = truck_events_avro
 - Description = Schema for the Kafka topic named 'truck_events_avro'
 - Type = Avro schema provider
 - Schema Group = truck-sensors-kafka
 - Compatibility: BACKWARD
 - Check the evolve checkbox
 - Copy the schema from [here](#) and paste it into the Schema Text area.
 - Click Save
5. Click the "+" button to add a schema, schema group (exists from previous step), and schema metadata for the Speed Event Sensor Kafka topic:
- Name = truck_speed_events_avro

- Description = Schema for the Kafka topic named 'truck_speed_events_avro'
- Type = Avro schema provider
- Schema Group = truck-sensors-kafka
- Compatibility: BACKWARD
- Check the evolve check box
- Copy the schema from [here](#) and paste it into the Schema Text area.
- Click **Save**.

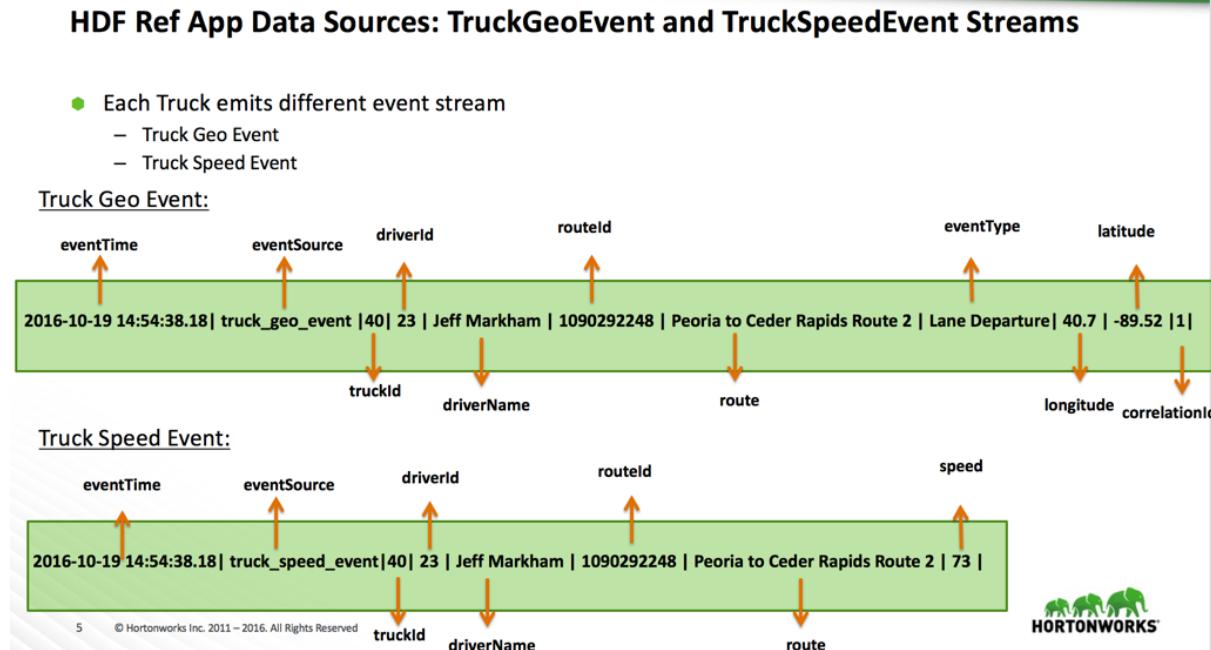
More Information

If you want to create these schemas programmatically using the Schema Registry client via REST rather than through the UI, you can find examples at this [Github location](#).

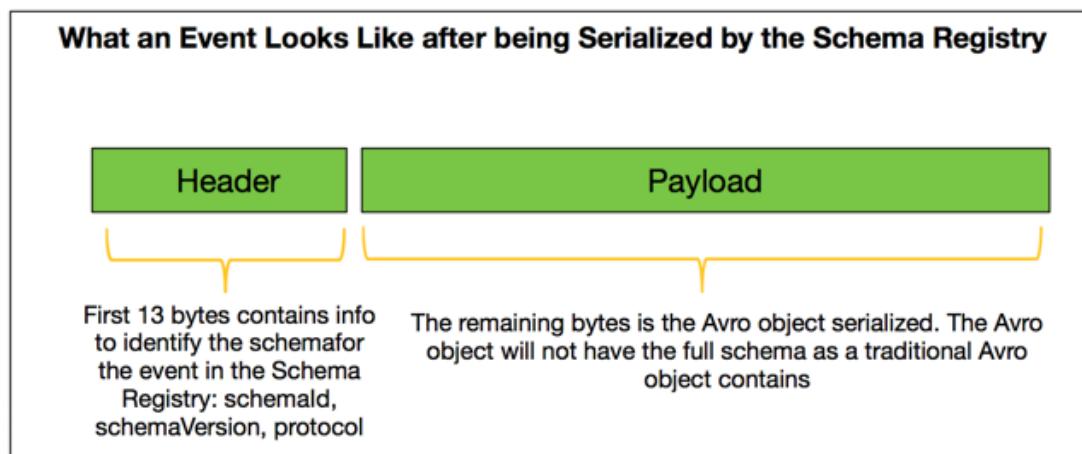
3. Creating a Dataflow Application

3.1. Data Producer Application Generates Events

The following is a sample of a raw truck event stream generated by the sensors.



The date producing application or data simulator publishes these serialized Avro raw events into Kafka topics. The following is what the raw event looks like serialized into Avro using the Schema Registry.



3.2. NiFi: Create a Dataflow Application

To make things easier to setup, import the NiFi Template for this flow by downloading it from this [Github location](#). After importing, select Use Case 1 process group. The below instructions are with respect to that flow.

3.2.1. NiFi Controller Services

Click on Flow Configuration Settings icon and select Controller Services tab.

Hortonworks Schema Registry Controller Service

1. Click on Flow Configuration Settings icon and select **Controller Services** tab.
2. You will see the HWX Schema Registry controller service. Edit the properties to configure the Schema Registry URL based on your environment. You can find this value in the Streaming Analytics Manager Service in Ambari for the configuration property called `registry.url`. An example of what the URL looks similar to `http://$REGISTRY_SERVER:7788/api/v1`.
3. Enable this controller service.

RecordReader and RecordWriter Controller Services

The RecordReader and RecordWriter controller services are new controller services that allows you convert events from one type (json, xml, csv, Avro) to another (json, xml, csv, Avro). These controller services use the Schema Registry to fetch the schema for the event to do this conversion. There are a number of different schema access strategies you can configure on the RecordReader and RecordWriter to tell the Record Reader/Writer how to look up the schema information. For example, if you are reading records serialized by the Hortonworks Schema Registry, the schema identifier required to look up the schema in the registry is embedded in the header of the payload. Hence, the RecordReader would use the schema access strategy called "HDX Content-Encoded Schema Reference". The following are the RecordReader and RecordWriter controller services used for the NiFi template imported:

- Avro Truck Events - Reads Avro events and looks up the schema id via the HDX Content-Encoded Schema Reference strategy. This schema id is then used to query the schema from the Hortonworks Schema Registry.
- CSV Truck Events - Reads csv events and looks up the schema name from the value of the "Schema Name" attribute. This schema lookup strategy is called the "Use 'Schema Name' Property" access strategy. This value of this schema name property is then used to query the schema from the Hortonworks Schema Registry.
- AvroRecordSetWriter - Writes events into Avro and looks up the schema identifier info using the HDX Schema Reference Attribute strategy. This controller also uses a write strategy of HDX Content-Encoded Schema Reference where the Avro object will have schema identifier information pre-appended on the header.
- AvroRecordSetWriter-Read-Schema-From-HDX-Via-Schema-Name - Writes events into Avro and looks up the schema using the Schema Name access strategy. This controller

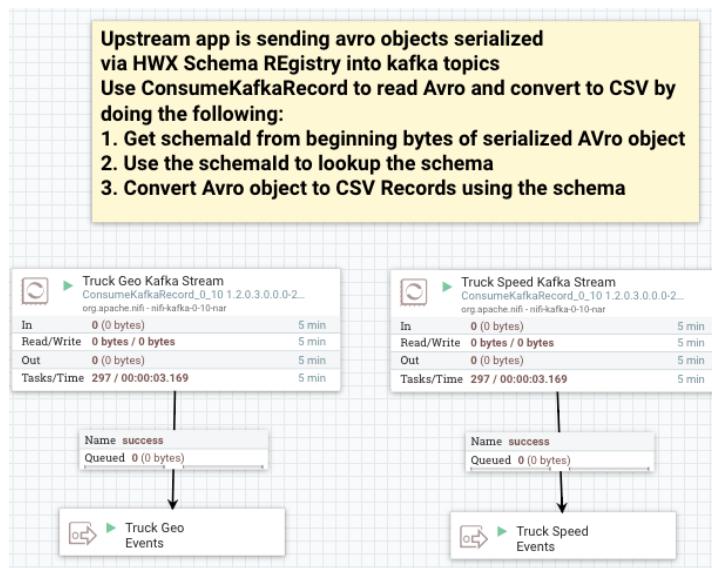
also uses a write strategy of HWX Content-Encoded Schema Reference where the Avro object will have schema identifier information pre-appended on the header.

- CSVRecordSetWriter - Writes events into CSV and looks up the schema identifier using the HWX Schema Reference Attribute strategy. The write schema strategy is also HWX Schema Reference attributes. This means when the csv is written the schema identifier information is stored in named attributes of the flow file.
- CSVRecordSetWriter-Read-Schema-From-HWX-Embedded - Similar to the previous csv writer but the schema identifier is looked using the HWX Content-Encoded Schema Reference strategy.

Enable all of these controller services.

3.2.2. NiFi Ingests the Raw Sensor Events

In the Use Case 1 process group, go into the "Acquire Events" process group. The first step in the NiFi flow is to ingest the raw serialized Avro events from the two Kafka topics. We will use the new ConsumerKafkaRecord processor for this.

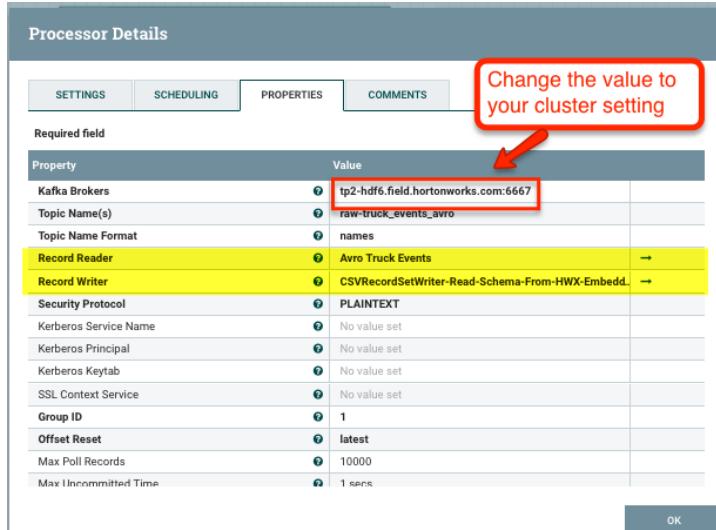


Both ConsumerKafkaRecord processors are configured with an AvroReader controller service and the CSVRecordSetWriter-Read-Schema-From-HWX-Embedded controller service to convert from Avro to CSV using a schema.



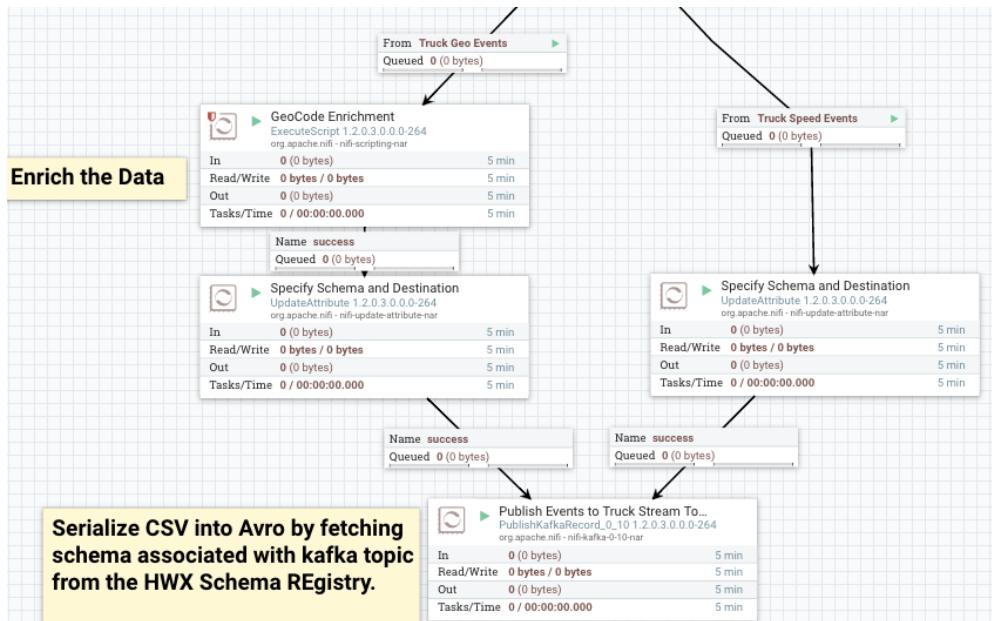
Note

Make sure for both processors, you change the Kafka Brokers property value to your cluster settings.



3.2.3. Publish Enriched Events to Kafka for Consumption by Analytics Applications

After NiFi has done the routing, transforms, and enrichment, NiFi will publish the enriched events into Kafka topics. These topics have a schema registered for it in the Schema Registry and we will store the schema identifier for the schema in the FlowFile attributes (UpdateAttribute processors) and use the PublishKafkaRecord processor to push the events into Kafka.



The PublishKafkaRecord processor is configured with the controller service 'CSV Truck Events' for the Record Reader and uses the AvroRecordSetWriter to write the events into Avro. It is a serialized Avro object with the schema identifier in the header that gets published to Kafka for consumption by SAM.



Note

Make sure for the PublishKafkaRecord, you change the Kafka Brokers property value to your cluster settings.

Processor Details

SETTINGS SCHEDULING PROPERTIES COMMENTS

Required field

Property	Value
Kafka Brokers	tp2-hdf6.field.hortonworks.com:6667
Topic Name	\$(kafka.topic)
Record Reader	CSV Truck Events
Record Writer	AvroRecordSetWriter
Security Protocol	PLAINTEXT
Kerberos Service Name	No value set
Kerberos Principal	No value set
Kerberos Keytab	No value set
SSL Context Service	No value set
Delivery Guarantee	Best Effort
Message Key Field	No value set
Max Request Size	1 MB
Acknowledgment Wait Time	5 secs
Max Metadata Wait Time	5 sec

Change the value to your cluster setting

OK

3.2.4. Start the NiFi Flow

Start the Process Grouped called "Use Case 1".

4. Creating a Stream Analytics Application

1. Create a Service Pool and Environment [12]
2. Create Your First Application [12]
3. Creating and Configuring the Kafka Source Stream [13]
4. Connecting Components [15]
5. Joining Multiple Streams [16]
6. Filtering Events in a Stream using Rules [17]
7. Using Aggregate Functions over Windows [18]
8. Implementing Business Rules on the Stream [19]
9. Transforming Data using a Projection Processor [21]
10. Creating Alerts with Notifications Sink [23]
11. Streaming Alerts to an Analytics Engine for Dashboarding [24]
12. Streaming Violation Events to an Analytics Engine for Descriptive Analytics [25]
13. Streaming Violation Events into a Data Lake and Operational Data Store [27]

4.1. Create a Service Pool and Environment

Before you create an application, you have to create a Service Pool and then an Environment that you associate with an application. Refer to the *Streaming Analytics Manager User Guide* sections on [Streaming Analytics Manager Environment Setup](#) and [Managing Stream Applications](#).

4.2. Create Your First Application

About This Task

The Streaming Analytics Manager provides capabilities to the application developer for building streaming applications. You can go to the Stream Builder UI by select the **Streaming Analytics Manager** service in Ambari and under **Quick Links** select **SAM UI**.

Creating a new stream application requires two steps: clicking the + icon, and then providing a unique name for the stream application and associating the application with an Environment.

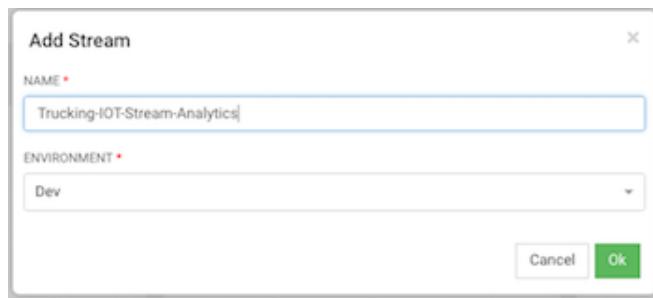
Steps

1. Click the + icon on the **My Applications** dashboard and choose **New Application**.
2. Specify the name of the stream application and the environment that you want it to use stream.



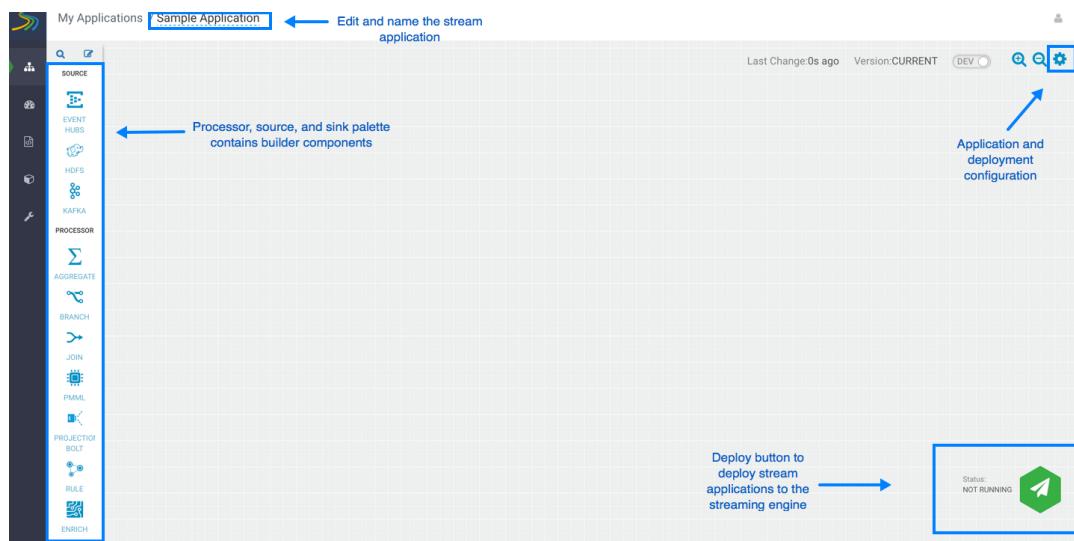
Note

The name of the stream application cannot have any spaces.



Result

SAM displays the Stream Builder canvas. Builder components on the canvas palette are the building blocks used to build stream applications. Now you are ready to start building the streaming application.



4.3. Creating and Configuring the Kafka Source Stream

About This Task

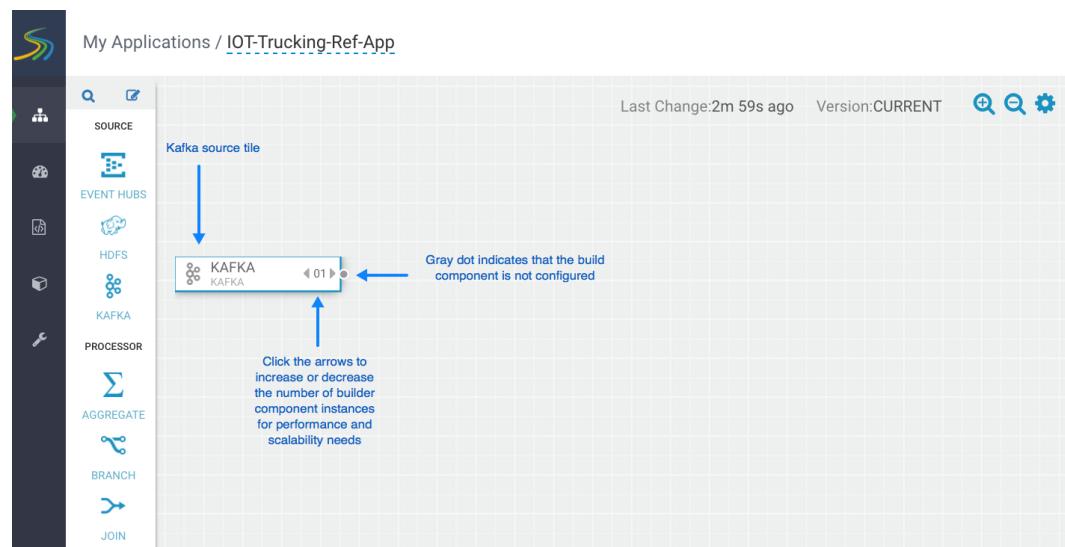
The first step in building a stream application is to drag builder components to the canvas. As described in the *Hortonworks DataFlow Overview*, Stream Builder offers four types of builder components: sources, processors, sinks, and custom components.

Every stream application must start with a source.

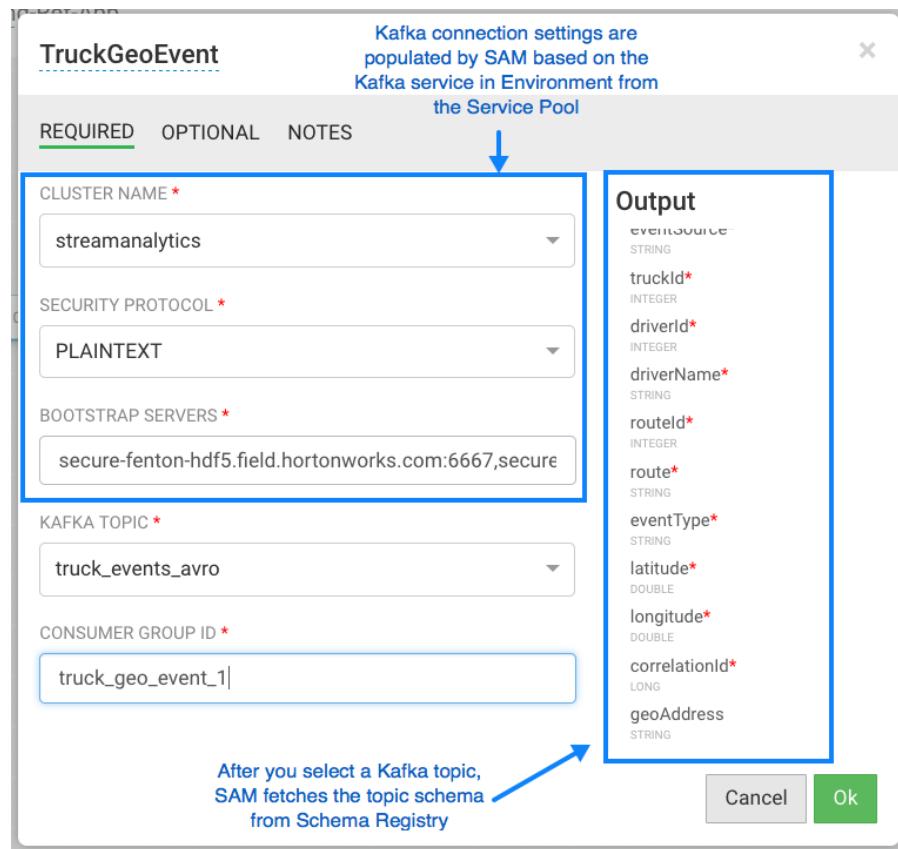
Complete the following instructions to start building a stream application. Use these steps to implement Requirement 4 of the use case.

Steps

1. Drag the Kafka builder component onto the canvas, creating a Kafka tile:



2. Set the number of run-time instances for your Kafka tile component by clicking the up arrow on the tile.
3. Double-click on the tile to begin configuring Kafka. After you specify a Kafka topic name, SAM communicates with the Schema Registry and displays the schema:



Result

Once you have configured your Kafka component correctly, the tile component displays a green dot.

More Information

[Hortonworks DataFlow Overview](#)

4.4. Connecting Components

About This Task

To pass a stream of events from one component to the next, create a connection between the two components. In addition to defining data flow, connections allow you to pass a schema from one component to another.

Steps

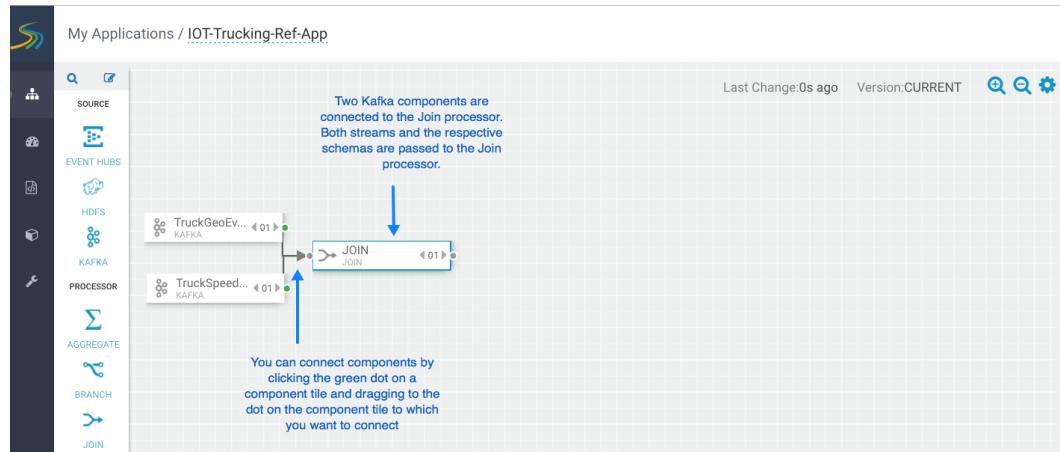
1. Click the green dot to the left of your source component.



2. Drag your cursor to the component tile to which you want to connect.

Example

The following example shows two connections: a connection from Kafka sink TruckGeoStream to the join processor, and a connection from the Kafka sink TruckSpeedStream to the same join processor.



4.5. Joining Multiple Streams

About This Task

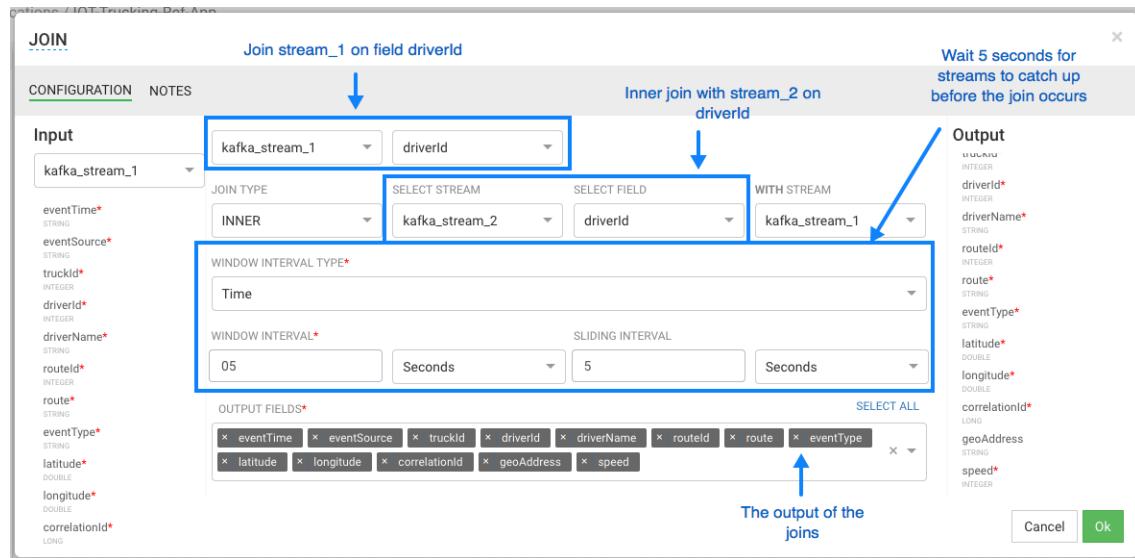
Joining multiple streams is an important SAM capability. You accomplish this by adding the Join processor to your strream application.

This section shows you how to configure a Join processor that joins the truck geo-event stream with the speed event stream, based on Requirement 5 of the use case.

Steps

1. Drag a Join processor onto your canvas and connect it to a source.
2. Double click the Join tile to open the **Configuration** dialog.
3. Configure the Join processors according to the example below.

Example



4.6. Filtering Events in a Stream using Rules

About This Task

SAM provides powerful capabilities to filter events in the stream. It uses a Rules Engine, which translates rules into SQL queries that operate on the stream of data.

The following steps demonstrate this, implementing Requirement 6 of the use case.

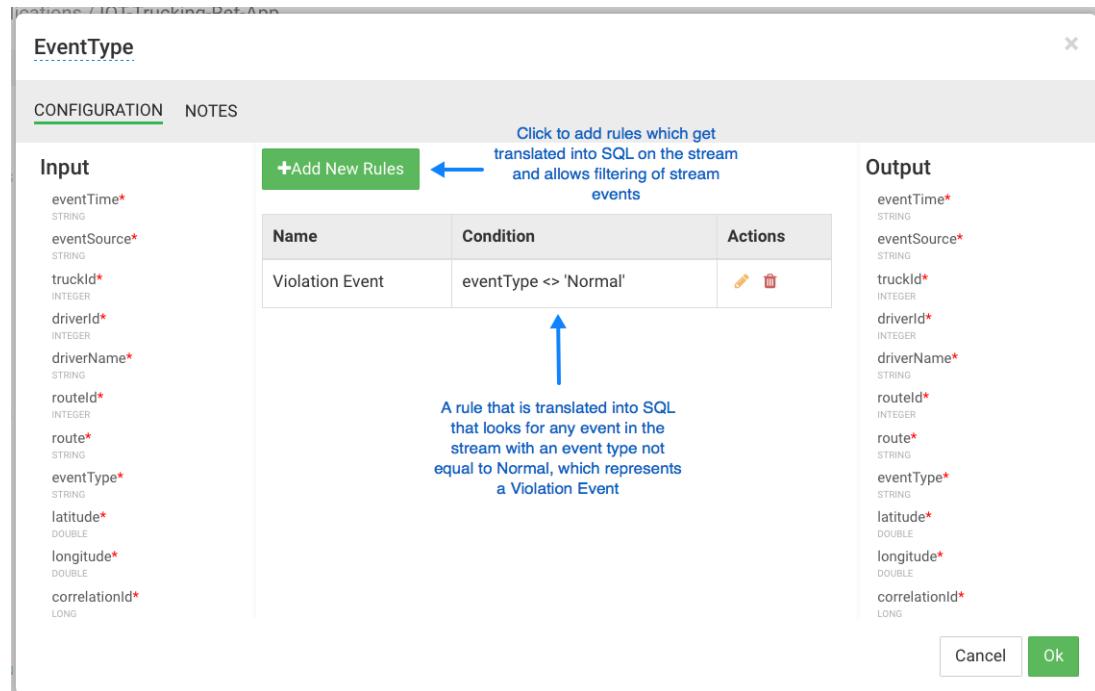
Steps

1. Drag the Rule processor to the canvas and connect it from the Join processors.



2. Double click the Rule processor, click the **Add new Rules** button, and create a new rule.
3. Click **OK** to save the new rule.

Example



4.7. Using Aggregate Functions over Windows

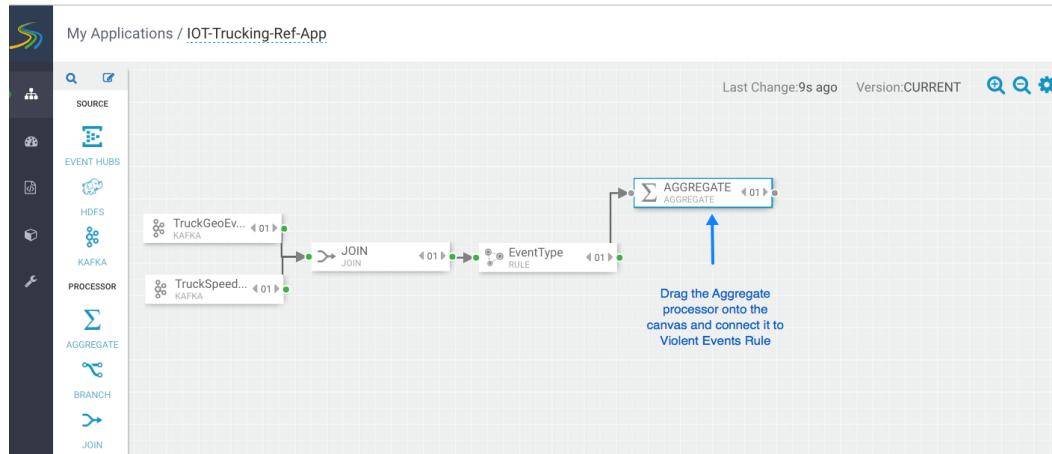
About This Task

Windowing is the ability to split an unbounded stream of data into finite sets based on specified criteria such as time or count, so that you can perform aggregate functions (such as sum or average) on the bounded set of events. SAM exposes these capabilities using the Aggregate processor. The Aggregate processor supports two window types, tumbling and sliding windows. The creation of a window can be based on time or count.

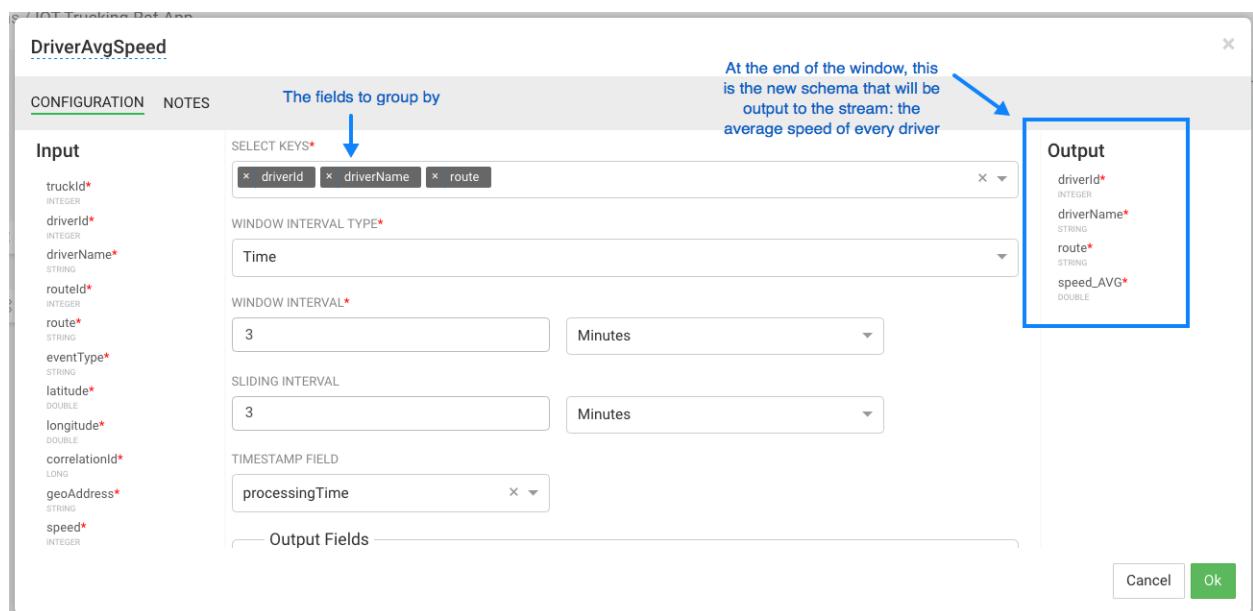
The following images show how to use the Aggregate processor to implement Requirement 8 of the use case.

Steps

1. Drag the Aggregate processor to the canvas and connect it to the Rule processor.



- Double-click on the Aggregate processor, and configure it to calculate the average speed of driver over a three-minute duration.



4.8. Implementing Business Rules on the Stream

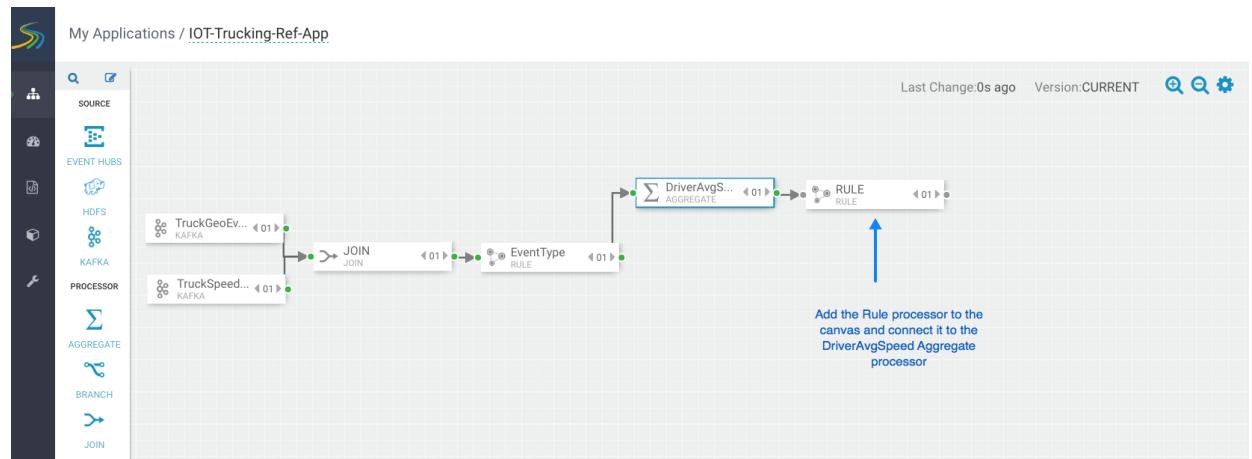
About This Task

This section shows you how to implement the business rule you created above to detect high speeding drivers. "High speed" is defined as greater than 80 miles per hour over a three-minute time window.

This step partially implements Requirement 8 of the use case.

Steps

- Drag the Rule processor onto the canvas and connect to it to the DriverAvgSpeed Aggregate processor:



2. Configure the business rule as follows:

Add New Rule

RULE NAME*
Speeding Driver

DESCRIPTION*
Driver who is speeding excessively

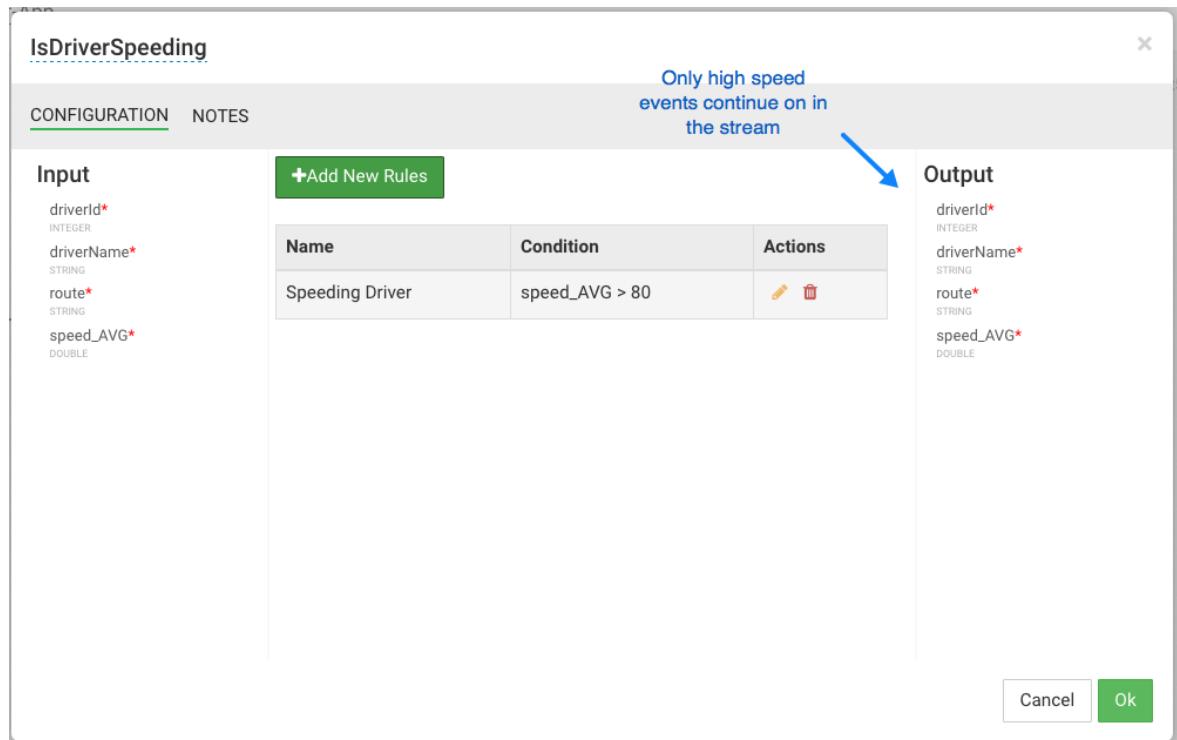
CREATE QUERY*
speed_AVG **GREATER_THAN** 80

QUERY PREVIEW:
speed_AVG > 80

Cancel **Ok**

Result

The fully configured business rule should look similar to the following. Only high speed events continue on in the stream.



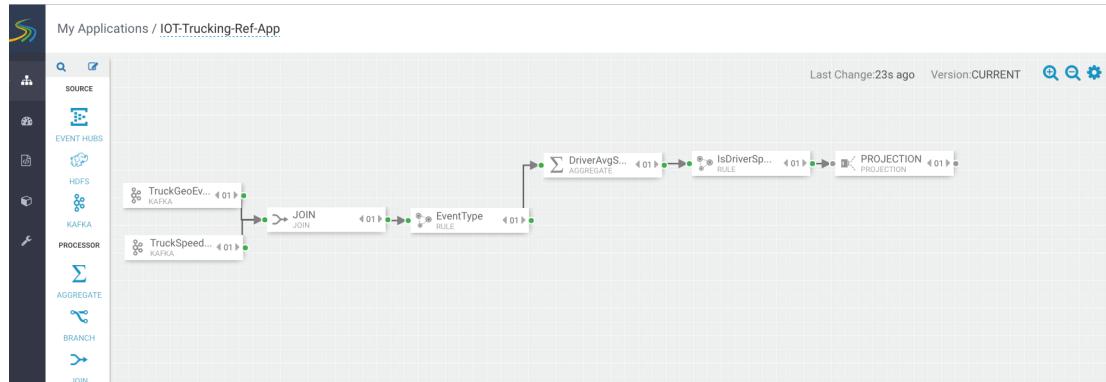
4.9. Transforming Data using a Projection Processor

About This Task

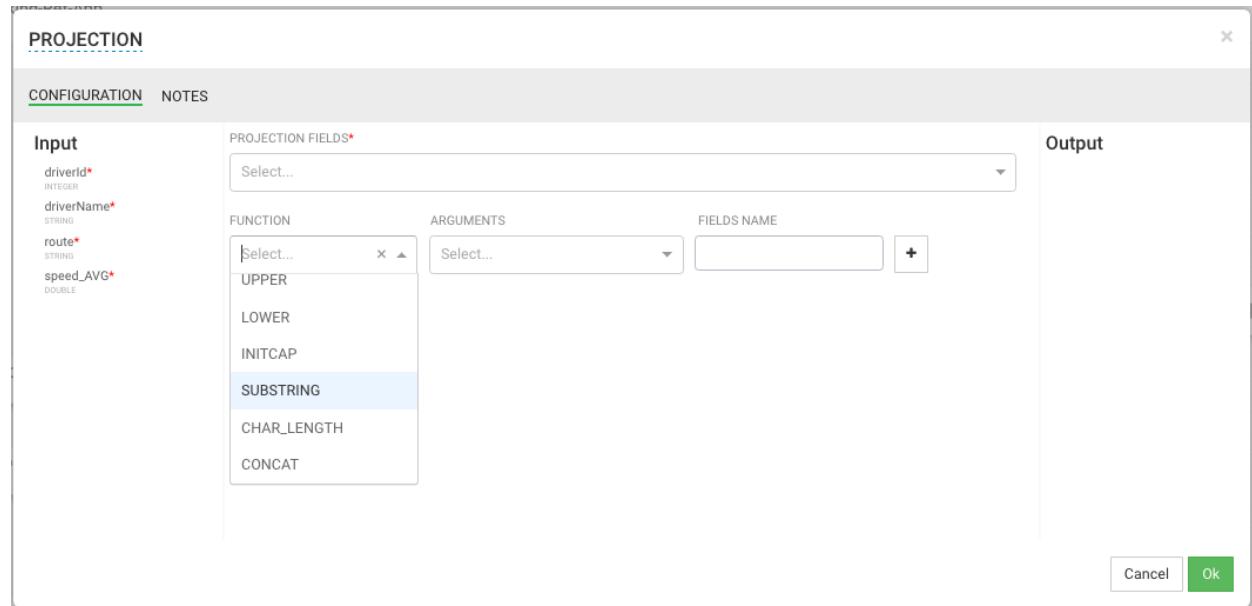
It is common to do transformations on the events in the stream. In our case, before we alert on the speeding driver, we want to convert the average speed we calculated in the aggregate processor into a integer from a double so it is easier to display in the alert. The projection processor allows you to do these transformations.

Steps

1. Drag the Projection processor onto the canvas and connect to it to the IsDriverSpeeding Rule processor:



2. When you double click on the projection processor, you will see a number of out of the box functions however a Round function does not exist.

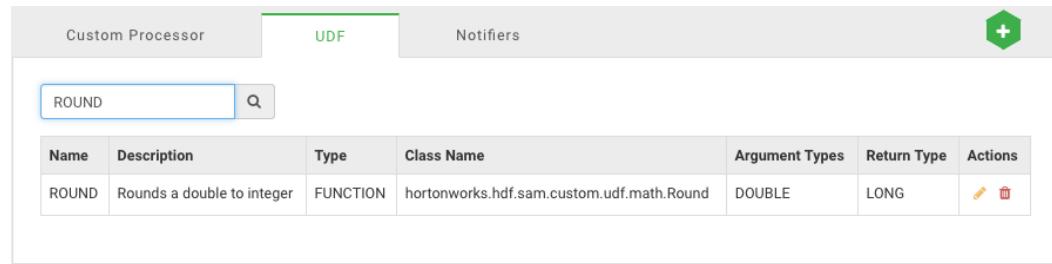


3. Adding UDFs (User Defined Functions) is easy to do within SAM. Follow the below steps to add Round UDF function to SAM.

- From the left-hand menu, click **Configuration**, then **Application Resources**.
- Select the **UDF** tab and click the + sign to create the ROUND UDF. The jar for this UDF can be downloaded from [here](#). The simple java class used to implement this Round function using the SAM SDK can be found [here](#). Unzip the downloaded artifact and use the jar called sam-custom-udf-0.0.5.jar. Configure the UDF with the following values:

NAME *	ROUND
DISPLAY NAME *	ROUND
DESCRIPTION *	Rounds a double to integer
TYPE *	FUNCTION
CLASSNAME *	hortonworks.hdf.sam.custom.udf.math.Round
UDF JAR *	<input type="button" value="Browse"/> sam-custom-udf-0.0.5.jar

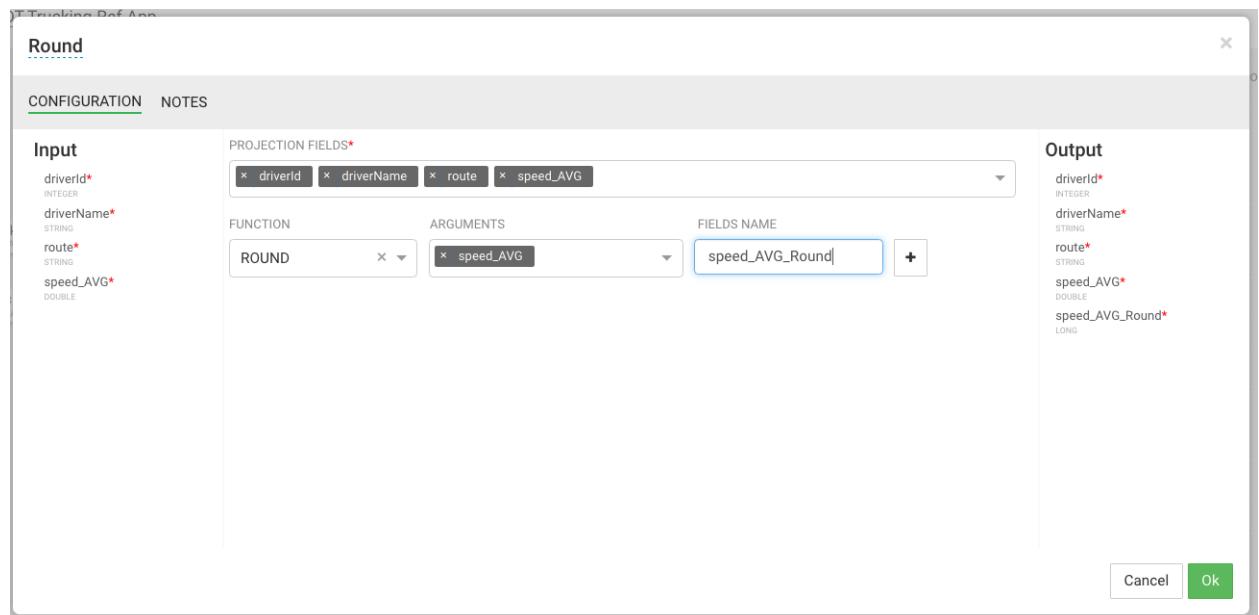
c. After uploading the UDF, you should see the new Round UDF created.



The screenshot shows the 'UDF' tab selected in the top navigation bar. A search bar contains the text 'ROUND'. Below it is a table with columns: Name, Description, Type, Class Name, Argument Types, Return Type, and Actions. One row is visible for the 'ROUND' function, which rounds a double to an integer.

Name	Description	Type	Class Name	Argument Types	Return Type	Actions
ROUND	Rounds a double to integer	FUNCTION	hortonworks.hdf.sam.custom.udf.math.Round	DOUBLE	LONG	

4. After creating the UDF, go back to your Application and double click on the on the Projection Processor you added to the canvas and you see ROUND in the **FUNCTION** drop down list. Configure the ROUND function as the following:



4.10. Creating Alerts with Notifications Sink

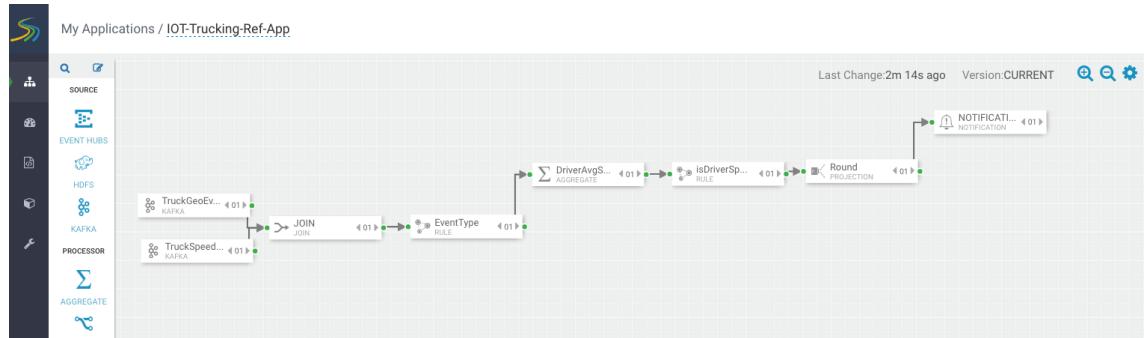
About This Task

The Notifications sink allows you to create alerts. The Notification sink supports email alerts, and it is extensible— you can plug in other types of notifications.

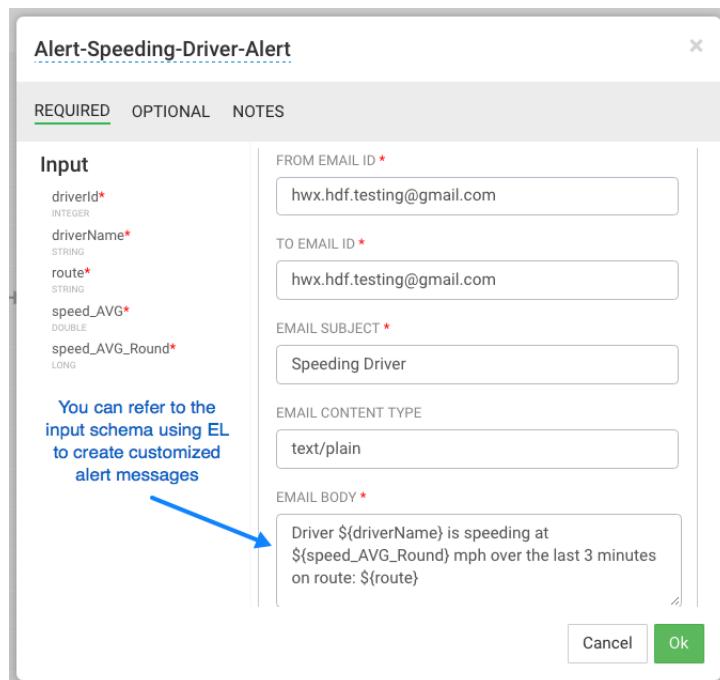
The following steps demonstrate how to create email alerts when drivers are speeding. Like custom UDFs, custom notifications can be added to SAM.

Steps

1. Drag the Notifications sink to the canvas and connect to it to the Round Projection.



2. Configure the Notifications sink to generate email alerts for high speeding drivers.



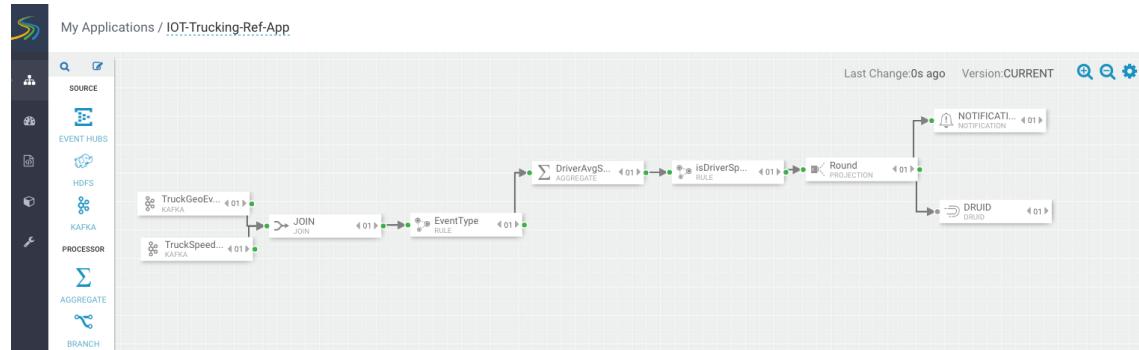
4.11. Streaming Alerts to an Analytics Engine for Dashboarding

About This Task

In addition creating notification alerts, a common use case requirement is to send these alerts to a dashboard so they can be displayed and visualized. SAM offers this capability by allowing you to stream data into DRUID and then using Superset to create dashboards and visualizations.

Steps

1. Drag the Druid sink to the canvas and connect to it to the Round Projection.



2. Stream these events into a Druid cube called **alerts-speeding-drivers-cube** by configuring the Druid processor like the following.

Alert-Speeding-Driver-Cube

REQUIRED	OPTIONAL	NOTES
Input	Datasource Name * <input type="text" value="alerts-speeding-drivers-cube"/> Zookeeper Connect String * <input type="text" value="secure-fenton-hdf1.field.hortonworks.com:2181,secure-f"/>	
	Dimensions * <div style="border: 1px solid #ccc; padding: 5px; display: inline-block;"> <input type="button" value="x driverId"/> <input type="button" value="x driverName"/> <input type="button" value="x route"/> <input type="button" value="x speed_AVG"/> <input type="button" value="x speed_AVG.Round"/> </div>	
	Timestamp Field Name * <input type="text" value="processingTime"/>	
	Window Period * <input type="text" value="PT10M"/>	
		<input type="button" value="Cancel"/> <input type="button" value="Ok"/>

3. In the **Creating Visualization Section**, describe how to create dashboards for the **alerts-speeding-drivers-cube**.

4.12. Streaming Violation Events to an Analytics Engine for Descriptive Analytics

About This Task

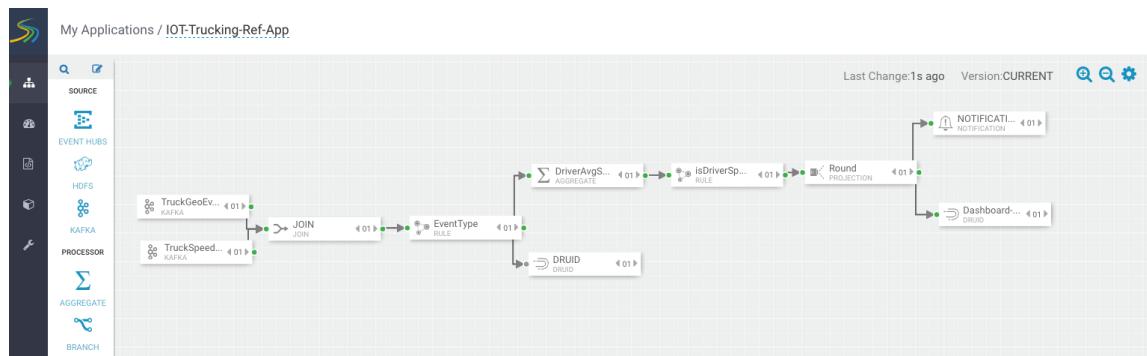
Now lets implement Requirement 7:

All infraction events need to be available for descriptive analytic (dash-boarding, visualizations, etc.) by a business analyst. The analyst needs the ability to do analysis on the streaming data.

The analytics engine in SAM is powered by Druid. The following steps show how to stream data into Druid, so that a business analyst can use the Stream Insight Superset module to generate descriptive analytics.

Steps

1. Drag the Druid processor to the canvas and connect it to the ViolationEvents Rule processor.



2. Configure the Druid processor. You can edit the ZooKeeper connect string in the advanced section of the Druid Service in Ambari, under the property `druid.zk.service.host`.

Violation-Events-Cube

The name of the insight data source/cube to which you want to stream data.

Business analysts use these data sources to query the data

REQUIRED **OPTIONAL** **NOTES**

Input

eventTime*	DATASOURCE NAME *
eventSource*	violation-events-cube
truckId*	ZOOKEEPER CONNECT STRING *
driverId*	secure-fenton-hdf1.field.hortonworks.com:2181,secure-f
driverName*	DIMENSIONS *
routeId*	x eventTime x eventSource x truckId
route*	x driverId x driverName x routeId x route
eventType*	x eventType x latitude x longitude
latitude*	x correlationId x geoAddress x speed
longitude*	
correlationId*	

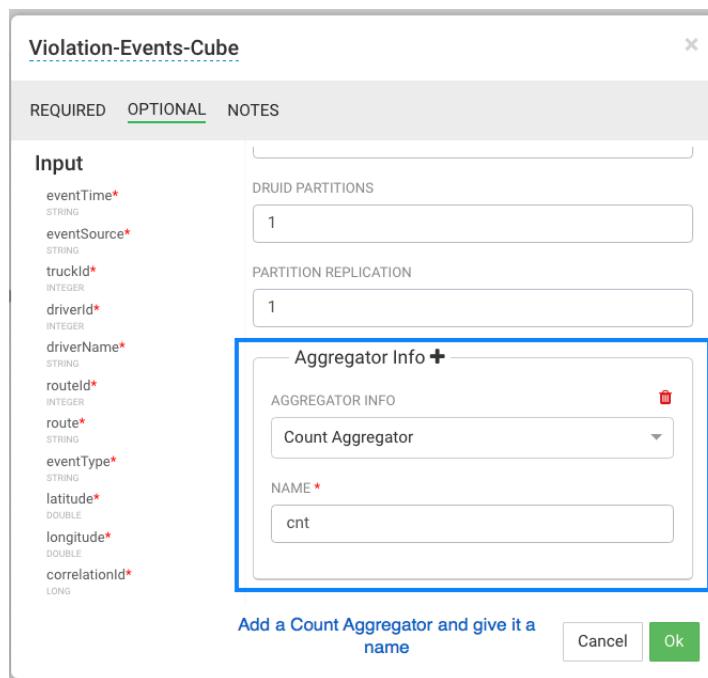
TIMESTAMP FIELD NAME *

processingTime

WINDOW PERIOD *

Cancel **Ok**

3. Configure the **Aggregator Info** settings, under the **OPTIONAL** menu



4.13. Streaming Violation Events into a Data Lake and Operational Data Store

About This Task

Another common requirement is to stream data into an operational data store like HBase to power real-time web apps as well as a data lake powered by HDFS for long term storage and batch etl and analytic.

Steps

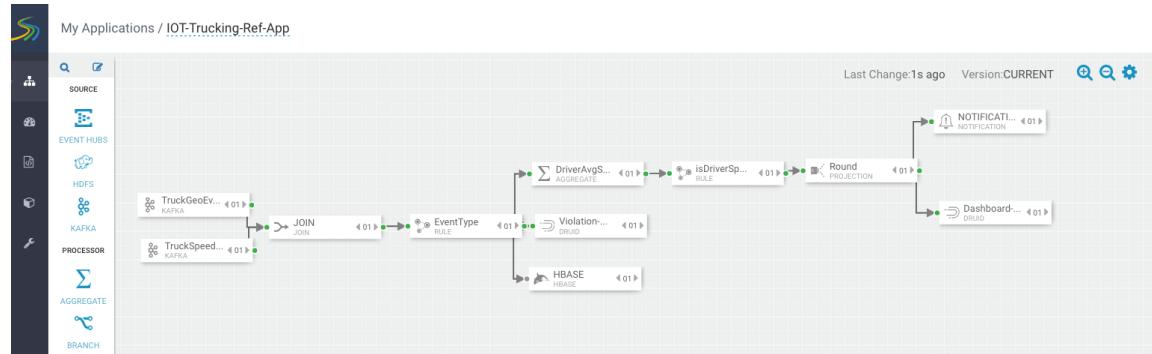
1. You will need ot have HBase service running. This can be easily done by adding the HDP HBase Service via Ambari. Create a new HBase table by logging into an node where Hbase client is installed then execute the below commands

```
cd /usr/hdp/current/hbase-client/bin
/hbase shell
create 'violation_events', {NAME=> 'events', VERSIONS => 3} ;
```

2. Create the following directory in HDFS and give it access to all users. Log into a node where HDFS client is installed and execute the below commands

```
su hdfs
hadoop fs -mkdir /apps/trucking-app
hadoop fs -chmod 777 /apps/trucking-app
```

3. Drag the HBase sink to the canvas and connect it to the ViolationEvents Rule processor.

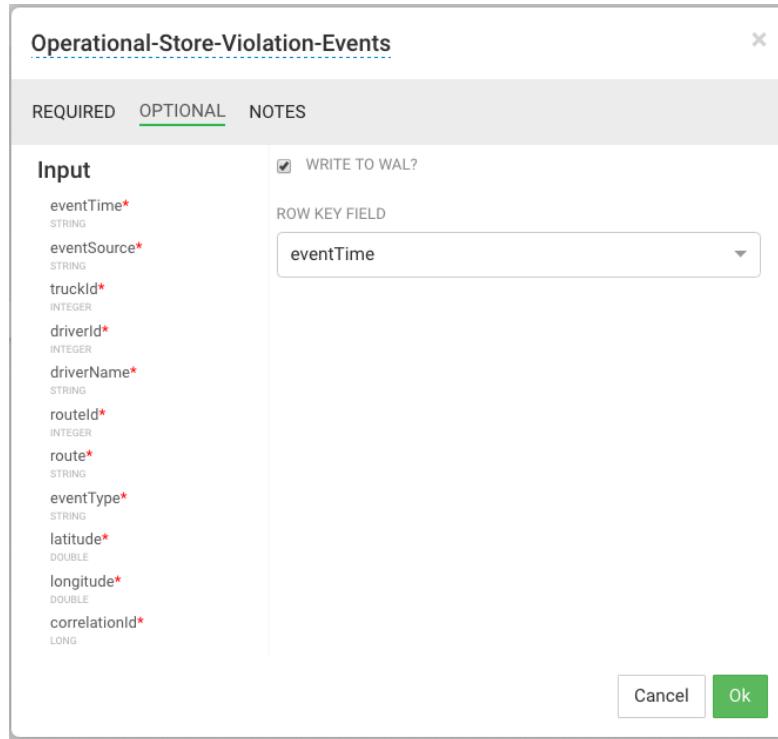


4. Configure the Hbase Sink as below.

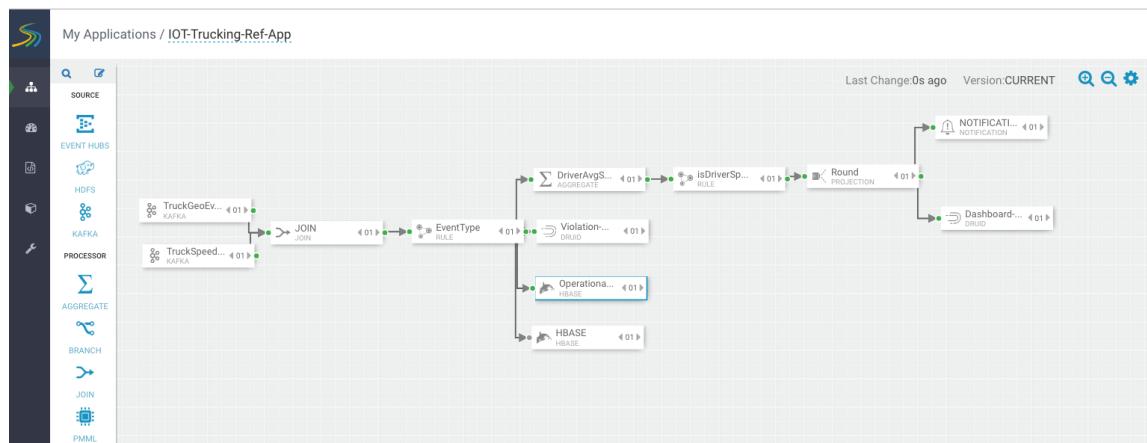
Operational-Store-Violation-Events

REQUIRED	OPTIONAL	NOTES
Input		
eventTime*	STRING	CLUSTER NAME *
eventSource*	STRING	streamanalytics
truckId*	INTEGER	HBASE TABLE *
driverId*	INTEGER	default:violation_events
driverName*	STRING	COLUMN FAMILY *
routeld*	INTEGER	events
route*	STRING	BATCH SIZE *
eventType*	STRING	100
latitude*	DOUBLE	
longitude*	DOUBLE	
correlationId*	LONG	

Cancel **Ok**



5. Drag the HDFS sink to the canvas and connect it to the ViolationEvents Rule processor.



6. Configure HDFS as below. Make sure you have permissions to write into the directory you have configured for HDFS path.

Data-Lake-HDFS

REQUIRED **OPTIONAL** **NOTES**

Input

eventTime*	STRING
eventSource*	STRING
truckId*	INTEGER
driverId*	INTEGER
driverName*	STRING
routId*	INTEGER
route*	STRING
eventType*	STRING
latitude*	DOUBLE
longitude*	DOUBLE
correlationId*	LONG

PATH *
/apps/trucking-app

FLUSH COUNT *
1000

ROTATION POLICY
Time Based Rotation

ROTATION INTERVAL MULTIPLIER *
3

ROTATION INTERVAL UNIT *
MINUTES

OUTPUT FIELDS *

Cancel **Ok**

5. Deploy an Application

5.1. Configure Deployment Settings

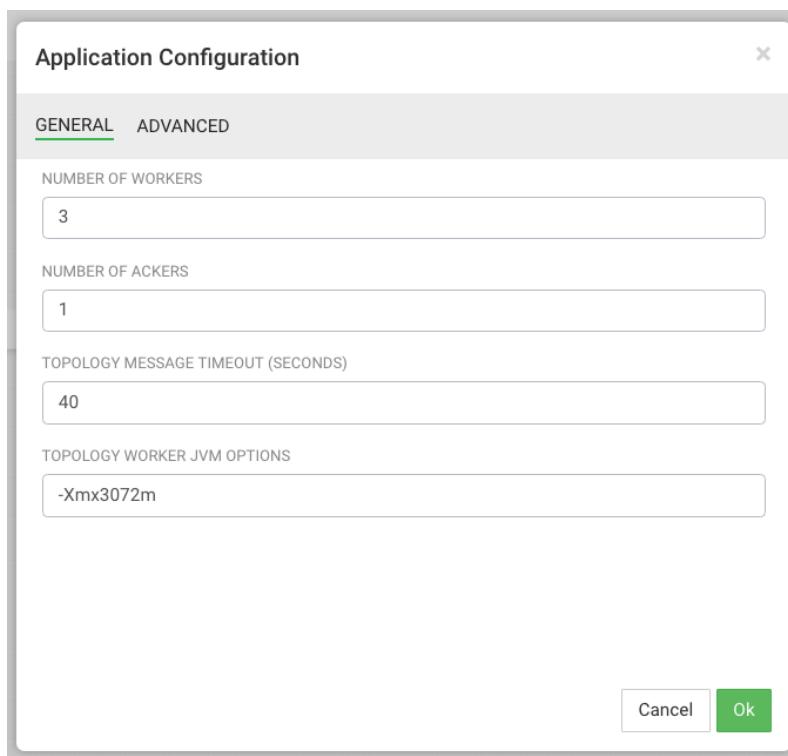
About This Task

Before deploying the application, you must configure deployment settings such as JVM size, number of ackers, and number of workers. Because this topology uses a number of joins and windows, you should increase the JVM heap size for the workers.

Steps

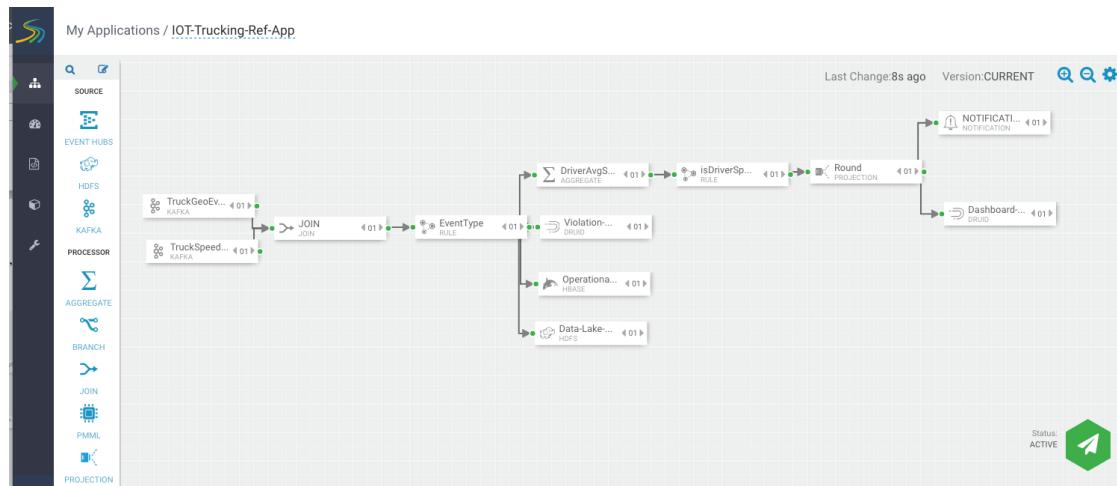
1. Click the gear icon on the top right corner of the canvas to display the **Application Configuration** dialog.
2. Increase **Number of Workers** to 5.
3. Set **Topology Worker JVM Options** to `-Xmx3072m`.

Example



5.2. Deploy the App

After you have configured the application's deployment settings, click the **Deploy** button on the lower right of the canvas.



During the deployment process, Streaming Analytics Manager completes the following tasks:

1. Construct the configurations for the different big data services used in the stream app.
2. Create a deployable jar of the streaming app.
3. Upload and deploy the app jar to streaming engine server.

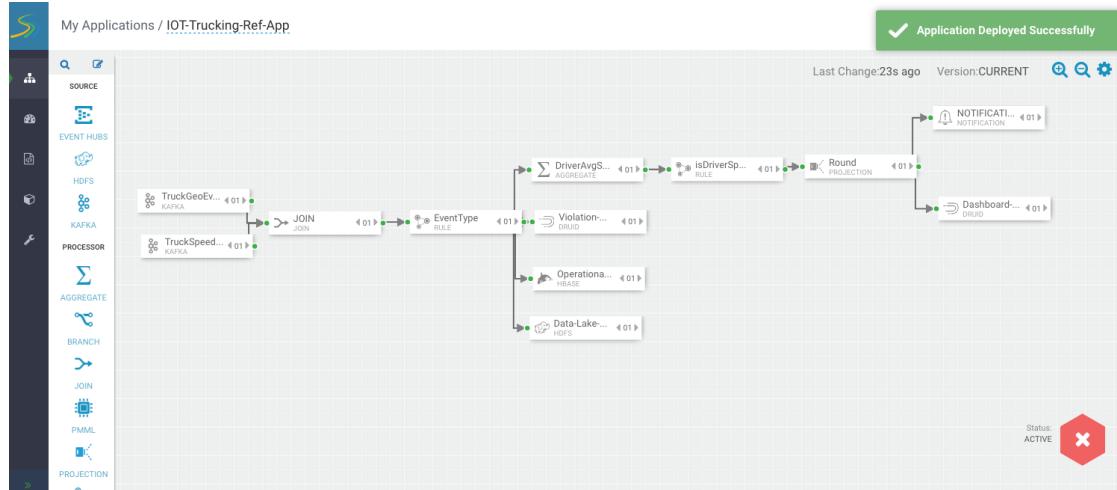
As SAM works through these tasks, it displays a progress bar.



Building Application Jars

The stream application is deployed to a Storm cluster based on the Storm Service defined in the Environment associated with the application.

After the application has been deployed successfully, SAM notifies you and updates the button to red to indicate it is deployed. Click the red button to kill/undeploy the app.



5.3. Running the Stream Simulator

Now that you have developed and deployed the NiFi Flow Application and the Stream Analytics Application, we are ready to run a data simulator that generates truck geo events and sensor events for the apps to process.

To generate the raw truck events serialized into Avro objects using the Schema registry and publish them into the raw Kafka topics, do the following:

1. Download the [Data-Loader](#) Unzip it and copy it to the node the cluster. Lets call the directory you unzip it to as: \$DATA_LOADER_HOME. Then execute the following. Make sure to replace variables below with your environment specific values (you can find the REST URL to schema registry in Ambari under SAM service for config value registry.url) . Make sure java (jdk 1.8) is on your classpath.

```
tar -zxvf $DATA_LOADER_HOME/routes.tar.gz

nohup java -cp \
stream-simulator-jar-with-dependencies.jar \
hortonworks.hdp.refapp.trucking.simulator.
SimulationRegistrySerializerRunnerApp \
20000 \
hortonworks.hdp.refapp.trucking.simulator.impl.domain.transport.Truck \
hortonworks.hdp.refapp.trucking.simulator.impl.collectors.
KafkaEventSerializedWithRegistryCollector \
1 \
$DATA_LOADER_HOME/routes/midwest/ \
10000 \
$KAFKA_BROKER_HOST:$KAFKA_PORT \
$REST_URL_TO_SCHEMA_REGISTRY \
ALL_STREAMS \
NONSECURE &
```

2. You should see events being published into the Kafka topics called: raw-truck_events_avro and raw-truck_speed_events_avro, NiFi should be consuming those, enriching them and then pushing into the truck_events_avro and truck_speed_events_avro kafka topics and then SAM consumes from those topics.

6. Stream Operations

The Stream Operation view provides management of the stream applications, including the following:

- Application life cycle management: start, stop, edit, delete
- Application performance metrics
- Troubleshooting, debugging
- Exporting and importing applications

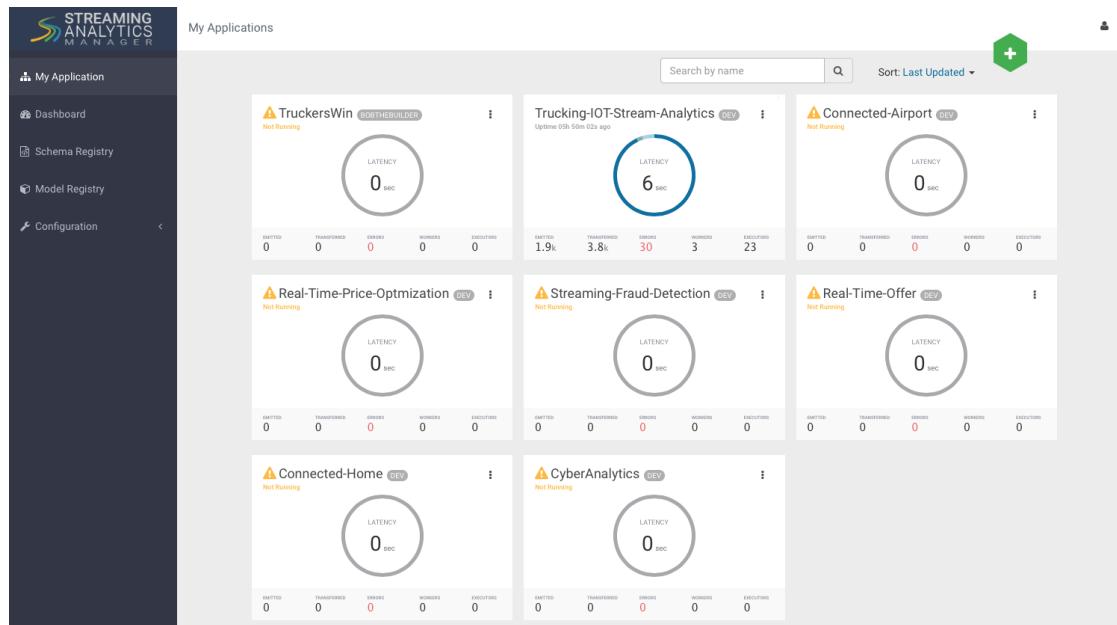
6.1. My Applications View

Once a stream application has been deployed, the Stream Operations displays operational views of the application.

One of these views is called **My Application** dashboard.

To access the application dashboard in SAM, click **My Application** tab (the hierarchy icon). The dashboard displays all applications built using Streaming Analytics Manager:

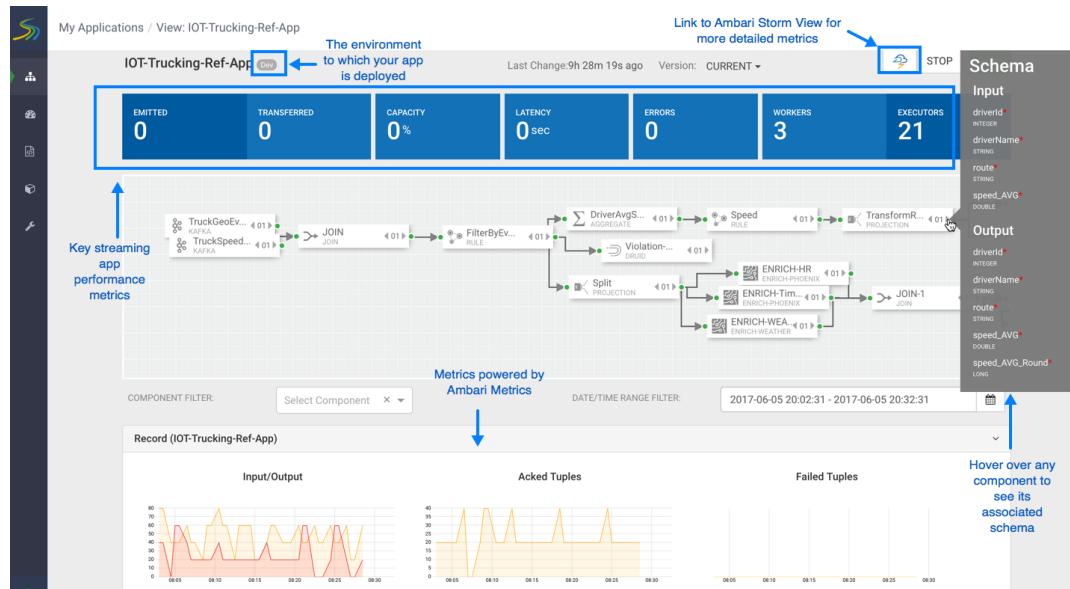
Each stream application is represented by an application tile. Hovering over the application tile provides status, metrics, and actions you can perform on the stream application.



6.2. Application Performance Monitoring

To view application performance metrics (APM) for the application, clicking on the application name on the application tile.

The following diagram describes elements of the APM view.



6.3. Troubleshooting and Debugging a Stream application

At the top right corner of the APM, there is a Storm icon that takes you to the Storm Ambari view.

The Storm Ambari View provides the following capabilities for deeper troubleshooting and debugging:

- Topology View and Metrics: shows a visual representation of the deployed topology and topology level Metrics.
- Distributed Log Search: allows users to search all logs across supervisor machines for a topology; results can include zipped logs.
- Dynamic Log Levels: allows Users and Administrators to dynamically change the log level settings for a running topology.
- Topology Event Inspector: allows viewing of tuples flowing through the topology along with the ability to turn on/off debug events without having to stop/restart the entire topology.
- Dynamic Worker Profiling: allows users to request worker profile data directly from the Storm UI (Heap Dumps, JStack Output, JProfile).

Use the first portion of the Ambari Storm View to review the topology summary and statistics, set event profiling, search logs, and dynamically change them.

TOPOLOGY SUMMARY

- ID: streamline-1-IOT-Trucking-Ref-App-3-1496685847
- Owner: storm
- Status: ACTIVE
- Uptime: 9h 42m 9s
- Workers: 3
- Executors: 21
- Tasks: 21
- Memory: 9408
- Worker-Host:Port: secure-sam-hdf5.field.hortonworks.com:6700, secure-sam-hdf6.field.hortonworks.com:6701, secure-sam-hdf6.field.hortonworks.com:6700

TOPOLOGY STATS

Window	Emitted	Transferred	Complete Latency (ms)	Acked	Failed
10m 0s	0	0	0	0	Change log levels dynamically
3h 0m 0s	0	0	0	0	
1d 0h 0m 0s	15360	17960	36160.785	3100	
All time	15360	17960	36160.785	3100	

Distributed log search across all logs → Search in Logs

Window: All time System Summary Debug OFF Turn on event profiling

Scroll down to review the deployed topology and see metrics about its components.

The streaming application that SAM deployed to the Storm engine

Kafka Spout Lag					
Id	Topic	Partition	Latest Offset	Spout Committed Offset	Lag
No Data Found.					

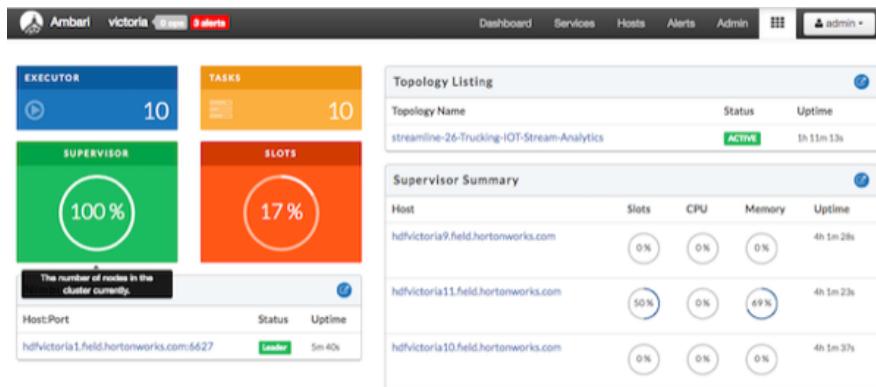
More detailed metrics on the individual components (source, sink, and processor) in the deployed application

Spouts										
Search by id										
Id	Executors	Tasks	Emitted	Transferred	Complete Latency (ms)	Acked	Failed	Error Host:Port	Last Error	Error Time
1-TruckGeoEvent	1	1	1340	1340	30596.000	1560	0			
2-TruckSpeedEvent	1	1	1380	1380	41797.844	1540	0			

Showing 1 to 2 of 2 entries.

6.3.1. Streaming Engine Infrastructure Metrics

The following dashboard shows infrastructure metrics for the streaming engine used; in this case, it shows details about the Storm cluster.



6.3.2. Changing Log Levels Dynamically and with Expiration Policies

When debugging a stream application, the ability to change the log dynamically is a powerful troubleshooting feature. However, since typical stream applications handle millions of events per second, changes to log levels can impact performance unless safeguards such as expiration policies are defined. The following diagram shows how to change log levels with expiration policies.

This screenshot shows the 'Change Log Level' interface within the Ambari topology listing. It includes fields for 'Logger' (set to 'com.your.organization.LoggerName'), 'Level' (set to 'ALL'), 'Timeout' (set to '30'), 'Expires At' (set to '30'), and an 'Action' button. A blue arrow points to the 'Logger' field with the text 'Click the log icon to configure the logger'.

6.3.3. Distributed Log Search

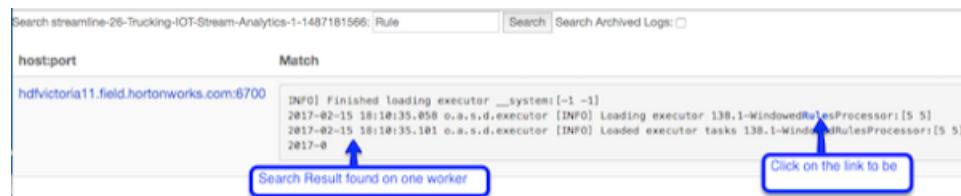
Storm is a distributed streaming engine, which means that many worker nodes can be used to power the streaming application. Because it has a distributed architecture, logs are distributed across the cluster on many worker nodes. Searching for log data across workers can be a painful process. With distributed log search, however, you can search across all logs located across all worker nodes.

The following steps describe how to use distributed log search.

1. Type your search string in the distributed log search text box:

This screenshot shows the distributed log search interface. It features a search bar at the top with the placeholder 'Type what you want to search across logs, and configure whether you want to also search archived logs'. Below the search bar are buttons for 'Window' (set to 'All time'), 'System Summary' (set to 'OFF'), 'Debug' (set to 'OFF'), and a search icon. To the right, there are checkboxes for 'Search archived logs' and 'Deep Search'. The bottom part of the interface is identical to the 'Change Log Level' screen shown above, with fields for 'Logger', 'Level', 'Timeout', 'Expires At', and an 'Action' button.

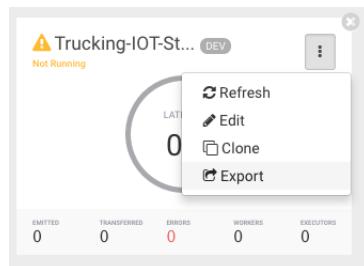
2. Review the results.
3. Click on the link to navigate to the exact location in the log file.



6.4. Exporting and Importing Stream applications

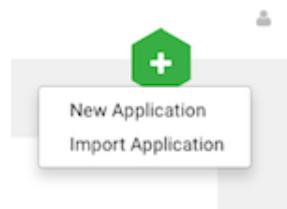
Service pool and environment abstractions combined with import and export capabilities allow you to move a stream application from one environment to another easily.

To export a stream application, click the Export icon on the **My Application** dashboard. This downloads a JSON file that represents your streaming application.

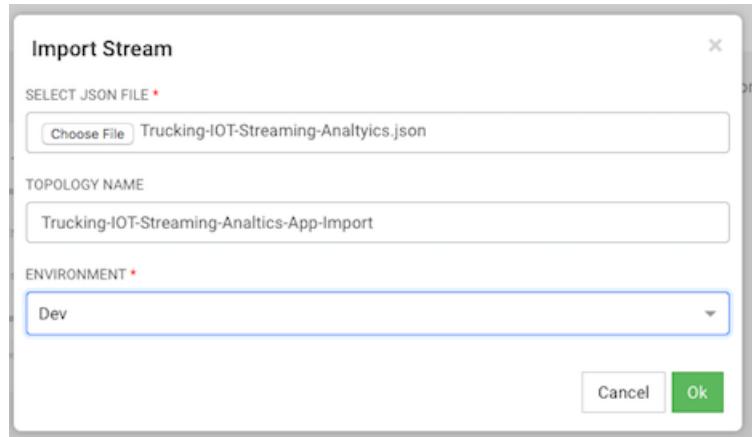


To import a stream application that was exported in JSON format:

1. Click on the + icon in **My Applications** View and select import application:



2. Select the JSON file that you want to import, provide a unique name for the application and specify which environment to use.

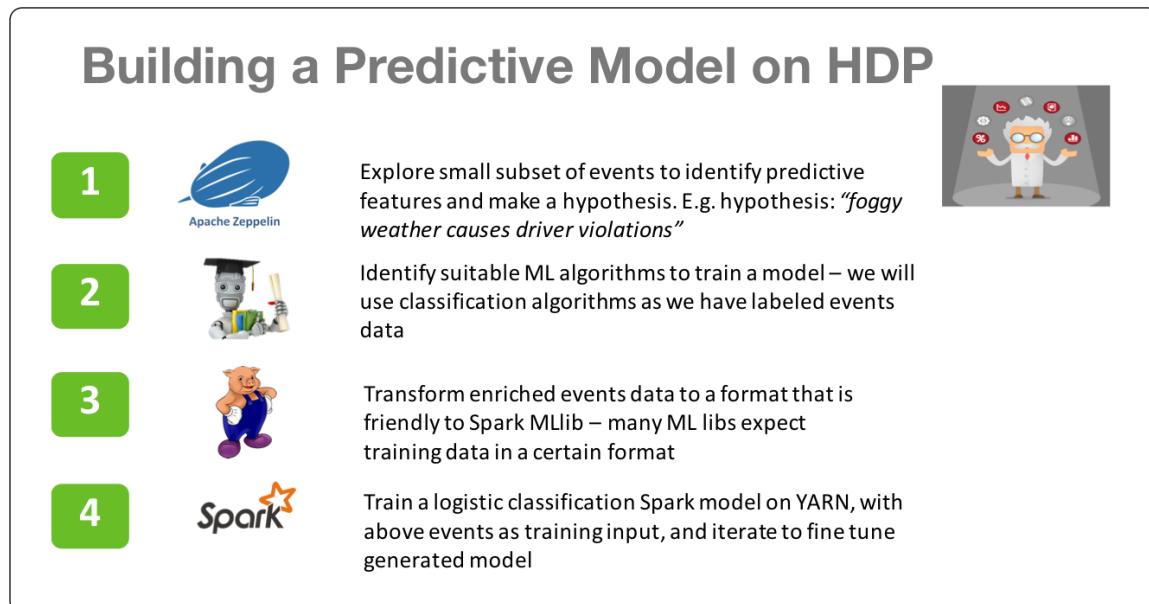


7. Advanced: Doing Predictive Analytics on the Stream

Requirement 10 of this use case states the following:

Execute a logistical regression Spark ML model on the events in the stream to predict if a driver is going to commit a violation. If violation is predicted, then alert on it.

HDP, the Hortonworks data at rest platform provides powerful set of tools for data engineers and scientists to build powerful analytics with data processing engines like Spark Streaming, Hive and Pig. The below diagram illustrates a typical analytics life cycle in HDP.



Once the model has been trained and optimized, insights can be created by scoring the model in real-time as events are coming in. The next set of steps in the life cycle has to do with scoring the model in real-time using HDF components.

Scoring a Predictive Model on HDF

- | | |
|---|---|
| 5

Model Registry | Export the Spark Mllib model and import into the HDF's Model Registry |
| 6

Enrich with Features | Use SAM's enrich/custom processors to enrich the event with the features required for the model |
| 7

Transform/Normalize | Use SAM's projection/custom processors to transform/normalize the streaming event and the features required for the model |
| 8

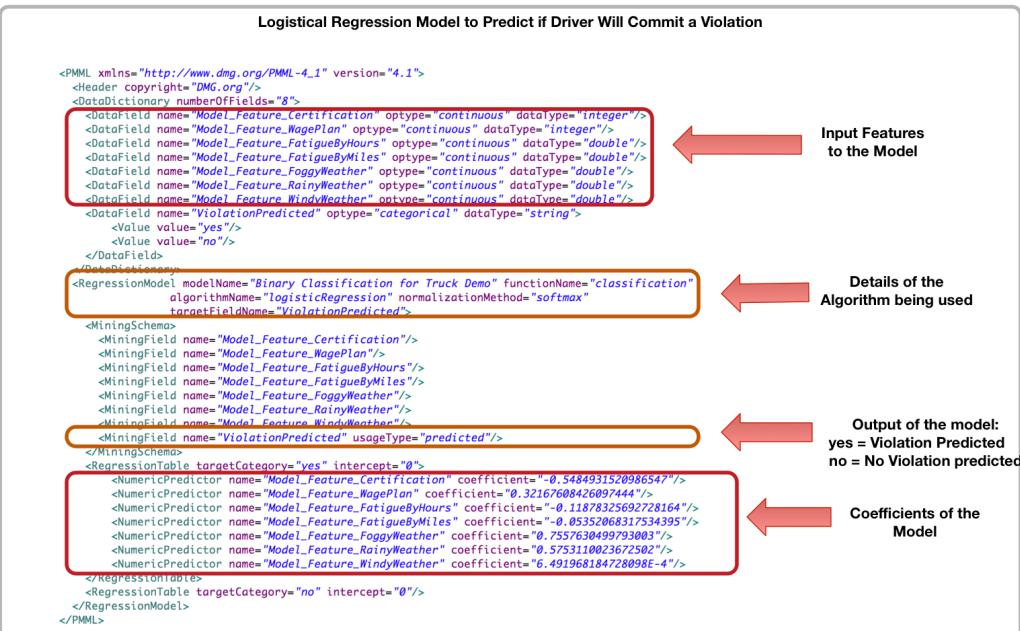
Score Model | Use SAM's PMML processor to score the model for each stream event with its required features |
| 9

Alert / Notify / Action | Use SAM's rule and notification processors to alert, notify and take action using the results of the model |

In the next few sections we will walk through how to do steps 5 through 9 in SAM.

7.1. Logistical Regression Model

With steps 1-4 with HDP, we were able to build a logistical regression model. The model was then exported into PMML. The below diagram illustrates the features, coefficients and output of the model.



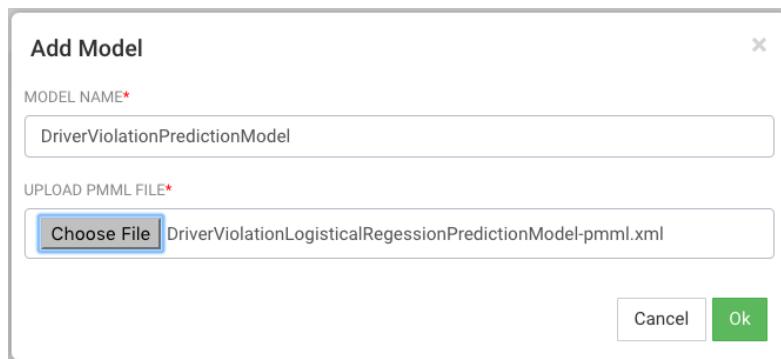
7.2. Export the Model into SAM's Model Registry

About This Task

SAM provides a registry where you can store PMML models. To get started with predictive analytics, upload this logistical regression model.

Steps

1. Download this [PMML model](#) and save it locally with an `.xml` extension.
2. Select the **Model Registry** menu item.
3. Click the `+` icon.
4. Give your PMML model a name.
5. From **Upload PMML File**, select the PMML file you just downloaded.

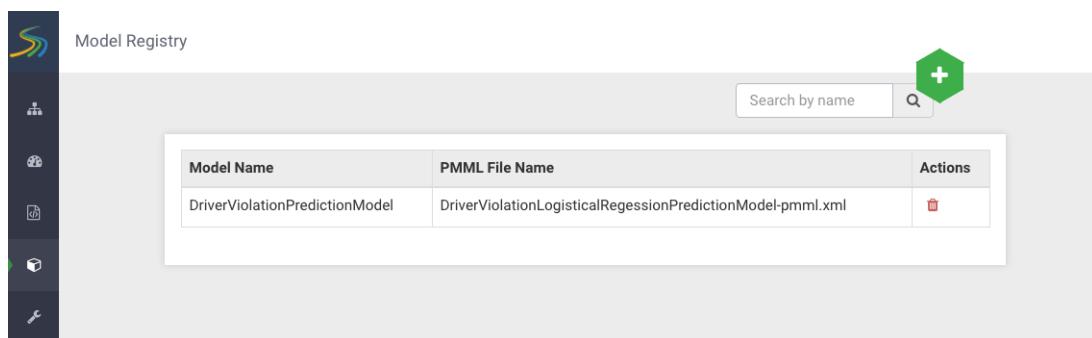


The screenshot shows the 'Add Model' dialog box. It has two main input fields: 'MODEL NAME*' which contains the value 'DriverViolationPredictionModel', and 'UPLOAD PMML FILE*' which contains the path 'DriverViolationLogisticalRegressionPredictionModel-pmml.xml'. Below these fields are two buttons: 'Cancel' and 'Ok'. The 'Ok' button is highlighted with a green background.

6. Click **Ok**.

Result

The model is saved in the **Model Registry**.



The screenshot shows the Model Registry interface. On the left is a sidebar with icons for Home, Model Registry, Predictions, and Help. The main area is titled 'Model Registry' and contains a table with one row. The table has three columns: 'Model Name' (DriverViolationPredictionModel), 'PMML File Name' (DriverViolationLogisticalRegressionPredictionModel-pmml.xml), and 'Actions' (with a red trash icon). At the top right of the main area are a search bar and a green '+' button.

7.3. Enrichment and Normalization of Model Features

Now that the model has been added to the model registry, you can use it in the streaming application by the PMML processor. Before the model can be executed, you must enrich and normalize the streaming events with the features required by the model. As the above diagram illustrates, there are 7 features in the model. None of these features come as part of the stream from the two sensors. So, based on the driverId and the latitude and longitude location, enrich the streaming event with these features and then normalize it required by the model. The table below describe each feature, enrichment store, and the normalization required.

Feature	Description	Enrichment Store	Normalization
Model_Feature_Certification	Identifies if the driver is certified or not	HBase/Phoenix table called drivers	"yes" # normalize to 1 "no" # normalize to 0
Model_Feature_WagePlan	Identifies if the driver is on an hourly or by miles wage plan	HBase/Phoenix table called drivers	"Hourly" # normalize to 1 "Miles" # normalize to 0
Model_Feature_FatigueByHour	The total number of hours driven by the driver in the last week	HBase/Phoenix table called timesheet	Scale by 100 to improve algorithm performance (e.g: hours/100)
Model_Feature_FatigueByMile	The total number of miles driven by the driver in the last week	HBase/Phoenix table called timesheet	Scale by 1000 to improve algorithm performance (e.g: miles/1000)
Model_Feature_FoggyWeather	Determines if for the given time and location, if the conditions are foggy	API to WeatherService	if (foggy) # normalize to 1 else 0
Model_Feature_RainyWeather	Determines if for the given time and location, if the conditions are rainy	API to WeatherService	if (rainy) -> normalize to 1 else 0
Model_Feature_WindyWeather	Determines if for the given time and location, if the conditions are windy	API to WeatherService	if (windy) # normalize to 1 else 0

7.4. Setting up your Enrichment Store and Building Custom UDFs and Processors

About This Task

As the table above indicate four of the seven features comes from HBase/Phoenix tables. This section gives you instructions on setting up the HBase/Phoenox tables timesheet and drivers, loading them with reference data, and downloading the custom UDFs and processors to do the enrichment and normalization.

Install HBase/Phoenix and download the sam-extensions

1. If HBase is not installed, install/add an HBase service.
2. Ensure that Phoenix is enabled on the HBase Cluster.
3. Download the [Sam-Custom-Extensions.zip](#) and save it to your local machine.

4. Unzip the contents. We will call the unzipped folder \$SAM_EXTENSIONS

Steps for Creating Phoenix Tables and Loading Reference Data

1. Copy the \$SAM_EXTENSIONS/scripts.tar.gz to a node where HBase/Phoenix client is installed.

2. On that node, untar the scripts.tar.gz. We will call this directory \$SCRIPTS

```
tar -zxvf scripts.tar.gz
```

3. Go to the directory where the phoenix script is located which will create the phoenix tables for enrichment and load it with reference data.

```
cd $SCRIPTS/phoenix
```

4. Open the file phoenix_create.sh and replace <ZK_HOST> with the FQDN of your ZooKeeper host.

5. Make the phoenix_create.sh script executable and execute it. Make sure you to JAVA_HOME

```
./phoenix_create.sh
```

Steps for Verifying Data has Populated Phoenix Tables

1. Start up sqlline Phoenix client.

```
cd /usr/hdp/current/phoenix-client/bin  
./sqlline.py $ZK_HOST:2181:/hbase-unsecure
```

2. List all the tables in Phoenix.

```
!tables
```

3. Query the drivers and timesheet tables.

```
select * from drivers;  
select * from timesheet;
```

7.5. Upload Custom Processors and UDFs for Enrichment and Normalization

To perform the above enrichment and normalization, upload the custom UDFs and processors you downloaded in the previous section.

7.5.1. Upload Custom UDFs

Steps for Uploading the Timestamp_Long UDF

1. From the left-hand menu, click **Configuration**, then **Application Resources**.

2. Select the **UDF** Tab and click the + sign to create the **TIMESTAMP_LONG** UDF. This udf will convert a string date time to a Timestamp long. The simple class for this UDF using the SAM SDK can be found [here](#).

3. The jar for this UDF is located in SAM_EXTENSIONS/sam-custom-udf-0.0.5.jar.

4. Configure the UDF with the following values:

NAME *
TIMESTAMP_LONG

DISPLAY NAME *
TIMESTAMP_LONG

DESCRIPTION *
Converts a string date time to a Timestamp Long

TYPE *
FUNCTION

CLASSNAME *
hortonworks.hdf.sam.custom.udf.time.ConvertToTimestampLong

UDF JAR *
Browse sam-custom-udf-0.0.5.jar

Cancel Ok

Steps for Configuring the Get_Week UDF

1. Select the UDF tab and click the + sign to create the GET_WEEK UDF.
2. The jar for this UDF is located in SAM_EXTENSIONS/sam-custom-udf-0.0.5.jar. This udf will convert a timestamp string into the week of the year which is required for an enrichment query. The simple class for this UDF using the SAM SDK can be found [here](#).
3. Configure the UDF with the following values:

NAME *
GET_WEEK

DISPLAY NAME *
GET_WEEK

DESCRIPTION *
For a given date time string, the functions returns the week of the date/time

TYPE *
FUNCTION

CLASSNAME *
hortonworks.hdf.sam.custom.udf.time.GetWeek

UDF JAR *
Browse sam-custom-udf-0.0.5.jar

Cancel Ok

7.5.2. Upload Custom Processors

Steps for Uploading the ENRICH-PHOENIX Custom Processor

1. From the left-hand menu, click **Configuration**, then **Application Resources**.
2. Select **Custom Processor** and click the + sign to create the ENRICH-PHOENIX processor. Configure the processor with the following values. This processor can be used to enriched streams with data from Phoenix based on a user supplied sql statement. The java class for this processor using the SAM SDK can be found [here](#).

Custom Processor

STREAMING ENGINE*: STORM

NAME*: ENRICH-PHOENIX

DESCRIPTION*: Enriches the input schema with data from Phoenix based on user supplied SQL.

CLASSNAME*: hortonworks.hdf.sam.custom.processor.enrich.phoenix.PhoenixEnrichmentSecureProcessor

UPLOAD JAR*: sam-custom-processor-0.0.5-jar-with-dependencies.jar

CONFIG FIELDS*:

Field Name	UI Name	Optional	Type	Default Value	Tooltip	Actions
zkServerUrl	Phoenix Zookeeper Connection URL	false	string		Zookeeper server url in the format of \$FQDN_ZK_HOST:\$ZK_PORT	<input type="button" value="Edit"/> <input type="button" value="Delete"/>
enrichmentSQL	Enrichment SQL	false	string		SQL to execute for the enriched values	<input type="button" value="Edit"/> <input type="button" value="Delete"/>
enrichedOutputFields	Enrichment Output Fields	false	string		The output field names to store new enriched values	<input type="button" value="Edit"/> <input type="button" value="Delete"/>
secureCluster	Secure Cluster	false	boolean	true	Check if connecting to a secure HBase/Phoenix Cluster	<input type="button" value="Edit"/> <input type="button" value="Delete"/>
kerberosClientPrincipal	Kerberos Client Principal	true	string		The principal uses to connect to secure HBase/PHoenix Cluster. Required if secureCluster is checked	<input type="button" value="Edit"/> <input type="button" value="Delete"/>
kerberosKeyTabFile	Kerberos Key Tab File	true	string		Kerberos Key Tab File location on each of the worker nodes for the principal configured	<input type="button" value="Edit"/> <input type="button" value="Delete"/>

INPUT SCHEMA*:

```

1 {
2   "name": "eventTime",
3   "type": "STRING",
4   "optional": false
5 },
6 {
7   "name": "eventSource",
8   "type": "STRING",
9   "optional": false
10 },
11 {
12   "name": "truckId",
13   "type": "INTEGER",
14   "optional": false
15 }
16

```

OUTPUT SCHEMA*:

```

1 {
2   "name": "eventTime",
3   "type": "STRING",
4   "optional": false
5 },
6 {
7   "name": "eventSource",
8   "type": "STRING",
9   "optional": false
10 },
11 {
12   "name": "truckId",
13   "type": "INTEGER",
14   "optional": false
15 }
16

```

ENRICH-PHOENIX Configuration Values

- **Streaming Engine – Storm**
- **Name – ENRICH-PHOENIX**
- **Description –** Enriches the input schema with data from Phoenix based on user supplied SQL
- **ClassName –**
hortonworks.hdf.sam.custom.processor.enrich.phoenix.PhoenixEnrichmentSecure

- **Upload Jar** – The jar for this custom processor can be found under SAM_EXTENSIONS/sam-custom-processor-0.0.5-jar-with-dependencies.jar

Click the **Add Config Fields** button and and the following 3 configuration fields

- Add a config field called **zkServerUrl** with the following values:
 - Field Name** – zkServerUrl
 - UI Name** – Phoenix ZooKeeper Connection URL
 - Optional** – false
 - Type** – string
 - ToolTip** – ZooKeeper server url in the format of \$FQDN_ZK_HOST:\$ZK_PORT
- Add a config field called **enrichmentSQL** with the following values:
 - Field Name** – enrichmentSQL
 - UI Name** – Enrichment SQL
 - Optional** – false
 - Type** – string
 - ToolTip** – SQL to execute for the enriched values
- Add a config field called **enrichedOutputFields** with the following values:
 - Field Name** – enrichedOutputFields
 - UI Name** – Enrichment Output Fields
 - Optional** – false
 - Type** – string
 - ToolTip** – The output field names to store new enriched values
- Add a config field called **secureCluster** with the following values:
 - Field Name** – secureCluster
 - UI Name** – Secure Cluster
 - Optional** – false
 - Type** – boolean
 - ToolTip** – Check if connecting to a secure HBase/Phoenix Cluster

- Add a config field called **kerberosClientPrincipal** with the following values:
 - a. **Field Name** – kerberosClientPrincipal
 - b. **UI Name** – Kerberos Client Principal
 - c. **Optional** – true
 - d. **Type** – string
 - e. **ToolTip** – The principal uses to connect to secure HBase/PHoenix Cluster. Required if secureCluster is checked
- Add a config field called **kerberosKeyTabFile** with the following values:
 - a. **Field Name** – kerberosKeyTabFile
 - b. **UI Name** – Kerberos Key Tab File
 - c. **Optional** – true
 - d. **Type** – string
 - e. **ToolTip** – Kerberos Key Tab File location on each of the worker nodes for the principal configured

Input and Output Schema for ENRICH-PHOENIX

- Copy this [input schema](#) and paste into the INPUT SCHEMA text area box
- Copy this [output schema](#) and paste into the OUTPUT SCHEMA text area box

Steps for Uploading the ENRICH-WEATHER Custom Processor

1. Select **Custom Processor** and click the + sign to create the ENRICH-WEATHER processor. This processor can be used to enrich streams with weather data based on time and lat/long location. The java class for this processor using the SAM SDK can be found [here](#).
2. Configure the processor with the following values.

STREAMING ENGINE*

NAME*

DESCRIPTION*

CLASSNAME*

UPLOAD JAR*

CONFIG FIELDS*

Field Name	UI Name	Optional	Type	Default Value	Tooltip	Actions
weatherServiceURL	Weather Web Service URL	false	string		The URL to the Weather Web Service	

INPUT SCHEMA*

```

40     "optional": false
41   },
42   {
43     "name": "latitude",
44     "type": "DOUBLE",
45     "optional": false
46   },
47   {
48     "name": "longitude",
49     "type": "DOUBLE",
50     "optional": false
51   },
52   {
53     "name": "correlationId",
54     "type": "LONG",
55     "optional": false
56   }
57 ]
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87 ]

```

OUTPUT SCHEMA*

```

73   {
74     "name": "Model_Feature_FoggyWeather",
75     "type": "DOUBLE",
76     "optional": false
77   },
78   {
79     "name": "Model_Feature_RainyWeather",
80     "type": "DOUBLE",
81     "optional": false
82   },
83   {
84     "name": "Model_Feature_WindyWeather",
85     "type": "DOUBLE",
86     "optional": false
87 ]

```

ENRICH-WEATHER Configuration Values

- **Streaming Engine** – Storm
- **Name** – ENRICH-WEATHER
- **Description** – Enrichment with normalized weather data for a geo location
- **ClassName** –
hortonworks.hdf.sam.custom.processor.enrich.weather.WeatherEnrichmentProcessor
- **Upload Jar** – The jar for this custom processor can be found under SAM_EXTENSIONS / sam-custom-processor-0.0.5.jar

Click the **Add Config Fields** button and and a configuration field with the following values:

- **Field Name** – weatherServiceURL
- **UI Name** – Weather Web Service URL
- **Optional** – false
- **Type** – string
- **Tooltip** – The URL to the Weather Web Service

Input and Output Schema for ENRICH-WEATHER

- Copy this **input schema** and paste into the INPUT SCHEMA text area box
- Copy this **output schema** and paste into the OUTPUT SCHEMA text area box

Steps for Uploading the NORMALIZE-MODEL-FEATURES Custom Processor

1. Select the Custom Processor Tab and click the + sign to create the NORMALIZE-MODEL-FEATURES processor. This processor is to normalized the enriched fields to format that the model is expecting.
2. Configure the processor with the following values.

The screenshot shows the 'Custom Processor' configuration interface. The 'STREAMING ENGINE' is set to 'STORM'. The 'NAME' field contains 'NORMALIZE-MODEL-FEATURES'. The 'DESCRIPTION' field contains 'Normalize the features of the model before passing it to model'. The 'CLASSNAME' field contains 'hortonworks.hdf.sam.custom.processor.enrich.driver.predictivemodel.FeatureNormal'. The 'UPLOAD JAR' field shows 'sam-custom-processor-0.0.5a.jar'. The 'CONFIG FIELDS' section is empty. The 'INPUT SCHEMA' and 'OUTPUT SCHEMA' sections display JSON code for the schema definitions.

```

    "INPUT SCHEMA": [
        {
            "name": "driverCertification",
            "type": "STRING",
            "optional": false
        },
        {
            "name": "driverWagePlan",
            "type": "STRING",
            "optional": false
        },
        {
            "name": "driverFatigueByHours",
            "type": "STRING",
            "optional": false
        }
    ],
    "OUTPUT SCHEMA": [
        {
            "name": "Model_Feature_WagePlan",
            "type": "INTEGER",
            "optional": false
        },
        {
            "name": "Model_Feature_FatigueByHours",
            "type": "DOUBLE",
            "optional": false
        },
        {
            "name": "Model_Feature_FatigueByMiles",
            "type": "DOUBLE",
            "optional": false
        }
    ]
  
```

NORMALIZE-MODEL-FEATURES Configuration Values

- **Streaming Engine – Storm**
- **Name – NORMALIZE-MODEL-FEATURES**
- **Description – Normalize the features of the model before passing it to model**
- **ClassName –**
hortonworks.hdf.sam.custom.processor.enrich.driver.predictivemodel.FeatureNormalizationProcessor
- **Upload Jar –** The jar for this custom processor can be found under SAM_EXTENSIONS / sam-custom-processor-0.0.5a.jar

Input and Output Schema for NORMALIZE-MODEL-FEATURES

- Copy this [input schema](#) and paste into the INPUT SCHEMA text area box
- Copy this [output schema](#) and paste into the OUTPUT SCHEMA text area box

Result

We have uploaded three custom processors required to do enrichment of the stream and normalization of the enriched values to feed into the model.

Name	Description	Jar File Name	Actions
ENRICH-PHOENIX	Enriches the input schema with data from Phoenix based on user supplied SQL	sam-custom-processor-0.0.5.jar-with-dependencies.jar	
ENRICH-WEATHER	Enrichment with normalized weather data for a geo location	sam-custom-processor-0.0.5.jar	
NORMALIZE-MODEL-FEATURES	Normalize the features of the model before passing it to model	sam-custom-processor-0.0.5a.jar	

If you go back to the Stream Builder, you will see three new custom processors on palette.



7.6. Scoring the Model in the Stream using a Streaming Split Join Pattern

About This Task

Now that you have created the enrichment store, loaded the enrichment data and uploaded the custom UDFs and processors to SAM, build the stream flow to score the model in real-time. In this case, you want to predict violations for events that are not blatant infractions.

Steps

1. Click into to the Trucking IOT application you built.
2. Double click the Event Type rule processor to display the **Add New Rule** dialog.
3. Configure the new rule with the following values:

Add New Rule

RULE NAME*
Non Violation Events

DESCRIPTION*
Events that are not violations that we want to do predictions on

CREATE QUERY*

eventType	EQUALS	'Normal'	+
-----------	--------	----------	----------

QUERY PREVIEW:

```
eventType = 'Normal'
```

Cancel **Ok**

Result

Your new rule is added to the Event Type processor.

EventType

CONFIGURATION **NOTES**

Input	+Add New Rules		Output
	Name	Condition	Actions
eventTime*	ViolationEvents	eventType <> 'Normal'	
eventSource*	Non Violation Events	eventType = 'Normal'	
truckId*			
driverId*			
driverName*			
routeld*			
route*			
eventType*			
latitude*			
longitude*			
correlationId*			

Cancel **Ok**

7.7. Streaming Split Join Pattern

About This Task

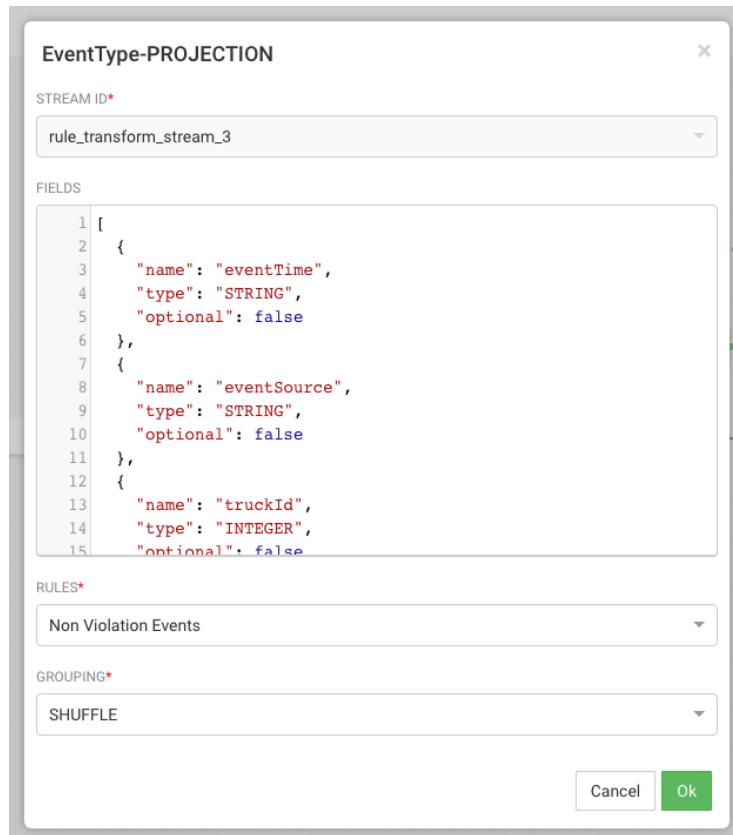
Your objective is to perform three enrichments:

- Retrieve a driver's certification and wage plan from the driver's table
- Retrieve the driver's hours and miles logged from the timesheet table
- Query weather information for a specific time and location.

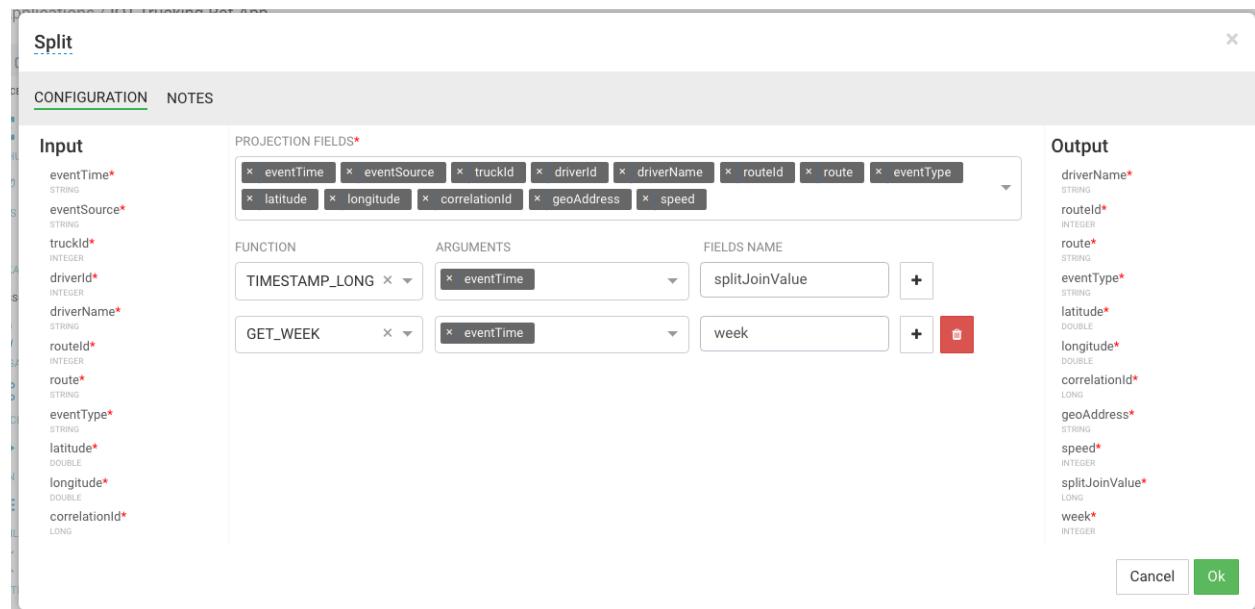
To do this, use the split join pattern to split the stream into 3, perform the enrichment in parallel, and then re-join the three streams.

Steps for Creating a Split Join Key

1. Create a new split key in the stream which allows you to join in a common field when you join the three stream. To do this, drag the projection processor to the canvas and create a connection from the EventType rule processor to this projectioin processor. When configuring the connection, select the Non Violation Events Rule which tells SAM to only send non violation events to this project processor.

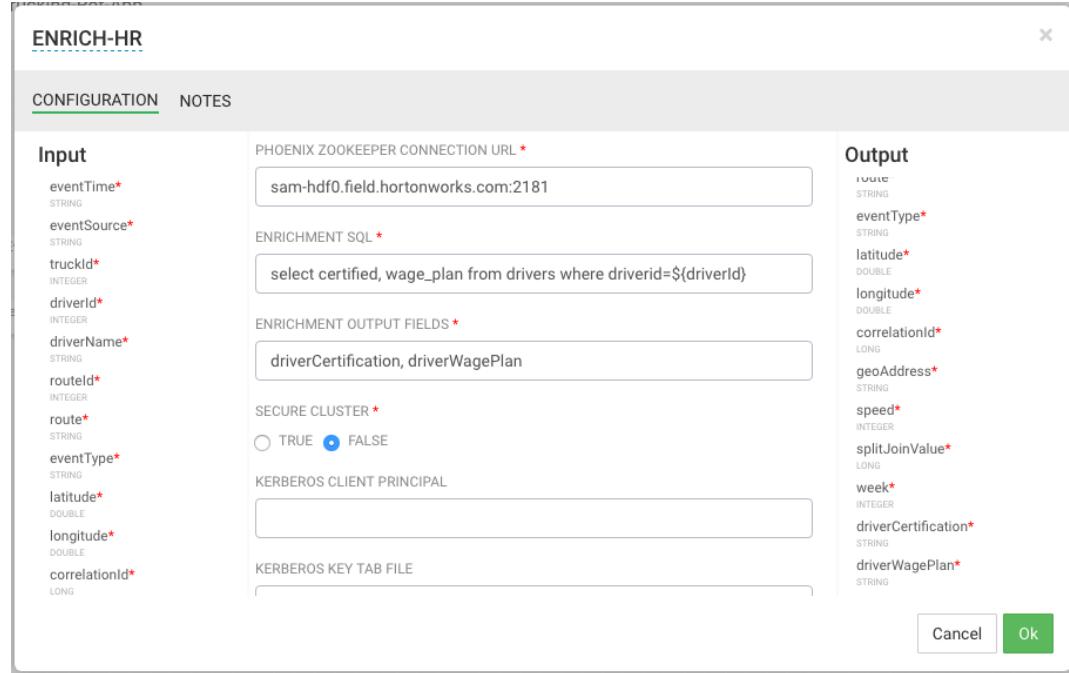


2. Configure the projection processor to create our split join key called splitJoinValue using the custom udf we uploaded earlier called "TIMESTAMP_LONG". We will also do a transformation which calculates the week based on the event time which is required for one of the enrichments downstream. Configure the processor like the following:

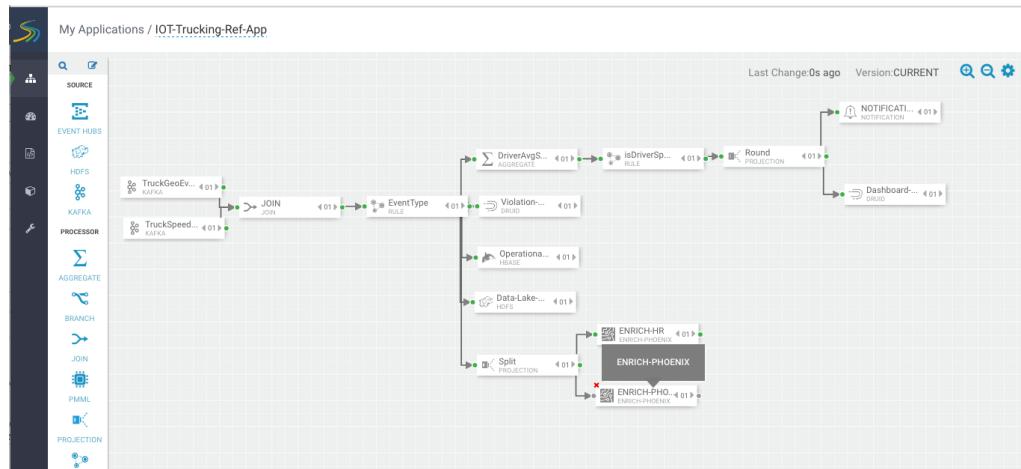


Steps for Splitting the Stream into Three to Perform Enrichments in Parallel

1. With the split join key created, we can split the stream into three to perform the enrichments in parallel. To do the first split to do the enrichment of the wage and certification status of driver, drag the "ENRICH-PHOENIX" processor the canvas and connect it from the Split project processor
2. Configure the enrich processor like below. After this processor executes, the output schema will have two fields populated called driverCertification and driverWagePlan.
 - a. ENRICHMENT SQL: select certified, wage_plan from drivers where driverid= \${driverId}
 - b. ENRICHMENT OUTPUT FIELDS: driverCertification, driverWagePlan
 - c. SECURE CLUSTER: false
 - d. INPUT SCHEMA MAPPINGS: Leave defaults
 - e. OUTPUT FIELDS: select all fields except for driverFatigueByHours and driverFatigueByMiles



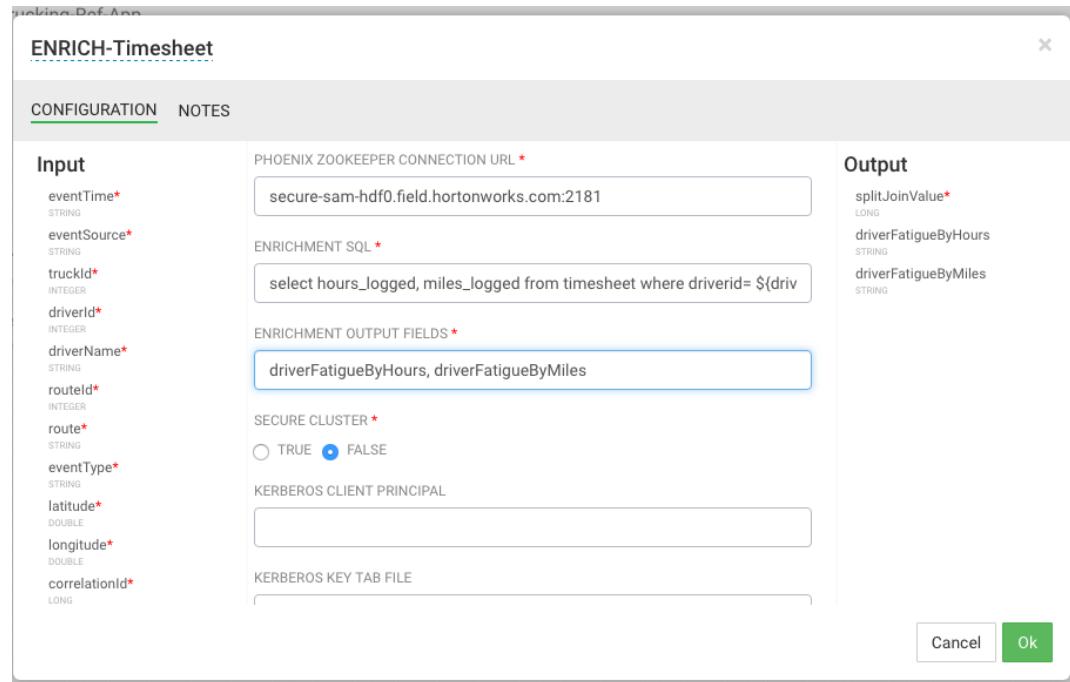
3. Create the second stream to do enrichment of the drivers hours and miles logged in last week by dragging another "ENRICH-PHOENIX" processor to the canvas and connect it from the Split projection processor.



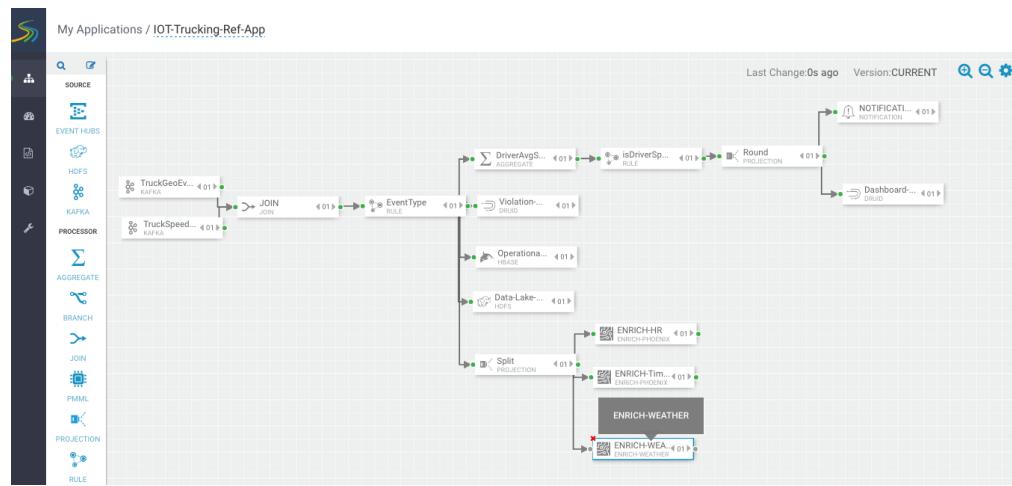
4. Configure the enrich processor like below. After this processor executes, the output schema will have two fields populated called driverFatigueByHours, driverFatigueByMiles.
- ENRICHMENT SQL: select hours_logged, miles_logged from timesheet where driverid= \${driverId} and week=\${week}
 - ENRICHMENT OUTPUT FIELDS: driverFatigueByHours, driverFatigueByMiles
 - SECURE CLUSTER: false

d. INPUT SCHEMA MAPPINGS: Leave defaults

e. OUTPUT FIELDS: select splitJoinValue, driverFatigueByHours, driverFatigueMiles

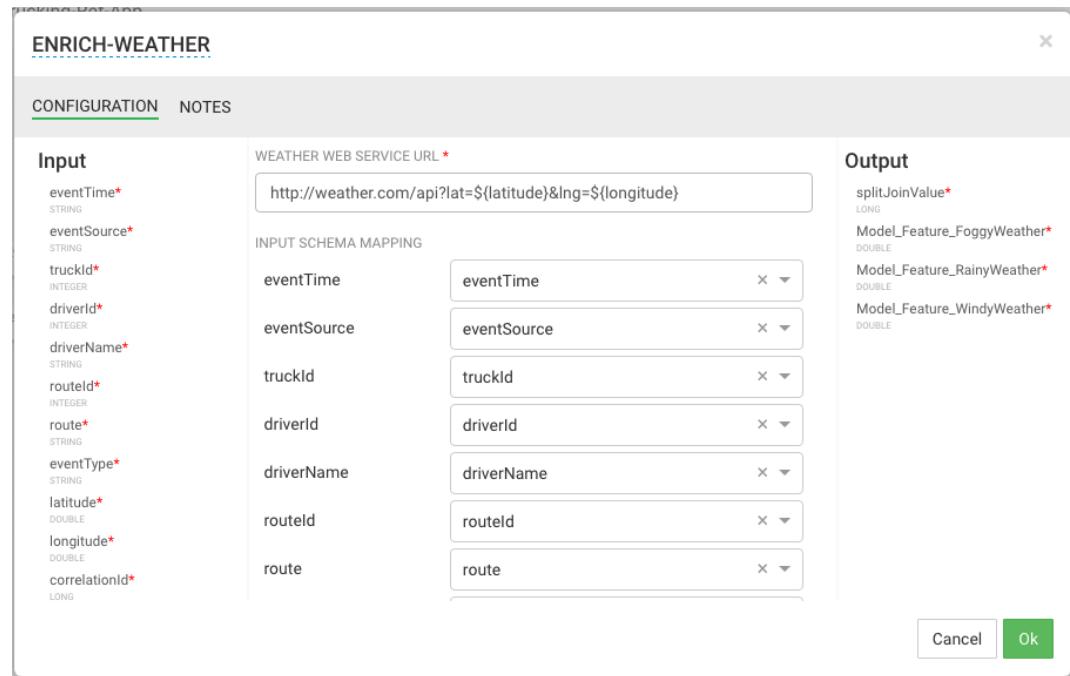


5. Create the third stream to do weather enrichment by dragging the custom processor we uploaded called "ENRICH-WEATHER" processor to the canvas and connect it from the Split project processor.



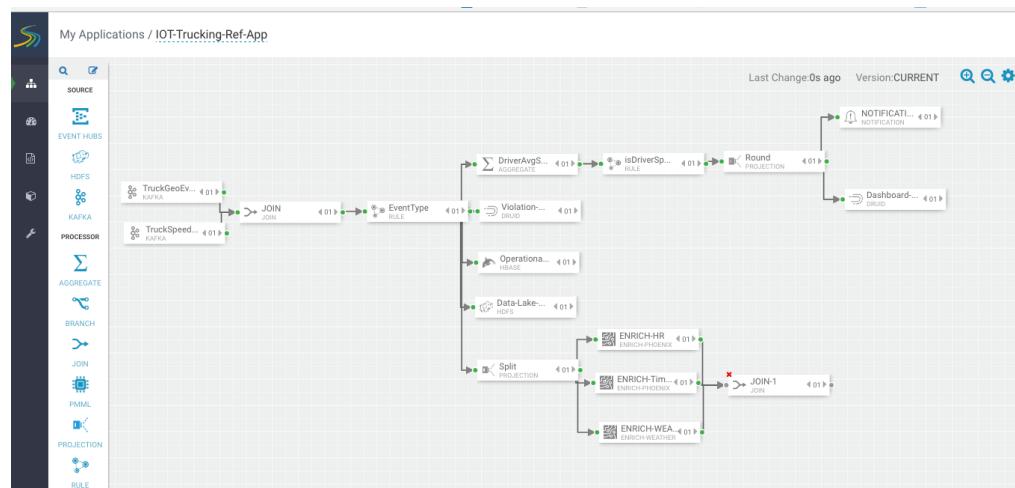
6. Configure the weather process like the following (currently the weather processor is just a stub that generates random normalized weather info). After this processor executes, the output schema will have three fields populated called Model_Feature_FoggyWeather, Model_Feature_RainyWeather, Model_Feature_WindyWeather.

- WEATHER WEB SERVICE URL: `http://weather.com/api?lat=${latitude}&lng=${longitude}`
- INPUT SCHEMA MAPPINGS: Leave defaults
- OUTPUT FIELDS: Select the splitJoinValue and the three model enriched features



Steps for Rejoining the Three Enriched Streams

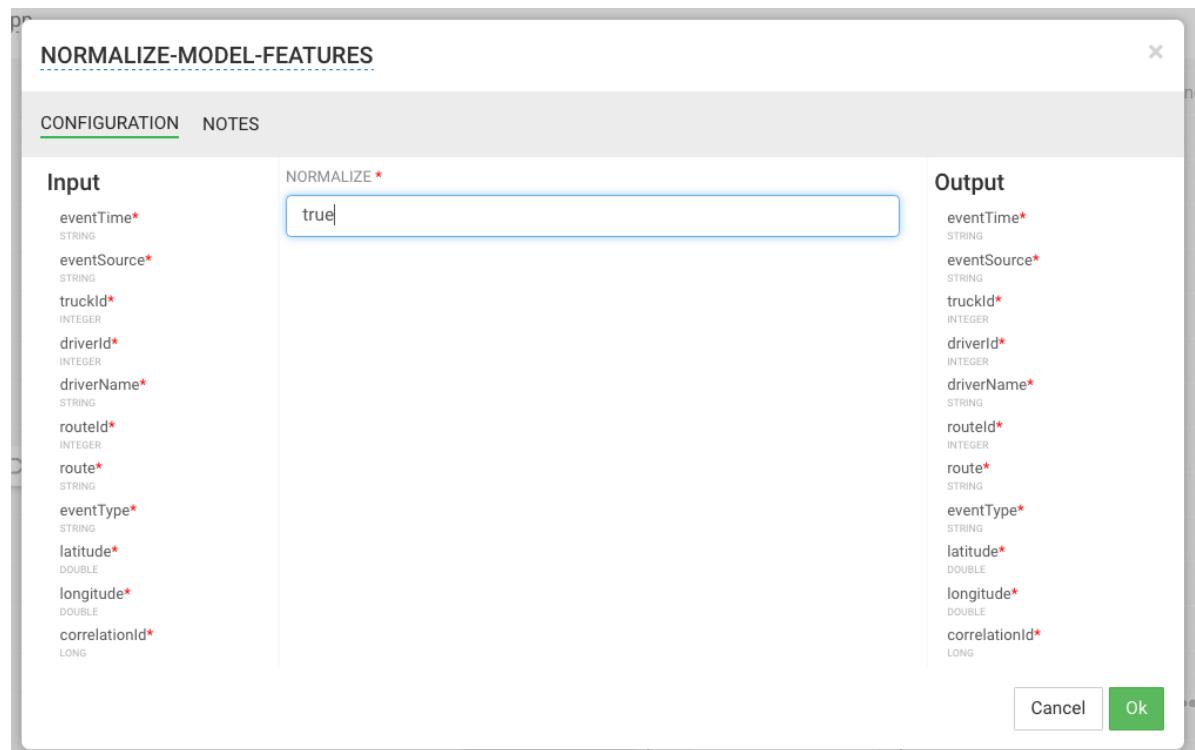
- Now that we have done the enrichment in parallel by splitting the stream into 3, we can now join the 3 streams by dragging the join processor to the canvas and connecting the join from the 3 streams.



- Configure the join processor like the following where we use the joinSplitValue to join all three streams.

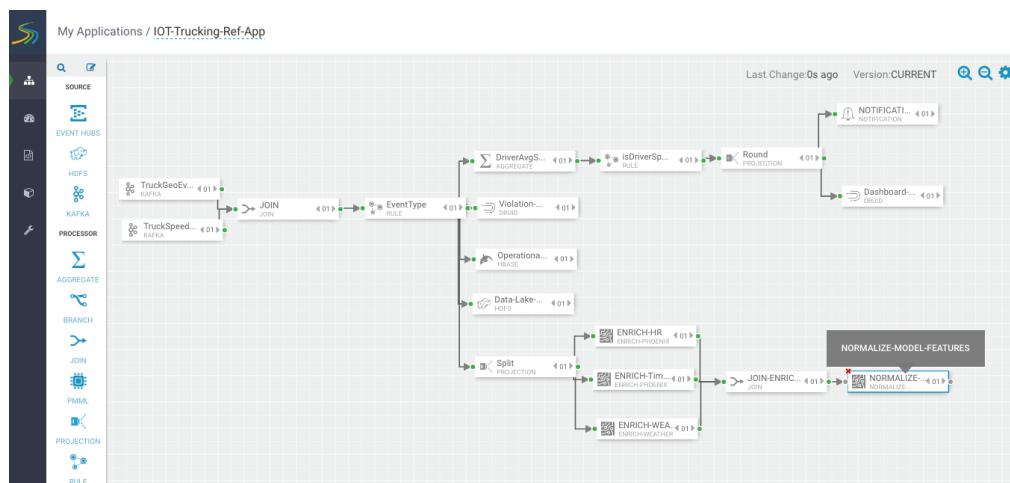
For the Output field, just click **SELECT ALL** to get all the fields across the three streams.

- Now that we have joined three enriched streams, lets normalize the data into the format that the model expects by dragging to the canvas the "NORMALIZE-MODEL-FEATURES" custom processor that we added. For the output fields select all the fields and leave the the mapping as defaults.



Result

Your flow looks similar to the following.



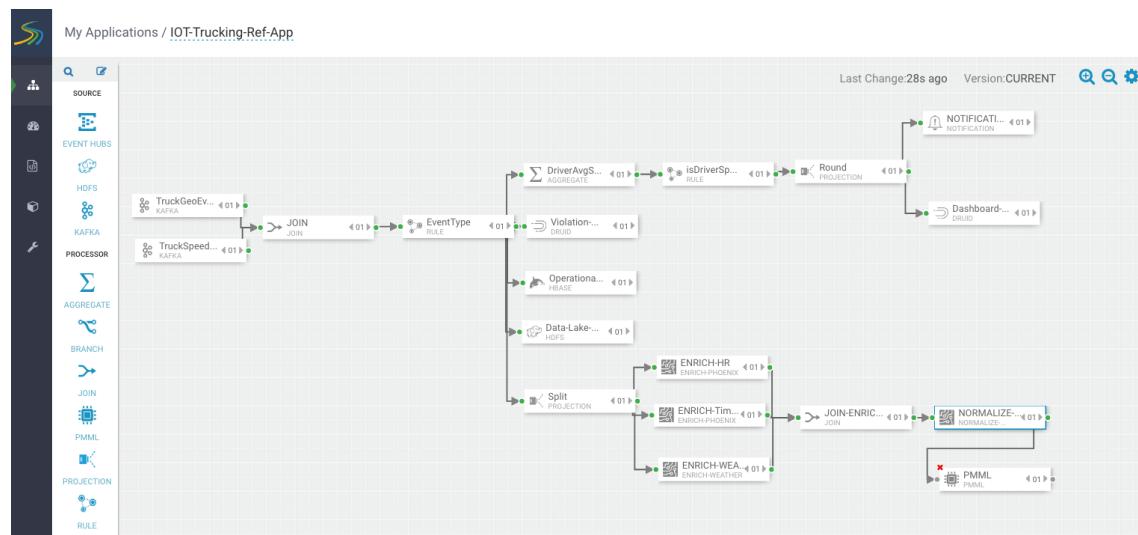
7.8. Score the Model using the PMML Processor and Alert

About This Task

Now you are ready to score the logistical regression model.

Steps

1. Drag the PMML processor the canvas and connect it to the Normalize processor.



2. Configure the PMML processor like the following by selecting the **DriverViolationPredictionModel** that you uploaded to the **Model Registry** earlier. After this processor executes, a new field called **ViolationPredicted** is added to stream for the result of the prediction. In output fields, select all the contextual fields you want to pass on including the model value result.

Predict

CONFIGURATION NOTES

Input

- eventTime*
- STRING
- eventSource*
- STRING
- truckId*
- INTEGER
- driverId*
- INTEGER
- driverName*
- STRING
- routId*
- INTEGER
- route*
- STRING
- eventType*
- STRING
- latitude*
- DOUBLE
- longitude*
- DOUBLE
- correlationId*
- LONG

MODEL NAME*

DriverViolationPredictionModel

OUTPUT FIELDS*

SELECT ALL

x eventTime x eventSource x truckId x driverId
x driverName x routId x route x latitude x longitude
x geoAddress x speed

Output

- ViolationPredicted*
- STRING
- eventTime*
- STRING
- eventSource*
- STRING
- truckId*
- INTEGER
- driverId*
- INTEGER
- driverName*
- STRING
- routId*
- INTEGER
- route*
- STRING
- latitude*
- DOUBLE
- longitude*
- DOUBLE
- geoAddress*
- STRING

Cancel **Ok**

3. Determine if the model predicted if the driver will commit a violation by dragging a rule processor to the canvas and configuring a rule like the following:

Edit Rule

RULE NAME*

Violation Predicted

DESCRIPTION*

model returned a prediction

CREATE QUERY*

ViolationPredicted **X** EQUALS **X** 'yes' **X** **+**

QUERY PREVIEW:

`ViolationPredicted = 'yes'`

Cancel **Ok**

4. If a violation is predicted, send it to a Druid to display on a dashboard. Drag the Druid processor to canvas and configure. Stream the events into a cube called **alerts-violation-predictions-cube**.

Dashboard-Predictions

REQUIRED OPTIONAL NOTES

Input

ViolationPredicted*
STRING
eventTime*
STRING
eventSource*
STRING
truckId*
INTEGER
driverId*
INTEGER
driverName*
STRING
routeId*
INTEGER
route*
STRING
eventType*
STRING
latitude*
DOUBLE
longitude*
DOUBLE

DATASOURCE NAME *

alerts-violation-predictions-cube

ZOOKEEPER CONNECT STRING *

secure-sam-hdf2.field.hortonworks.com:2181,secure-sar

DIMENSIONS *

x ViolationPredicted x eventTime
x eventSource x truckId x driverId
x driverName x routeId x route
x eventType x latitude x longitude
x correlationId x geoAddress x speed

x ▾

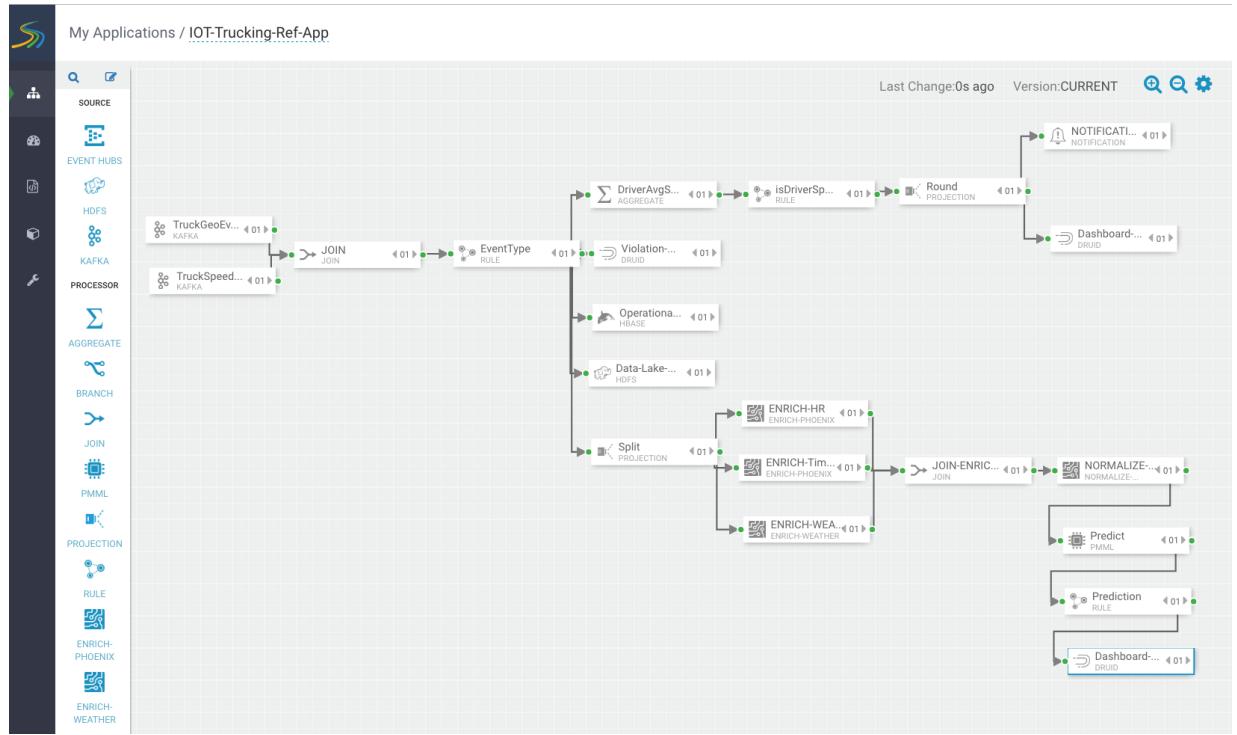
TIMESTAMP FIELD NAME *

processingTime

Cancel Ok

Result

The final flow looks like the following:



8. Creating Visualizations Using Superset

A business analyst can create a wide array of visualizations to gather insights on streaming data. The platform supports over 30+ visualizations the business analyst can create. For visualization examples, see the [Gallery of Superset Visualizations](#).

The general process for creating and viewing visualizations is as follows:

1. Whenever you add new data sources to Druid via a Stream App, perform the **Refresh Druid Metadata** action on the Superset menu.
2. Using the Superset Stream Insight UI, create one or more "slices". A slice is one business visualization associated with a data source (e.g: Druid cube).
3. Using the Dashboard menu, add the slices to your dashboard and organize their layout.



Note

Note that when a SAM app streams data to a new cube using the Druid processor, it will take about 30 minutes for the cube to appear in Superset. This is because Superset has to wait for the first segment to be created in Druid. After the cube appears, users can analyze the streaming data immediately as it is streaming in.

8.1. Creating Insight Slices

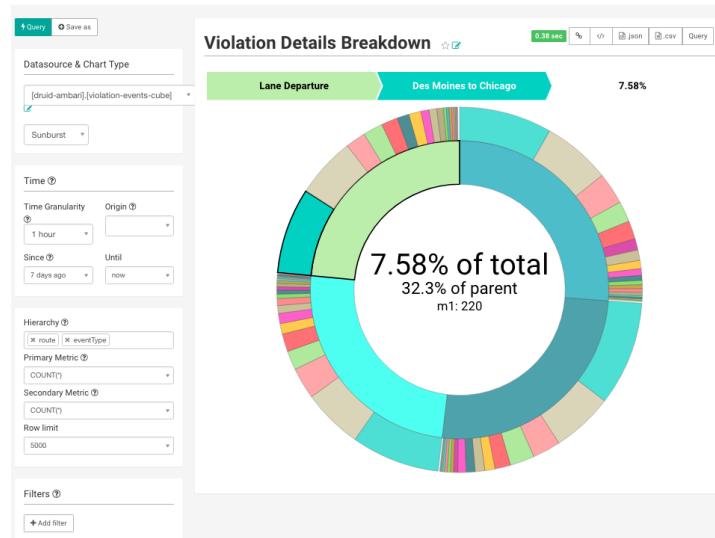
The following steps demonstrate a typical flow for creating a slice:

1. Choose **Slices** on the Menu.
2. Click + to create a new Slice.
3. Select the Druid Data Source that you want to use for the new visualization:

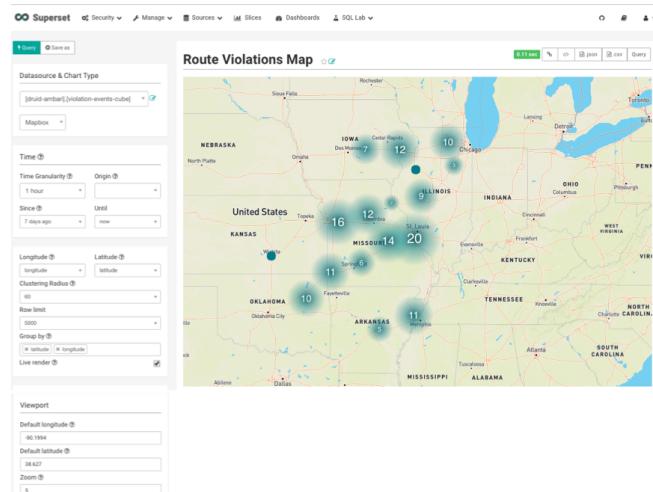
Action	Data Source	Cluster	Changed By	Changed On	Time Offset
Edit	Alerts-High-Speed-Cube-V2	Streaming Analytics Manager - Stream Insight	George Vetticaden	2017-02-07 15:50:59.807995	0
Edit	driver-violations-cube-2	Streaming Analytics Manager - Stream Insight	George Vetticaden	2017-02-07 17:29:22.544088	0

4. Select a Chart Type from the menu.

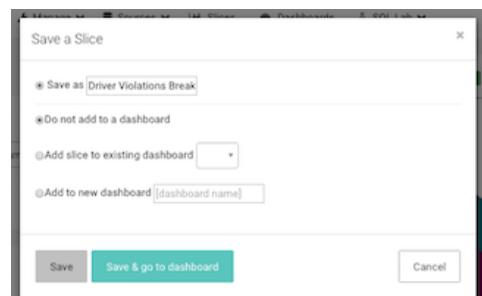
This example creates a "Sunburst" visualization where we are rolling up multiple dimensions like route, eventType and driver info. Configure the chart and click **Execute Query**.



5. Another visualization could be integration with [MapBox](#). Here we are mapping where violations are occurring the most based on the lat/long location of the event



6. To save the slice, specify a name and click **Save**.



8.2. Adding Insight Slices to a Dashboard

After you create slices, you can organize them into a dashboards:

1. Click the Dashboard menu item.
2. Click + to create a new Dashboard.
3. Configure the dashboard: specify a name and the slices to include in the Dashboard.

The screenshot shows the 'Add Dashboard' interface in Superset. The 'Title' field contains 'Trucking IOT Dashboard'. The 'Slug' field contains 'truckin-giot-dashboard'. The 'Slices' section lists five slices: 'Total Violations in Last Hour', 'Top Violation Drivers', 'Driver Violations Breakdown', 'Direction Infraction Details', and 'Routes with Infractions'. The 'Owners' section lists 'George Vetticaden'. The 'Position JSON' section contains a JSON object for widget positioning. The 'CSS' section contains CSS styles. The 'JSON Metadata' section contains dynamically generated JSON metadata. At the bottom are 'Save' and 'Cancel' buttons.

4. Arrange the slices on the dashboard as desired, and then click **Save**.

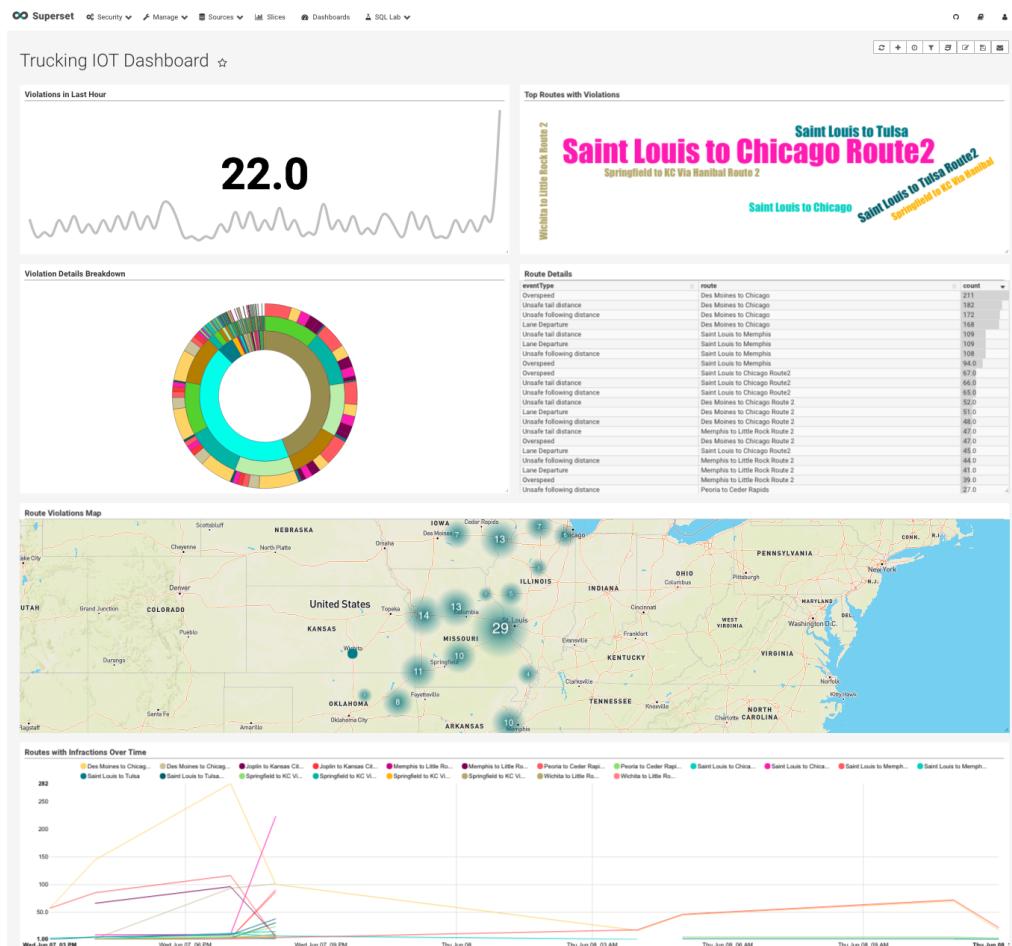
8.2.1. Dashboards for the Trucking IOT App

The IOT Trucking app that we implementing using the Stream Builder was streaming violation events, alerts and predictions into three cubes:

- violation-events-cube
- alerts-speeding-drivers-cube
- alerts-violation-predictions-cube

Based on the powerful visualizations that SuperSet offers, you can create the below powerful dashboards in minutes.

IoT Dashboard



Alerts Dashboard

