

Hortonworks Data Platform

Command Line Upgrade

(Aug 29, 2016)

Hortonworks Data Platform: Command Line Upgrade

Copyright © 2012-2016 Hortonworks, Inc. Some rights reserved.

The Hortonworks Data Platform, powered by Apache Hadoop, is a massively scalable and 100% open source platform for storing, processing and analyzing large volumes of data. It is designed to deal with data from many sources and formats in a very quick, easy and cost-effective manner. The Hortonworks Data Platform consists of the essential set of Apache Software Foundation projects that focus on the storage and processing of Big Data, along with operations, security, and governance for the resulting system. This includes Apache Hadoop – which includes MapReduce, Hadoop Distributed File System (HDFS), and Yet Another Resource Negotiator (YARN) – along with Ambari, Falcon, Flume, HBase, Hive, Kafka, Knox, Oozie, Phoenix, Pig, Ranger, Slider, Spark, Sqoop, Storm, Tez, and ZooKeeper. Hortonworks is the major contributor of code and patches to many of these projects. These projects have been integrated and tested as part of the Hortonworks Data Platform release process and installation and configuration tools have also been included.

Unlike other providers of platforms built using Apache Hadoop, Hortonworks contributes 100% of our code back to the Apache Software Foundation. The Hortonworks Data Platform is Apache-licensed and completely open source. We sell only expert technical support, [training](#) and partner-enablement services. All of our technology is, and will remain, free and open source.

Please visit the [Hortonworks Data Platform](#) page for more information on Hortonworks technology. For more information on Hortonworks services, please visit either the [Support](#) or [Training](#) page. Feel free to [contact us](#) directly to discuss your specific needs.



Except where otherwise noted, this document is licensed under
Creative Commons Attribution ShareAlike 3.0 License.
<http://creativecommons.org/licenses/by-sa/3.0/legalcode>

Table of Contents

1. Upgrade from HDP 2.4 to HDP 2.5.0 Manually	1
1.1. Getting Ready to Upgrade	2
1.2. Upgrade HDP 2.4 Components	11
1.3. Symlink Directories with hdp-select	17
1.4. Configure and Start Apache ZooKeeper	17
1.5. Configure Hadoop	17
1.6. Start Hadoop Core	18
1.7. Verify HDFS Filesystem Health	21
1.8. Configure YARN and MapReduce	22
1.9. Start YARN/MapReduce Services	25
1.10. Run Hadoop Smoke Tests	26
1.11. Configure and Start Apache HBase	27
1.12. Configure Apache Phoenix	27
1.13. Configure and Start Apache Accumulo	29
1.14. Configure and Start Apache Tez	29
1.15. Configure and Start Apache Hive and Apache HCatalog	31
1.16. Configure and Start Apache Oozie	33
1.17. Configure and Start Apache WebHCat	37
1.18. Configure Apache Pig	40
1.19. Configure and Start Apache Sqoop	40
1.20. Configure, Start, and Validate Apache Flume	40
1.21. Configure, Start, and Validate Apache Mahout	41
1.22. Configure and Start Hue	41
1.23. Configure and Start Apache Knox	42
1.23.1. Upgrade the Knox Gateway	43
1.23.2. Verify the Knox Upgrade	44
1.23.3. Downgrade the Knox Gateway to the Previous Version	44
1.23.4. Verify the Knox Downgrade Was Successful	45
1.24. Configure and Validate Apache Falcon	45
1.25. Configure and Start Apache Storm	45
1.26. Configure and Start Apache Ranger	48
1.26.1. Preparing Your Cluster to Upgrade Ranger	48
1.26.2. Stop the Ranger Services	49
1.26.3. Preparing the Cluster for Upgrade	49
1.26.4. Registering the HDP 2.5.0 Repo	49
1.26.5. Install the Ranger Components	51
1.26.6. Restart the Ranger Services	52
1.26.7. Enable Ranger Plugins	52
1.26.8. Enable KMS Configuration	53
1.26.9. Configure and Start Apache Ranger on a Kerberized Cluster	54
1.27. Configuring and Upgrading Apache Spark	62
1.28. Upgrade Apache Slider	63
1.29. Upgrade Apache Kafka	63
1.29.1. Downgrading Kafka	63
1.30. Finalize the Upgrade	63
1.31. Migrate the Audit Logs from DB to Solr	64
1.32. Install New HDP 2.5.0 Services	64
2. Upgrade from HDP 2.3 to HDP 2.5.0 Manually	65

2.1. Getting Ready to Upgrade	66
2.2. Upgrade HDP 2.3 Components	75
2.3. Symlink Directories with hdp-select	81
2.4. Configure and Start Apache ZooKeeper	81
2.5. Configure Hadoop	82
2.6. Start Hadoop Core	82
2.7. Verify HDFS Filesystem Health	85
2.8. Configure YARN and MapReduce	86
2.9. Start YARN/MapReduce Services	89
2.10. Run Hadoop Smoke Tests	90
2.11. Configure and Start Apache HBase	91
2.12. Configure Apache Phoenix	91
2.13. Configure and Start Apache Accumulo	93
2.14. Configure and Start Apache Tez	94
2.15. Configure and Start Apache Hive and Apache HCatalog	95
2.16. Configure and Start Apache Oozie	98
2.17. Configure and Start Apache WebHCat	102
2.18. Configure Apache Pig	103
2.19. Configure and Start Apache Sqoop	104
2.20. Configure, Start, and Validate Apache Flume	104
2.21. Configure, Start, and Validate Apache Mahout	105
2.22. Configure and Start Hue	106
2.23. Configure and Start Apache Knox	106
2.23.1. Upgrade the Knox Gateway	107
2.23.2. Verify the Knox Upgrade	108
2.23.3. Downgrade the Knox Gateway to the Previous Version	108
2.23.4. Verify the Knox Downgrade Was Successful	109
2.24. Configure and Validate Apache Falcon	109
2.25. Configure and Start Apache Storm	110
2.26. Configure and Start Apache Ranger	111
2.26.1. Preparing Your Cluster to Upgrade Ranger	111
2.26.2. Stop the Ranger Services	113
2.26.3. Preparing the Cluster for Upgrade	113
2.26.4. Registering the HDP 2.5.0 Repo	113
2.26.5. Install the Ranger Components	115
2.26.6. Restart the Ranger Services	116
2.26.7. Enable Ranger Plugins	116
2.26.8. Enable KMS Configuration	117
2.26.9. Configure and Start Apache Ranger on a Kerberized Cluster	118
2.27. Configuring and Upgrading Apache Spark	126
2.28. Upgrade Apache Slider	127
2.29. Upgrade Apache Kafka	127
2.29.1. Downgrading Kafka	127
2.30. Finalize the Upgrade	127
2.31. Migrate the Audit Logs from DB to Solr	128
2.32. Install New HDP 2.5.0 Services	128
3. Upgrade from HDP 2.2 to HDP 2.5.0 Manually	129
3.1. Getting Ready to Upgrade	130
3.2. Upgrade HDP 2.2 Components	140
3.3. Symlink Directories with hdp-select	145
3.4. Configure and Start Apache ZooKeeper	145

3.5. Configure Hadoop	145
3.6. Start Hadoop Core	146
3.7. Verify HDFS Filesystem Health	149
3.8. Configure YARN and MapReduce	150
3.9. Start YARN/MapReduce Services	153
3.10. Run Hadoop Smoke Tests	154
3.11. Configure and Start Apache HBase	155
3.12. Configure Apache Phoenix	155
3.13. Configure and Start Apache Accumulo	157
3.14. Configure and Start Apache Tez	158
3.15. Configure and Start Apache Hive and Apache HCatalog	159
3.16. Configure and Start Apache Oozie	162
3.17. Configure and Start Apache WebHCat	165
3.18. Configure Apache Pig	168
3.19. Configure and Start Apache Sqoop	169
3.20. Configure, Start, and Validate Apache Flume	169
3.21. Configure, Start, and Validate Apache Mahout	170
3.22. Configure and Start Hue	171
3.23. Configure and Start Apache Knox	172
3.23.1. Upgrade the Knox Gateway	172
3.23.2. Verify the Knox Upgrade	173
3.24. Configure and Validate Apache Falcon	173
3.25. Configure and Start Apache Storm	174
3.26. Configure and Start Apache Ranger	175
3.26.1. Preparing Your Cluster to Upgrade Ranger	175
3.26.2. Stop the Ranger Services	176
3.26.3. Preparing the Cluster for Upgrade	176
3.26.4. Registering the HDP 2.4 Repo	176
3.26.5. Install the Ranger Components	178
3.26.6. Restart the Ranger Services	179
3.26.7. Enable Ranger Plugins	180
3.26.8. Enable KMS Configuration	180
3.26.9. Configure and Start Apache Ranger on a Kerberized Cluster	181
3.27. Configuring and Upgrading Apache Spark	189
3.28. Upgrade Apache Slider	190
3.29. Upgrade Apache Kafka	191
3.29.1. Downgrading Kafka	191
3.30. Finalize the Upgrade	191
3.31. Migrate the Audit Logs from DB to Solr	192
3.32. Install New HDP 2.5.0 Services	192
4. Upgrade from HDP 2.1 to HDP 2.5.0 Manually	193
4.1. Getting Ready to Upgrade	194
4.2. Upgrade HDP 2.1 Components	201
4.3. Symlink Directories with hdp-select	208
4.4. Configure and Start Apache ZooKeeper	209
4.5. Configure Hadoop	209
4.6. Start Hadoop Core	210
4.7. Verify HDFS Filesystem Health	212
4.8. Configure YARN and MapReduce	213
4.9. Start YARN/MapReduce Services	217
4.10. Run Hadoop Smoke Tests	218

4.11. Configure and Start Apache HBase	219
4.12. Configure Apache Phoenix	219
4.13. Configure and Start Apache Accumulo	221
4.14. Configure and Start Apache Tez	222
4.15. Configure and Start Apache Hive and Apache HCatalog	224
4.16. Configure and Start Apache Oozie	227
4.17. Configure and Start Apache WebHCat	231
4.18. Configure Apache Pig	233
4.19. Configure and Start Apache Sqoop	234
4.20. Configure, Start, and Validate Apache Flume	234
4.21. Configure and Validate Apache Mahout	236
4.22. Configure and Start Hue	236
4.23. Configure and Start Apache Knox	237
4.23.1. Upgrade the Knox Gateway	238
4.23.2. Verify the Knox Upgrade	238
4.23.3. Downgrade the Knox Gateway to the Previous Version	239
4.23.4. Verify the Knox Downgrade Was Successful	239
4.24. Configure and Validate Apache Falcon	240
4.25. Configure and Start Apache Storm	245
4.26. Configure and Start Apache Ranger	247
4.26.1. Preparing Your Cluster to Upgrade Ranger	247
4.26.2. Stop the Ranger Services	248
4.26.3. Install the Ranger Components	249
4.26.4. Restart the Ranger Services	249
4.26.5. Remove Existing Startup Files and Symbolic Links	249
4.26.6. Enable Ranger Plugins	250
4.27. Configuring and Upgrading Apache Spark	251
4.28. Upgrade Apache Slider	252
4.29. Upgrade Apache Kafka	252
4.29.1. Downgrading Kafka	253
4.30. Finalize the Upgrade	253
4.31. Install New HDP 2.5 Services	253
5. Upgrade from HDP 2.0 to HDP 2.5.0 Manually	254
5.1. Getting Ready to Upgrade	255
5.2. Upgrade HDP 2.0 Components	260
5.3. Symlink Directories with hdp-select	265
5.4. Configure and Start Apache ZooKeeper	266
5.5. Configure Hadoop	266
5.6. Set RPC Authentication	267
5.7. Start Hadoop Core	267
5.8. Verify HDFS Filesystem Health	269
5.9. Configure YARN and MapReduce	270
5.10. Start YARN/MapReduce Services	272
5.11. Run Hadoop Smoke Tests	273
5.12. Configure and Start Apache HBase	274
5.13. Configure and Start Apache Hive and Apache HCatalog	274
5.14. Configure and Start Apache Oozie	278
5.15. Configure and Start Apache WebHCat (Templeton)	280
5.16. Configure and Start Apache Pig	282
5.17. Configure and Start Apache Sqoop	282
5.18. Configure, Start, and Validate Apache Flume	283

5.19. Configure, Start, and Validate Apache Mahout	284
5.20. Configure and Start Hue	285
5.21. Finalize the Upgrade	286
5.22. Install New HDP 2.5 Services	286
6. Upgrade from HDP 1.3 to HDP 2.5.0 Manually	287
6.1. Getting Ready to Upgrade	288
6.2. Upgrade HDP 1.3 Components	295
6.3. Symlink Directories with hdp-select	298
6.4. Configure and Start Apache ZooKeeper	299
6.5. Configure and Start Hadoop	300
6.6. Migrate the HDP Configurations	300
6.7. Create Local Directories	310
6.8. Start Hadoop Core	311
6.9. Verify HDFS Filesystem Health	313
6.10. Configure YARN and MapReduce	314
6.11. Start YARN/MapReduce Services	317
6.12. Run Hadoop Smoke Tests	318
6.13. Configure and Start Apache HBase	319
6.14. Configure and Start Apache Hive and Apache HCatalog	319
6.15. Configure and Start Apache Oozie	324
6.16. Configure and Start Apache WebHCat (Templeton)	326
6.17. Configure and Start Apache Pig	330
6.18. Configure and Start Apache Sqoop	331
6.19. Configure, Start, and Validate Apache Flume	331
6.20. Configure, Start, and Validate Apache Mahout	332
6.21. Configure and Start Hue	333
6.22. Finalize the Upgrade	334
6.23. Install New HDP 2.5 Services	334

List of Tables

1.1. Hive Metastore Database Backup and Restore	5
1.2. Oozie Metastore Database Backup and Restore	6
1.3. Hue Database Backup and Restore	6
1.4. Ranger_Admin install.properties names and values	50
1.5. Ranger_Usersync install.properties names and values	51
1.6. KMS install.properties names and values	53
1.7. Properties for the /etc/hadoop/conf/core-site.xml file	55
1.8. Ranger-admin install.properties	56
1.9.	57
1.10. Ranger-tagsync install.properties and values	57
1.11. Ranger-kms install.properties and values	57
1.12. hdfs-site.xml Property Names and Values	58
1.13. Ranger Tagsync_install.properties and values	64
2.1. Hive Metastore Database Backup and Restore	70
2.2. Oozie Metastore Database Backup and Restore	70
2.3. Hue Database Backup and Restore	71
2.4. Ranger_Admin install.properties names and values	114
2.5. Ranger_Usersync install.properties names and values	115
2.6. KMS install.properties names and values	117
2.7. Properties for the /etc/hadoop/conf/core-site.xml file	119
2.8. Ranger-admin install.properties	120
2.9.	121
2.10. Ranger-tagsync install.properties and values	121
2.11. Ranger-kms install.properties and values	121
2.12. hdfs-site.xml Property Names and Values	122
2.13. Ranger Tagsync_install.properties and values	128
3.1. Hive Metastore Database Backup and Restore	134
3.2. Oozie Metastore Database Backup and Restore	134
3.3. Hue Database Backup and Restore	135
3.4. Ranger_Admin install.properties names and values	177
3.5. Ranger_Usersync install.properties names and values	178
3.6. KMS install.properties names and values	180
3.7. Properties for the /etc/hadoop/conf/core-site.xml file	182
3.8. Ranger-admin install.properties	183
3.9.	184
3.10. Ranger-tagsync install.properties and values	184
3.11. Ranger-kms install.properties and values	185
3.12. hdfs-site.xml Property Names and Values	185
3.13. Ranger Tagsync_install.properties and values	192
4.1. Hive Metastore Database Backup and Restore	197
4.2. Oozie Metastore Database Backup and Restore	198
4.3. Hue Database Backup and Restore	198
5.1. Hive Metastore Database Backup and Restore	257
5.2. Oozie Metastore Database Backup and Restore	258
5.3. Hue Database Backup and Restore	258
6.1. Hive Metastore Database Backup and Restore	292
6.2. Oozie Metastore Database Backup and Restore	293
6.3. Hue Database Backup and Restore	293

6.4. HDP 1.3.2 Hadoop Core Site (core-site.xml)	301
6.5. HDP 1.3.2 Hadoop Core Site (hdfs-site.xml)	302
6.6. HDP 1.3.2 Configs now in Capacity Scheduler for HDP 2.x (mapred-site.xml)	304
6.7. HDP 1.3.2 Configs now in capacity scheduler for HDP 2.x (capacity-scheduler.xml)	304
6.8. HDP 1.3.2 Configs and HDP 2.x for hadoop-env.sh	305

1. Upgrade from HDP 2.4 to HDP 2.5.0 Manually

This chapter provides instructions on how to manually upgrade to HDP 2.5.0 from the HDP 2.4 release. It assumes the existing HDP 2.4.x was also installed manually.



Important

If you installed and manage HDP 2.4 with Ambari, **you must use the [Ambari Upgrade Guide](#)** to perform the HDP 2.4 to HDP 2.5.0 upgrade.

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

Rolling Upgrade involves complex orchestration as well as side-by-side installation. It is too complex for a manual procedure, and is therefore supported only as an Ambari feature. If you wish to perform a Rolling Upgrade, refer to the Ambari Install instructions to install Ambari, then follow the Ambari Rolling Upgrade instructions, see [Ambari Upgrade Guide](#).

The HDP packages for a complete installation of HDP 2.5.0 occupies about 6.5 GB of disk space.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.

The following provides an overview of steps for upgrading to the latest release of HDP 2.5.x from HDP 2.4:

1. Get ready to upgrade
2. Upgrade HDP 2.4 Components and stop all services
3. Use hdp-select to symlink the HDP 2.5.0 components into "current," in place of the former HDP 2.4 components
4. Configure and Start Apache ZooKeeper
5. Configure and Start Hadoop
6. Start HDFS
7. Configure and start YARN/MapReduce
8. Configure and Start Apache HBase
9. Configure and Start Apache Phoenix

10. Configure and Start Apache Accumulo
11. Configure and Start Apache Tez
12. Configure and Start Apache Hive and Apache HCatalog
13. Configure and Start Apache Oozie
14. Configure and Start Apache WebHCat (Templeton)
15. Configure and Start Apache Pig
16. Configure and Start Apache Sqoop
17. Configure and Start Apache Flume
18. Configure and Start Apache Mahout
19. Configure and Start Apache Hue
20. Configure and Start Apache Knox
21. Configure and Start Apache Falcon
22. Configure and Start Apache Storm
23. Configure and Start Apache Ranger
24. Configure and Start Apache Spark
25. Upgrade Apache Slider
26. Upgrade Apache Kafka
27. Finalize the Upgrade
28. Install new HDP 2.5 services

1.1. Getting Ready to Upgrade

HDP Stack upgrade involves upgrading from HDP 2.4 to HDP 2.5 versions and adding the new HDP 2.5 services. These instructions change your configurations.



Note

You must use **kinit** before running the commands as any particular user.

Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP 2.5.0 consumes about 6.5 GB of disk space.

The first step is to ensure you keep a backup copy of your HDP 2.4 configurations.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Back up the HDP directories for any hadoop components you have installed.

The following is a list of all HDP directories:

- `/etc/hadoop/conf`
- `/etc/hbase/conf`
- `/etc/hive-hcatalog/conf`
- `/etc/hive-webhcat/conf`
- `/etc/accumulo/conf`
- `/etc/hive/conf`
- `/etc/pig/conf`
- `/etc/sqoop/conf`
- `/etc/flume/conf`
- `/etc/mahout/conf`
- `/etc/oozie/conf`
- `/etc/hue/conf`
- `/etc/knox/conf`
- `/etc/zookeeper/conf`
- `/etc/tez/conf`
- `/etc/storm/conf`
- `/etc/falcon/conf`
- `/etc/slider/conf/`
- `/etc/ranger/admin/conf`, `/etc/ranger/usersync/conf` (If Ranger is installed, also take a backup of `install.properties` for all the plugins, ranger admin & ranger usersync.)
- Optional - Back up your userlogs directories, `${mapred.local.dir}/userlogs`.

2. Oozie runs a periodic purge on the shared library directory. The purge can delete libraries that are needed by jobs that started before the upgrade began and

that finish after the upgrade. To minimize the chance of job failures, you should extend the `oozie.service.ShareLibService.purge.interval` and `oozie.service.ShareLibService.temp.sharelib.retention.days` settings.

Add the following content to the `oozie-site.xml` file prior to performing the upgrade:

```
<property>
<name>oozie.service.ShareLibService.purge.interval</name>
<value>1000</value><description>
How often, in days, Oozie should check for old ShareLibs and LauncherLibs to
purge from HDFS.
</description>
</property>

<property>
<name>oozie.service.ShareLibService.temp.sharelib.retention.days</name>
<value>1000</value>
<description>
ShareLib retention time in days.</description>
</property>
```

3. Stop all long-running applications deployed using Slider:

```
su - yarn "usr/hdp/current/slider-client/bin/slider list"
```

For all applications returned in previous command, run `su - yarn "/usr/hdp/current/slider-client/bin/slider stop <app_name>"`

4. Stop all services (including MapReduce) except HDFS, ZooKeeper, and Ranger, and client applications deployed on HDFS.

See [Stopping HDP Services](#) for more information.

Component	Command
Accumulo	<code>/usr/hdp/current/accumulo-client/bin\$ /usr/hdp/current/accumulo-client/bin/stop-all.sh</code>
Knox	<code>cd \$GATEWAY_HOME su - knox -c "bin/gateway.sh stop"</code>
Falcon	<code>su - falcon "/usr/hdp/current/falcon-server/bin/falcon-stop"</code>
Oozie	<code>su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-stop.sh"</code>
WebHCat	<code>su - webhcat -c "/usr/hdp/hive-webhcat/sbin/webhcat_server.sh stop"</code>
Hive	<p>Run this command on the Hive Metastore and Hive Server2 host machine:</p> <pre>ps aux awk '{print \$1,\$2}' grep hive awk '{print \$2}' xargs kill >/dev/null 2>&1</pre> <p>Or you can use the following:</p> <pre>Killall -u hive -s 15 java</pre>

Component	Command
HBase RegionServers	<code>su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-daemon.sh --config /etc/hbase/conf stop regionserver"</code>
HBase Master host machine	<code>su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh --config /etc/hbase/conf stop master"</code>
YARN & Mapred Histro	<p>Run this command on all NodeManagers:</p> <pre>su - yarn -c "export /usr/hdp/current/hadoop-client/libexec && /usr/hdp/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop nodemanager"</pre> <p>Run this command on the History Server host machine:</p> <pre>su - mapred -c "export /usr/hdp/current/hadoop-client/libexec && && /usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-jobhistory-daemon.sh --config /etc/hadoop/conf stop historyserver"</pre> <p>Run this command on the ResourceManager host machine(s):</p> <pre>su - yarn -c "export /usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop resourcemanager"</pre> <p>Run this command on the ResourceManager host machine:</p> <pre>su -yarn -c "export /usr/hdp/current/hadoop-client/libex && /usr/hdp/current/hadoop-yarn-timelineserver/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"</pre> <p>Run this command on the YARN Timeline Server node:</p> <pre>su -l yarn -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-yarn/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"</pre>
Storm	<pre>storm kill topology-name</pre> <pre>sudo service supervisord stop</pre>
Spark (History server)	<code>su - spark -c "/usr/hdp/current/spark-client/sbin/stop-history-server.sh"</code>

5. If you have the Hive component installed, back up the Hive Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, see your database documentation.

Table 1.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<code>mysqldump \$dbname > \$outputfilename.sqlsbr</code>	<code>mysql \$dbname < \$inputfilename.sqlsbr</code>

Database Type	Backup	Restore
	For example: mysqldump hive > /tmp/mydir/ backup_hive.sql	For example: mysql hive < /tmp/mydir/ backup_hive.sql
PostgreSQL	sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr For example: sudo -u postgres pg_dump hive > / tmp/mydir/backup_hive.sql	sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr For example: sudo -u postgres psql hive < /tmp/ mydir/backup_hive.sql
Oracle	Export the database: exp username/password@database full=yes file=output_file.dmp	Import the database: imp username/password@database file=input_file.dmp

6. If you have the Oozie component installed, back up the Oozie metastore database.

These instructions are provided for your convenience. Check your database documentation for the latest backup instructions.

Table 1.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	mysqldump \$dbname > \$outputfilename.sql For example: mysqldump oozie > /tmp/mydir/ backup_hive.sql	mysql \$dbname < \$inputfilename.sql For example: mysql oozie < /tmp/mydir/ backup_oozie.sql
PostgreSQL	sudo -u \$username pg_dump \$databasename > \$outputfilename.sql For example: sudo -u postgres pg_dump oozie > /tmp/mydir/ backup_oozie.sql	sudo -u \$username psql \$databasename < \$inputfilename.sql For example: sudo -u postgres psql oozie < /tmp/mydir/ backup_oozie.sql
Oracle	Export the database: exp username/password@database full=yes file=output_file.dmp	Import the database: imp username/password@database file=input_file.dmp

7. **Optional:** Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 1.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	mysqldump \$dbname > \$outputfilename.sqlsbr For example:	mysql \$dbname < \$inputfilename.sqlsbr For example:

Database Type	Backup	Restore
	<code>mysqldump hue > /tmp/mydir/backup_hue.sql</code>	<code>mysql hue < /tmp/mydir/backup_hue.sql</code>
PostgreSQL	<code>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</code> For example: <code>sudo -u postgres pg_dump hue > / tmp/mydir/backup_hue.sql</code>	<code>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</code> For example: <code>sudo -u postgres psql hue < /tmp/ mydir/backup_hue.sql</code>
Oracle	Connect to the Oracle database using sqlplus. Export the database. For example: <code>exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr</code>	Import the database: For example: <code>imp username/password@database file=input_file.dmp</code>
SQLite	<code>/etc/init.d/hue stop</code> <code>su \$HUE_USER</code> <code>mkdir ~/hue_backup</code> <code>sqlite3 desktop.db .dump > ~/ hue_backup/desktop.bak</code> <code>/etc/init.d/hue start</code>	<code>/etc/init.d/hue stop</code> <code>cd /var/lib/hue</code> <code>mv desktop.db desktop.db.old</code> <code>sqlite3 desktop.db < ~/hue_backup/ desktop.bak</code> <code>/etc/init.d/hue start</code>

8. Back up the Knox data/security directory.

```
cp -RL /etc/knox/data/security ~/knox_backup
```

9. Save the namespace by executing the following commands:

```
su - hdfs  
  
hdfs dfsadmin -safemode enter  
  
hdfs dfsadmin -saveNamespace
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

10. Run the `fsck` command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.)

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-old-  
fsck-1.log"
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

11. Use the following instructions to compare status before and after the upgrade.

The following commands must be executed by the user running the HDFS service (by default, the user is hdfs).

- a. Capture the complete namespace of the file system. (The following command does a recursive listing of the root file system.)



Important

Make sure the namenode is started.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-old-lsr-1.log"
```



Note

In secure mode you must have Kerberos credentials for the hdfs user.

- b. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs dfsadmin -c "-report > dfs-old-report-1.log"
```

- c. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-old-report-1.log"
```

- d. **Optional:** You can copy all or unrecoverable only data storelibext-customer directory in HDFS to a local file system or to a backup instance of HDFS.

- e. **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.

12.Finalize any prior HDFS upgrade, if you have not done so already.

```
su - hdfs -c "hdfs dfsadmin -finalizeUpgrade"
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

13Stop remaining services (HDFS, Zookeeper, and Ranger).

See [Stopping HDP Services](#) for more information.

Component	Command
HDFS	<p>On all DataNodes:</p> <p>If you are running secure cluster, run following command as root:</p> <pre>/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode</pre> <p>Else:</p>

Component	Command
	<pre>su - hdfs -c "usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode"</pre> <p>If you are not running a highly available HDFS cluster, stop the Secondary NameNode by executing this command on the Secondary NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop secondarynamenode"</pre> <p>On the NameNode host machine(s)</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop namenode"</pre> <p>If you are running NameNode HA, stop the ZooKeeper Failover Controllers (ZKFC) by executing this command on the NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop zkfc"</pre> <p>If you are running NameNode HA, stop the JournalNodes by executing these command on the JournalNode host machines:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh ==config /etc/hadoop/conf stop journalnode"</pre>
Zookeeper Host machines	<pre>su - zookeeper "/usr/hdp/current/zookeeper-server/bin/zookeeper-server stop"</pre>
Ranger (XA Secure)	<pre>service ranger-admin stop</pre> <pre>service ranger-usersync stop</pre>

14. Back up your NameNode metadata.



Note

It's recommended to take a backup of the full `/hadoop.hdfs/namenode` path.

- a. Copy the following checkpoint files into a backup directory.

The NameNode metadata is stored in a directory specified in the `hdfs-site.xml` configuration file under the configuration value `"dfs.namenode.dir"`.

For example, if the configuration value is:

```
<property>
  <name>dfs.namenode.name.dir</name>
  <value>/hadoop/hdfs/namenode</value>
</property>
```

Then, the NameNode metadata files are all housed inside the directory `/hadoop.hdfs/namenode`.

- b. Store the layoutVersion of the namenode.

```
${dfs.namenode.name.dir}/current/VERSION
```

15. Verify that edit logs in `${dfs.namenode.name.dir}/current/edits*` are empty.

- a. Run: `hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out`

- b. Verify the edits.out file. It should only have OP_START_LOG_SEGMENT transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

- c. If edits.out has transactions other than OP_START_LOG_SEGMENT, run the following steps and then verify edit logs are empty.

- Start the existing version NameNode.
- Ensure there is a new FS image file.
- Shut the NameNode down:

```
hdfs dfsadmin - saveNamespace
```

16. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path
component in this version of HDFS.

Please rollback and delete or rename this path, or upgrade with the
-renameReserved key-value pairs option to automatically rename these
paths during upgrade.
```

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup.

For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, specify `-upgrade -renameReserved .snapshot=.my-snapshot, .reserved=.my-reserved`.

If no key-value pairs are specified with `-renameReserved`, the NameNode then suffixes reserved paths with:

```
.<LAYOUT-VERSION>.UPGRADE_RENAMED
```

For example: `.snapshot.-51.UPGRADE_RENAMED`.



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

17. Upgrade the JDK on all nodes to JDK 7 or JDK 8 before upgrading HDP.

1.2. Upgrade HDP 2.4 Components

The upgrade process to HDP 2.5.0 involves the following steps. See the [Release Notes](#) for the HDP 2.5.0.0 repo information. Select your OS:

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.4 components. This command uninstalls the HDP 2.4 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"spark*" "slider*" "hdp_mon_nagios_addons" "bigtop"
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
yum list installed | grep @HDP2.4
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP 2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo namestatus
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "hive-webhcat" "hive-hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "tez"
"storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark"
"slider" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

RHEL/CentOS/Oracle

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.4 components. This command uninstalls the HDP 2.4 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
yum list installed | grep @HDP2.4
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP 2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos5/2.x/updates/2.5.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "hue" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.4 components. This command uninstalls the HDP 2.4 components. It leaves the user data, and metadata, but removes your configurations:

```
zypper rm "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
yum list installed | grep @HDP2.4
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "hue" "tez"
"storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark"
"slider*" "hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5.0
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component files names should appear in the returned list.

SLES

1. On all hosts, clean the zypper repository.

```
zypper clean -all
```

2. Remove your old HDP 2.4 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*"
"pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*"
"falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common"
"hue-shell" "knox*" "ranger*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
zypper search --installed-only --repo HDP-2.4.0.0
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
http://public-repo-1.hortonworks.com/HDP/susellsp3/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "oozie" "collectd" "gccxml"
"pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "tez" "storm" "falcon"
"flume" "phoenix" "accumulo" "mahout" "knox" "spark" "spark-python"
"hdp_mon_nagios_addons" "slider*" "hive-webcat" "hive-hcatalog"
```

```
zypper up -r HDP-2.5.0
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component file names should appear in the returned list.

Ubuntu 12

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.4 components. This command uninstalls the HDP 2.4 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez.*"
"storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue.*"
"knox*" "spark*" "slider*" "hdp_mon_nagios_addons" --purge
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
yum list installed | grep @HDP2.4
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2.5.0.0/hdp.list -O /etc/apt/sources.list.d/hdp.list
```


6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm"
"falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*"
"hdp_mon_nagios_addons"
```

**Note**

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

Debian

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.4 components. This command uninstalls the HDP 2.4 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*"
"storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue*" "knox*"
"spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.4 component binaries are uninstalled:

```
yum list installed | grep @HDP2.4
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
debian<version>/2.x/updates/2.5.0.0/hdp.list -O /etc/apt/
sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.4 components:

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm"
"falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*"
"hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

1.3. Symlink Directories with hdp-select



Warning

HDP 2.5.0 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

1. Before you run hdp-select, remove one link:

```
rm /usr/bin/oozie
```

2. Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5.0.0-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5.0.0-2800
```

1.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available HDFS cluster, configure and restart Apache ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.

1. Replace your configuration after upgrading on all the ZooKeeper nodes. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "/usr/hdp/current/zookeeper-server/bin/zookeeper-server start"
```

1.5. Configure Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed in HDP 2.4. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed since HDP 2.4. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

Ubuntu/Debian

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

1.6. Start Hadoop Core



Warning

Before you start HDFS on a highly available HDFS cluster, you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

To start HDFS, run commands as the `$HDFS_USER`.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using

another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading on all the HDFS nodes. Replace the HDFS template configuration in `/etc/hdfs/conf`.
2. If you are upgrading from a highly available HDFS cluster configuration, start all JournalNodes. On each JournalNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start journalnode"
```



Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

3. If you are running HDFS on a highly available namenode, you must first start the Zookeeper service



Note

Perform this step only if you are on a highly available HDFS cluster.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc"
```

4. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually.

On the active NameNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the `-upgrade` option only once. After you have completed this step, you can bring up the NameNode using this command without including the `-upgrade` option.

To check if the Upgrade is in progress, check that the `"\previous"` directory has been created in the `\NameNode` and `\JournalNode` directories. The `"\previous"` directory contains a snapshot of the data before upgrade.

In a highly available HDFS cluster configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the

shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode.

To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the '-bootstrapStandby' flag. Do NOT start this standby NameNode with the '-upgrade' flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The bootstrapStandby command downloads the most recent fsimage from the active NameNode into the \$dfs.name.dir directory of the standby NameNode. You can enter that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

5. Verify that the NameNode is up and running:

```
ps -ef|grep -i NameNode
```

6. If you do not have a highly available HDFS cluster configuration (non_HA namenode), start the Secondary NameNode.



Note

Do not perform this step if you have a highly available HDFS cluster configuration.

On the Secondary NameNode host machine, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

7. Verify that the Secondary NameNode is up and running.



Note

Do not perform this step if you have a highly available HDFS cluster environment.

```
ps -ef|grep SecondaryNameNode
```

8. Start DataNodes.

On each of the DataNodes, enter the following command. Note: If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start datanode"
```

9. Verify that the DataNode process is up and running:

```
ps -ef|grep DataNode
```

10. Verify that NameNode can go out of safe mode.

```
>su - hdfs -c "hdfs dfsadmin -safemode wait"
```

You should see the following result: Safe mode is OFF

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

1.7. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

1. Run the `fsck` command on namenode as `$HDFS_USER`:

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-new-fsck-1.log"
```

You should see feedback that the filesystem under path `/` is HEALTHY.

2. Run `hdfs` namespace and report.

- a. List directories.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-new-lsr-1.log"
```

- b. Open the `dfs-new-lsr-1.log` and confirm that you can see the file and directory listing in the namespace.

- c. Run report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-new-report-1.log"
```

- d. Open the `dfs-new-report` file and validate the admin report.

3. Compare the namespace report before the upgrade and after the upgrade. Verify that user files exist after upgrade.

The file names are listed below:

```
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log
```

```
dfs-old-lsr-1.log < -- > dfs-new-lsr-1.log
```



Note

You must do this comparison manually to catch all errors.

4. From the NameNode WebUI, determine if all DataNodes are up and running.

```
http://<namenode>:<namenodeport>
```

5. If you are on a highly available HDFS cluster, go to the StandbyNameNode web UI to see if all DataNodes are up and running:

```
http://<standbynamenode>:<namenodeport>
```

6. If you are **not** on a highly available HDFS cluster, go to the SecondaryNameNode web UI to see if the secondary node is up and running:

```
http://<secondarynamenode>:<secondarynamenodeport>
```

7. Verify that read and write to hdfs works successfully.

```
hdfs dfs -put [input file] [output file]
```

```
hdfs dfs -cat [output file]
```

1.8. Configure YARN and MapReduce

After you upgrade Hadoop, complete the following steps to update your configs.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

In secure mode, you must have Kerberos credentials for the hdfs user.

1. Upload the MapReduce tarball to HDFS. As the HDFS user, for example 'hdfs':

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/mapreduce.tar.gz /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

```
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/mapreduce"
```

```
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<
$version>/mapreduce/mapreduce.tar.gz"
```

2. Make sure that the following properties are in `/etc/hadoop/conf/mapred-site.xml`:

- Make sure `mapreduce.application.framework.path` exists in `mapred-site.xml`:

```
<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}/mapreduce/mapreduce.tar.gz#mr-framework</
value>
</property>

<property>
  <name>yarn.app.mapreduce.am.admin-command-opts</name>
  <value>-Dhdp.version=${hdp.version}</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
  <value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}/hadoop/lib/native:/usr/
hdp/${hdp.version}/hadoop/
  lib/native/Linux-amd64-64</value>
</property>

<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.admin.reduce.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.application.classpath</name>
  <value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
  $PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/common/*,
  $PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/yarn/*,
  $PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
  $PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
  /usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
  /etc/hadoop/conf/secure</value>
```



```
</property>
```



Note

You do not need to modify `${hdp.version}`.



Note

If you are planning to use Spark in yarn-client mode, make Spark work in yarn-client mode 2.5.0.0-<\$version>.

3. Make sure the following property is in `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.application.classpath</name>
  <value>$HADOOP_CONF_DIR,/usr/hdp/${hdp.version}/hadoop-client/*,
    /usr/hdp/${hdp.version}/hadoop-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/lib/*</value>
</property>
```

4. On secure clusters only, add the following properties to `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.timeline-service.recovery.enabled:</name>
  <value>TRUE</value>
</property>

<property>
  <name>yarn.timeline-service.state-store.class: org.apache.hadoop.yarn.
server.timeline.recovery:</name>
  <value>LevelDbTimelineStateStore</value>
</property>

<property>
  <name>yarn.timeline-service.leveldb-state-store.path:</name>
  <value><the same as the default of "yarn.timeline-service-leveldb-
timeline-store.path</value>
</property>
```

5. For secure clusters, you must create and configure the `container-executor.cfg` configuration file:

- Create the `container-executor.cfg` file in `/usr/hdp/2.5.0.0-<version>/hadoop-yarn/bin/container-executor`.
- Insert the following properties:

```
<property>
  yarn.nodemanager.linux-container-executor.group=hadoop
  banned.users=hdfs,yarn,mapred
  min.user.id=1000
</property>
```

- `yarn.nodemanager.linux-container-executor.group` - Configured value of `yarn.nodemanager.linux-container-executor.group`. This must match the value of `yarn.nodemanager.linux-container-executor.group` in `yarn-site.xml`.
- `banned.users` - Comma-separated list of users who can not run container-executor.
- `min.user.id` - Minimum value of user id. This prevents system users from running container-executor.
- `allowed.system.users` - Comma-separated list of allowed system users.
- Set the file `/etc/hadoop/conf/container-executor.cfg` file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/bin/container-executor
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/bin/container-executor
```

1.9. Start YARN/MapReduce Services

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.



Note

The `su` commands in this section use "yarn" to represent the YARN Service user and `mapreduce` to represent the MAPREDUCE Service user. If you are using another name for these Service users, you need to substitute your Service user name for "yarn" or "mapreduce" in each of the `su` commands.

1. Manually clear the ResourceManager state store.

```
su - yarn -c "yarn resourcemanager -format-state-store"
```

2. Start the ResourceManager on all your ResourceManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh start resourcemanager"
ps -ef | grep -i resourcemanager
```

3. Start the TimelineServer on your TimelineServer host.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-timeline-servers/sbin/yarn-daemon.sh start timelineserver"
ps -ef | grep -i timelineserver
```

4. Start the NodeManager on all your NodeManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh
start nodemanager"

ps -ef | grep -i nodemanager
```

5. To start MapReduce, run the following commands:

```
su - mapreduce -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-
jobhistory-daemon.sh start historyserver"

ps -ef | grep -i jobhistoryserver
```

1.10. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user.

The job uses MapReduce to write 100MB of data into HDFS with RandomWriter

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.
jar
    randomwriter -Dtest.randomwrite.total_bytes=10000000 test-after-
upgrade.
```

You should see messages similar to:

```
map 0% reduce 0%
...map 100% reduce 100%
Job ... completed successfully
```

MapReduce upgraded successfully. You can now upgrade your other components.

Basic Troubleshooting

To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

`http://<resource manager host>:8088/cluster/nodes`

The number of active nodes should be equal to the number of nodemanagers.

Accessing error messages:

1. Access the ApplicationMaster WebUI to view the container logs.
2. At your console logs for MapReduce job, look for a line with this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track the job: http://<resource
manager host>:8088/proxy/application_1380673658357_0007/
```

3. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

1.11. Configure and Start Apache HBase



Note

The `su` commands in this section use "hbase" to represent the HBASE Service user. If you are using another name for your HBASE Service user, you need to substitute your HBASE Service user name for "hbase" in each of the `su` commands.

The `hbase.bucketcache.percentage.in.combinedcache` is removed in HDP 2.5.0. This simplifies the configuration of block cache. BucketCache configurations from HDP 2.4 needs to be recalculated to attain identical memory allotments in HDP 2.5.0. The L1 LruBlockCache is whatever `hfile.block.cache.size` is set to and the L2 BucketCache is whatever `hbase.bucketcache.size` is set to.

1. Replace your configuration after upgrading. Replace the Apache HBase template configuration in `/etc/hbase/conf`.
2. Start services. From root, assuming that `$HBASE_USER=hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

3. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

1.12. Configure Apache Phoenix

To configure Apache Phoenix, complete the following steps:

1. Add the following property to the `/etc/hbase/hbase-site.xml` file on all HBase nodes, the MasterServer, and all RegionServers to prevent deadlocks from occurring during maintenance on global indexes:

```
<property>
  <name>hbase.regionserver.wal.codec</name>
  <value>org.apache.hadoop.hbase.regionserver.wal.IndexedWALEditCodec</
value>
</property>
```

2. To enable user-defined functions, configure the following property in `/etc/hbase/conf` on all Hbase nodes.

```
<property>
  <name>phoenix.functions.allowUserDefinedFunctions</name>
  <value>true</value>
  <description>enable UDF functions</description>
</property>
```

3. Ensure the client side hbase-site.xml matches the server side configuration.
4. **(Optional)** To use local indexing, set the following three properties. These properties ensure collocation of data table and local index regions:



Note

The local indexing feature is a technical preview and considered under development. Do not use this feature in your production systems. If you have questions regarding this feature, contact Support by logging a case on our Hortonworks Support Portal at [//hortonworks.com/support/](http://hortonworks.com/support/).

Add the following two properties, if they do not already exist, to the master side configurations:

```
<property>
  <name>hbase.master.loadbalancer.class</name>
  <value>org.apache.phoenix.hbase.index.balancer.IndexLoadBalancer</value>
</property>

<property>
  <name>hbase.coprocessor.master.classes</name>
  <value>org.apache.phoenix.hbase.index.master.IndexMasterObserver</value>
</property>
```

Add the following, if it does not already exist, to the Region server side configurations:

```
<property>
  <name>hbase.coprocessor.regionserver.classes</name>
  <value>org.apache.hadoop.hbase.regionserver.LocalIndexMerger</value>
</property>
```

5. If the folder specified in hbase.tmp.dir property on hbase-site.xml does not exist, create that directory with adequate permissions.
6. Restart the HBase Master and Region Servers.

Configuring Phoenix to Run in a Secure Cluster

Perform the following additional steps to configure Phoenix to run in a secure Hadoop cluster:

1. To link the HBase configuration file with the Phoenix libraries:

```
ln -sf HBASE_CONFIG_DIR/hbase-site.xml PHOENIX_HOME/bin/hbase-site.xml
```

2. To link the Hadoop configuration file with the Phoenix libraries:

```
ln -sf HADOOP_CONFIG_DIR/core-site.xml PHOENIX_HOME/bin/core-site.xml
```



Note

When running the pssql.py and sqlline.py Phoenix scripts in secure mode, you can safely ignore the following warnings.

```
14/04/19 00:56:24 WARN util.NativeCodeLoader:  
Unable to load native-hadoop library for your platform...  
using builtin-java classes where applicable  
  
14/04/19 00:56:24 WARN util.DynamicClassLoader: Failed to identify the fs of  
dir hdfs://<HOSTNAME>:8020/apps/hbase/data/lib, ignored java.io.IOException:  
No FileSystem for scheme: hdfs
```

1.13. Configure and Start Apache Accumulo



Note

The `su` commands in this section use "accumulo" to represent the Accumulo Service user. If you are using another name for your Apache Accumulo Service user, you need to substitute your Accumulo Service user name for "accumulo" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/accumulo/conf` from the template to the conf directory in Accumulo hosts.
2. Start the services:

```
su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh  
`hostname` master"  
  
su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh  
`hostname` tserver"  
  
su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh  
`hostname` gc"  
  
su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh  
`hostname` tracer"  
  
su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh  
`hostname` monitor"
```

3. Check that the processes are running

```
ps -ef | grep accumulo
```

or visit `http://<hostname>:50095` in your browser

1.14. Configure and Start Apache Tez



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

To upgrade Apache Tez:

1. Copy your previously backed-up copy of `tez-site.xml` into the `/etc/tez/conf` directory.
2. Upload the Tez tarball to HDFS.

```
su - hdfs
hdfs dfs -mkdir -p /hdp/apps/<hdp_version>/tez/
hdfs dfs -put /usr/hdp/<hdp_version>/tez/lib/tez.tar.gz /hdp/apps/
<hdp_version>/tez/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/<hdp_version>/tez
hdfs dfs -chmod -R 444 /hdp/apps/<hdp_version>/tez/tez.tar.gz
```

Where `<hdp_version>` is the current HDP version, for example `2.5.0.0-2800`.

3. Edit the `tez.lib.uris` property in the `tez-site.xml` file to point to `/hdp/apps/<hdp_version>/tez/tez.tar.gz`

```
...
<property>
  <name>tez.lib.uris</name>
  <value>/hdp/apps/<hdp_version>/tez/tez.tar.gz</value>
</property>
...
```

Where `<hdp_version>` is the current HDP version, for example `2.5.0.0-2800`.

4. **Optional** Earlier releases of Tez did not have access control. In the current version of Tez, the default behavior restricts the ability to view the Tez history to only the owner of the job. To retain unrestricted access for non-secure clusters, set `tez.am.view-acls` set to `"*"`.
5. Change the value of the `tez.tez-ui.history-url.base` property to the url for the upgraded Tez View. For information on setting up the Tez view, see [Deploying the Tez View](#) in the HDP Ambari Views Guide.
6. **Run Tez Smoke Test**

To smoke test your Tez upgrade, you can run the following Tez Example job as a regular user.

MRRSleep Tez Example job sleeps for a defined period of time in mapper and reducer

```
hadoop jar /usr/hdp/current/tez-client/tez-tests-<version>.jar mrrsleep -m 1
-r 1 -mt 100 -rt 100
```

You should see messages similar to:

```
...DAG: State: SUCCEEDED Progress: 100%
DAG completed. FinalState=SUCCEEDED
```

Tez upgraded successfully. You can now upgrade your other components.

1.15. Configure and Start Apache Hive and Apache HCatalog



Note

The `su` commands in this section use "hive" to represent the Hive Service user. If you are using another name for your Hive Service user, you need to substitute your Hive Service user name for "hive" in each of the `su` commands.

1. Prior to starting the upgrade process, set the following in your hive configuration file:

```
datanucleus.autoCreateSchema=false
```

2. Copy the jdbc connector jar from `OLD_HIVE_HOME/lib` to `CURRENT_HIVE_HOME/lib`.
3. Upgrade the Apache Hive Metastore database schema. Restart the Hive Metastore database and run:

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -upgradeSchema -dbType" <$databaseType>
```

The value for `$databaseType` can be `derby`, `mysql`, `oracle`, or `postgres`.



Note

If you are using PostgreSQL, you should reset the Hive Metastore database owner to `<HIVE_USER>`:

```
psql -U <POSTGRES_USER> -c
ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER TO <HIVE_USER>
```



Note

If you are using Oracle 11, you may see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
optimize.mapjoin.mapreduce does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
heapsize does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
server2.enable.impersonation does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
semantic.analyzer.factory.impl does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.
convert.sortmerge.join.noconditionaltask does not exist
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.
internal:1521/XE
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver
Metastore connection User: hiveuser
Starting upgrade metastore schema from version 0.13.0 to 0.14.0
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql
Error: ORA-00955: name is already used by an existing object
(state=42000,code=955)
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.
oracle.sql: Schema script failed, errorcode 2
```



```
Completed upgrade-0.13.0-to-0.14.0.oracle.sql
schemaTool completed
```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.

4. Edit the hive-site.xml file and modify the properties based on your environment. Search for TODO in the file for the properties to replace.

- a. Edit the following properties in the hive-site.xml file:

```
<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache
  </description>
</property>
```

- b. **Optional:** To enable the Hive builtin authorization mode, make the following changes. If you want to use the advanced authorization provided by Ranger, refer to the [Ranger](#) instructions.

Set the following Hive authorization parameters in the hive-site.xml file:

```
<property>
  <name>hive.server2.enable.doAs</name>
  <value>>false</value>
</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
    StorageBasedAuthorizationProvider,org.apache.hadoop.hive.ql.security.
    authorization.MetaStoreAuthzAPIAuthorizeEmbedOnly</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
    SQLStdConfOnlyAuthorizeFactory</value>
</property>
```

Also set hive.users.in.admin.role to the list of comma-separated users who need to be added to admin role. A user who belongs to the admin role needs to run the "set role" command before getting the privileges of the admin role, as this role is not in the current roles by default.

Set the following in the hiveserver2-site.xml file.

```
<property>
  <name>hive.security.authenticator.manager</name>
```

```

<value>org.apache.hadoop.hive ql.security.
SessionStateUserAuthenticator</value>
</property>

<property>
  <name>hive.security.authorization.enabled</name>
  <value>true</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive ql.security.authorization.plugin.sqlstd.
SQLStdHiveAuthorizeFactory</value>
</property>

```

- c. For a remote Hive metastore database, set the IP address (or fully-qualified domain name) and port of the metastore host using the following hive-site.xml property value.

```

<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server.
    To enable HiveServer2, leave the property value empty.
  </description>
</property>

```

You can further fine-tune your configuration settings based on node hardware specifications, using the HDP utility script.

5. Start Hive Metastore.

On the Hive Metastore host machine, run the following command:

```

su - hive -c "nohup /usr/hdp/current/hive-metastore/bin/hive
--service metastore -hiveconf hive.log.file=hivemetastore.log
>/var/log/hive/hivemetastore.out 2>/var/log/hive/
hivemetastoreerr.log &"

```

6. Start Hive Server2.

On the Hive Server2 host machine, run the following command:

```

su - hive

nohup /usr/hdp/current/hive-server2/bin/hiveserver2 -hiveconf
hive.metastore.uris=" " -hiveconf hive.log.file=hiveserver2.log
>/var/log/hive/hiveserver2.out 2> /var/log/hive/
hiveserver2err.log &

```

1.16. Configure and Start Apache Oozie



Note

The duration of the Oozie upgrade is dependent on the amount of job history stored in ooziedb. This history must be backed up and restored during the

upgrade process. Best practice when planning for upgrade is to backup ooziedb from your production oozie server and restore it to a test or development oozie server. This can help you estimate the time that is required to upgrade Oozie during your production upgrade.



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and "oozie" to represent the Oozie Service user. If you are using another name for your HDFS Service user or your Oozie Service user, you need to substitute your Service user names for "hdfs" or "oozie" in each of the `su` commands.

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore oozie-site.xml from your backup to the conf directory on each oozie server and client.
2. Copy the JDBC jar to libext-customer:

- a. Create the /usr/hdp/current/oozie/libext-customer directory.

```
cd /usr/hdp/current/oozie-server
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the libext-customer directory.

```
chmod -R 777 /usr/hdp/current/oozie-server/libext-customer
```

3. Copy these files to the libext-customer directory

```
cp /usr/hdp/current/hadoop-client/lib/hadoop*lzo*.jar /usr/hdp/
current/oozie-server/libext-customer

cp /usr/share/HDP-oozie/ext.zip /usr/hdp/current/oozie-server/
libext-customer/
```

Also, copy Oozie db jar in libext-customer.

4. If Falcon was also installed and configured before upgrade in HDP 2.4.x, then after upgrade you might also need to do the following:

```
cp /usr/hdp/current/falcon-server/oozie/ext/falcon-oozie-el-
extension-*.jar /usr/hdp/current/oozie-server/libext-customer
```

5. Extract share-lib.

```
/usr/hdp/current/oozie/bin/oozie-setup.sh sharelib create -fs
hdfs://<namenode>:8020
```

To verify that the sharelibs extracted correctly, run the following command:

```
oozie admin -oozie http://<oozie server host address>:11000/  
oozie -shareliblist
```

There should be:

- Available ShareLib
- oozie
- hive
- distcp
- hcatalog
- sqoop
- mapreduce-streaming
- pig

Change the ownership and permissions of the oozie directory:

```
su -l hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie"  
su -l hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

Add the two zip files from spark-client phone to spark sharelib:

```
hdfs dfs -put /usr/hdp/current/spark-client/python/lib/py4j-0.9-src.zip /  
user/oozie/share/lib/lib_<timestamp>/spark
```

```
hdfs dfs -put /usr/hdp/current/spark-client/python/lib/pyspark.zip /user/  
oozie/share/lib/lib_<timestamp>/spark
```

6. If a previous version of Oozie was created using auto schema creation, run the following SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

7. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh  
upgrade -run"
```

8. As root, prepare the Oozie WAR file.

```
chown oozie:oozie /usr/hdp/current/oozie-server/oozie-server/  
conf/server.xml
```

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-setup.sh  
prepare-war -d /usr/hdp/current/oozie-server/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar
New Oozie WAR file with added 'JARs' at /var/lib/oozie/oozie-server/webapps/
oozie.war
```

9. Make sure that following property is added in `oozie-log4j.properties`:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where `${oozie.instance.id}` is determined by oozie, automatically.

10.If you have custom Oozie actions, you must define them in `oozie-site.xml`. Edit the `/etc/oozie/conf/oozie-site.xml` file and add the following properties:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>[Comma separated list of custom actions]</value>
</property>
<property>
  <name>oozie.service.SparkConfigurationService.spark.configuration</name>
  <value>*=/etc/spark/conf/</value>
</property>
```

For example, if you have added Spark Action, enter the following:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>spark-action-0.1.xsd</value>
</property>
```

11.Configure HTTPS for the Oozie server.

- a. Create a self signed certificate or get certificate from a trusted CA for the Oozie Server
- b. Import the certificate to the client JDK trust store on all client nodes.
- c. In the Ambari Oozie configuration, set the following environment variables in `oozie-env.sh`, adding them if it does not exist:

```
export OOZIE_HTTPS_PORT=11443
export OOZIE_HTTPS_KEYSTORE_FILE=/home/oozie/.keystore
export OOZIE_HTTPS_KEYSTORE_PASS=password
```

- d. Change `OOZIE_HTTP_PORT={{oozie_server_port}}` to `OOZIE_HTTP_PORT=11000`.

- e. Set the `oozie.base.url` to the HTTPS address.

- f. Save the configuration, and restart the Oozie components.

12. Use the `/usr/hdp/current/oozie-server/bin/oozied.sh` script with the `start` parameter to start the Oozie server.

13. Check processes.

```
ps -ef | grep -i oozie
```

14. When needed, the `/usr/hdp/current/oozie-server/bin/oozied.sh` script with the `stop` parameter stops the Oozie server.

1.17. Configure and Start Apache WebHCat



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and webhcat to represent the WebHCat Service user. If you are using another name for these Service users, you need to substitute your Service user name for "hdfs" or "webhcat" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/webhcat/conf` from the template to the conf directory in webhcat hosts.
2. Modify the Apache WebHCat configuration files.
 - a. Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User` (in this example, hdfs):

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/"

su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/"

su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /
hdp/apps/2.5.0.0-<$version>/pig/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.
gz /hdp/apps/2.5.0.0-<$version>/hive/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.
gz /hdp/apps/2.5.0.0-<$version>/sqoop/"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/pig.
tar.gz"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/
hive.tar.gz"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/
sqoop.tar.gz"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

- b. Update the following properties in the webhcat-site.xml configuration file, as their values have changed:

```
<property>
  <name>templeton.pig.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/pig/pig.tar.gz</value>
</property>

<property>
  <name>templeton.hive.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/hive/hive.tar.gz</value>
</property>

<property>
  <name>templeton.streaming.jar</name>
  <value>hdfs:///hdp/apps/${hdp.version}/mapreduce/
    hadoop-streaming.jar</value>
  <description>The hdfs path to the Hadoop streaming jar file.</
description>
</property>

<property>
  <name>templeton.sqoop.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/sqoop/sqoop.tar.gz</value>
  <description>The path to the Sqoop archive.</description>
</property>

<property>
  <name>templeton.sqoop.path</name>
  <value>sqoop.tar.gz/sqoop/bin/sqoop</value>
  <description>The path to the Sqoop executable.</description>
</property>

<property>
  <name>templeton.sqoop.home</name>
  <value>sqoop.tar.gz/sqoop</value>
  <description>The path to the Sqoop home in the exploded archive.
    </description>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- c. Add the following property if it is not present in webhcat-site.xml:

```
<property>
  <name>templeton.libjars</name>
  <value>/usr/hdp/current/zookeeper-client/zookeeper.jar,/usr/hdp/current/
hive-client/lib/hive-common.jar</value>
  <description>Jars to add the classpath.</description>
</property>
```

- d. Remove the following obsolete properties from webhcat-site.xml:

```
<property>
  <name>templeton.controller.map.mem</name>
```

```
<value>1600</value>
<description>Total virtual memory available to map tasks.</description>
</property>

<property>
  <name>hive.metastore.warehouse.dir</name>
  <value>/path/to/warehouse/dir</value>
</property>
```

- e. Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHCat to impersonate your additional HDP 2.5.0 groups and hosts:

```
<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>
</property>

<property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>
</property>
```

Where:

`hadoop.proxyuser.hcat.group`

Is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

`hadoop.proxyuser.hcat.hosts`

A comma-separated list of the hosts which are allowed to submit requests by 'hcat'.

3. Start WebHCat:

```
su - webhcat -c "/usr/hdp/current/hive-webhcat/sbin/
webhcat_server.sh start"
```

4. Smoke test WebHCat.

- a. If you have a non-secure cluster, on the WebHCat host machine, run the following command to check the status of WebHCat server:

```
curl http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

You should see the following return status:

```
"status": "ok", "version": "v1"
```

- b. If you are using a Kerberos secure cluster, run the following command:

```
curl --negotiate -u: http://$WEBHCAT_HOST_MACHINE:50111/
templeton/v1/status
```

You should see the following return status


```
{ "status": "ok", "version": "v1" } [machine@acme]$
```

1.18. Configure Apache Pig

Replace your configuration after upgrading. Copy `/etc/pig/conf` from the template to the `conf` directory in pig hosts.

1.19. Configure and Start Apache Sqoop



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Copy `/etc/sqoop/conf` from the template to the `conf` directory in sqoop hosts.
2. As the HDFS Service user, upload the Apache Sqoop tarball to HDFS.

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"

su - hdfs -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
```

3. If you are using the MySQL database as a source or target, then the MySQL connector jar must be updated to 5.1.29 or later.

Refer to the MySQL web site for information on updating the MySQL connector jar.

1.20. Configure, Start, and Validate Apache Flume

1. Replace your configuration after upgrading. Copy `/etc/flume/conf` from the template to the `conf` directory in Flume hosts.
2. By default, Apache Flume does not start running immediately upon installation. To validate your Flume upgrade, replace your default `conf/flume.conf` with the provided `flume.conf` file, restart Flume, and see if the data is flowing by examining the destination.

Use this `flume.conf` file:

```
#1. Name the components on this agent
a1.sources = r1
```

```

a1.sinks = k1
a1.channels = c1

#2.Describe/configure the source
a1.sources.r1.type = seq

#3. Describe the sink
a1.sinks.k1.type = file_roll
a1.sinks.k1.channel = c1
a1.sinks.k1.sink.directory = /tmp/flume

#4. Use a channel which buffers events in memory
a1.channels.c1.type = memory

#5. Bind the source and sink to the channel
a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1

```

3. After starting Flume, check /tmp/flume to see if there are any files there. The files should contain simple sequential numbers.
4. After validating, stop Flume and revert changes to flume.conf to prevent your disk from filling up.

1.21. Configure, Start, and Validate Apache Mahout



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Replace your configuration after upgrading. Copy /etc/mahout/conf from the template to the conf directory in mahout hosts.

To validate Apache Mahout:

1. Create a test user:

```
su - hdfs -c "dfs -put /tmp/sample-test.txt /user/testuser"
```

2. Set up mahout to convert the plain text file sample-test.txt into a sequence file that is in the output directory mahouttest.

```
mahout seqdirectory --input /user/testuser/sample-test.txt --output /user/testuser/mahouttest --charset utf-8
```

1.22. Configure and Start Hue



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using

another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

For HDP 2.5.0, use the Hue version shipped with HDP 2.5.0. If you have a previous version of Hue, use the following steps to upgrade Hue.

1. Migrate hue.ini setting from your old hue.ini configuration file to new hue.ini configuration file.
2. If you are using the embedded SQLite database, remember to restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database. For example:

```
su - hue
cd /var/lib/hue
mv desktop.db desktop.db.old
sqlite3 desktop.db < ~/hue_backup/desktop.bak
exit
```

3. Synchronize Database

```
cd /usr/lib/hue
source ./build/env/bin/activate
hue syncdb
hue migrate
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

1.23. Configure and Start Apache Knox

When working with the Apache Knox Gateway in your Hadoop cluster, it is important you have the latest version of Knox installed so you can take advantage of new features and enhancements, in addition to ensuring your instance of Knox is in sync with other Hadoop components (e.g. Ranger, Spark, Hive, Hue, etc.) for stability and performance. For example, if you need to upgrade your Hadoop cluster from 2.4 to 2.5, you should also make sure that your individual Hadoop components are also upgraded to the latest version.

HDP enables you to perform a rolling upgrade in 2.2.x and onward. A rolling upgrade means that you can upgrade a component, or the entire Hadoop stack, without losing service, and your users can continue to use the cluster and run jobs with no application or server downtime. The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple instances of the gateway and a load balancer serving each gateway instance from a single URL.

The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple

instances of the gateway and a load balancer serving each gateway instance from a single URL. Once the upgrade process is completed, you are up and running with the latest version of Knox on each server you have designated as a Knox server.



Note

In this document, whenever you see a `{ }` with a value inside, this denotes a value you must define.

1.23.1. Upgrade the Knox Gateway

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Knox manually to the latest version. Because “rolling upgrades” are now supported in HDP 2.5.0, it is not important which version of Knox you are currently running, only that you have an instance of the Knox Gateway running.



Note

If you have not already installed Knox, refer to the "Install the Knox RPMs on the Knox Server" section of the *Non-Ambari Cluster Installation Guide* for instructions on how to install and configure the Knox Gateway.

Before upgrading the Knox Gateway, there are a several steps you must follow to ensure your configuration files, settings, and topology files can be copied to the new Knox Gateway instance when the upgrade is complete, which are described below.

1. Back up your existing `conf` directory if you have not already done so.
2. Stop each Knox server if you have not already done so.

```
su -l Knox /usr/hdp/{the current Knox version}/knox/bin/gateway.sh stop
```

3. Select the HDP server version you are upgrading to after you have stopped each Knox server if you have not already done so.

```
hdp-select set Knox-server {the HDP server version}
```

4.



Note

The `su` commands in this section use "Knox" to represent the Knox Service user. If you are using another name for your Knox Service user, you need to substitute your Knox Service user name for "Knox" in each of the `su` commands.

For HDP 2.5.0, the default paths for Knox change. Upgrade Knox in order to update these paths.

- a. Restore the backed up security directory. This places the master secret and keystores back in place for the new deployment.

- b. Start the Gateway:

```
su -Knox -c "/usr/hdp/current/Knox-server/bin/gateway.sh start"
```

- c. Unzip your previously saved configuration directory (the `conf` directory you backed up in step 1) into the new `/var/log/knox/gateway.conf` directory to import these files.
- d. Restart the Knox server to complete the upgrade.

```
su -l knox /usr/hdp/{the new HDP server version}/knox/bin/gateway.sh
start
```

1.23.2. Verify the Knox Upgrade

To verify the upgrade was successful, follow the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful.
2. Verify you have cluster access using the `LISTSTATUS` WebHDFS API call.

```
curl -ivk -u {user}:{password} https://{knox host}:8443 /gateway/webhdfs/v1/
tmp?op=LISTSTATUS
```

3. Verify the Knox version using the Knox Admin service and Version API.

```
curl -ivk -u {adminuser}:{adminpassword} https://{knox host}:8443 /gateway/
admin/v1/version
```



Note

The Admin API requires you to be a member of an Admin group, as specified in the `admin.xml` authorization provider.

When you have verified the Knox upgrade was successful, you can begin using Knox. If, however, the upgrade was unsuccessful, you need to downgrade the Knox Gateway to the previous version. The steps to downgrade the Knox Gateway are described in the next section.

1.23.3. Downgrade the Knox Gateway to the Previous Version

If the Knox Gateway upgrade was unsuccessful, you need to downgrade Knox to the previous version to ensure you have a working Knox Gateway for your cluster. To downgrade Knox, follow the steps listed below.

1. For each server running Knox, stop the server.

```
su -l knox /usr/hdp/{current HDP server version}/knox/bin/gateway.sh stop
```

2. Select the HDP server version you want to use to downgrade your Knox Gateway.

```
hdp-select set knox-server {previous HDP server version}
```

3. Restart the server.

```
su -l knox /usr/hdp/{previous HDP server version}/knox/bin/gateway.sh start
```

1.23.4. Verify the Knox Downgrade Was Successful

When the restart is complete, verify you are running an older version of Knox by following the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful
2. Verify you have cluster access using the `LISTSTATUSWebHDFS` API call.
3. Check the Knox version using the Knox Admin service and Version API using the following command:

```
curl -ivk -u {adminuser}"{adminpassword} https://{knox host}:8443 /gateway/admin/v1/version
```

1.24. Configure and Validate Apache Falcon



Note

In HDP 2.5.0, if authorization is enabled (for example, in the properties file with `*.falcon.security.authorization.enabled=true`) then Access Control List (ACL) is mandated for all entities.

Upgrade Apache Falcon after you have upgraded HDFS, Hive, Oozie, and Pig. Stop Oozie jobs while running Falcon.

1. Replace your configuration after upgrading. Copy `/etc/falcon/conf` from the template to the conf directory in falcon hosts.
2. Check your Falcon entities. There should be no changes, but in some cases you may need to update your entities post-upgrade.
3. In HDP 2.5.0 for Falcon, TLS is enabled by default. When TLS is enabled, Falcon starts on `https://<falcon_host>.15443/`. You can disable TLS by adding the following line to the `startup.properties` file.

```
"*.falcon.enableTLS=false
```

4. If Transport Layer Security (TLS) is disabled, check the `client.properties` file to make sure the property `"falcon.uri"` is set as follows:

```
falcon.uri=http://<falcon_host>:15000/
```

1.25. Configure and Start Apache Storm



Note

The `su` commands in this section use `"zookeeper"` to represent the Zookeeper Service user. If you are using another name for your Zookeeper Service user, you need to substitute your Zookeeper Service user name for `"zookeeper"` in each of the `su` commands.

Apache Storm is fairly independent of changes to the HDP cluster:

1. Deactivate all running topologies.

2. Delete all states under zookeeper:

```
/usr/hdp/current/zookeeper-client/bin/zkCli.sh (optionally in  
secure environment specify -server zk.server:port)
```

3. `rmr /storm`

4. Delete all states under the storm-local directory:

```
rm -rf <value of stormlocal.dir>
```

5. Stop Storm Services on the storm node.

6. Update the following configs in `storm.yaml`:

- `storm.thrift.transport=org.apache.storm.security.auth.SimpleTransportPlugin`
- `storm.messaging.transport=org.apache.storm.messaging.netty.Context`
- `nimbus.topology=org.apache.storm.nimbus.DefaultTopologyValidator`
- `topology.spout.wait.strategy=org.apache.storm.spout.SleepSpoutWaitStrategy`
- `topology.kryo.factory=org.apache.storm.serialization.DefaultKryoFactory`
- `topology.tuple.serializer=org.apache.storm.serialization.types.ListDelegator`
- `nimbus.authorizer=org.apache.storm.security.suth.authorizer.SimpleACLAuthorizer`
(applicable only in a secure cluster)
- `drpc.authorizer=org.apache.storm.security.auth.authorizer.DRPCSimpleACLAuthorizer`
(applicable only in a secure cluster)
- `ui.filter=org.apache.storm.secuity.auth.KerberosPrincipalToLocal`
(applicable only in a secure cluster)

7. Stop ZooKeeper Services on the storm node.

```
su - zookeeper -c "export ZOOCFGDIR=/etc/zookeeper/conf ; export  
ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-env.sh ; /  
usr/lib/zookeeper/bin/zkServer.sh stop"
```

8. Remove Storm and zookeeper from the storm node and install the HDP 2.5.0 version:

- For **RHEL/CentOS/Oracle Linux**:

```
yum erase storm
```

```
yum erase zookeeper
```

```
yum install storm
```

```
yum install zookeeper
```

- For **SLES**:

```
zypper rm storm
zypper rm zookeeper
zypper install storm
zypper install zookeeper
```

- For **Ubuntu/Debian**:

```
apt-get remove storm --purge
apt-get remove zookeeper --purge
apt-get install storm
apt-get install zookeeper
```

9. Replace your configuration after upgrading. Copy `/etc/storm/conf` from the template to the conf directory .

10. Replace your ZooKeeper configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.

11. Ensure ZooKeeper is running. On the storm node, run the following command:

```
su - zookeeper -c "source /etc/zookeeper/conf/zookeeper-env.sh; export
ZOO_CFG_DIR=/etc/zookeeper/conf; /usr/hdp/current/zookeeper-server/bin/
zkServer.sh start >> $ZOO_LOG_DIR/zoo.out\""
```

where

- `$ZOO_LOG_DIR` is the directory where ZooKeeper server logs are stored. For example, `/var/log/zookeeper`.

12. Start nimbus, then supervisor/ui/drpc/logviewer:

```
/usr/hdp/current/storm-nimbus/bin/storm nimbus.
```

13. Start Storm, using a process controller, such as supervisor:

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm
supervisor
```

You can use the same command syntax to start Storm using nimbus/ui and logviewer.

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm nimbus
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm ui
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm logviewer
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm drpc
```


1.26. Configure and Start Apache Ranger

This section describes how to upgrade the Ranger service.

1.26.1. Preparing Your Cluster to Upgrade Ranger

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Apache Ranger manually to the latest version. This section describes the steps you need to follow to prepare your cluster for the Ranger upgrade.

1. Back up the following Ranger configuration directories:

- Ranger Policy Administration Service

```
/etc/ranger/admin/conf
```

- Ranger UserSync

```
/etc/ranger/usersync/conf
```

- Ranger Plugins:

- Hadoop

```
/etc/hadoop/conf
```

- Hive

```
/etc/hive/conf
```

- HBase

```
/etc/hbase/conf
```

- Knox

```
/etc/knox/conf
```

- Storm

```
/etc/storm/conf
```

2. Backup the Ranger Policy and Audit databases. Make sure to take note of the following details in the `install.properties` file:

- db_host
- db_name
- db_user
- db_password
- policy manager configuration
- LDAP directory configuration

- LDAP settings
- LDAP AD domain
- LDAP URL

```
mysqldump -u root -p root ranger > dest_dir/filename.sql  
mysqldump -u root -p root ranger_audit > dest_dir/audit_filename.sql
```

1.26.2. Stop the Ranger Services

Now that you have prepared your cluster for the Ranger upgrade, you need to stop the Ranger Admin and Ranger UserSync services. To stop the Ranger services, perform the steps described below.

1. Stop the Ranger Policy Admin service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-admin stop
```

2. Stop the Ranger UserSync service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-usersync stop
```

3. Stop the applicable services using the Ranger plugin (HDFS, HBase, Hive, Knox, Storm).

See [Stopping HDP Services](#) for more information.

1.26.3. Preparing the Cluster for Upgrade

Before you begin the upgrade process, you need to perform a series of steps to prepare the cluster for upgrade. These steps are described in the "*Getting Ready To Upgrade*" section of this guide, which you need to follow before continuing to upgrade Ranger. Some of these steps include:

- Backing up HDP directories
- Stopping all long-running applications and services.
- Backing up the Hive and Oozie metastore databases.
- Backing up Hue
- Backing up specific directories and configurations

1.26.4. Registering the HDP 2.5.0 Repo

After you have prepared your cluster for the upgrade, you need to register the HDP 2.5.0 repo. This requires you to perform the following steps:

1. (Optional) The Ranger components should already have installed in the at the beginning of the HDP upgrade process, but you can use the following commands to confirm that the Ranger packages have been installed:

```
hdp-select status ranger-admin
hdp-select status ranger-usersync
```

If the packages have not been installed, you can use the install commands specific to your OS. For example, for RHEL/CentOS you would use the following commands to install the packages.

```
yum install ranger_2_5_*-admin
yum install ranger_2_5_*-usersync
```

2. Select the Ranger Admin and Ranger UserSync versions you want to use.

```
hdp-select set ranger-admin <HDP_server_version>
hdp-select set ranger-usersync <HDP_server_version>
```

3. Change ownership of `/etc/ranger/admin/conf/` and `/etc/ranger/usersync/conf/` to `ranger:ranger`:

```
chown -R ranger:ranger /etc/ranger/admin/conf/
```

4. Solr must be installed and configured before installing RangerAdmin or any of the Ranger component plugins.

For information regarding installation and configuration of Solr, see [Using Apache Solr for Ranger Audits](#).

5. Update the `install.properties` file to migrate the database credentials properties and `POLICYMGR_EXTERNAL-URL` property from HDP 2.4. to HDP 2.5.0.

Table 1.4. Ranger_Admin install.properties names and values

Property Name	Property Value
DB_FLAVOR	MySQL(ORACLE POSTGRES MSSQL SQLA)
db_root_user	root
db_root_password	Password of db (eg: vagrant)
db_host	Hostname : where your db does exist
polycmgr_external_url	http://<hostname>:6080
polycmgr_http_enabled	true
authentication_method	UNIX(LDAP ACTIVE_DIRECTORY UNIX NONE)
audit_store	solr
audit_solr_urls	http://<solr_host>:6083/solr/ranger_audits



Note

When you migrate to new version, you have to remove `/usr/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

6. Install the Ranger Admin component. Be sure to set the `JAVA_HOME` environment variable if it is not already set.

```
cd /usr/hdp/current/ranger-admin/
```

```
cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-site.xml ews/webapp/WEB-INF/classes/conf/

cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-default-site.xml ews/webapp/WEB-INF/classes/conf/

cp ews/webapp/WEB-INF/classes/conf.dist/security-applicationContext.xml ews/webapp/WEB-INF/classes/conf/

./setup.sh
```

This should successfully install the Ranger Admin component.

7. Start the Ranger Admin component.

```
service ranger-admin start
```

8. Configure and setup the Ranger UserSync component by migrating the properties from the HDP 2.4 `install.properties` file (POLICY_MGR_URL, SYNC_SOURCE and LDAP/AD properties).

Table 1.5. Ranger_Usersync install.properties names and values

Property Name	Property Value
POLICY_MGR_URL	http://<hostname>:6080
SYNC_SOURCE	unix
SYNC_INTERVAL	5



Note

When you migrate to new version, you have to remove `/user/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

9. Install the Ranger UserSync component. Be sure to set the JAVA_HOME component if it is not already set.

```
cd /usr/hdp/current/ranger-usersync/

./setup.sh
```

10 Start the Ranger UserSync component.

```
service ranger-usersync start
```

1.26.5. Install the Ranger Components

Next, you need to re-install each Ranger component again to ensure you have the latest version. Because you have already upgraded your HDP stack, you only need to follow the instructions in the *Non-Ambari Cluster Installation Guide* to install each Ranger component. You must install the following Ranger components:

- Ranger Policy Admin
- Ranger UserSync

- Ranger Plugins:
 - HDFS
 - HBase
 - Hive
 - Knox
 - Storm
 - Solr
 - Kafka
 - YARN



Note

When installing each Ranger component, you also need to make sure you upgrade each individual component to version 2.5.0 before restarting each service.

1.26.6. Restart the Ranger Services

Once you have re-installed each Ranger component, you then need to restart these components to ensure the new configurations are loaded in your cluster. This includes restarting the Policy Admin and UserSync components, NameNode, and each Ranger plugin.



Note

Before restarting the NameNode, make sure to remove the `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable the NameNode after finishing the upgrade.

The *Non-Ambari Cluster Installation Guide* describes how you can start the following Ranger services:

- Ranger Policy Admin service

```
service ranger-admin start
```

- Ranger UserSync service

```
service ranger-usersync start
```

1.26.7. Enable Ranger Plugins

The final step in the Ranger upgrade process requires you to re-enable the Ranger plugins. Although you are only required to enable HDFS in your cluster, you should re-enable all of the Ranger plugins because class names have changed for the 2.5.0 release, and to ensure smooth operation of Ranger services in your cluster.



Note

When you enable each Ranger plugin, be sure to remove all 2.4 class name values.



Note

Re-enabling a Ranger plugin does not affect policies you have already created. As long as you use the same database as the Policy store, all of your data remains intact.

To re-enable the Ranger plugins, use the links listed below to access instructions in the *Non-Ambari Cluster Installation* guide that describe editing the `install.properties` file and enabling the Ranger plugins:



Important

Before enabling the HDFS plugin, remove `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable this plugin after the upgrade is complete.

- [HDFS Plugin](#)
- [YARN Plugin](#)
- [Kafka Plugin](#)
- [HBase Plugin](#)
- [Hive Plugin](#)
- [Knox Plugin](#)
- [Storm Plugin](#)

1.26.8. Enable KMS Configuration

Configure and setup the Ranger KMS configuration by editing the KMS `install.properties` file:

Table 1.6. KMS `install.properties` names and values

Property Name	Property Name
DB_FLAVOR	MYSQL
db_root_user	foot
db_root_password	<db password>
db_host	<db hostname>
db_name	rangerkms (default)
db_user	rangerkms (default)
db_password	<kms db password>
hadoop_conf	/etc/hadoop/conf
POLICY_MGR_URL	http://<hostname>:6080

Property Name	Property Name
REPOSITORY_NAME	kmsdev
XAAUDIT.SOLR.ENABLE	false (default), change to true if desired
XAAUDIT.SOLR.URL	http://<solr_host>:6083/solr/ranger_audits

1.26.9. Configure and Start Apache Ranger on a Kerberized Cluster

Beginning with HDP 2.5.0, kerberos authentication is supported for Ranger and its plugins. Use the this section if you have an HDP 2.4 cluster in a kerberized environment with Ranger in simple authentication mode that you want to upgrade.

For additional information regarding Kerberos, refer to [Kerberos Overview](#) in the *Hadoop Security Guide*.

1.26.9.1. Create Keytabs and Principals

1. Follow these steps to create keytabs and principals:

For Ranger Admin:

- a. Create rangeradmin/<FQDN of Ranger Admin>@<REALM>.

b.

```
> kadmin.local
> addprinc -randkey rangeradmin/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangeradmin/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

For Ranger Lookup:

- a. Create rangerlookup/<FQDN of Ranger Admin>@<REALM>.

b.

```
> kadmin.local
> addprinc -randkey rangerlookup/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerlookup/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/ranger-
upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Usersync:

- a. Create rangerusersync/<FQDN>@<REALM>.

b.

```
> kadmin.local
> addprinc -randkey rangerusersync/<FQDN of Ranger usersync>
Eg: addprinc -randkey rangerusersync/ranger-upgrade-0707-2.openstacklocal
```

```
> xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Tagsync

a. Create rangertagsync/<FQDN>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangertagsync/<FQDN of Ranger tagsync>
Eg: addprinc -randkey rangertagsync/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger KMS:

a. Create rangerkms/<FQDN of Ranger Admin>@<REALM>

```
b.
> kadmin.local
> addprinc -randkey rangerkms/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerkms/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/<FQDN of Ranger
Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

2. If the Kerberos server and admin are on different hosts, copy the keytab on the admin host, assign permission to user `ranger`, and change the permissions.

```
scp the <rangeradmin_keytab_file> to <new_path>
chown ranger <rangeradmin_keytab_path>
chmod 400 <rangeradmin_keytab_path>
```

3. Use `kdestroy` to delete the Kerberos credentials cache file.
4. Set the following properties and values in the `/etc/hadoop/conf/core-site.xml` file:

Table 1.7. Properties for the `/etc/hadoop/conf/core-site.xml` file

Property Name	Property Value
fs.defaultFS	hdfs://ranger-upgrade-0707-2.openstacklocal:8020
hadoop.security.authentication	kerberos
hadoop.security.authorization	true
hadoop.security.auth_to_local	RULE:[1:\$1@\$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/ RULE:[1:\$1@\$0](.*@EXAMPLE.COM)s/@.*/ RULE:[2:\$1@\$0](dn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0](nn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0]

Property Name	Property Value
	(rangeradmin@EXAMPLE.COM)s/.*/ranger/ RULE: [2:\$1@\$0](rangerusersync@EXAMPLE.COM)s/.*/ rangerusersync/ DEFAULT

See [Creating Mappings Between Principals and UNIX Usernames](#) in the *Hadoop Security Guide*.

The following is an example of a `core-site.xml` file with the properties set for Kerberos:

```
<configuration>

  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://ranger-upgrade-0707-2.openstacklocal:8020</value>
    <final>true</final>
  </property>

  <property>
    <name>hadoop.security.authentication</name>
    <value>kerberos</value>
  </property>

  <property>
    <name>hadoop.security.authorization</name>
    <value>true</value>
  </property>

  <property>
    <name>hadoop.security.auth_to_local</name>
    <value>RULE:[1:$1@$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/
RULE:[1:$1@$0](.*@EXAMPLE.COM)s/@.*/ /
RULE:[2:$1@$0](dn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](nn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](rangeradmin@EXAMPLE.COM)s/.*/ranger/
RULE:[2:$1@$0](rangerusersync@EXAMPLE.COM)s/.*/rangerusersync/
DEFAULT</value>
  </property>

</configuration>
```

1.26.9.2. Run Setup Again for Ranger Admin

1. Stop the ranger-admin service.
2. Add the following properties and values to the ranger-admin `install.properties` file:

Table 1.8. Ranger-admin install.properties

Property Name	Property Value
spnego_principal	HTTP/<FQDN_OF_Ranger_Admin_Cluster>@<REALM>
spnego_keytab	<HTTP keytab path>
token_valid	30
cookie_domain	<FQDN_OF_Ranger_Admin_Cluster>

Property Name	Property Value
cookie_path	/
admin_principal	rangeradmin/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
admin_keytab	<rangeradmin keytab path>
lookup_principal	rangerlookup/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
lookup_keytab	<rangerlookup keytab path>
hadoop_conf	/etc/hadoop/conf

3. Execute the `setup.sh` script.
4. Start the ranger-admin service.
5. Stop the ranger-usersync service.
6. Add the following properties and values to the ranger-sync `install.properties` file:

Table 1.9.

Property Name	Property Value
usersync_principal	rangerusersync/<FQDN>@<REALM>
usersyn_keytab	<rangerusersync keytab path>
hadoop_conf	/etc/hadoop/conf

7. Execute the `setup.sh` script.
8. Start the ranger-sync service.
9. Stop the ranger-tagsync service.
10. Add the following properties and values to the ranger-tagsync `install.properties` file:

Table 1.10. Ranger-tagsync install.properties and values

Property Name	Property Value
tagsync_principal	rangertagsync/<FQDN>@<REALM>
tagsync_keytab	<rangertagsync keytab path>
hadoop_conf	/etc/hadoop/conf

11. Execute the `setup.sh` script.
12. Start the ranger-tagsync service.
13. Stop the ranger-kms service.
14. Add the following properties and values to the ranger-kms `install.properties` file:

Table 1.11. Ranger-kms install.properties and values

Property Name	Property Value
kms_principal	rangerkms/<FQDN>@<REALM>
kms_keytab	<rangerkms keytab path>

Property Name	Property Value
hadoop_conf	/etc/hadoop/conf

1.26.9.3. Install and Enable the Ranger HDFS Plugin

This section documents how to install and enable the Ranger HDFS plugin. You might want to consider making similar changes for the other Ranger plugins that you are using.

The Ranger HDFS plugin is located at `/usr/hdp/<version>/ranger-hdfs-plugin`.

Follow these steps to install and enable the Ranger HDFS Plugin:

1. Edit the relevant lines in the `install.properties` file:

```
POLICY_MGR_URL=http://<FQDN of ranger admin host>:6080
REPOSITORY_NAME=hadoopdev
Audit info (Solr/HDFS options available)
```

2. Enter the following commands to enable the HDFS plugin:

```
export JAVA_HOME=<JAVA Path>
./enable-hdfs-plugin.sh
```

3. Start and stop the namenode:

```
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh stop
namenode"
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start
namenode"
```

4. In the custom repo configuration file, add the component user, for example, `hdfs`, as a value for each of the following properties:

- `policy.download.auth.users` or `policy.grantrevoke.auth.users`
- `tag.download.auth.users`

5. Verify that the plugin communicates with Ranger admin using the **Audit # plugins** tab.

6. Set the following properties in the `hdfs-site.xml` file:

Table 1.12. hdfs-site.xml Property Names and Values

Property Name	Property Value
<code>dfs.permissions.enabled</code>	<code>true</code>
<code>dfs.permissions.supergroup</code>	<code>hdfs</code>
<code>dfs.block.access.token.enable</code>	<code>true</code>
<code>dfs.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.secondary.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.principal</code>	<code>HTTP/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.keytab</code>	<code>/etc/security/keytabs/spnego.service.keytab</code>
<code>dfs.datanode.kerberos.principal</code>	<code>dn/_HOST@EXAMPLE.COM</code>
<code>dfs.namenode.keytab.file</code>	<code>/etc/security/keytabs/nn.service.keytab</code>
<code>dfs.secondary.namenode.keytab.file</code>	<code>/etc/security/keytabs/nn.service.keytab</code>
<code>dfs.datanode.keytab.file</code>	<code>/etc/security/keytabs/dn.service.keytab</code>

Property Name	Property Value
dfs.https.port	50470
dfs.namenode.https-address	Example:ip-10-111-59-170.ec2.internal:50470
dfs.datanode.data.dir.perm	750
dfs.cluster administrators	hdfs
dfs.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}
dfs.secondary.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}

The following is an example of a `hdfs-site.xml` file with the properties set for Kerberos:

```
<property>
  <name>dfs.permissions</name>
  <value>true</value>
  <description> If "true", enable permission checking in
HDFS. If "false", permission checking is turned
off, but all other behavior is
unchanged. Switching from one parameter value to the other does
not change the mode, owner or group of files or
directories. </description>
</property>

<property>
  <name>dfs.permissions.supergroup</name>
  <value>hdfs</value>
  <description>The name of the group of
super-users.</description>
</property>

<property>
  <name>dfs.namenode.handler.count</name>
  <value>100</value>
  <description>Added to grow Queue size so that more
client connections are allowed</description>
</property>

<property>
  <name>ipc.server.max.response.size</name>
  <value>5242880</value>
</property>

<property>
  <name>dfs.block.access.token.enable</name>
  <value>true</value>
  <description> If "true", access tokens are used as capabilities
for accessing datanodes. If "false", no access tokens are checked on
accessing datanodes. </description>
</property>

<property>
  <name>dfs.namenode.kerberos.principal</name>
  <value>nn/_HOST@EXAMPLE.COM</value>
  <description> Kerberos principal name for the
NameNode </description>
</property>
```

```
<property>
  <name>dfs.secondary.namenode.kerberos.principal</name>
  <value>nn/_HOST@EXAMPLE.COM</value>
  <description>Kerberos principal name for the secondary NameNode.
</description>
</property>

<property>
  <!--cluster variant -->
  <name>dfs.secondary.http.address</name>
  <value>ip-10-72-235-178.ec2.internal:50090</value>
  <description>Address of secondary namenode web server</description>
</property>

<property>
  <name>dfs.secondary.https.port</name>
  <value>50490</value>
  <description>The https port where secondary-namenode
  binds</description>
</property>

<property>
  <name>dfs.web.authentication.kerberos.principal</name>
  <value>HTTP/_HOST@EXAMPLE.COM</value>
  <description>The HTTP Kerberos principal used by Hadoop-Auth in the
  HTTP endpoint.
  The HTTP Kerberos principal MUST start with 'HTTP/' per Kerberos HTTP
  SPNEGO specification.
  </description>
</property>

<property>
  <name>dfs.web.authentication.kerberos.keytab</name>
  <value>/etc/security/keytabs/spnego.service.keytab</value>
  <description>The Kerberos keytab file with the credentials for the HTTP
  Kerberos principal used by Hadoop-Auth in the HTTP endpoint.
  </description>
</property>

<property>
  <name>dfs.datanode.kerberos.principal</name>
  <value>dn/_HOST@EXAMPLE.COM</value>
  <description>
    The Kerberos principal that the DataNode runs as. "_HOST" is replaced
    by the real
    host name.
  </description>
</property>

<property>
  <name>dfs.namenode.keytab.file</name>
  <value>/etc/security/keytabs/nn.service.keytab</value>
  <description>
    Combined keytab file containing the namenode service and host
    principals.
  </description>
</property>

<property>
  <name>dfs.secondary.namenode.keytab.file</name>
```

```
<value>/etc/security/keytabs/nn.service.keytab</value>
<description>
Combined keytab file containing the namenode service and host
principals.
</description>
</property>

<property>
  <name>dfs.datanode.keytab.file</name>
  <value>/etc/security/keytabs/dn.service.keytab</value>
  <description>
The filename of the keytab file for the DataNode.
  </description>
</property>

<property>
  <name>dfs.https.port</name>
  <value>50470</value>
  <description>The https port where namenode
binds</description>
</property>

<property>
  <name>dfs.https.address</name>
  <value>ip-10-111-59-170.ec2.internal:50470</value>
  <description>The https address where namenode binds</description>
</property>

<property>
  <name>dfs.datanode.data.dir.perm</name>
  <value>750</value>
  <description>The permissions that should be there on
dfs.data.dir directories. The datanode will not come up if the
permissions are different on existing dfs.data.dir directories. If
the directories don't exist, they will be created with this
permission.</description>
</property>

<property>
  <name>dfs.access.time.precision</name>
  <value>0</value>
  <description>The access time for HDFS file is precise upto this
value.The default value is 1 hour. Setting a value of 0
disables access times for HDFS.
  </description>
</property>

<property>
  <name>dfs.cluster.administrators</name>
  <value> hdfs</value>
  <description>ACL for who all can view the default
servlets in the HDFS</description>
</property>

<property>
  <name>ipc.server.read.threadpool.size</name>
  <value>5</value>
  <description></description>
</property>
```

```
<property>
  <name>dfs.namenode.kerberos.internal.spnego.principal</name>
  <value>${dfs.web.authentication.kerberos.principal}</value>
</property>

<property>
  <name>dfs.secondary.namenode.kerberos.internal.spnego.principal</name>
  <value>${dfs.web.authentication.kerberos.principal}</value>
</property>
```

7. For Download Policy to be successful, use the Ranger UI to update the service configuration with the following custom properties:

```
policy.download.auth.users=<Component service user>
tag.download.auth.users=<Component service user>(if tag download)
```

8. For Grant/Revoke for Hive and Hbase to be successful, use the Ranger UI to update the service configuration with the following custom property:

```
policy.grantrevoke.auth.users = <Component service user>
```

9. For Test Connection and Resource Lookup to be successful, use the Ranger UI to add lookup user in the permission list of the policies.

1.27. Configuring and Upgrading Apache Spark

Before you can upgrade Apache Spark, you must have first upgraded your HDP components to the latest version (in this case, 2.5.0). This section assumes that you have already upgraded your components for HDP 2.5.0. If you have not already completed these steps, return to [Getting Ready to Upgrade](#) and [Upgrade 2.4 Components](#) for instructions on how to upgrade your HDP components to 2.5.0.

To upgrade Spark, start the service and update configurations.

1. Stop the Spark history-server. If you are using the Spark thrift-server, stop the thrift-server.

```
su - spark -c "$SPARK_HOME/sbin/stop-history-server.sh"
su - spark -c "$SPARK_HOME/sbin/stop-thriftserver.sh"
```

2. Remove any reference to hdp.version from the Spark configuration files.

Remove spark.yarn.services property from spark-defaults.conf.

Make sure that spark.history.provider is set to org.apache.spark.deploy.history.FsHistoryProvider.

3. Restart the history-server:

```
su - spark -c "$SPARK_HOME/sbin/start-history-server.sh"
```

4. If you are using the Spark thrift-server, restart the thrift-server. See [\(Optional\) Starting the Spark Thrift Server](#).
5. Validate the Spark installation. As user spark, run the examples in the [Running Spark Applications](#) in the *Spark Guide*.

For additional configuration information, see the [Spark Guide](#).

1.28. Upgrade Apache Slider

To upgrade Apache Slider, simply upgrade the Slider client.

1. Upgrade Slider client:

```
hdp-select set slider-client 2.5.0.0-<version>  
slider version
```

1.29. Upgrade Apache Kafka

Upgrade each Apache Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the new version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"  
hdp-select set kafka-broker 2.4.4.0-2633  
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.
3. If the upgrade process fails, follow the steps in "Downgrading Kafka" to return to your previous version of Kafka.

1.29.1. Downgrading Kafka

Downgrade each Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the previous version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"  
hdp-select set kafka-broker 2.5.0.0-2041  
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.

1.30. Finalize the Upgrade



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.

1. Verify your file system health before finalizing the upgrade. After you finalize an upgrade, your backup is discarded!
2. As the \$HDFS_USER, enter:

```
su - hdfs -c "dfsadmin -finalizeUpgrade"
```

1.31. Migrate the Audit Logs from DB to Solr

Beginning with HDP 2.5.0, audit log to DB support has been removed. If your logs were previously stored on DB, you can migrate the logs to Solr. Refer to [Migrating Audit Logs from DB to Solr](#).

1.32. Install New HDP 2.5.0 Services

1. Install new HDP 2.5.0 Services (see the [Non-Ambari Cluster Installation Guide](#)):
 - Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
 - SmartSense – a next generation subscription model that features upgrade and configuration recommendations.
2. As Ranger has upgraded to a new version, you can Tag-sync, using the following Tagsync > install.properties, if required.

Table 1.13. Ranger Tagsync_install.properties and values

Property Name	Property Value
TAG_DEST_RANGER_ENDPOINT	http://<hostname>:6080
If TAG_SOURCE_ATLAS_ENABLED=true	
TAG_SOURCE_ATLAS_KAFKA_BOOTSTRAP_SERVERS	<hostname>:6667
TAG_SOURCE_ATLAS_KAFKA_ZOOKEEPER_CONNECT	<hostname>:2181
TAG_SOURCE_ATLAS_KAFKA_ENTITIES_GROUP_ID	ranger_entities_consumer
TAG_SOURCE_ATLAS_KAFKA_SERVICE_NAME	kafka
TAG_SOURCE_ATLAS_KAFKA_SECURITY_PROTOCOL	PLAINTEXTSASL
If TAG_SOURCE_ATLASREST_ENABLED = true	
TAG_SOURCE_ATLASREST_ENDPOINT	http://<hostname>:21000
TAG_SOURCE_ATLASREST_DOWNLOAD_INTERVAL_IN_MINUTES	100000
TAG_SOURCE_ATLASREST_USERNAME	
TAG_SOURCE_ATLASREST_PASSWORD	
If TAG_SOURCE_FILE_ENABLED = true	
TAG_SOURCE_FILE_ENABLED	true
TAG_SOURCE_FILE_FILENAME	/etc/ranger/data/tags.json (where the tag-sync is installed)

2. Upgrade from HDP 2.3 to HDP 2.5.0 Manually

This chapter provides instructions on how to manually upgrade to HDP 2.5.0 from the HDP 2.3.0 release. It assumes the existing HDP 2.3 was also installed manually.



Important

If you installed and manage HDP 2.3 with Ambari, **you must use the [Ambari Upgrade Guide](#)** to perform the the HDP 2.3 to HDP 2.5.0 upgrade.

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

Rolling Upgrade involves complex orchestration as well as side-by-side installation. It is too complex for a manual procedure, and is therefore supported only as an Ambari feature. If you wish to perform a Rolling Upgrade, refer to the Ambari Install instructions to install Ambari, then follow the Ambari Rolling Upgrade instructions, see [Ambari Upgrade Guide](#).

The HDP packages for a complete installation of HDP 2.5.0 consumes about 6.5 GB of disk space.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.

The following provides an overview of steps for upgrading to the latest release of HDP 2.3.2 from HDP 2.3:

1. Get ready to upgrade
2. Upgrade HDP 2.3 Components and stop all services
3. Use hdp-select to symlink the HDP 2.3.2 components into "current," in place of the former HDP 2.3 components
4. Configure and Start Apache ZooKeeper
5. Configure and Start Hadoop
6. Start HDFS
7. Configure and start YARN/MapReduce
8. Configure and Start Apache HBase
9. Configure and Start Apache Phoenix

10. Configure and Start Apache Accumulo
11. Configure and Start Apache Tez
12. Configure and Start Apache Hive and Apache HCatalog
13. Configure and Start Apache Oozie
14. Configure and Start Apache WebHCat (Templeton)
15. Configure and Start Apache Pig
16. Configure and Start Apache Sqoop
17. Configure and Start Apache Flume
18. Configure and Start Apache Mahout
19. Configure and Start Apache Hue
20. Configure and Start Apache Knox
21. Configure and Start Apache Falcon
22. Configure and Start Apache Storm
23. Configure and Start Apache Ranger
24. Configure and Start Apache Spark
25. Upgrade Apache Slider
26. Upgrade Apache Kafka
27. Finalize the Upgrade
28. Install new HDP 2.5.0 services

2.1. Getting Ready to Upgrade

HDP Stack upgrade involves upgrading from HDP 2.3 to HDP 2.5.0 versions and adding the new HDP 2.5.0 services. These instructions change your configurations.



Note

You must use **kinit** before running the commands as any particular user.

Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP 2.5 consumes about 6.5 GB of disk space.

The first step is to ensure you keep a backup copy of your HDP 2.3 configurations.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Back up the HDP directories for any hadoop components you have installed.

The following is a list of all HDP directories:

- `/etc/hadoop/conf`
- `/etc/hbase/conf`
- `/etc/hive-hcatalog/conf`
- `/etc/hive-webhcat/conf`
- `/etc/accumulo/conf`
- `/etc/hive/conf`
- `/etc/pig/conf`
- `/etc/sqoop/conf`
- `/etc/flume/conf`
- `/etc/mahout/conf`
- `/etc/oozie/conf`
- `/etc/hue/conf`
- `/etc/knox/conf`
- `/etc/zookeeper/conf`
- `/etc/tez/conf`
- `/etc/storm/conf`
- `/etc/falcon/conf`
- `/etc/slider/conf/`
- `/etc/ranger/admin/conf`, `/etc/ranger/usersync/conf` (If Ranger is installed, also take a backup of `install.properties` for all the plugins, ranger admin & ranger usersync.)
- Optional - Back up your userlogs directories, `${mapred.local.dir}/userlogs`.

2. Oozie runs a periodic purge on the shared library directory. The purge can delete libraries that are needed by jobs that started before the upgrade began and

that finish after the upgrade. To minimize the chance of job failures, you should extend the `oozie.service.ShareLibService.purge.interval` and `oozie.service.ShareLibService.temp.sharelib.retention.days` settings.

Add the following content to the `oozie-site.xml` file prior to performing the upgrade:

```
<property>
<name>oozie.service.ShareLibService.purge.interval</name>
<value>1000</value><description>
How often, in days, Oozie should check for old ShareLibs and LauncherLibs to
purge from HDFS.
</description>
</property>

<property>
<name>oozie.service.ShareLibService.temp.sharelib.retention.days</name>
<value>1000</value>
<description>
ShareLib retention time in days.</description>
</property>
```

3. Stop all long-running applications deployed using Slider:

```
su - yarn -c "/usr/hdp/current/slider-client/bin/slider list"
```

For all applications returned in previous command, run `su - yarn "/usr/hdp/current/slider-client/bin/slider stop <app_name>"`

4. Stop all services (including MapReduce) except HDFS, ZooKeeper, and Ranger, and client applications deployed on HDFS.

See [Stopping HDP Services](#) for more information.

Component	Command
Accumulo	<code>/usr/hdp/current/accumulo-client/bin/stop-all.sh</code>
Knox	<code>cd \$GATEWAY_HOME su - knox -c "bin/gateway.sh stop"</code>
Falcon	<code>su - falcon "/usr/hdp/current/falcon-server/bin/falcon-stop"</code>
Oozie	<code>su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-stop.sh"</code>
WebHCat	<code>su - webhcat -c "/usr/hdp/hive-webhcat/sbin/webhcat_server.sh stop"</code>
Hive	<p>Run this command on the Hive Metastore and Hive Server2 host machine:</p> <pre>ps aux awk '{print \$1,\$2}' grep hive awk '{print \$2}' xargs kill >/dev/null 2>&1</pre> <p>Or you can use the following:</p> <pre>Killall -u hive -s 15 java</pre>

Component	Command
HBase RegionServers	<pre>su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-daemon.sh --config /etc/hbase/conf stop regionserver"</pre>
HBase Master host machine	<pre>su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh --config /etc/hbase/conf stop master"</pre>
YARN & Mapred History	<p>Run this command on all NodeManagers:</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop nodemanager"</pre> <p>Run this command on the History Server host machine:</p> <pre>su - mapred -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-jobhistory-daemon.sh --config /etc/hadoop/conf stop historyserver"</pre> <p>Run this command on the ResourceManager host machine(s):</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop resourcemanager"</pre> <p>Run this command on the ResourceManager host machine:</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-timelineserver/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"</pre> <p>Run this command on the YARN Timeline Server node:</p> <pre>su -l yarn -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-yarn/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"</pre>
Storm	<pre>storm kill topology-name</pre> <pre>sudo service supervisord stop</pre>
Spark (History server)	<pre>su - spark -c "/usr/hdp/current/spark-client/sbin/stop-history-server.sh"</pre>

5. If you have the Hive component installed, back up the Hive Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, see your database documentation.

Table 2.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hive < /tmp/mydir/backup_hive.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hive < /tmp/mydir/backup_hive.sql</pre>
Oracle	<p>Export the database:</p> <pre>exp username/password@database full=yes file=output_file.dmp</pre>	<p>Import the database:</p> <pre>imp username/password@database file=input_file.dmp</pre>

6. If you have the Oozie component installed, back up the Oozie metastore database.

These instructions are provided for your convenience. Check your database documentation for the latest backup instructions.

Table 2.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sql</pre> <p>For example:</p> <pre>mysqldump oozie > /tmp/mydir/backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sql</pre> <p>For example:</p> <pre>mysql oozie < /tmp/mydir/backup_oozie.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump oozie > /tmp/mydir/backup_oozie.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres psql oozie < /tmp/mydir/backup_oozie.sql</pre>
Oracle	<p>Export the database:</p> <pre>exp username/password@database full=yes file=output_file.dmp</pre>	<p>Import the database:</p> <pre>imp username/password@database file=input_file.dmp</pre>

7. **Optional:** Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 2.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hue > /tmp/mydir/backup_hue.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hue < /tmp/mydir/backup_hue.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hue > /tmp/mydir/backup_hue.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hue < /tmp/mydir/backup_hue.sql</pre>
Oracle	<p>Connect to the Oracle database using sqlplus. Export the database.</p> <p>For example:</p> <pre>exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr</pre>	<p>Import the database:</p> <p>For example:</p> <pre>imp username/password@database file=input_file.dmp</pre>
SQLite	<pre>/etc/init.d/hue stop</pre> <pre>su \$HUE_USER</pre> <pre>mkdir ~/hue_backup</pre> <pre>sqlite3 desktop.db .dump > ~/hue_backup/desktop.bak</pre> <pre>/etc/init.d/hue start</pre>	<pre>/etc/init.d/hue stop</pre> <pre>cd /var/lib/hue</pre> <pre>mv desktop.db desktop.db.old</pre> <pre>sqlite3 desktop.db < ~/hue_backup/desktop.bak</pre> <pre>/etc/init.d/hue start</pre>

8. Back up the Knox data/security directory.

```
cp -RL /etc/knox/data/security ~/knox_backup
```

9. Save the namespace by executing the following commands:

```
su - hdfs
```

```
hdfs dfsadmin -safemode enter
```

```
hdfs dfsadmin -saveNamespace
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

10. Run the `fsck` command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.)

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-old-fsck-1.log"
```




Note

In secure mode, you must have Kerberos credentials for the hdfs user.

11. Use the following instructions to compare status before and after the upgrade.

The following commands must be executed by the user running the HDFS service (by default, the user is hdfs).

- a. Capture the complete namespace of the file system. (The following command does a recursive listing of the root file system.)



Important

Make sure the namenode is started.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-old-lsr-1.log"
```



Note

In secure mode you must have Kerberos credentials for the hdfs user.

- b. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs dfsadmin -c "-report > dfs-old-report-1.log"
```

- c. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-old-report-1.log"
```

- d. **Optional:** You can copy all or unrecoverable only data store libext-customer directory in HDFS to a local file system or to a backup instance of HDFS.

- e. **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.

12. Finalize any prior HDFS upgrade, if you have not done so already.

```
su - hdfs -c "hdfs dfsadmin -finalizeUpgrade"
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

13. Stop remaining services (HDFS, Zookeeper, and Ranger).

See [Stopping HDP Services](#) for more information.

Component	Command
HDFS	On all DataNodes: If you are running secure cluster, run following command as root:

Component	Command
	<pre>/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode</pre> <p>Else:</p> <pre>su - hdfs -c "usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode"</pre> <p>If you are not running a highly available HDFS cluster, stop the Secondary NameNode by executing this command on the Secondary NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop secondarynamenode"</pre> <p>On the NameNode host machine(s)</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop namenode"</pre> <p>If you are running NameNode HA, stop the ZooKeeper Failover Controllers (ZKFC) by executing this command on the NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop zkfc"</pre> <p>If you are running NameNode HA, stop the JournalNodes by executing these command on the JournalNode host machines:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh ==config /etc/hadoop/conf stop journalnode"</pre>
Zookeeper Host machines	<pre>su - zookeeper -c "/usr/hdp/current/zookeeper-server/bin/zookeeper-server stop"</pre>
Ranger (XA Secure)	<pre>service ranger-admin stop</pre> <pre>service ranger-usersync stop</pre>

14. Back up your NameNode metadata.



Note

It's recommended to take a backup of the full `/hadoop.hdfs/namenode` path.

- a. Copy the following checkpoint files into a backup directory.

The NameNode metadata is stored in a directory specified in the `hdfs-site.xml` configuration file under the configuration value `"dfs.namenode.name.dir"`.

For example, if the configuration value is:

```
<property>
  <name>dfs.namenode.name.dir</name>
```

```
<value>/hadoop/hdfs/namenode</value>
</property>
```

Then, the NameNode metadata files are all housed inside the directory /
hadoop.hdfs/namenode.

- b. Store the layoutVersion of the namenode.

```
${dfs.namenode.name.dir}/current/VERSION
```

15. Verify that edit logs in `${dfs.namenode.name.dir}/current/edits*` are empty.

- a. Run: `hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out`

- b. Verify the edits.out file. It should only have OP_START_LOG_SEGMENT transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

- c. If edits.out has transactions other than OP_START_LOG_SEGMENT, run the following steps and then verify edit logs are empty.

- Start the existing version NameNode.
- Ensure there is a new FS image file.
- Shut the NameNode down:

```
hdfs dfsadmin - saveNamespace
```

16. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path
component in this version of HDFS.
```

```
Please rollback and delete or rename this path, or upgrade with the
-renameReserved key-value pairs option to automatically rename these
paths during upgrade.
```

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup.

For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, specify `-upgrade -renameReserved .snapshot=.my-snapshot,.reserved=.my-reserved`.

If no key-value pairs are specified with `-renameReserved`, the NameNode then suffixes reserved paths with:

```
.<LAYOUT-VERSION>.UPGRADE_RENAMED
```

For example: `.snapshot.-51.UPGRADE_RENAMED`.



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

17. Upgrade the JDK on all nodes to JDK 7 or JDK 8 before upgrading HDP.

2.2. Upgrade HDP 2.3 Components

The upgrade process to HDP 2.5 involves the following steps. See the [Release Notes](#) for the HDP 2.5.0.0 repo information. Select your OS:

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.3 components. This command uninstalls the HDP 2.3 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"spark*" "slider*" "hdp_mon_nagios_addons" "bigtop"
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
yum list installed | grep @HDP2.3
```

4. Remove your old `hdp.repo` file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP 2.5 repo:

- Download the `hdp.repo` file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

RHEL/CentOS/Oracle

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.3 components. This command uninstalls the HDP 2.3 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
yum list installed | grep @HDP2.3
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP 2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos5/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "hue" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.3 components. This command uninstalls the HDP 2.3 components. It leaves the user data, and metadata, but removes your configurations:

```
zypper rm "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
yum list installed | grep @HDP2.3
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "hue" "tez"
"storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark"
"slider*" "hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5.0
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component files names should appear in the returned list.

SLES

1. On all hosts, clean the zypper repository.

```
zypper clean -all
```

2. Remove your old HDP 2.3 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*"
"pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*"
"falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common"
"hue-shell" "knox*" "ranger*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
zypper search --installed-only --repo HDP-2.3.6.0
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/susel1sp3/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "oozie" "collectd" "gccxml"
"pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "tez" "storm" "falcon"
"flume" "phoenix" "accumulo" "mahout" "knox" "spark" "spark-python"
"hdp_mon_nagios_addons" "slider*" "hive-webcat" "hive-hcatalog"
```

```
zypper up -r HDP-2.5.0
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component files names should appear in the returned list.

Ubuntu 12

1. On all hosts, clean the apt-get repository.

```
apt-get clean
```

2. Remove your old HDP 2.3 components. This command uninstalls the HDP 2.3 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop-2-3-*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig-2-3-*" "hdfs*" "sqoop-2-3-*" "zookeeper-2-3-*" "hbase-2-3-*"
"hive-2-3-*" "tez-2-3-*" "storm-2-3-*" "falcon-2-3-*" "flume-2-3-*"
"phoenix-2-3-*" "accumulo-2-3-*" "mahout" "hue" "knox-2-3-*" "spark*"
"slider*" "hdp_mon_nagios_addons" --purge
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
apt --installed list |grep 2.3.
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```


5. Download the HDP 2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2.5.0.0/hdp.list - O /etc/apt/sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

Debian

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.3 components. This command uninstalls the HDP 2.3 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue*" "knox*" "spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.3 component binaries are uninstalled:

```
yum list installed | grep @HDP2.3
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/debian<version>/2.x/updates/2.5.0.0/hdp.list - O /etc/apt/sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.3 components:

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm"
"falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "spark" "slider*"
"hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

2.3. Symlink Directories with hdp-select



Warning

HDP 2.5.0 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

1. Before you run hdp-select, remove one link:

```
rm /usr/bin/oozie
```

2. Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5.0.0-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5.0.0-2800
```

2.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available HDFS cluster, configure and restart Apache ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.

1. Replace your configuration after upgrading on all the ZooKeeper nodes. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "/usr/hdp/current/zookeeper-server/bin/zookeeper-server
start"
```

2.5. Configure Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed in HDP 2.3. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code do not load, as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed since HDP 2.3. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code do not load, as this is not where lzo is installed.

Ubuntu/Debian

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code do not load, as this is not where lzo is installed.

2.6. Start Hadoop Core



Warning

Before you start HDFS on a highly available HDFS cluster, you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

To start HDFS, run commands as the `$HDFS_USER`.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading on all the HDFS nodes. Replace the HDFS template configuration in `/etc/hdfs/conf`.
2. If you are upgrading from a highly available HDFS cluster configuration, start all JournalNodes. On each JournalNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start journalnode"
```



Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

3. If you are running HDFS on a highly available namenode, you must first start the Zookeeper service



Note

Perform this step only if you are on a highly available HDFS cluster.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc"
```

4. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually.

On the active NameNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the `-upgrade` option only once. After you have completed this step, you can bring up the NameNode using this command without including the `-upgrade` option.

To check if the Upgrade is in progress, check that the "`\previous`" directory has been created in the `\NameNode` and `\JournalNode` directories. The "`\previous`" directory contains a snapshot of the data before upgrade.

In a highly available HDFS cluster configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode.

To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the '-bootstrapStandby' flag. Do NOT start this standby NameNode with the '-upgrade' flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The bootstrapStandby command downloads the most recent fsimage from the active NameNode into the \$dfs.name.dir directory of the standby NameNode. You can enter that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

5. Verify that the NameNode is up and running:

```
ps -ef|grep -i NameNode
```

6. If you do not have a highly available HDFS cluster configuration (non_HA namenode), start the Secondary NameNode.



Note

Do not perform this step if you have a highly available HDFS cluster configuration.

On the Secondary NameNode host machine, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

7. Verify that the Secondary NameNode is up and running.



Note

Do not perform this step if you have a highly available HDFS cluster environment.

```
ps -ef|grep SecondaryNameNode
```

8. Start DataNodes.

On each of the DataNodes, enter the following command. Note: If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start datanode"
```

9. Verify that the DataNode process is up and running:

```
ps -ef | grep DataNode
```

10. Verify that NameNode can go out of safe mode.

```
> su - hdfs -c "hdfs dfsadmin -safemode wait"
```

You should see the following result: `Safe mode is OFF`

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

2.7. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

1. Run the `fsck` command on namenode as `$HDFS_USER`:

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-new-fsck-1.log"
```

You should see feedback that the filesystem under path `/` is `HEALTHY`.

2. Run `hdfs` namespace and report.

a. List directories.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-new-lsr-1.log"
```

b. Open the `dfs-new-lsr-1.log` and confirm that you can see the file and directory listing in the namespace.

c. Run `report` command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-new-report-1.log"
```

d. Open the `dfs-new-report` file and validate the admin report.

3. Compare the namespace report before the upgrade and after the upgrade. Verify that user files exist after upgrade.

The file names are listed below:

```
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log
```

```
dfs-old-lsr-1.log < -- > dfs-new-lsr-1.log
```



Note

You must do this comparison manually to catch all errors.

4. From the NameNode WebUI, determine if all DataNodes are up and running.

```
http://<namenode>:<namenodeport>
```

5. If you are on a highly available HDFS cluster, go to the StandbyNameNode web UI to see if all DataNodes are up and running:

```
http://<standbynamenode>:<namenodeport>
```

6. If you are **not** on a highly available HDFS cluster, go to the SecondaryNameNode web UI to see if the secondary node is up and running:

```
http://<secondarynamenode>:<secondarynamenodeport>
```

7. Verify that read and write to hdfs works successfully.

```
hdfs dfs -put [input file] [output file]
```

```
hdfs dfs -cat [output file]
```

2.8. Configure YARN and MapReduce

After you upgrade Hadoop, complete the following steps to update your configs.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

In secure mode, you must have Kerberos credentials for the hdfs user.

1. Upload the MapReduce tarball to HDFS. As the HDFS user, for example 'hdfs':

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/mapreduce.tar.gz /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

```
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<
$version>/mapreduce"
```

```
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<
$version>/mapreduce/mapreduce.tar.gz"
```

2. Make sure that the following properties are in `/etc/hadoop/conf/mapred-site.xml`:

- Make sure `mapreduce.application.framework.path` exists in `mapred-site.xml`:

```
<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}/mapreduce/mapreduce.tar.gz#mr-framework</
value>
</property>

<property>
  <name>yarn.app.mapreduce.am.admin-command-opts</name>
  <value>-Dhdp.version=${hdp.version}</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
  <value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}/hadoop/lib/native:/usr/
hdp/${hdp.version}/hadoop/
  lib/native/Linux-amd64-64</value>
</property>

<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.admin.reduce.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.application.classpath</name>
  <value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
  $PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/common/*,
  $PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/yarn/*,
  $PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
  $PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
```



```
$PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
/usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
/etc/hadoop/conf/secure</value>
</property>
```

**Note**

You do not need to modify `${hdp.version}`.

**Note**

If you are planning to use Spark in yarn-client mode, make Spark work in yarn-client mode 2.5.0.0-<\$version>.

3. Make sure the following property is in `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.application.classpath</name>
  <value>$HADOOP_CONF_DIR,/usr/hdp/${hdp.version}/hadoop-client/*,
  /usr/hdp/${hdp.version}/hadoop-client/lib/*,
  /usr/hdp/${hdp.version}/hadoop-hdfs-client/*,
  /usr/hdp/${hdp.version}/hadoop-hdfs-client/lib/*,
  /usr/hdp/${hdp.version}/hadoop-yarn-client/*,
  /usr/hdp/${hdp.version}/hadoop-yarn-client/lib/*</value>
</property>
```

4. On secure clusters only, add the following properties to `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.timeline-service.recovery.enabled</name>
  <value>TRUE</value>
</property>

<property>
  <name>yarn.timeline-service.state-store.class</name>
  <value>org.apache.hadoop.yarn.server.timeline.recovery.
LevelDbTimelineStateStore</value>
</property>

<property>
  <name>yarn.timeline-service.leveldb-state-store.path</name>
  <value><the same as the default of "yarn.timeline-service-leveldb-
timeline-store.path"></value>
</property>
```

5. For secure clusters, you must create and configure the container-executor.cfg configuration file:

- Create the container-executor.cfg file in `/etc/hadoop/conf/container-executor.cfg`.
- Insert the following properties:

```
<property>
yarn.nodemanager.linux-container-executor.group=hadoop
banned.users=hdfs,yarn,mapred
min.user.id=1000
</property>
```

- `yarn.nodemanager.linux-container-executor.group` - Configured value of `yarn.nodemanager.linux-container-executor.group`. This must match the value of `yarn.nodemanager.linux-container-executor.group` in `yarn-site.xml`.
- `banned.users` - Comma-separated list of users who can not run container-executor.
- `min.user.id` - Minimum value of user id. This prevents system users from running container-executor.
- `allowed.system.users` - Comma-separated list of allowed system users.
- Set the file `/etc/hadoop/conf/container-executor.cfg` file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn/bin/container-executor
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn/bin/container-executor
```

2.9. Start YARN/MapReduce Services

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.



Note

The `su` commands in this section use "yarn" to represent the YARN Service user and `mapreduce` to represent the MAPREDUCE Service user. If you are using another name for these Service users, you need to substitute your Service user name for "yarn" or "mapreduce" in each of the `su` commands.

1. Manually clear the ResourceManager state store.

```
su - yarn -c "yarn resourcemanager -format-state-store"
```

2. Start the ResourceManager on all your ResourceManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh start resourcemanager"
ps -ef | grep -i resourcemanager
```

3. Start the TimelineServer on your TimelineServer host.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-timelineserver/sbin/yarn-daemon.sh start timelineserver"
ps -ef | grep -i timelineserver
```

4. Start the NodeManager on all your NodeManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh
start nodemanager"

ps -ef | grep -i nodemanager
```

5. To start MapReduce, run the following commands:

```
su - mapreduce -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-
jobhistory-daemon.sh start historyserver"

ps -ef | grep -i jobhistoryserver
```

2.10. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user.

The job uses MapReduce to write 100MB of data into HDFS with RandomWriter

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.
jar
    randomwriter -Dtest.randomwrite.total_bytes=10000000 test-after-
upgrade.
```

You should see messages similar to:

```
map 0% reduce 0%
...map 100% reduce 100%
Job ... completed successfully
```

MapReduce upgraded successfully. You can now upgrade your other components.

Basic Troubleshooting

To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

`http://<resource manager host>:8088/cluster/nodes`

The number of active nodes should be equal to the number of nodemanagers.

Accessing error messages:

1. Access the ApplicationMaster WebUI to view the container logs.
2. At your console logs for MapReduce job, look for a line with this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track the job: http://<resource
manager host>:8088/proxy/application_1380673658357_0007/
```

3. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

2.11. Configure and Start Apache HBase



Note

The `su` commands in this section use "hbase" to represent the HBASE Service user. If you are using another name for your HBASE Service user, you need to substitute your HBASE Service user name for "hbase" in each of the `su` commands.

To enable Kerberos for clusters with dual home network setting, each HBase RegionServer must have its own key. See [Installing HDP Manually](#).

The `hbase.bucketcache.percentage.in.combinedcache` is removed in HDP 2.5.0. This simplifies the configuration of block cache. BucketCache configurations from HDP 2.3 needs to be recalculated to attain identical memory allotments in HDP 2.5.0. The L1 LruBlockCache is whatever `hfile.block.cache.size` is set to and the L2 BucketCache is whatever `hbase.bucketcache.size` is set to.

1. Replace your configuration after upgrading. Replace the Apache HBase template configuration in `/etc/hbase/conf`.
2. Start services. From root, assuming that `$HBASE_USER=hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

3. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

4. Run HBase smoke test. Verify that following commands can be executed successfully from HBase shell.

```
create 'testtable', 'cf1'
```

```
put 'testtable', 'r1', 'cf1:c1', 'v1'
```

```
disable 'testtable'
```

```
drop 'testtable'
```

2.12. Configure Apache Phoenix

To configure Apache Phoenix, complete the following steps:

1. Add the following property to the `/etc/hbase/hbase-site.xml` file on all HBase nodes, the MasterServer, and all RegionServers to prevent deadlocks from occurring during maintenance on global indexes:

```
<property>
  <name>hbase.regionserver.wal.codec</name>
  <value>org.apache.hadoop.hbase.regionserver.wal.IndexedWALEditCodec</value>
</property>
```

2. To enable user-defined functions, configure the following property in `/etc/hbase/conf` on all Hbase nodes.

```
<property>
  <name>phoenix.functions.allowUserDefinedFunctions</name>
  <value>true</value>
  <description>enable UDF functions</description>
</property>
```

3. Ensure the client side `hbase-site.xml` matches the server side configuration.
4. **(Optional)** To use local indexing, set the following three properties. These properties ensure collocation of data table and local index regions:



Note

The local indexing feature is a technical preview and considered under development. Do not use this feature in your production systems. If you have questions regarding this feature, contact Support by logging a case on our Hortonworks Support Portal at [//hortonworks.com/support/](http://hortonworks.com/support/).

Add the following two properties, if they do not already exist, to the master side configurations:

```
<property>
  <name>hbase.master.loadbalancer.class</name>
  <value>org.apache.phoenix.hbase.index.balancer.IndexLoadBalancer</value>
</property>

<property>
  <name>hbase.coprocessor.master.classes</name>
  <value>org.apache.phoenix.hbase.index.master.IndexMasterObserver</value>
</property>
```

Add the following, if it does not already exist, to the Region server side configurations:

```
<property>
  <name>hbase.coprocessor.regionserver.classes</name>
  <value>org.apache.hadoop.hbase.regionserver.LocalIndexMerger</value>
</property>
```

5. If the folder specified in `hbase.tmp.dir` property on `hbase-site.xml` does not exist, create that directory with adequate permissions.
6. Restart the HBase Master and Region Servers.

Configuring Phoenix to Run in a Secure Cluster

Perform the following additional steps to configure Phoenix to run in a secure Hadoop cluster:

1. To link the HBase configuration file with the Phoenix libraries:

```
ln -sf HBASE_CONFIG_DIR/hbase-site.xml PHOENIX_HOME/bin/hbase-site.xml
```

2. To link the Hadoop configuration file with the Phoenix libraries:

```
ln -sf HADOOP_CONFIG_DIR/core-site.xml PHOENIX_HOME/bin/core-site.xml
```



Note

When running the `pssql.py` and `slline.py` Phoenix scripts in secure mode, you can safely ignore the following warnings.

```
14/04/19 00:56:24 WARN util.NativeCodeLoader:
Unable to load native-hadoop library for your platform...
    using builtin-java classes where applicable

14/04/19 00:56:24 WARN util.DynamicClassLoader: Failed to identify the fs of
dir hdfs://<HOSTNAME>:8020/apps/hbase/data/lib, ignored java.io.IOException:
No FileSystem for scheme: hdfs
```

2.13. Configure and Start Apache Accumulo



Note

The `su` commands in this section use "accumulo" to represent the Accumulo Service user. If you are using another name for your Apache Accumulo Service user, you need to substitute your Accumulo Service user name for "accumulo" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/accumulo/conf` from the template to the `conf` directory in Accumulo hosts.
2. Start the services:

```
ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
master\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
tserver\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname` gc\"
accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
tracer\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
monitor\" accumulo"
```

3. Check that the processes are running

```
ps -ef | grep accumulo
```

or visit <http://<hostname>:50095> in your browser

4. Run Accumulo smoke test. Verify that following commands can be executed successfully from Accumulo shell.

```
createtable testtable
```

```
insert row1 colf colq value1
```

```
scan
```

```
deletetable testtable
```

```
tables
```

2.14. Configure and Start Apache Tez



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

To upgrade Apache Tez:

1. Copy your previously backed-up copy of `tez-site.xml` into the `/etc/tez/conf` directory.
2. Upload the Tez tarball to HDFS.

```
su - hdfs
hdfs dfs -mkdir -p /hdp/apps/<hdp_version>/tez/
hdfs dfs -put /usr/hdp/<hdp_version>/tez/lib/tez.tar.gz /hdp/apps/
<hdp_version>/tez/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/<hdp_version>/tez
hdfs dfs -chmod -R 444 /hdp/apps/<hdp_version>/tez/tez.tar.gz
```

Where `<hdp_version>` is the current HDP version, for example 2.5.0.0-2800.

3. Edit the `tez.lib.uris` property in the `tez-site.xml` file to point to `/hdp/apps/<hdp_version>/tez/tez.tar.gz`

```
...
<property>
  <name>tez.lib.uris</name>
  <value>/hdp/apps/<hdp_version>/tez/tez.tar.gz</value>
</property>
...
```

Where `<hdp_version>` is the current HDP version, for example 2.5.0.0-2800.

4. **Optional** Earlier releases of Tez did not have access control. In the current version of Tez, the default behavior restricts the ability to view the Tez history to only the owner of the job. To retain unrestricted access for non-secure clusters, set `tez.am.view-acls` set to `"*"`.
5. Change the value of the `tez.tez-ui.history-url.base` property to the url for the upgraded Tez View. For information on setting up the Tez view, see [Deploying the Tez View](#) in the HDP Ambari Views Guide.
6. **Run Tez Smoke Test**

To smoke test your Tez upgrade, you can run the following Tez Example job as a regular user.

MRRSleep Tez Example job sleeps for a defined period of time in mapper and reducer

```
hadoop jar /usr/hdp/current/tez-client/tez-tests-<version>.jar mrrsleep -m 1
-r 1 -mt 100 -rt 100
```

You should see messages similar to:

```
...DAG: State: SUCCEEDED Progress: 100%
DAG completed. FinalState=SUCCEEDED
```

Tez upgraded successfully. You can now upgrade your other components.

2.15. Configure and Start Apache Hive and Apache HCatalog



Note

The `su` commands in this section use "hive" to represent the Hive Service user. If you are using another name for your Hive Service user, you need to substitute your Hive Service user name for "hive" in each of the `su` commands.

1. Restore backup copy of `hive-site.xml` to `/etc/hive/conf`, on all nodes where `hivemetastore/hiveserver2` should be running. Prior to starting the upgrade process, set the following in your hive configuration file:
2. Copy the jdbc connector jar from `OLD_HIVE_HOME/lib` to `CURRENT_HIVE_HOME/lib`.
3. Upgrade the Apache Hive Metastore database schema. Restart the Hive Metastore database and run:

```
datanucleus.autoCreateSchema=false
```

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -upgradeSchema
-dbType <$databaseType>"
```

The value for `$databaseType` can be `derby`, `mysql`, `oracle`, or `postgres`.



Note

If you are using PostgreSQL, you should reset the Hive Metastore database owner to <HIVE_USER>:

```
psql -U <POSTGRES_USER> -c
ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER TO <HIVE_USER>
```



Note

If you are using Oracle 11, you may see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
optimize.mapjoin.mapreduce does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
heapsize does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
server2.enable.impersonation does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
semantic.analyzer.factory.impl does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.
convert.sortmerge.join.noconditionaltask does not exist
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.
internal:1521/XE
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver
Metastore connection User: hiveuser
Starting upgrade metastore schema from version 0.13.0 to 0.14.0
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql
Error: ORA-00955: name is already used by an existing object
(state=42000,code=955)
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.
oracle.sql: Schema script failed, errorcode 2
Completed upgrade-0.13.0-to-0.14.0.oracle.sql
schemaTool completed
```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.

4. Edit the hive-site.xml file and modify the properties based on your environment. Search for TODO in the file for the properties to replace.

- a. Edit the following properties in the hive-site.xml file:

```
<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache
  </description>
</property>
```

- b. Remove the following property in hive-site.xml, if present:

```
<property>
  <name>hive.exec.post.hooks</name>
  <value>org.apache.hadoop.hive.ql.hooks.ATSHook,org.apache.atlas.hive.
hook.HiveHook</value>
</property>
```

- c. **Optional:** To enable the Hive builtin authorization mode, make the following changes. If you want to use the advanced authorization provided by Ranger, refer to the [Ranger](#) instructions.

Set the following Hive authorization parameters in the hive-site.xml file:

```
<property>
  <name>hive.server2.enable.doAs</name>
  <value>>false</value>
</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
StorageBasedAuthorizationProvider,org.apache.hadoop.hive.ql.security.
authorization.MetaStoreAuthzAPIAuthorizeEmbedOnly</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
SQLStdConfOnlyAuthorizeFactory</value>
</property>
```

Also set hive.users.in.admin.role to the list of comma-separated users who need to be added to admin role. A user who belongs to the admin role needs to run the "set role" command before getting the privileges of the admin role, as this role is not in the current roles by default.

Set the following in the hiveserver2-site.xml file.

```
<property>
  <name>hive.security.authenticator.manager</name>
  <value>org.apache.hadoop.hive.ql.security.
SessionStateUserAuthenticator</value>
</property>

<property>
  <name>hive.security..authorization.enabled</name>
  <value>>true</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
SQLStdHiveAuthorizeFactory</value>
</property>
```

- d. For a remote Hive metastore database, set the IP address (or fully-qualified domain name) and port of the metastore host using the following hive-site.xml property value.

```
<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server.
    To enable HiveServer2, leave the property value empty.
  </description>
</property>
```

You can further fine-tune your configuration settings based on node hardware specifications, using the HDP utility script.

5. Start Hive Metastore.

On the Hive Metastore host machine, run the following command:

```
su - hive -c "nohup /usr/hdp/current/hive-metastore/bin/hive
--service metastore -hiveconf hive.log.file=hivemetastore.log
>/var/log/hive/hivemetastore.out 2>/var/log/hive/
hivemetastoreerr.log &"
```

6. Start Hive Server2.

On the Hive Server2 host machine, run the following command:

```
su - hive

nohup /usr/hdp/current/hive-server2/bin/hiveserver2 -hiveconf
hive.metastore.uris=" " -hiveconf hive.log.file=hiveserver2.log
>/var/log/hive/hiveserver2.out 2> /var/log/hive/
hiveserver2err.log &
```

2.16. Configure and Start Apache Oozie



Note

The duration of the Oozie upgrade is dependent on the amount of job history stored in ooziedb. This history must be backed up and restored during the upgrade process. Best practice when planning for upgrade is to backup ooziedb from your production oozie server and restore it to a test or development oozie server. This can help you estimate the time that is required to upgrade Oozie during your production upgrade.



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and "oozie" to represent the Oozie Service user. If you are using another name for your HDFS Service user or your Oozie Service user, you need to substitute your Service user names for "hdfs" or "oozie" in each of the `su` commands.

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore oozie-site.xml from your backup to the conf directory on each oozie server and client.
2. Copy the JDBC jar to libext-customer:

- a. Create the /usr/hdp/current/oozie/libext-customer directory.

```
cd /usr/hdp/current/oozie-server  
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the libext-customer directory.

```
chmod -R 777 /usr/hdp/current/oozie-server/libext-customer
```

3. Copy these files to the libext-customer directory

```
cp /usr/hdp/current/hadoop-client/lib/hadoop*lzo*.jar /usr/hdp/  
current/oozie-server/libext-customer  
  
cp /usr/share/HDP-oozie/ext.zip /usr/hdp/current/oozie-server/  
libext-customer/
```

Also, copy Oozie db jar in libext-customer.

4. If Falcon was also installed and configured before upgrade in HDP 2.3.x, then after upgrade you might also need to do the following:

```
cp /usr/hdp/current/falcon-server/oozie/ext/falcon-oozie-el-  
extension-*jar /usr/hdp/current/oozie-server/libext-customer
```

5. Extract share-lib.

```
/usr/hdp/current/oozie/bin/oozie-setup.sh sharelib create -fs  
hdfs://<namenode>:8020
```

To verify that the sharelibs extracted correctly, run the following command:

```
oozie admin -oozie http://<oozie server host address>:11000/  
oozie -shareliblist
```

There should be:

- Available ShareLib
- oozie
- hive
- distcp

- hcatalog
- sqoop
- mapreduce-streaming
- pig

Change the ownership and permissions of the oozie directory:

```
su -l hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie"
```

```
su -l hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

Add the two zip files from spark-client phone to spark sharelib:

```
hdfs dfs -put /usr/hdp/current/spark-client/python/lib/py4j-0.9-src.zip /
user/oozie/share/lib/lib_<timestamp>/spark
```

```
hdfs dfs -put /usr/hdp/current/spark-client/python/lib/pyspark.zip /user/
oozie/share/lib/lib_<timestamp>/spark
```

6. If a previous version of Oozie was created using auto schema creation, run the following SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

7. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh
upgrade -run"
```

8. As root, prepare the Oozie WAR file.

```
chown oozie:oozie /usr/hdp/current/oozie-server/oozie-server/
conf/server.xml
```

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-setup.sh
prepare-war -d /usr/hdp/current/oozie-server/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar
New Oozie WAR file with added 'JARS' at /var/lib/oozie/oozie-server/webapps/
oozie.war
```

9. Make sure that following property is added in oozie-log4j.properties:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where \${oozie.instance.id} is determined by oozie, automatically.

10. If you have custom Oozie actions, you must define them in oozie-site.xml. Edit the /etc/oozie/conf/oozie-site.xml file and add the following properties:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>[Comma seperated list of custom actions]</value>
</property>
<property>
  <name>oozie.service.SparkConfigurationService.spark.configuration</name>
  <value>*/etc/spark/conf/</value>
</property>
```

For example, if you have added Spark Action, enter the following:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>spark-action-0.1.xsd</value>
</property>
```

11. Configure HTTPS for the Oozie server.

- a. Create a self signed certificate or get certificate from a trusted CA for the Oozie Server
- b. Import the certificate to the client JDK trust store on all client nodes.
- c. In the Ambari Oozie configuration, set the following environment variables in `oozie-env.sh`, adding them if it does not exist:

```
export OOZIE_HTTPS_PORT=11443
export OOZIE_HTTPS_KEYSTORE_FILE=/home/oozie/.keystore
export OOZIE_HTTPS_KEYSTORE_PASS=password
```

- d. Change `OOZIE_HTTP_PORT={{oozie_server_port}}` to `OOZIE_HTTP_PORT=11000`.
- e. Set the `oozie.base.url` to the HTTPS address.
- f. Save the configuration, and restart the Oozie components.

12. Use the `/usr/hdp/current/oozie-server/bin/oozied.sh` script with the `start` parameter to start the Oozie server.

13. Check processes.

```
ps -ef | grep -i oozie
```

14. When needed, the `/usr/hdp/current/oozie-server/bin/oozied.sh` script with the `stop` parameter stops the Oozie server.

2.17. Configure and Start Apache WebHCat



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and webhcat to represent the WebHCat Service user. If you are using another name for these Service users, you need to substitute your Service user name for "hdfs" or "webhcat" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/hive-webhcat/conf` from the template to the `conf` directory in webhcat hosts.
2. Modify the Apache WebHCat configuration files.
 - a. Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User` (in this example, `hdfs`):

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/"
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/"
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /
hdp/apps/2.5.0.0-<$version>/pig/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.
gz /hdp/apps/2.5.0.0-<$version>/hive/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.
gz /hdp/apps/2.5.0.0-<$version>/sqoop/"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/pig.
tar.gz"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/
hive.tar.gz"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/
sqoop.tar.gz"
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

- b. Add the following property if it is not present in `webhcat-site.xml`:

```
<property>
  <name>templeton.libjars</name>
  <value>/usr/hdp/current/zookeeper-client/zookeeper.jar,/usr/hdp/current/
hive-client/lib/hive-common.jar</value>
  <description>Jars to add the classpath.</description>
</property>
```

- c. Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHCat to impersonate your additional HDP 2.5.0 groups and hosts:

```
<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>
</property>

<property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>
</property>
```

Where:

`hadoop.proxyuser.hcat.group`

Is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

`hadoop.proxyuser.hcat.hosts`

A comma-separated list of the hosts which are allowed to submit requests by 'hcat'.

3. Start WebHCat:

```
su - hcat -c "/usr/hdp/current/hive-webhcat/sbin/
webhcat_server.sh start"
```

4. Smoke test WebHCat.

- a. If you have a non-secure cluster, on the WebHCat host machine, run the following command to check the status of WebHCat server:

```
curl http://$WEBHCAT_HOST_MACHINE:20111/templeton/v1/status
```

You should see the following return status:

```
"status": "ok", "version": "v1"
```

- b. If you are using a Kerberos secure cluster, run the following command:

```
curl --negotiate -u: http://$WEBHCAT_HOST_MACHINE:20111/
templeton/v1/status
```

You should see the following return status

```
{"status": "ok", "version": "v1"}[machine@acme]$
```

2.18. Configure Apache Pig

Replace your configuration after upgrading. Copy `/etc/pig/conf` from the template to the `conf` directory in pig hosts.

2.19. Configure and Start Apache Sqoop



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Copy `/etc/sqoop/conf` from the template to the conf directory in sqoop hosts.
2. As the HDFS Service user, upload the Apache Sqoop tarball to HDFS.

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"

su - hdfs -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
```

3. If you are using the MySQL database as a source or target, then the MySQL connector jar must be updated to 5.1.29 or later.

Refer to the MySQL web site for information on updating the MySQL connector jar.

2.20. Configure, Start, and Validate Apache Flume

1. Replace your configuration after upgrading. Copy `/etc/flume/conf` from the template to the conf directory in Flume hosts.
2. By default, Apache Flume does not start running immediately upon installation. To validate your Flume upgrade, replace your default `conf/flume.conf` with the provided `flume.conf` file, restart Flume, and see if the data is flowing by examining the destination.

Use this `flume.conf` file:

```
#1. Name the components on this agent
a1.sources = r1
a1.sinks = k1
a1.channels = c1

#2. Describe/configure the source
a1.sources.r1.type = seq

#3. Describe the sink
a1.sinks.k1.type = file_roll
a1.sinks.k1.channel = c1
```

```

a1.sinks.k1.sink.directory = /tmp/flume

#4. Use a channel which buffers events in memory
a1.channels.c1.type = memory

#5. Bind the source and sink to the channel
a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1

```

3. After starting Flume, check /tmp/flume to see if there are any files there. The files should contain simple sequential numbers.
4. After validating, stop Flume and revert changes to flume.conf to prevent your disk from filling up.

2.21. Configure, Start, and Validate Apache Mahout



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Replace your configuration after upgrading. Copy /etc/mahout/conf from the template to the conf directory in mahout hosts.

To validate Apache Mahout:

1. Create a test user named "testuser" in the Linux cluster and in HDFS, and log in as that user.
2. Export the required environment variables for Mahout:

```

export JAVA_HOME="your_jdk_home_install_location_here"
export HADOOP_HOME=/usr/hdp/current/hadoop-client
export MAHOUT_HOME=/usr/hdp/current/mahout-client
export PATH="$PATH":$HADOOP_HOME/bin:$MAHOUT_HOME/bin
export CLASSPATH="$CLASSPATH":$MAHOUT_HOME

```

3. Upload a few megabytes of natural-language plain text to the Linux server as /tmp/sample-test.txt.
4. Transfer the sample-test.txt file to a subdirectory of the testusers's HDFS home directory.

```

hdfs dfs -mkdir /user/testuser/testdata
hdfs dfs -put /tmp/sample-test.txt /user/testuser/testdata

```

5. Set up mahout to convert the plain text file sample-test.txt into a sequence file that is in the output directory mahouttest.

```

mahout seqdirectory --input /user/testuser/sample-test.txt --output /user/
testuser/mahouttest --charset utf-8

```

2.22. Configure and Start Hue



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

For HDP 2.5.0, use the Hue version shipped with HDP 2.5.0. If you have a previous version of Hue, use the following steps to upgrade Hue.

1. Migrate hue.ini setting from your old hue.ini configuration file to new hue.ini configuration file.
2. If you are using the embedded SQLite database, remember to restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database. For example:

```
su - hue
cd /var/lib/hue
mv desktop.db desktop.db.old
sqlite3 desktop.db < ~/hue_backup/desktop.bak
exit
```

3. Synchronize Database

```
cd /usr/lib/hue
source ./build/env/bin/activate
hue syncdb
hue migrate
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

2.23. Configure and Start Apache Knox

When working with the Apache Knox Gateway in your Hadoop cluster, it is important you have the latest version of Knox installed so you can take advantage of new features and enhancements, in addition to ensuring your instance of Knox is in sync with other Hadoop components (e.g. Ranger, Spark, Hive, Hue, etc.) for stability and performance. For example, if you need to upgrade your Hadoop cluster from 2.2 to 2.3.x, you should also make sure that your individual Hadoop components are also upgraded to the latest version.

HDP enables you to perform a rolling upgrade in 2.2.x and onward. A rolling upgrade means that you can upgrade a component, or the entire Hadoop stack, without losing service, and your users can continue to use the cluster and run jobs with no application

or server downtime. The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple instances of the gateway and a load balancer serving each gateway instance from a single URL.

The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple instances of the gateway and a load balancer serving each gateway instance from a single URL. Once the upgrade process is completed, you are up and running with the latest version of Knox on each server you have designated as a Knox server.



Note

In this document, whenever you see a `{ }` with a value inside, this denotes a value you must define.

2.23.1. Upgrade the Knox Gateway

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Knox manually to the latest version. Because "rolling upgrades" are now supported in HDP 2.5.0, it is not important which version of Knox you are currently running, only that you have an instance of the Knox Gateway running.



Note

If you have not already installed Knox, refer to the "Install the Knox RPMs on the Knox Server" section of the *Non-Ambari Cluster Installation Guide* for instructions on how to install and configure the Knox Gateway.

Before upgrading the Knox Gateway, there are a several steps you must follow to ensure your configuration files, settings, and topology files can be copied to the new Knox Gateway instance when the upgrade is complete, which are described below.

1. Back up your existing `conf` directory if you have not already done so.
2. Stop each Knox server if you have not already done so.

```
su -l Knox /usr/hdp/{the current Knox version}/knox/bin/gateway.sh stop
```

3. Select the HDP server version you are upgrading to after you have stopped each Knox server if you have not already done so.

```
hdp-select set Knox-server {the HDP server version}
```

4.



Note

The `su` commands in this section use "Knox" to represent the Knox Service user. If you are using another name for your Knox Service user, you need to substitute your Knox Service user name for "Knox" in each of the `su` commands.

For HDP 2.5.0, the default paths for Knox change. Upgrade Knox in order to update these paths.

- a. Restore the backed up security directory. This places the master secret and keystores back in place for the new deployment.

- b. Start the Gateway:

```
su -knox -c "/usr/hdp/current/knox-server/bin/gateway.sh start"
```

- c. Unzip your previously saved configuration directory (the `conf` directory you backed up in step 1) into the new `/var/log/knox/gateway.conf` directory to import these files.

- d. Restart the Knox server to complete the upgrade.

```
su -l knox /usr/hdp/{the new HDP server version}/knox/bin/gateway.sh  
start
```

2.23.2. Verify the Knox Upgrade

To verify the upgrade was successful, follow the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful.
2. Verify you have cluster access using the `LISTSTATUS` WebHDFS API call.

```
curl -ivk -u {user}:{password} https://{knox host}:8443 /gateway/webhdfs/v1/  
tmp?op=LISTSTATUS
```

3. Verify the Knox version using the Knox Admin service and Version API.

```
curl -ivk -u {adminuser}:{adminpassword} https://{knox host}:8443 /gateway/  
admin/v1/version
```



Note

The Admin API requires you to be a member of an Admin group, as specified in the `admin.xml` authorization provider.

When you have verified the Knox upgrade was successful, you can begin using Knox. If, however, the upgrade was unsuccessful, you need to downgrade the Knox Gateway to the previous version. The steps to downgrade the Knox Gateway are described in the next section.

2.23.3. Downgrade the Knox Gateway to the Previous Version

If the Knox Gateway upgrade was unsuccessful, you need to downgrade Knox to the previous version to ensure you have a working Knox Gateway for your cluster. To downgrade Knox, follow the steps listed below.

1. For each server running Knox, stop the server.

```
su -l knox /usr/hdp/{current HDP server version}/knox/bin/gateway.sh stop
```

2. Select the HDP server version you want to use to downgrade your Knox Gateway.

```
hdp-select set Knox-server {previous HDP server version}
```

3. Restart the server.

```
su -l Knox /usr/hdp/{previous HDP server version}/Knox/bin/gateway.sh start
```

2.23.4. Verify the Knox Downgrade Was Successful

When the restart is complete, verify you are running an older version of Knox by following the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful
2. Verify you have cluster access using the `LISTSTATUSWebHDFS` API call.
3. Check the Knox version using the Knox Admin service and Version API using the following command:

```
curl -ivk -u {adminuser}:{adminpassword} https://{Knox host}:8443 /gateway/admin/v1/version
```

2.24. Configure and Validate Apache Falcon



Note

In HDP 2.5.0, if authorization is enabled (for example, in the properties file with `*.falcon.security.authorization.enabled=true`) then Access Control List (ACL) is mandated for all entities.

Upgrade Apache Falcon after you have upgraded HDFS, Hive, Oozie, and Pig. Stop Oozie jobs while running Falcon.

1. Replace your configuration after upgrading. Copy `/etc/falcon/conf` from the template to the conf directory in falcon hosts.
2. Check your Falcon entities. There should be no changes, but in some cases you may need to update your entities post-upgrade.
3. In HDP 2.3.2 for Falcon, TLS is enabled by default. When TLS is enabled, Falcon starts on `https://<falcon_host>.15443/`. You can disable TLS by adding the following line to the `startup.properties` file.

```
"*.falcon.enableTLS=false
```

4. If Transport Layer Security (TLS) is disabled, check the `client.properties` file to make sure the property `"falcon.uri"` is set as follows:

```
falcon.uri=http://<falcon_host>:15000/
```

2.25. Configure and Start Apache Storm



Note

The `su` commands in this section use "zookeeper" to represent the Zookeeper Service user. If you are using another name for your Zookeeper Service user, you need to substitute your Zookeeper Service user name for "zookeeper" in each of the `su` commands.

Apache Storm is fairly independent of changes to the HDP cluster, but you must upgrade Storm for rolling upgrade support in HDP 2.5.0 and be on the latest version of Storm.

1. Deactivate all running topologies.

2. Delete all states under zookeeper:

```
/usr/hdp/current/zookeeper-client/bin/zkCli.sh (optionally in  
secure environment specify -server zk.server:port)
```

```
rmr /storm
```

3. Delete all states under the storm-local directory:

```
rm -rf <value of storm.local.dir>/*
```

4. Stop Storm Services on the storm node.

5. Stop ZooKeeper Services on the storm node.

```
su - zookeeper -c "export ZOOCFGDIR=/etc/zookeeper/conf ; export  
ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-env.sh ; /  
usr/lib/zookeeper/bin/zkServer.sh stop"
```

6. Remove Storm and zookeeper from the storm node and install the HDP 2.5.0 version:

- For **RHEL/CentOS/Oracle Linux**:

```
yum erase storm
```

```
yum erase zookeeper
```

```
yum install storm
```

```
yum install zookeeper
```

- For **SLES**:

```
zypper rm storm
```

```
zypper rm zookeeper
```

```
zypper install storm
```

```
zypper install zookeeper
```

- For **Ubuntu/Debian**:

```
apt-get remove storm --purge

apt-get remove zookeeper --purge

apt-get install storm

apt-get install zookeeper
```

7. After upgrading Storm, replace your configuration in `/etc/storm/conf/storm.yaml` change references `backtype.storm` to `org.apache.storm`. This is needed because namespace migration in apache storm.
8. Replace your ZooKeeper configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
9. Ensure ZooKeeper is running. On the storm node, run the following command:

```
su - zookeeper -c "source /etc/zookeeper/conf/zookeeper-env.sh; export
ZOO_CFG_DIR=/etc/zookeeper/conf; /usr/hdp/current/zookeeper-server/bin/
zkServer.sh start >> $ZOO_LOG_DIR/zoo.out\""
```

where

- `$ZOO_LOG_DIR` is the directory where ZooKeeper server logs are stored. For example, `/var/log/zookeeper`.

- 10 Start Storm processes, using a process controller (like supervisord) or manually. Start nimbus, then supervisor/ui/drpc/logviewer:

```
su - storm /usr/hdp/current/storm-nimbus/bin/storm nimbus
```

You can use the same command syntax to start supervisor, ui, logviewer and drpc.

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm
supervisor
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm ui
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm logviewer
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm drpc
```

2.26. Configure and Start Apache Ranger

This section describes how to upgrade the Ranger service.

2.26.1. Preparing Your Cluster to Upgrade Ranger

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Apache Ranger manually to the latest version. This section describes the steps you need to follow to prepare your cluster for the Ranger upgrade.

1. Back up the following Ranger configuration directories(optional, since it has already been backed up in Getting ready to upgrade step):

- Ranger Policy Administration Service

```
/etc/ranger/admin/conf
```

- Ranger UserSync

```
/etc/ranger/usersync/conf
```

- Ranger Plugins:

- Hadoop

```
/etc/hadoop/conf
```

- Hive

```
/etc/hive/conf
```

- HBase

```
/etc/hbase/conf
```

- Knox

```
/etc/knox/conf
```

- Storm

```
/etc/storm/conf
```

- Kafka

```
/etc/kafka/conf
```

2. Backup the Ranger Policy and Audit databases. Make sure to take note of the following details in the `install.properties` file:

- db_host
- db_name
- db_user
- db_password
- policy manager configuration
- LDAP directory configuration
- LDAP settings
- LDAP AD domain

- LDAP URL

```
mysqldump -u root -p root ranger > dest_dir/filename.sql  
mysqldump -u root -p root ranger_audit > dest_dir/audit_filename.sql
```

2.26.2. Stop the Ranger Services

Now that you have prepared your cluster for the Ranger upgrade, you need to stop the Ranger Admin and Ranger UserSync services. To stop the Ranger services, perform the steps described below.

1. Stop the Ranger Policy Admin service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-admin stop
```

2. Stop the Ranger UserSync service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-usersync stop
```

3. Stop the applicable services using the Ranger plugin (HDFS, HBase, Hive, Knox, Storm, YARN, Kafka).

See [Stopping HDP Services](#) for more information.

2.26.3. Preparing the Cluster for Upgrade

Before you begin the upgrade process, you need to perform a series of steps to prepare the cluster for upgrade. These steps are described in the *"Getting Ready To Upgrade"* section of this guide, which you need to follow before continuing to upgrade Ranger. Some of these steps include:

- Backing up HDP directories
- Stopping all long-running applications and services.
- Backing up the Hive and Oozie metastore databases.
- Backing up Hue
- Backing up specific directories and configurations

2.26.4. Registering the HDP 2.5.0 Repo

After you have prepared your cluster for the upgrade, you need to register the HDP 2.5.0 repo. This requires you to perform the following steps:

1. (Optional) The Ranger components should already have installed in the at the beginning of the HDP upgrade process, but you can use the following commands to confirm that the Ranger packages have been installed:

```
hdp-select status ranger-admin  
hdp-select status ranger-usersync
```

If the packages have not been installed, you can use the install commands specific to your OS. For example, for RHEL/CentOS you would use the following commands to install the packages.

```
yum install ranger_2_5_*-admin
yum install ranger_2_5_*-usersync
```

2. Select the Ranger Admin and Ranger UserSync versions you want to use.

```
hdp-select set ranger-admin <HDP_server_version>
hdp-select set ranger-usersync <HDP_server_version>
```

3. Change ownership of `/etc/ranger/admin/conf/` and `/etc/ranger/usersync/conf/` to `ranger:ranger`:

```
chown -R ranger:ranger /etc/ranger/admin/conf/
```

4. Solr must be installed and configured before installing RangerAdmin or any of the Ranger component plugins.

For information regarding installation and configuration of Solr, see [Using Apache Solr for Ranger Audits](#).

5. Update the `install.properties` file to migrate the database credentials properties and `POLICYMGR_EXTERNAL-URL` property from HDP 2.3. to HDP 2.5.0.

Table 2.4. Ranger_Admin install.properties names and values

Property Name	Property Value
DB_FLAVOR	MySQL(ORACLE POSTGRES MSSQL SQLA)
db_root_user	root
db_root_password	Password of db (eg: vagrant)
db_host	Hostname : where your db does exist
polycmgr_external_url	http://<hostname>:6080
polycmgr_http_enabled	true
authentication_method	UNIX(LDAP ACTIVE_DIRECTORY UNIX NONE)
audit_store	solr
audit_solr_urls	http://<solr_host>:6083/solr/ranger_audits



Note

When you migrate to new version, you have to remove `/usr/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

6. Install the Ranger Admin component. Be sure to set the `JAVA_HOME` environment variable if it is not already set.

```
cd /usr/hdp/current/ranger-admin/

cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-site.xml ews/webapp/WEB-INF/classes/conf/
```

```
cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-default-site.xml ews/
webapp/WEB-INF/classes/conf/

cp ews/webapp/WEB-INF/classes/conf.dist/security-applicationContext.xml
ews/webapp/WEB-INF/classes/conf/

./setup.sh
```

This should successfully install the Ranger Admin component.

7. Start the Ranger Admin component.

```
service ranger-admin start
```

8. Configure and setup the Ranger UserSync component by migrating the properties from the HDP 2.3 `install.properties` file (POLICY_MGR_URL, SYNC_SOURCE and LDAP/AD properties).

Table 2.5. Ranger_Usersync install.properties names and values

Property Name	Property Value
POLICY_MGR_URL	http://<hostname>:6080
SYNC_SOURCE	unix
SYNC_INTERVAL	5



Note

When you migrate to new version, you have to remove `/user/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

9. Install the Ranger UserSync component. Be sure to set the JAVA_HOME component if it is not already set.

```
cd /usr/hdp/current/ranger-usersync/

./setup.sh
```

10 Start the Ranger UserSync component.

```
service ranger-usersync start
```

2.26.5. Install the Ranger Components

Next, you need to re-install each Ranger component again to ensure you have the latest version. Because you have already upgraded your HDP stack, you only need to follow the instructions in the *Non-Ambari Cluster Installation Guide* to install each Ranger component. You must install the following Ranger components:

- Ranger Policy Admin
- Ranger UserSync
- Ranger Plugins:

- HDFS
- HBase
- Hive
- Knox
- Storm
- YARN
- Kafka
- Solr



Note

When installing each Ranger component, you also need to make sure you upgrade each individual component to version 2.5.0 before restarting each service.

With this release, Ranger has also added support for the following components:

- Atlas

2.26.6. Restart the Ranger Services

Once you have re-installed each Ranger component, you then need to restart these components to ensure the new configurations are loaded in your cluster. This includes restarting the Policy Admin and UserSync components, NameNode, and each Ranger plugin.



Note

Before restarting the NameNode, make sure to remove the `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable the NameNode after finishing the upgrade.

The *Non-Ambari Cluster Installation Guide* describes how you can start the following Ranger services:

- Ranger Policy Admin service

```
service ranger-admin start
```

- Ranger UserSync service

```
service ranger-usersync start
```

2.26.7. Enable Ranger Plugins

The final step in the Ranger upgrade process requires you to re-enable the Ranger plugins. Although you are only required to enable HDFS in your cluster, you should re-enable all of

the Ranger plugins because class names have changed for the 2.5.0 release, and to ensure smooth operation of Ranger services in your cluster.



Note

When you enable each Ranger plugin, be sure to remove all 2.3 class name values.



Note

Re-enabling a Ranger plugin does not affect policies you have already created. As long as you use the same database as the Policy store, all of your data remains intact.

To re-enable the Ranger plugins, use the links listed below to access instructions in the *Non-Ambari Cluster Installation* guide that describe editing the `install.properties` file and enabling the Ranger plugins:



Important

Before enabling the HDFS plugin, remove `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable this plugin after the upgrade is complete.

- [HDFS Plugin](#)
- [YARN Plugin](#)
- [Kafka Plugin](#)
- [HBase Plugin](#)
- [Hive Plugin](#)
- [Knox Plugin](#)
- [Storm Plugin](#)

2.26.8. Enable KMS Configuration

Configure and setup the Ranger KMS configuration by editing the KMS `install.properties` file:

Table 2.6. KMS `install.properties` names and values

Property Name	Property Name
DB_FLAVOR	MYSQL
db_root_user	foot
db_root_password	<db password>
db_host	<db hostname>
db_name	rangerkms (default)
db_user	rangerkms (default)
db_password	<kms db password>

Property Name	Property Name
hadoop_conf	/etc/hadoop/conf
POLICY_MGR_URL	http://<hostname>:6080
REPOSITORY_NAME	kmsdev
XAAUDIT.SOLR.ENABLE	false (default), change to true if desired
XAAUDIT.SOLR.URL	http://<solr_host>:6083/solr/ranger_audits

2.26.9. Configure and Start Apache Ranger on a Kerberized Cluster

Beginning with HDP 2.5.0, kerberos authentication is supported for Ranger and its plugins. Use the this section if you have an HDP 2.3 cluster in a kerberized environment with Ranger in simple authentication mode that you want to upgrade.

For additional information regarding Kerberos, refer to [Kerberos Overview](#) in the *Hadoop Security Guide*.

2.26.9.1. Create Keytabs and Principals

1. Follow these steps to create keytabs and principals:

For Ranger Admin:

- a. Create rangeradmin/<FQDN of Ranger Admin>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangeradmin/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangeradmin/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

For Ranger Lookup:

- a. Create rangerlookup/<FQDN of Ranger Admin>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangerlookup/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerlookup/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/ranger-
upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Usersync:

- a. Create rangerusersync/<FQDN>@<REALM>.

```
b.
> kadmin.local
```

```
> addprinc -randkey rangerusersync/<FQDN of Ranger usersync>
Eg: addprinc -randkey rangerusersync/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Tagsync

a. Create rangertagsync/<FQDN>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangertagsync/<FQDN of Ranger tagsync>
Eg: addprinc -randkey rangertagsync/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger KMS:

a. Create rangerkms/<FQDN of Ranger Admin>@<REALM>

```
b.
> kadmin.local
> addprinc -randkey rangerkms/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerkms/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/<FQDN of Ranger
Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

- If the Kerberos server and admin are on different hosts, copy the keytab on the admin host, assign permission to user ranger, and change the permissions.

```
scp the <rangeradmin_keytab_file> to <new_path>
chown ranger <rangeradmin_keytab_path>
chmod 400 <rangeradmin_keytab_path>
```

- Use **kdestroy** to delete the Kerberos credentials cache file.
- Set the following properties and values in the `/etc/hadoop/conf/core-site.xml` file:

Table 2.7. Properties for the `/etc/hadoop/conf/core-site.xml` file

Property Name	Property Value
fs.defaultFS	hdfs://ranger-upgrade-0707-2.openstacklocal:8020
hadoop.security.authentication	kerberos
hadoop.security.authorization	true
hadoop.security.auth_to_local	RULE:[1:\$1@\$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/ RULE:[1:\$1@\$0](.*@EXAMPLE.COM)s/@.*/ RULE:[2:\$1@

Property Name	Property Value
	\$0](dn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0](nn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0](rangeradmin@EXAMPLE.COM)s/.*/ranger/ RULE:[2:\$1@\$0](rangerusersync@EXAMPLE.COM)s/.*/rangerusersync/ DEFAULT

The following is an example of a `core-site.xml` file with the properties set for Kerberos:

```
<configuration>

  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://ranger-upgrade-0707-2.openstacklocal:8020</value>
    <final>true</final>
  </property>

  <property>
    <name>hadoop.security.authentication</name>
    <value>kerberos</value>
  </property>

  <property>
    <name>hadoop.security.authorization</name>
    <value>true</value>
  </property>

  <property>
    <name>hadoop.security.auth_to_local</name>
    <value>RULE:[1:$1@$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/
RULE:[1:$1@$0](.*@EXAMPLE.COM)s/@.*/ /
RULE:[2:$1@$0](dn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](nn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](rangeradmin@EXAMPLE.COM)s/.*/ranger/
RULE:[2:$1@$0](rangerusersync@EXAMPLE.COM)s/.*/rangerusersync/
DEFAULT</value>
  </property>

</configuration>
```

2.26.9.2. Run Setup Again for Ranger Admin

1. Stop the ranger-admin service.
2. Add the following properties and values to the ranger-admin `install.properties` file:

Table 2.8. Ranger-admin `install.properties`

Property Name	Property Value
spnego_principal	HTTP/<FQDN_OF_Ranger_Admin_Cluster>@<REALM>
spnego_keytab	<HTTP keytab path>
token_valid	30
cookie_domain	<FQDN_OF_Ranger_Admin_Cluster>
cookie_path	/

Property Name	Property Value
admin_principal	rangeradmin/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
admin_keytab	<rangeradmin keytab path>
lookup_principal	rangerlookup/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
lookup_keytab	<rangerlookup keytab path>
hadoop_conf	/etc/hadoop/conf

3. Execute the `setup.sh` script.
4. Start the ranger-admin service.
5. Stop the ranger-usersync service.
6. Add the following properties and values to the ranger-sync `install.properties` file:

Table 2.9.

Property Name	Property Value
usersync_principal	rangerusersync/<FQDN>@<REALM>
usersyn_keytab	<rangerusersync keytab path>
hadoop_conf	/etc/hadoop/conf

7. Execute the `setup.sh` script.
8. Start the ranger-sync service.
9. Stop the ranger-tagsync service.
- 10 Add the following properties and values to the ranger-tagsync `install.properties` file:

Table 2.10. Ranger-tagsync install.properties and values

Property Name	Property Value
tagsync_principal	rangertagsync/<FQDN>@<REALM>
tagsync_keytab	<rangertagsync keytab path>
hadoop_conf	/etc/hadoop/conf

- 11 Execute the `setup.sh` script.
- 12 Start the ranger-tagsync service.
- 13 Stop the ranger-kms service.
- 14 Add the following properties and values to the ranger-kms `install.properties` file:

Table 2.11. Ranger-kms install.properties and values

Property Name	Property Value
kms_principal	rangerkms/<FQDN>@<REALM>
kms_keytab	<rangerkms keytab path>
hadoop_conf	/etc/hadoop/conf

2.26.9.3. Install and Enable the Ranger HDFS Plugin

This section documents how to install and enable the Ranger HDFS plugin. You might want to consider making similar changes for the other Ranger plugins that you are using.

The Ranger HDFS plugin is located at `/usr/hdp/<version>/ranger-hdfs-plugin`.

Follow these steps to install and enable the Ranger HDFS Plugin:

1. Edit the relevant lines in the `install.properties` file:

```
POLICY_MGR_URL=http://<FQDN of ranger admin host>:6080
REPOSITORY_NAME=hadoopdev
Audit info (Solr/HDFS options available)
```

2. Enter the following commands to enable the HDFS plugin:

```
export JAVA_HOME=<JAVA Path>
./enable-hdfs-plugin.sh
```

3. Start and stop the namenode:

```
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh stop
namenode"
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start
namenode"
```

4. In the custom repo configuration file, add the component user, for example, `hdfs`, as a value for each of the following properties:

- `policy.download.auth.users` or `policy.grantrevoke.auth.users`
- `tag.download.auth.users`

5. Verify that the plugin communicates with Ranger admin using the **Audit # plugins** tab.

6. Set the following properties in the `hdfs-site.xml` file:

Table 2.12. hdfs-site.xml Property Names and Values

Property Name	Property Value
<code>dfs.permissions.enabled</code>	<code>true</code>
<code>dfs.permissions.supergroup</code>	<code>hdfs</code>
<code>dfs.block.access.token.enable</code>	<code>true</code>
<code>dfs.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.secondary.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.principal</code>	<code>HTTP/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.keytab</code>	<code>/etc/security/keytabs/spnego.service.keytab</code>
<code>dfs.datanode.kerberos.principal</code>	<code>dn/_HOST@EXAMPLE.COM</code>
<code>dfs.namenode.keytab.file</code>	<code>/etc/security/keytabs/nn.service.keytab</code>
<code>dfs.secondary.namenode.keytab.file</code>	<code>/etc/security/keytabs/nn.service.keytab</code>
<code>dfs.datanode.keytab.file</code>	<code>/etc/security/keytabs/dn.service.keytab</code>
<code>dfs.https.port</code>	<code>50470</code>
<code>dfs.namenode.https-address</code>	<code>Example:ip-10-111-59-170.ec2.internal:50470</code>
<code>dfs.datanode.data.dir.perm</code>	<code>750</code>

Property Name	Property Value
dfs.cluster administrators	hdfs
dfs.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}
dfs.secondary.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}

The following is an example of a `hdfs-site.xml` file with the properties set for Kerberos:

```
<property>
  <name>dfs.permissions</name>
  <value>true</value>
  <description> If "true", enable permission checking in
HDFS. If "false", permission checking is turned
off, but all other behavior is
unchanged. Switching from one parameter value to the other does
not change the mode, owner or group of files or
directories. </description>
</property>

<property>
  <name>dfs.permissions.supergroup</name>
  <value>hdfs</value>
  <description>The name of the group of
super-users.</description>
</property>

<property>
  <name>dfs.namenode.handler.count</name>
  <value>100</value>
  <description>Added to grow Queue size so that more
client connections are allowed</description>
</property>

<property>
  <name>ipc.server.max.response.size</name>
  <value>5242880</value>
</property>

<property>
  <name>dfs.block.access.token.enable</name>
  <value>true</value>
  <description> If "true", access tokens are used as capabilities
for accessing datanodes. If "false", no access tokens are checked on
accessing datanodes. </description>
</property>

<property>
  <name>dfs.namenode.kerberos.principal</name>
  <value>nn/_HOST@EXAMPLE.COM</value>
  <description> Kerberos principal name for the
NameNode </description>
</property>

<property>
  <name>dfs.secondary.namenode.kerberos.principal</name>
  <value>nn/_HOST@EXAMPLE.COM</value>
  <description>Kerberos principal name for the secondary NameNode.
```

```
    </description>
  </property>

  <property>
    <!--cluster variant -->
    <name>dfs.secondary.http.address</name>
    <value>ip-10-72-235-178.ec2.internal:50090</value>
    <description>Address of secondary namenode web server</description>
  </property>

  <property>
    <name>dfs.secondary.https.port</name>
    <value>50490</value>
    <description>The https port where secondary-namenode
    binds</description>
  </property>

  <property>
    <name>dfs.web.authentication.kerberos.principal</name>
    <value>HTTP/_HOST@EXAMPLE.COM</value>
    <description> The HTTP Kerberos principal used by Hadoop-Auth in the
    HTTP endpoint.
    The HTTP Kerberos principal MUST start with 'HTTP/' per Kerberos HTTP
    SPNEGO specification.
    </description>
  </property>

  <property>
    <name>dfs.web.authentication.kerberos.keytab</name>
    <value>/etc/security/keytabs/spnego.service.keytab</value>
    <description>The Kerberos keytab file with the credentials for the HTTP
    Kerberos principal used by Hadoop-Auth in the HTTP endpoint.
    </description>
  </property>

  <property>
    <name>dfs.datanode.kerberos.principal</name>
    <value>dn/_HOST@EXAMPLE.COM</value>
    <description>
    The Kerberos principal that the DataNode runs as. "_HOST" is replaced
    by the real
    host name.
    </description>
  </property>

  <property>
    <name>dfs.namenode.keytab.file</name>
    <value>/etc/security/keytabs/nn.service.keytab</value>
    <description>
    Combined keytab file containing the namenode service and host
    principals.
    </description>
  </property>

  <property>
    <name>dfs.secondary.namenode.keytab.file</name>
    <value>/etc/security/keytabs/nn.service.keytab</value>
    <description>
    Combined keytab file containing the namenode service and host
    principals.
```

```
    </description>
  </property>

  <property>
    <name>dfs.datanode.keytab.file</name>
    <value>/etc/security/keytabs/dn.service.keytab</value>
    <description>
      The filename of the keytab file for the DataNode.
    </description>
  </property>

  <property>
    <name>dfs.https.port</name>
    <value>50470</value>
    <description>The https port where namenode
      binds</description>
  </property>

  <property>
    <name>dfs.https.address</name>
    <value>ip-10-111-59-170.ec2.internal:50470</value>
    <description>The https address where namenode binds</description>
  </property>

  <property>
    <name>dfs.datanode.data.dir.perm</name>
    <value>750</value>
    <description>The permissions that should be there on
      dfs.data.dir directories. The datanode will not come up if the
      permissions are different on existing dfs.data.dir directories. If
      the directories don't exist, they will be created with this
      permission.</description>
  </property>

  <property>
    <name>dfs.access.time.precision</name>
    <value>0</value>
    <description>The access time for HDFS file is precise upto this
      value.The default value is 1 hour. Setting a value of 0
      disables access times for HDFS.
    </description>
  </property>

  <property>
    <name>dfs.cluster.administrators</name>
    <value> hdfs</value>
    <description>ACL for who all can view the default
      servlets in the HDFS</description>
  </property>

  <property>
    <name>ipc.server.read.threadpool.size</name>
    <value>5</value>
    <description></description>
  </property>

  <property>
    <name>dfs.namenode.kerberos.internal.spnego.principal</name>
    <value>${dfs.web.authentication.kerberos.principal}</value>
  </property>
```

```
<property>
  <name>dfs.secondary.namenode.kerberos.internal.spnego.principal</name>
  <value>${dfs.web.authentication.kerberos.principal}</value>
</property>
```

7. For Download Policy to be successful, use the Ranger UI to update the service configuration with the following custom properties for each supported component:

```
policy.download.auth.users=<Component service user>
tag.download.auth.users=<Component service user>(if tag download)
```

8. For Grant/Revoke for Hive and Hbase to be successful, use the Ranger UI to update the service configuration with the following custom property:

```
policy.grantrevoke.auth.users = <Component service user>
```

9. For Test Connection and Resource Lookup to be successful, use the Ranger UI to add lookup user in the permission list of the policies.

2.27. Configuring and Upgrading Apache Spark

Before you can upgrade Apache Spark, you must have first upgraded your HDP components to the latest version (in this case, 2.5.0). This section assumes that you have already upgraded your components for HDP 2.5.0. If you have not already completed these steps, return to [Getting Ready to Upgrade](#) and [Upgrade 2.3 Components](#) for instructions on how to upgrade your HDP components to 2.5.0.

To upgrade Spark, start the service and update configurations.

1. Stop the Spark history-server. If you are using the Spark thrift-server, stop the thrift-server.

```
su - spark -c "$SPARK_HOME/sbin/stop-history-server.sh"
su - spark -c "$SPARK_HOME/sbin/stop-thriftserver.sh"
```

2. Remove any reference to hdp.version from the Spark configuration files.

Remove spark.yarn.services property from spark-defaults.conf.

Make sure that spark.history.provider is set to org.apache.spark.deploy.history.FsHistoryProvider.

3. Restart the history-server:

```
su - spark -c "$SPARK_HOME/sbin/start-history-server.sh"
```

4. If you are using the Spark thrift-server, restart the thrift-server. See [\(Optional\) Starting the Spark Thrift Server](#).
5. Validate the Spark installation. As user spark, run the examples in [Running Spark Applications](#) in the *Spark Guide*.

For additional configuration information, see the [Spark Guide](#).

2.28. Upgrade Apache Slider

To upgrade Apache Slider, simply upgrade the Slider client.

1. Upgrade Slider client:

```
hdp-select set slider-client 2.5.0.0-<version>  
slider version
```

2.29. Upgrade Apache Kafka

Upgrade each Apache Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the new version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"  
hdp-select set kafka-broker 2.3.4.0-2633  
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.
3. If the upgrade process fails, follow the steps in "Downgrading Kafka" to return to your previous version of Kafka.

2.29.1. Downgrading Kafka

Downgrade each Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the previous version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"  
hdp-select set kafka-broker 2.5.0.0-2041  
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.

2.30. Finalize the Upgrade



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.

1. Verify your file system health before finalizing the upgrade. After you finalize an upgrade, your backup is discarded!
2. As the \$HDFS_USER, enter:

```
su - hdfs -c "dfsadmin -finalizeUpgrade"
```

2.31. Migrate the Audit Logs from DB to Solr

Beginning with HDP 2.5.0, audit log to DB support has been removed. If your logs were previously stored on DB, you can migrate the logs to Solr. Refer to [Migrating Audit Logs from DB to Solr](#).

2.32. Install New HDP 2.5.0 Services

1. Install new HDP 2.5.0 Services (see the [Non-Ambari Cluster Installation Guide](#)):
 - Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
 - SmartSense – a next generation subscription model that features upgrade and configuration recommendations.
2. As Ranger has upgraded to a new version, you can Tag-sync, using the following Tagsync > install.properties, if required.

Table 2.13. Ranger Tagsync_install.properties and values

Property Name	Property Value
TAG_DEST_RANGER_ENDPOINT	http://<hostname>:6080
If TAG_SOURCE_ATLAS_ENABLED=true	
TAG_SOURCE_ATLAS_KAFKA_BOOTSTRAP_SERVERS	<hostname>:6667
TAG_SOURCE_ATLAS_KAFKA_ZOOKEEPER_CONNECT	<hostname>:2181
TAG_SOURCE_ATLAS_KAFKA_ENTITIES_GROUP_ID	ranger_entities_consumer
TAG_SOURCE_ATLAS_KAFKA_SERVICE_NAME	kafka
TAG_SOURCE_ATLAS_KAFKA_SECURITY_PROTOCOL	PLAINTEXTSASL
If TAG_SOURCE_ATLASREST_ENABLED = true	
TAG_SOURCE_ATLASREST_ENDPOINT	http://<hostname>:21000
TAG_SOURCE_ATLASREST_DOWNLOAD_INTERVAL_IN_MINUTES	60000
TAG_SOURCE_ATLASREST_USERNAME	
TAG_SOURCE_ATLASREST_PASSWORD	
If TAG_SOURCE_FILE_ENABLED = true	
TAG_SOURCE_FILE_ENABLED	true
TAG_SOURCE_FILE_FILENAME	/etc/ranger/data/tags.json (where the tag-sync is installed)

3. Upgrade from HDP 2.2 to HDP 2.5.0 Manually

This chapter provides instructions on how to manually upgrade to HDP 2.5.0 from the HDP 2.2 release. It assumes the existing HDP 2.2 was also installed manually.



Important

If you installed and manage HDP 2.2 with Ambari, **you must use the [Ambari Upgrade Guide](#)** to perform the the HDP 2.2 to HDP 2.5.0 upgrade.

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

Rolling Upgrade involves complex orchestration as well as side-by-side installation. It is too complex for a manual procedure, and is therefore supported only as an Ambari feature. If you wish to perform a Rolling Upgrade, refer to the Ambari Install instructions to install Ambari, then follow the Ambari Rolling Upgrade instructions, see [Ambari Upgrade Guide](#).

The HDP packages for a complete installation of HDP 2.5.0 consumes about 6.5 GB of disk space.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.

The following provides an overview of steps for upgrading to the latest release of HDP 2.5.0 from HDP 2.2:

1. Get ready to upgrade
2. Upgrade HDP 2.2 components
3. Use hdp-select to symlink the HDP 2.5.0 components into "current," in place of the former HDP 2.2 components
4. Configure and Start Apache ZooKeeper
5. Configure and Start Hadoop
6. Start HDFS
7. Configure and start YARN/MapReduce
8. Configure and Start Apache HBase
9. Configure and Start Apache Phoenix
10. Configure and Start Apache Accumulo
11. Configure and Start Apache Tez

- 12.Configure and Start Apache Hive and Apache HCatalog
- 13.Configure and Start Apache Oozie
- 14.Configure and Start Apache WebHCat (Templeton)
- 15.Configure and Start Apache Pig
- 16.Configure and Start Apache Sqoop
- 17.Configure and Start Apache Flume
- 18.Configure and Start Apache Mahout
- 19.Configure and Start Hue
- 20.Configure and Start Apache Knox
- 21.Configure and Start Apache Falcon
- 22.Configure and Start Apache Storm
- 23.Configure and Start Apache Ranger
- 24.Configure and Start Apache Spark
- 25.Upgrade Apache Slider
- 26.Upgrade Apache Kafka
- 27.Finalize the Upgrade
- 28.Install new HDP 2.5.0 services

3.1. Getting Ready to Upgrade

HDP Stack upgrade involves upgrading from HDP 2.2 to HDP 2.5.0 versions and adding the new HDP 2.5.0 services. These instructions change your configurations.



Note

You must use **kinit** before running the commands as any particular user.

Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP 2.5.0 consumes about 6.5 GB of disk space.

The first step is to ensure you keep a backup copy of your HDP 2.2 configurations.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Back up the HDP directories for any hadoop components you have installed.

The following is a list of all HDP directories:

- `/etc/hadoop/conf`
- `/etc/hbase/conf`
- `/etc/phoenix/conf`
- `/etc/hive-hcatalog/conf`
- `/etc/hive-webhcat/conf`
- `/etc/accumulo/conf`
- `/etc/hive/conf`
- `/etc/pig/conf`
- `/etc/sqoop/conf`
- `/etc/flume/conf`
- `/etc/mahout/conf`
- `/etc/oozie/conf`
- `/etc/hue/conf`
- `/etc/knox/conf`
- `/etc/zookeeper/conf`
- `/etc/tez/conf`
- `/etc/falcon/conf`
- `/etc/slider/conf/`
- `/etc/storm/conf/`
- `/etc/storm-slider-client/conf/`
- `/etc/spark/conf/`
- `/etc/ranger/admin/conf`, `/etc/ranger/usersync/conf` (If Ranger is installed, also take a backup of `install.properties` for all the plugins, ranger admin & ranger usersync.)
- Optional - Back up your userlogs directories, `${mapred.local.dir}/userlogs`.

2. Oozie runs a periodic purge on the shared library directory. The purge can delete libraries that are needed by jobs that started before the upgrade began and that finish after the upgrade. To minimize the chance of job failures, you should

extend the `oozie.service.ShareLibService.purge.interval` and `oozie.service.ShareLibService.temp.sharelib.retention.days` settings.

Add the following content to the `oozie-site.xml` file prior to performing the upgrade:

```
<property>
<name>oozie.service.ShareLibService.purge.interval</name>
<value>1000</value><description>
How often, in days, Oozie should check for old ShareLibs and LauncherLibs to
purge from HDFS.
</description>
</property>

<property>
<name>oozie.service.ShareLibService.temp.sharelib.retention.days</name>
<value>1000</value>
<description>
ShareLib retention time in days.</description>
</property>
```

3. Stop all long-running applications deployed using Slider:

```
su - yarn -c "/usr/hdp/current/slider-client/bin/slider list"
```

For all applications returned in previous command, run `su - yarn "/usr/hdp/current/slider-client/bin/slider stop <app_name>"`

4. Stop all services (including MapReduce) except HDFS, ZooKeeper, and Ranger, and client applications deployed on HDFS.

See [Stopping HDP Services](#) for more information.

Component	Command
Accumulo	<code>/usr/hdp/current/accumulo-client/bin/stop-all.sh</code>
Knox	<code>cd \$GATEWAY_HOME su - knox -c "bin/gateway.sh stop"</code>
Falcon	<code>su - falcon "/usr/hdp/current/falcon-server/bin/falcon-stop"</code>
Oozie	<code>su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-stop.sh"</code>
WebHCat	<code>su - webhcat -c "/usr/hdp/current/hive-webhcat/sbin/webhcat_server.sh stop"</code>
Hive	<p>Run this command on the Hive Metastore and Hive Server2 host machine:</p> <pre>ps aux awk '{print \$1,\$2}' grep hive awk '{print \$2}' xargs kill >/dev/null 2>&1</pre> <p>Or you can use the following:</p> <pre>Killall -u hive -s 15 java</pre>
HBase RegionServers	<code>su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-daemon.sh --config /etc/hbase/conf stop regionserver"</code>

Component	Command
HBase Master host machine	<pre>su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh --config /etc/hbase/conf stop master"</pre>
YARN & Mapred History	<p>Run this command on all NodeManagers:</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop nodemanager"</pre> <p>Run this command on the History Server host machine:</p> <pre>su - mapred -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-jobhistory-daemon.sh --config /etc/hadoop/conf stop historyserver"</pre> <p>Run this command on the ResourceManager host machine(s):</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop resourcemanager"</pre> <p>Run this command on the ResourceManager host machine:</p> <pre>su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-timelineserver/sbin/yarn-daemon.sh --config /etc/hadoop/conf"</pre> <p>Run this command on the YARN Timeline Server node:</p> <pre>su -l yarn -c "export HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec && /usr/hdp/current/hadoop-yarn-timelineserver/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"</pre>
Storm	<p>Deactivate all running topologies:</p> <pre>storm kill topology-name</pre> <p>Delete all states under zookeeper:</p> <pre>/usr/hdp/current/zookeeper-client/bin/zkCli.sh (optionally in secure environment specify -server zk.server:port)</pre> <pre>rmr /storm</pre> <p>Delete all states under the storm-local directory:</p> <pre>rm -rf <value of storm.local.dir></pre> <p>Stop Storm Serves on the storm node. If you are using Process Control System supervisord then use:</p>

Component	Command
	<pre>sudo service supervisord stop</pre> <p>Otherwise you can find storm processes and kill them:</p> <pre>ps -ef grep -i storm.home</pre> <p>Stop ZooKeeper Services on the storm node:</p> <pre>su - zookeeper -c "export ZOOCFGDIR=/etc/zookeeper/conf ; export ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-env.sh ; /usr/hdp/current//zookeeper-server/bin/zkServer.sh stop"</pre>
Spark (History server)	<pre>su - spark -c "/usr/hdp/current/spark-client/sbin/stop-history-server.sh"</pre>

5. If you have the Hive component installed, back up the Hive Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, see your database documentation.

Table 3.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hive < /tmp/mydir/backup_hive.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>su - postgres -c "pg_dump hive > /tmp/mydir/backup_hive.sql"</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hive < /tmp/mydir/backup_hive.sql</pre>
Oracle	<p>Export the database:</p> <pre>exp username/password@database full=yes file=output_file.dmp</pre>	<p>Import the database:</p> <pre>imp username/password@database file=input_file.dmp</pre>

6. If you have the Oozie component installed, back up the Oozie metastore database.

These instructions are provided for your convenience. Check your database documentation for the latest backup instructions.

Table 3.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sql</pre> <p>For example:</p>	<pre>mysql \$dbname < \$inputfilename.sql</pre> <p>For example:</p> <pre>mysql oozie < /tmp/mydir/backup_oozie.sql</pre>

Database Type	Backup	Restore
	<code>mysqldump oozie > /tmp/mydir/backup_oozie.sql</code>	
PostgreSQL	<code>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql</code> For example: <code>su - postgres -c "pg_dump oozie > /tmp/mydir/ backup_oozie.sql"</code>	<code>sudo -u \$username psql \$databasename < \$inputfilename.sql</code> For example: <code>sudo -u postgres psql oozie < /tmp/mydir/ backup_oozie.sql</code>
Oracle	Export the database: <code>exp username/password@database full=yes file=output_file.dmp</code>	Import the database: <code>imp username/password@database file=input_file.dmp</code>

7. Optional: Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 3.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	<code>mysqldump \$dbname > \$outputfilename.sqlsbr</code> For example: <code>mysqldump hue > /tmp/mydir/ backup_hue.sql</code>	<code>mysql \$dbname < \$inputfilename.sqlsbr</code> For example: <code>mysql hue < /tmp/mydir/ backup_hue.sql</code>
PostgreSQL	<code>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</code> For example: <code>sudo -u postgres pg_dump hue > / tmp/mydir/backup_hue.sql</code>	<code>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</code> For example: <code>sudo -u postgres psql hue < /tmp/ mydir/backup_hue.sql</code>
Oracle	Connect to the Oracle database using sqlplus. Export the database. For example: <code>exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr</code>	Import the database: For example: <code>imp username/password@database file=input_file.dmp</code>
SQLite	<code>/etc/init.d/hue stop</code> <code>su \$HUE_USER</code> <code>mkdir ~/hue_backup</code> <code>sqlite3 desktop.db .dump > ~/ hue_backup/desktop.bak</code> <code>/etc/init.d/hue start</code>	<code>/etc/init.d/hue stop</code> <code>cd /var/lib/hue</code> <code>mv desktop.db desktop.db.old</code> <code>sqlite3 desktop.db < ~/hue_backup/ desktop.bak</code> <code>/etc/init.d/hue start</code>

8. Back up the Knox data/security directory.

```
cp -RL /etc/knox/data/security ~/knox_backup
```


9. Save the namespace by executing the following commands:

```
su - hdfs

hdfs dfsadmin -safemode enter

hdfs dfsadmin -saveNamespace
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

10. Run the `fsck` command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.)

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-old-fsck-1.log"
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

11. Use the following instructions to compare status before and after the upgrade.

The following commands must be executed by the user running the HDFS service (by default, the user is hdfs).

a. Capture the complete namespace of the file system. (The following command does a recursive listing of the root file system.)



Important

Make sure the namenode is started.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-old-lsr-1.log"
```



Note

In secure mode you must have Kerberos credentials for the hdfs user.

b. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-old-report-1.log"
```

c. **Optional:** You can copy all or unrecoverable only data store libext-customer directory in HDFS to a local file system or to a backup instance of HDFS.

d. **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.

12. Finalize any prior HDFS upgrade, if you have not done so already.

```
su - hdfs -c "hdfs dfsadmin136finalizeUpgrade"
```



Note

In secure mode, you must have Kerberos credentials for the hdfs user.

13 Stop remaining services (HDFS, Zookeeper, and Ranger).

See [Stopping HDP Services](#) for more information.

Component	Command
HDFS	<p>On all DataNodes:</p> <p>If you are running secure cluster, run following command as root:</p> <pre>/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode</pre> <p>Else:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode"</pre> <p>If you are not running a highly available HDFS cluster, stop the Secondary NameNode by executing this command on the Secondary NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop secondarynamenode"</pre> <p>On the NameNode host machine(s)</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop namenode"</pre> <p>If you are running NameNode HA, stop the ZooKeeper Failover Controllers (ZKFC) by executing this command on the NameNode host machine:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop zkfc"</pre> <p>If you are running NameNode HA, stop the JournalNodes by executing these command on the JournalNode host machines:</p> <pre>su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop journalnode"</pre>
Zookeeper Host machines	<pre>su - zookeeper -c "/usr/hdp/current/zookeeper-server/bin/zookeeper-server stop"</pre>
Ranger (XA Secure)	<pre>service ranger-admin stop</pre> <pre>service ranger-usersync stop</pre>

14 Back up your NameNode metadata.



Note

It's recommended to take a backup of the full `/hadoop.hdfs/namenode` path.

- a. Copy the following checkpoint files into a backup directory.

The NameNode metadata is stored in a directory specified in the `hdfs-site.xml` configuration file under the configuration value `"dfs.namenode.name.dir"`.

For example, if the configuration value is:

```
<property>
  <name>dfs.namenode.name.dir</name>
  <value>/hadoop/hdfs/namenode</value>
</property>
```

Then, the NameNode metadata files are all housed inside the directory `/hadoop.hdfs/namenode`.

- b. Store the layoutVersion of the namenode.

```
${dfs.namenode.name.dir}/current/VERSION
```

15. Verify that the edit logs in `${dfs.namenode.name.dir}/current/edits*` are empty.

- a. Run the following command on the active NameNode machine:

```
hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out
```

- b. Verify the `edits.out` file. It should only have `OP_START_LOG_SEGMENT` transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

- c. If `edits.out` has transactions other than `OP_START_LOG_SEGMENT`, run the following steps and then verify edit logs are empty.
 - Start the existing version NameNode.
 - Ensure there is a new FS image file.
 - Shut the NameNode down:

```
hdfs dfsadmin - saveNamespace
```

16. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path
component in this version of HDFS.

Please rollback and delete or rename this path, or upgrade with the
-renameReserved key-value pairs option to automatically rename these
paths during upgrade.
```

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup.

For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, specify `-upgrade -renameReserved .snapshot=.my-snapshot,.reserved=.my-reserved`.

If no key-value pairs are specified with `-renameReserved`, the NameNode suffixes reserved paths with:

```
.<LAYOUT-VERSION>.UPGRADE_RENAMED
```

For example: `.snapshot.-51.UPGRADE_RENAMED`.



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

17. Upgrade the JDK on all nodes to JDK 7 before upgrading HDP. If you are running JDK 7, no action is required.



Important

If you want to upgrade to from JDK 7 to JDK 8, you must update the HDP stack before upgrading to JDK 8. For example, the high-level process should follow:

- a. Run HDP 2.2 with JDK 7.
- b. Perform the stack upgrade to HDP 2.5.0. See [Upgrade HDP 2.2 Components](#).
- c. Upgrade JDK from 7 to 8.

3.2. Upgrade HDP 2.2 Components

The upgrade process to HDP 2.5.0 involves the following steps. See the [Release Notes](#) for the HDP 2.5.0.0 repo information. Select your OS:

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.2 components. This command uninstalls the HDP 2.2 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"spark" "ranger" "slider" "hdp_mon_nagios_addons" "bigtop"
```

3. Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP 2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.2 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "hadoop-lzo" "openssl" "hive-webhcat" "hive-
catalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase"
"hue" "hive" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout"
"knox" "kafka" "spark" "ranger" "slider" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.2 components. This command uninstalls the HDP 2.2 components. It leaves the user data, and metadata, but removes your configurations:

```
zypper rm "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"spark" "slider" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.5.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.2 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "hue" "tez"
"storm" "falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "kafka"
"spark" "slider" "hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component files names should appear in the returned list.

SLES

1. On all hosts, clean the zypper repository.

```
zypper clean -a
```

2. Remove your old HDP 2.2 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*"
"pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*"
"falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common"
"hue-shell" "knox*" "ranger*" "spark" "slider" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.2 component binaries are uninstalled:

```
zypper search --installed-only --repo HDP-2.2
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/susel1sp3/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.2 components:

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "oozie" "collectd" "gccxml"
"pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "tez" "storm" "falcon"
"flume" "phoenix" "accumulo" "mahout" "knox" "ranger" "kafka" "spark"
"spark-python" "hdp_mon_nagios_addons" "slider" "hive-webhcat" "hive-hcatalog"
```

```
zypper up -r HDP-2.5
```

```
zypper install oozie-client
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep
hcatalog
```

No component files names should appear in the returned list.

Ubuntu 12

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.2 components. This command uninstalls the HDP 2.2 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez.
*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue.*"
"knox*" "spark" "slider" "hdp_mon_nagios_addons" --purge
```

3. Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/
updates/2.5.0.0/hdp.list - O /etc/apt/sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.2 components:

```
apt-get install "hadoop" "hadoop-lzo" "hadoop-hdfs" "libhdfs0" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm"
"falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "kafka" "spark"
"slider" "hdp_mon_nagios_addons"
```




Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

Debian

1. On all hosts, clean the apt-get repository.

```
apt-get clean
```

2. Remove your old HDP 2.2 components. This command uninstalls the HDP 2.2 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "hive-webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*"
"storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue*" "knox*"
"spark*" "slider*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.2 component binaries are uninstalled:

```
dpkg -l | grep "^ii" | grep hadoop
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
debian<version>/2.x/updates/2.5.0.0/hdp.list -O /etc/apt/
sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.2 components:

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "hive-webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm"
"falcon" "flume" "phoenix" "accumulo" "mahout" "knox" "kafka" "spark"
"slider" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Ranger, see [Upgrade Ranger](#) for more information on the upgrade path.

3.3. Symlink Directories with hdp-select



Warning

HDP 2.5.0 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

1. Before you run hdp-select, remove one link:

```
rm /usr/bin/oozie
```

2. Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5.0.0-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5.0.0-2800
```

3.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available HDFS cluster, configure and restart Apache ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.

1. Replace your configuration after upgrading on all the ZooKeeper nodes. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "/usr/hdp/current/zookeeper-server/bin/zookeeper-server start"
```

3.5. Configure Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed in HDP 2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed since HDP 2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

Ubuntu/Debian

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

3.6. Start Hadoop Core



Warning

Before you start HDFS on a highly available HDFS cluster, you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

To start HDFS, run commands as the `$HDFS_USER`.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading on all the HDFS nodes. Replace the HDFS template configuration in `/etc/hdfs/conf`.

2. If you are upgrading from a highly available HDFS cluster configuration, start all JournalNodes. On each JournalNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start journalnode"
```



Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

3. If you are running HDFS on a highly available namenode, you must first start the Zookeeper service



Note

Perform this step only if you are on a highly available HDFS cluster.

Run the following command on all NameNode hosts.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc"
```

4. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually.

On the active NameNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the -upgrade option only once. After you have completed this step, you can bring up the NameNode using this command without including the -upgrade option.



Note

If you receive the error:

```
Failed to find Premain-Class manifest attribute in /usr/hdp/<HDP-version>/hadoop/lib/ranger-hdfs-plugin-shim-0.5.0.***.jar
Error occurred during initialization of VM
agent library failed to init: instrument
```

after an upgrade to HDP-2.5, remove set-hdfs-plugin-env.sh from the /usr/hdp/<hdp-version>/hadoop/conf/ directory.

To check if the Upgrade is in progress, check that the "\previous" directory has been created in the \NameNode and \JournalNode directories. The "\previous" directory contains a snapshot of the data before upgrade.

In a highly available HDFS cluster configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode.

To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the '-bootstrapStandby' flag. Do NOT start this standby NameNode with the '-upgrade' flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The bootstrapStandby command downloads the most recent fsimage from the active NameNode into the \$dfs.name.dir directory of the standby NameNode. You can enter that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

5. Verify that the NameNode is up and running:

```
ps -ef | grep -i NameNode
```

6. If you do not have a highly available HDFS cluster configuration (non_HA namenode), start the Secondary NameNode.



Note

Do not perform this step if you have a highly available HDFS cluster configuration.

On the Secondary NameNode host machine, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

7. Verify that the Secondary NameNode is up and running.



Note

Do not perform this step if you have a highly available HDFS cluster environment.

```
ps -ef | grep SecondaryNameNode
```

8. Start DataNodes.

On each of the DataNodes, enter the following command. Note: If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-  
daemon.sh start datanode"
```

9. Verify that the DataNode process is up and running:

```
ps -ef|grep DataNode
```

10. Verify that NameNode can go out of safe mode.

```
>su - hdfs -c "hdfs dfsadmin -safemode wait"
```

You should see the following result: Safe mode is OFF

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

3.7. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

1. Run the `fsck` command on namenode as `$HDFS_USER`:

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-new-  
fsck-1.log"
```

Open `dfs-new-fsck-1.log` to see that the filesystem under path `/` is HEALTHY.

2. Run `hdfs` namespace and report.

a. List directories.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-new-lsr-1.log"
```

b. Open the `dfs-new-lsr-1.log` and confirm that you can see the file and directory listing in the namespace.

c. Run report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-new-report-1.log"
```

d. Open the `dfs-new-report` file and validate the admin report.

3. Compare the namespace report before the upgrade and after the upgrade. Verify that user files exist after upgrade.

The file names are listed below:

```
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log
```

```
dfs-old-lsr-1.log < -- > dfs-new-lsr-1.log
```



Note

You must do this comparison manually to catch all errors.

4. From the NameNode WebUI, determine if all DataNodes are up and running.

```
http://<namenode>:<namenodeport>
```

5. If you are on a highly available HDFS cluster, go to the StandbyNameNode web UI to see if all DataNodes are up and running:

```
http://<standbynamenode>:<namenodeport>
```

6. If you are **not** on a highly available HDFS cluster, go to the SecondaryNameNode web UI to see if the secondary node is up and running:

```
http://<secondarynamenode>:<secondarynamenodeport>
```

7. Verify that read and write to hdfs works successfully.

```
hdfs dfs -put [input file] [output file]
```

```
hdfs dfs -cat [output file]
```

3.8. Configure YARN and MapReduce

After you upgrade Hadoop, complete the following steps to update your configs.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

In secure mode, you must have Kerberos credentials for the hdfs user.

1. Upload the MapReduce tarball to HDFS. As the HDFS user, for example 'hdfs':

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/  
mapreduce/"
```

```
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/
mapreduce.tar.gz /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

```
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<
$version>/mapreduce"
```

```
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<
$version>/mapreduce/mapreduce.tar.gz"
```

2. Make sure that the following properties are in /etc/hadoop/conf/mapred-site.xml:

- Make sure `mapreduce.application.framework.path` exists in `mapred-site.xml`:

```
<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}/mapreduce/mapreduce.tar.gz#mr-framework</
value>
</property>

<property>
  <name>yarn.app.mapreduce.am.admin-command-opts</name>
  <value>-Dhdp.version=${hdp.version}</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
  <value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}/hadoop/lib/native:/usr/
hdp/${hdp.version}/hadoop/
  lib/native/Linux-amd64-64</value>
</property>

<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.admin.reduce.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.application.classpath</name>
  <value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
```



```
$PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/common/*,
$PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/yarn/*,
$PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
$PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
/usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
/etc/hadoop/conf/secure</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.



Note

If you are planning to use Spark in yarn-client mode, make Spark work in yarn-client mode 2.5.0.0-<\$version>.

3. Make sure the following property is in `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.application.classpath</name>
  <value>$HADOOP_CONF_DIR,/usr/hdp/${hdp.version}/hadoop-client/*,
    /usr/hdp/${hdp.version}/hadoop-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/lib/*</value>
</property>
```

4. On secure clusters only, add the following properties to `/etc/hadoop/conf/yarn-site.xml`:

```
<property>
  <name>yarn.timeline-service.recovery.enabled</name>
  <value>TRUE</value>
</property>

<property>
  <name>yarn.timeline-service.state-store.class</name>
  <value>org.apache.hadoop.yarn.server.timeline.recovery.
    LevelDbTimelineStateStore</value>
</property>

<property>
  <name>yarn.timeline-service.leveldb-state-store.path</name>
  <value><the same as the default of "yarn.timeline-service-leveldb-
    timeline-store.path">
  </value>
</property>
```

5. For secure clusters, you must create and configure the container-executor.cfg configuration file:

- Create the container-executor.cfg file in `/etc/hadoop/conf/container-executor.cfg`.
- Insert the following properties:

```
yarn.nodemanager.linux-container-executor.group=hadoop
banned.users=hdfs,yarn,mapred
min.user.id=1000
```

- `yarn.nodemanager.linux-container-executor.group` - Configured value of `yarn.nodemanager.linux-container-executor.group`. This must match the value of `yarn.nodemanager.linux-container-executor.group` in `yarn-site.xml`.
- `banned.users` - Comma-separated list of users who can not run container-executor.
- `min.user.id` - Minimum value of user id. This prevents system users from running container-executor.
- `allowed.system.users` - Comma-separated list of allowed system users.
- Set the file `/etc/hadoop/conf/container-executor.cfg` file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn/bin/container-executor
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn/bin/container-executor
```

3.9. Start YARN/MapReduce Services

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.



Note

The `su` commands in this section use "yarn" to represent the YARN Service user and `mapreduce` to represent the MAPREDUCE Service user. If you are using another name for these Service users, you need to substitute your Service user name for "yarn" or "mapreduce" in each of the `su` commands.

1. Manually clear the ResourceManager state store.

```
su - yarn -c "yarn resourcemanager -format-state-store"
```

2. Start the ResourceManager on all your ResourceManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh start resourcemanager"

ps -ef | grep -i resourcemanager
```

3. Start the TimelineServer on your TimelineServer host.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-timelineviewer/sbin/yarn-daemon.sh start timelineserver"
```

```
ps -ef | grep -i timelineserver
```

4. Start the NodeManager on all your NodeManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh  
start nodemanager"
```

```
ps -ef | grep -i nodemanager
```

5. To start MapReduce, run the following commands:

```
su - mapred -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-  
jobhistory-daemon.sh start historyserver"
```

```
ps -ef | grep -i jobhistoryserver
```

3.10. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user.

The job uses MapReduce to write 100MB of data into HDFS with RandomWriter

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.  
jar  
randomwriter -Dtest.randomwrite.total_bytes=10000000 test-after-  
upgrade.
```

You should see messages similar to:

```
map 0% reduce 0%  
...map 100% reduce 100%  
Job ... completed successfully
```

MapReduce upgraded successfully. You can now upgrade your other components.

Basic Troubleshooting

To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

```
http://<resource manager host>:8088/cluster/nodes
```

The number of active nodes should be equal to the number of nodemanagers.

Accessing error messages:

1. Access the ApplicationMaster WebUI to view the container logs.
2. At your console logs for MapReduce job, look for a line with this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track the job: http://<resource  
manager host>:8088/proxy/application_1380673658357_0007/
```

3. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

3.11. Configure and Start Apache HBase



Note

The `su` commands in this section use "hbase" to represent the HBASE Service user. If you are using another name for your Apache HBASE Service user, you need to substitute your HBASE Service user name for "hbase" in each of the `su` commands.

To enable Kerberos for clusters with dual home network setting, each HBase RegionServer must have its own key. See [Installing HDP Manually](#).

The `hbase.bucketcache.percentage.in.combinedcache` is removed in HDP 2.5.0. This simplifies the configuration of block cache. BucketCache configurations from HDP 2.2 need to be recalculated to attain identical memory allotments in HDP 2.5.0. The L1 LruBlockCache is whatever `hfile.block.cache.size` is set to and the L2 BucketCache is whatever `hbase.bucketcache.size` is set to.

1. Replace your configuration after upgrading. Replace the HBase template configuration in `/etc/hbase/conf`.
2. Start services. From root, assuming that `$HBASE_USER=hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

3. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

4. Run HBase smoke test. Verify that following commands can be executed successfully from HBase shell.

```
create 'testtable', 'cf1'
```

```
put 'testtable', 'r1', 'cf1:c1', 'v1'
```

```
disable 'testtable'
```

```
drop 'testtable'
```

3.12. Configure Apache Phoenix

To configure Apache Phoenix, complete the following steps:

1. Add the following property to the `/etc/hbase/hbase-site.xml` file on all HBase nodes, the MasterServer, and all RegionServers to prevent deadlocks from occurring during maintenance on global indexes:

```
<property>
  <name>hbase.regionserver.wal.codec</name>
  <value>org.apache.hadoop.hbase.regionserver.wal.IndexedWALEditCodec</value>
</property>
```

2. To enable user-defined functions, configure the following property in `/etc/hbase/conf` on all Hbase nodes.

```
<property>
  <name>phoenix.functions.allowUserDefinedFunctions</name>
  <value>true</value>
  <description>enable UDF functions</description>
</property>
```

3. Ensure the client side `hbase-site.xml` matches the server side configuration.
4. **(Optional)** To use local indexing, set the following three properties. These properties ensure collocation of data table and local index regions:



Note

The local indexing feature is a technical preview and considered under development. Do not use this feature in your production systems. If you have questions regarding this feature, contact Support by logging a case on our Hortonworks Support Portal at <http://hortonworks.com/support/>.

Add the following two properties, if they do not already exist, to the master side configurations:

```
<property>
  <name>hbase.master.loadbalancer.class</name>
  <value>org.apache.phoenix.hbase.index.balancer.IndexLoadBalancer</value>
</property>

<property>
  <name>hbase.coprocessor.master.classes</name>
  <value>org.apache.phoenix.hbase.index.master.IndexMasterObserver</value>
</property>
```

Add the following, if it does not already exist, to the Region server side configurations:

```
<property>
  <name>hbase.coprocessor.regionserver.classes</name>
  <value>org.apache.hadoop.hbase.regionserver.LocalIndexMerger</value>
</property>
```

5. If the folder specified in `hbase.tmp.dir` property on `hbase-site.xml` does not exist, create that directory with adequate permissions.
6. Restart the HBase Master and Region Servers.

Configuring Phoenix to Run in a Secure Cluster

Perform the following additional steps to configure Phoenix to run in a secure Hadoop cluster:

1. To link the HBase configuration file with the Phoenix libraries:

```
ln -sf HBASE_CONFIG_DIR/hbase-site.xml PHOENIX_HOME/bin/hbase-site.xml
```

2. To link the Hadoop configuration file with the Phoenix libraries:

```
ln -sf HADOOP_CONFIG_DIR/core-site.xml PHOENIX_HOME/bin/core-site.xml
```



Note

When running the `pssql.py` and `slline.py` Phoenix scripts in secure mode, you can safely ignore the following warnings.

```
14/04/19 00:56:24 WARN util.NativeCodeLoader:
Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable

14/04/19 00:56:24 WARN util.DynamicClassLoader: Failed to identify the fs of
dir hdfs://<HOSTNAME>:8020/apps/hbase/data/lib, ignored java.io.IOException:
No FileSystem for scheme: hdfs
```

3.13. Configure and Start Apache Accumulo



Note

The `su` commands in this section use "accumulo" to represent the Apache Accumulo Service user. If you are using another name for your Accumulo Service user, you need to substitute your Accumulo Service user name for "accumulo" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/accumulo/conf` from the template to the `conf` directory in Accumulo hosts.
2. Start the services:

```
ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
master\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
tserver\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname` gc\"
accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
tracer\" accumulo"

ssh `hostname` "sudo su - -c \"export ACCUMULO_CONF_DIR=/etc/accumulo/conf/
server;/usr/hdp/current/accumulo-client/bin/start-server.sh `hostname`
monitor\" accumulo"
```

3. Check that the processes are running

```
ps -ef | grep accumulo
```

or visit <http://<hostname>:50095> in your browser

4. Run Accumulo smoke test. Verify that following commands can be executed successfully from Accumulo shell.

```
createtable testtable
```

```
insert row1 colf colq value1
```

```
scan
```

```
deletetable testtable
```

```
tables
```

3.14. Configure and Start Apache Tez



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

To upgrade Apache Tez:

1. Copy your previously backed-up copy of `tez-site.xml` into the `/etc/tez/conf` directory.
2. Upload the Tez tarball to HDFS.

```
su - hdfs
hdfs dfs -mkdir -p /hdp/apps/<hdp_version>/tez/
hdfs dfs -put /usr/hdp/<hdp_version>/tez/lib/tez.tar.gz /hdp/apps/
<hdp_version>/tez/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/<hdp_version>/tez
hdfs dfs -chmod -R 444 /hdp/apps/<hdp_version>/tez/tez.tar.gz
```

Where `<hdp_version>` is the current HDP version, for example 2.5.0.0-2800.

3. Edit the `tez.lib.uris` property in the `tez-site.xml` file to point to `/hdp/apps/<hdp_version>/tez/tez.tar.gz`

```
...
<property>
  <name>tez.lib.uris</name>
  <value>/hdp/apps/<hdp_version>/tez/tez.tar.gz</value>
</property>
...
```

Where `<hdp_version>` is the current HDP version, for example 2.5.0.0-2800.

4. **Optional** Earlier releases of Tez did not have access control. In the current version of Tez, the default behavior restricts the ability to view the Tez history to only the owner of the job. To retain unrestricted access for non-secure clusters, set `tez.am.view-acls` set to `"*"`.
5. Change the value of the `tez.tez-ui.history-url.base` property to the url for the upgraded Tez View. For information on setting up the Tez view, see [Using the Tez View](#) in the HDP Ambari Views Guide.

6. Run Tez Smoke Test

To smoke test your Tez upgrade, you can run the following Tez Example job as a regular user.

MRRSleep Tez Example job sleeps for a defined period of time in mapper and reducer

```
hadoop jar /usr/hdp/current/tez-client/tez-tests-<version>.jar mrrsleep -m 1
-r 1 -mt 100 -rt 100
```

You should see messages similar to:

```
...DAG: State: SUCCEEDED Progress: 100%
DAG completed. FinalState=SUCCEEDED
```

Tez upgraded successfully. You can now upgrade your other components.

3.15. Configure and Start Apache Hive and Apache HCatalog



Note

The `su` commands in this section use "hive" to represent the Apache Hive Service user. If you are using another name for your Hive Service user, you need to substitute your Hive Service user name for "hive" in each of the `su` commands.



Note

When using HiveServer2 in HTTP mode, you must configure the mapping from Kerberos Principals to short names in the "hadoop.security.auth_to_local" property setting in the `core-site.xml` file.

1. Prior to starting the upgrade process, set the following in your hive configuration file:

```
datanucleus.autoCreateSchema=false
```

2. Copy the jdbc connector jar from `OLD_HIVE_HOME/lib` to `CURRENT_HIVE_HOME/lib`.
3. Upgrade the Hive Metastore database schema. Restart the Hive Metastore database and run:

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -upgradeSchema
-dbType <$databaseType>"
```


The value for \$databaseType can be derby, mysql, oracle, or postgres.



Note

If you are using PostgreSQL, you should reset the Hive Metastore database owner to <HIVE_USER>:

```
psql -U <POSTGRES_USER> -c "
ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER TO <HIVE_USER>"
```



Note

If you are using Oracle 11, you may see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
optimize.mapjoin.mapreduce does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
heapspace does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
server2.enable.impersonation does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
semantic.analyzer.factory.impl does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.
convert.sortmerge.join.noconditionaltask does not exist
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.
internal:1521/XE
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver
Metastore connection User: hiveuser
Starting upgrade metastore schema from version 0.13.0 to 0.14.0
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql
Error: ORA-00955: name is already used by an existing object
(state=42000,code=955)
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.
oracle.sql: Schema script failed, errorcode 2
Completed upgrade-0.13.0-to-0.14.0.oracle.sql
schemaTool completed
```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.

4. Edit the hive-site.xml file and modify the properties based on your environment. Search for TODO in the file for the properties to replace.

- a. Edit the following properties in the hive-site.xml file:

```
<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache
</description>
```

```
</property>
```

- b. **Optional:** To enable the Hive builtin authorization mode, make the following changes. If you want to use the advanced authorization provided by Ranger, refer to the [Ranger](#) instructions.

Set the following Hive authorization parameters in the hive-site.xml file:

```
<property>
  <name>hive.server2.enable.doAs</name>
  <value>>false</value>
</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
    StorageBasedAuthorizationProvider,org.apache.hadoop.hive.ql.security.
    authorization.MetaStoreAuthzAPIAuthorizeEmbedOnly</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
    SQLStdConfOnlyAuthorizeFactory</value>
</property>
```

Also set hive.users.in.admin.role to the list of comma-separated users who need to be added to admin role. A user who belongs to the admin role needs to run the "set role" command before getting the privileges of the admin role, as this role is not in the current roles by default.

Set the following in the hiveserver2-site.xml file.

```
<property>
  <name>hive.security.authenticator.manager</name>
  <value>org.apache.hadoop.hive.ql.security.
    SessionStateUserAuthenticator</value>
</property>

<property>
  <name>hive.security..authorization.enabled</name>
  <value>true</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
    SQLStdHiveAuthorizeFactory</value>
</property>
```

- c. For a remote Hive metastore database, set the IP address (or fully-qualified domain name) and port of the metastore host using the following hive-site.xml property value.

```
<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server.</description>
```

```
To enable HiveServer2, leave the property value empty.  
</description>  
</property>
```

You can further fine-tune your configuration settings based on node hardware specifications, using the HDP utility script.

5. Start Hive Metastore.

On the Hive Metastore host machine, run the following command:

```
su - hive -c "nohup /usr/hdp/current/hive-metastore/bin/hive  
--service metastore -hiveconf hive.log.file=hivemetastore.log  
>/var/log/hive/hivemetastore.out 2>/var/log/hive/  
hivemetastoreerr.log &"
```

By default hivemetastore.log gets populated into the directory specified in hive-log4j.properties available in hive conf directory (By default it is /tmp/hive).

[Optional]: you can specify your hive.log.dir while starting hive metastore process.

```
eg: -hiveconf hive.log.dir=/var/log/hive
```

6. Start Hive Server2.

On the Hive Server2 host machine, run the following command:

```
su - hive  
  
nohup /usr/hdp/current/hive-server2/bin/hiveserver2 -hiveconf  
hive.metastore.uris=" " -hiveconf hive.log.file=hiveserver2.log  
>/var/log/hive/hiveserver2.out 2> /var/log/hive/  
hiveserver2err.log &
```

By default hiveserver2.log gets populated into the directory specified in hive-log4j.properties available in hive conf directory (By default it is /tmp/hive).

[Optional]: you can specify your hive.log.dir while starting hiveserver2 process.

```
eg: -hiveconf hive.log.dir=/var/log/hive
```

3.16. Configure and Start Apache Oozie



Note

The duration of the Oozie upgrade is dependent on the amount of job history stored in ooziedb. This history must be backed up and restored during the upgrade process. Best practice when planning for upgrade is to backup ooziedb from your production oozie server and restore it to a test or development oozie server. This can help you estimate the time that is required to upgrade Oozie during your production upgrade.



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and "oozie" to represent the Oozie Service user. If you are using another name for your HDFS Service user or your Oozie Service user, you need to substitute your Service user names for "hdfs" or "oozie" in each of the `su` commands.

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore `oozie-site.xml` from your backup to the `conf` directory on each oozie server and client.

2. Copy the JDBC jar from `/usr/share/java` to `libext-customer`:

- a. Create the `/usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer` directory.

```
cd /usr/hdp/2.5.0.0-<version>/oozie-server
```

```
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the `libext-customer` directory.

```
chmod -R 777 /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer
```

3. Copy these files to the `libext-customer` directory

```
cp /usr/hdp/2.5.0.0-<version>/hadoop-client/lib/hadoop*lzo*.jar /usr/hdp/current/oozie-server/libext-customer
```

```
cp /usr/share/HDP-oozie/ext.zip /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer/
```

Also, copy Oozie db jar in `libext-customer`.

4. If Falcon was also installed and configured before upgrade in HDP 2.2.x, then after upgrade you might also need to do the following:

```
cp /usr/hdp/current/falcon-server/oozie/ext/falcon-oozie-el-extension-*.jar /usr/hdp/current/oozie-server/libext-customer
```

5. Extract share-lib.

```
/usr/hdp/current/oozie/bin/oozie-setup.sh sharelib create -fs hdfs://<namenode>:8020
```

To verify that the sharelibs extracted correctly, run the following command:

```
oozie admin -oozie http://<oozie server host address>:11000/oozie -shareliblist
```

There should be:

- Available ShareLib
- oozie
- hive
- distcp
- hcatalog
- sqoop
- mapreduce-streaming
- pig

Change the ownership and permissions of the oozie directory:

```
su -l hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie"
```

```
su -l hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

6. If a previous version of Oozie was created using auto schema creation, run the following SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

7. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh  
upgrade -run"
```

8. As root, prepare the Oozie WAR file.

```
chown oozie:oozie /usr/hdp/current/oozie-server/oozie-server/  
conf/server.xml
```

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-setup.sh  
prepare-war -d /usr/hdp/current/oozie-server/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar  
New Oozie WAR file with added 'JARS' at /var/lib/oozie/oozie-server/webapps/  
oozie.war
```

9. Make sure that following property is added in oozie-log4j.properties:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p  
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where `${oozie.instance.id}` is determined by oozie, automatically.

10.If you have custom Oozie actions, you must define them in `oozie-site.xml`. Edit the `/etc/oozie/conf/oozie-site.xml` file and add the following property:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>[Comma seperated list of custom actions]</value>
</property>
```

For example, if you have added Spark Action, enter the following:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>spark-action-0.1.xsd</value>
</property>
```

11.Configure HTTPS for the Oozie server.

- a. Create a self signed certificate or get certificate from a trusted CA for the Oozie Server
- b. Import the certificate to the client JDK trust store on all client nodes.
- c. In the Ambari Oozie configuration, set the following environment variables in `oozie-env.sh`, adding them if it does not exist:

```
export OOZIE_HTTPS_PORT=11443
export OOZIE_HTTPS_KEYSTORE_FILE=/home/oozie/.keystore
export OOZIE_HTTPS_KEYSTORE_PASS=password
```

- d. Change `OOZIE_HTTP_PORT={{oozie_server_port}}` to `OOZIE_HTTP_PORT=11000`.
- e. Set the `oozie.base.url` to the HTTPS address.
- f. Save the configuration, and restart the Oozie components.

12.Start Oozie as the Oozie user:

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-start.sh"
```

13.Check processes.

```
ps -ef | grep -i oozie
```

3.17. Configure and Start Apache WebHCat



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and webhcat to represent the Apache WebHCat Service user. If you are using another name for these Service users, you need to substitute your Service user name for "hdfs" or "webhcat" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/webhcat/conf` from the template to the conf directory in webhcat hosts.
2. Modify the Apache WebHCat configuration files.
 - a. Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User` (in this example, `hdfs`):

```

su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/"

su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/"

su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /
hdp/apps/2.5.0.0-<$version>/pig/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.
gz /hdp/apps/2.5.0.0-<$version>/hive/"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.
gz /hdp/apps/2.5.0.0-<$version>/sqoop/"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/
pig.tar.gz"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/
hive.tar.gz"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/
sqoop.tar.gz"

su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"

```

- b. Update the following properties in the `webhcat-site.xml` configuration file, as their values have changed:

```

<property>
  <name>templeton.pig.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/pig/pig.tar.gz</value>
</property>

<property>
  <name>templeton.hive.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/hive/hive.tar.gz</value>
</property>

<property>
  <name>templeton.streaming.jar</name>
  <value>hdfs:///hdp/apps/${hdp.version}/mapreduce/
    hadoop-streaming.jar</value>
  <description>The hdfs path to the Hadoop streaming jar file.</
description>

```

```

</property>

<property>
  <name>templeton.sqoop.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/sqoop/sqoop.tar.gz</value>
  <description>The path to the Sqoop archive.</description>
</property>

<property>
  <name>templeton.sqoop.path</name>
  <value>sqoop.tar.gz/sqoop/bin/sqoop</value>
  <description>The path to the Sqoop executable.</description>
</property>

<property>
  <name>templeton.sqoop.home</name>
  <value>sqoop.tar.gz/sqoop</value>
  <description>The path to the Sqoop home in the exploded archive.
    </description>
</property>

```



Note

You do not need to modify `${hdp.version}`.

- c. Add the following property if it is not present in `webhcat-site.xml`:

```

<property>
  <name>templeton.libjars</name>
  <value>/usr/hdp/current/zookeeper-client/zookeeper.jar,/usr/hdp/current/
hive-client/lib/hive-common.jar</value>
  <description>Jars to add the classpath.</description>
</property>

```

- d. Remove the following obsolete properties from `webhcat-site.xml`:

```

<property>
  <name>templeton.controller.map.mem</name>
  <value>1600</value>
  <description>Total virtual memory available to map tasks.</description>
</property>

<property>
  <name>hive.metastore.warehouse.dir</name>
  <value>/path/to/warehouse/dir</value>
</property>

```

- e. Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHCat to impersonate your additional HDP 2.5.0 groups and hosts:

```

<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>
</property>

<property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>

```



```
</property>
```

Where:

```
hadoop.proxyuser.hcat.group
```

Is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

```
hadoop.proxyuser.hcat.hosts
```

A comma-separated list of the hosts which are allowed to submit requests by 'hcat'.

3. Start WebHCat:

```
su - hcat
```

```
/usr/hdp/current/hive-webhcat/sbin/webhcat_server.sh start
```

4. Smoke test WebHCat.

- a. If you have a non-secure cluster, on the WebHCat host machine, run the following command to check the status of WebHCat server:

```
curl http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

You should see the following return status:

```
"status": "ok", "version": "v1"
```

- b. If you are using a Kerberos secure cluster, run the following command:

```
curl --negotiate -u: http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

You should see the following return status

```
{"status": "ok", "version": "v1"}[machine@acme]$
```

3.18. Configure Apache Pig

1. Replace your configuration after upgrading. Copy `/etc/pig/conf` from the template to the conf directory in pig hosts.
2. To validate the Apache Pig upgrade, complete the following steps:
 - a. On the host machine where Pig is installed, run the following commands:

```
sudo su -c "hadoop fs -copyFromLocal /etc/passwd passwd" $HDFS_USER
```

- b. Create a Pig script file named `/tmp/id.pig` that contains the following Pig Latin commands:

```
A = load 'passwd' using PigStorage(':'); B = foreach A generate $0 as id;
store B into '/tmp/id.out';
```

- c. Run the Pig script:

```
su - $HDFS_USER
pig -l /tmp/pig.log /tmp/id.pig
```

3.19. Configure and Start Apache Sqoop



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Copy `/etc/sqoop/conf` from the template to the conf directory in sqoop hosts.
2. As the HDFS Service user, upload the Apache Sqoop tarball to HDFS.

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"

su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"

su - hdfs -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
```

3. If you are using the MySQL database as a source or target, then the MySQL connector jar must be updated to 5.1.29 or later.

Refer to the MySQL web site for information on updating the MySQL connector jar.

4. Because Sqoop is a client tool with no server component, you need to run your own jobs to validate the upgrade.

3.20. Configure, Start, and Validate Apache Flume

1. To confirm that Apache Flume is working correctly, create an example configuration file. The following snippet is a sample configuration that can be set using the properties file. For more detailed information, see the "Flume User Guide."

```
a1.sources = pstream
a1.channels = memoryChannel
a1.channels.memoryChannel.type = memory

a1.sources.pstream.channels = memoryChannel
a1.sources.pstream.type = exec
a1.sources.pstream.command = tail -f /etc/passwd

a1.sinks = hdfsSink
```

```

a1.sinks.hdfsSink.type = hdfs
a1.sinks.hdfsSink.channel = memoryChannel
a1.sinks.hdfsSink.hdfs.path = hdfs://<NN host>:8020/tmp/flumetest
a1.sinks.hdfsSink.hdfs.fileType = SequenceFile
a1.sinks.hdfsSink.hdfs.writeFormat = Text

```

The source here is defined as an exec source. The agent runs a given command on startup, which streams data to stdout, where the source gets it. The channel is defined as an in-memory channel and the sink is an HDFS sink.

2. Given this configuration, you can start Flume as follows:

```

$ bin/flume-ng agent --conf ./conf --conf-file example.conf --name a1 -
Dflume.root.logger=INFO,console

```



Note

The directory specified for `--conf` argument would include a shell script `flume-env.sh` and potentially a `log4j` properties file. In this example, we pass a Java option to force Flume to log to the console and we go without a custom environment script.

3. After validating data in `hdfs://tmp/flumetest`, stop Flume and restore any backup files. Copy `/etc/flume/conf` to the `conf` directory in Flume hosts.

3.21. Configure, Start, and Validate Apache Mahout



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Replace your configuration after upgrading. Copy `/etc/mahout/conf` from the template to the `conf` directory in mahout hosts.

To validate Apache Mahout:

1. Create a test user named "testuser" in the Linux cluster and in HDFS, and log in as that user.
2. Export the required environment variables for Mahout:

```

export JAVA_HOME="your_jdk_home_install_location_here"
export HADOOP_HOME=/usr/hdp/current/hadoop-client
export MAHOUT_HOME=/usr/hdp.current/mahout-client
export PATH="$PATH":$HADOOP_HOME/bin:$MAHOUT_HOME/bin
export CLASSPATH="$CLASSPATH":$MAHOUT_HOME

```

3. Upload a few megabytes of natural-language plain text to the Linux server as `/tmp/sample-test.txt`.
4. Transfer the `sample-test.txt` file to a subdirectory of the testusers's HDFS home directory.

```
hdfs dfs -mkdir /user/testuser/testdata
hdfs dfs -put /tmp/sample-test.txt /user/testuser/testdata
```

5. Enter the mahout command to convert the plain text file sample-test.txt into a sequence file stored in the output directory mahouttest:

```
mahout seqdirectory --input /user/testuser/testdata --output /user/testuser/
mahouttest -ow --charset utf-8
```

3.22. Configure and Start Hue



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

For HDP 2.5.0, use the Hue version shipped with HDP 2.5.0. If you have a previous version of Hue, use the following steps to upgrade Hue.

Complete one of the following:

- **SQLite**

1. Copy the hue.ini settings from your old hue.ini configuration file to new hue.ini configuration file.
2. Restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database. For example:

```
su - hue
cd /var/lib/hue
mv desktop.db desktop.db.old
sqlite3 desktop.db < ~/hue_backup/desktop.bak
exit
```

3. Synchronize database.

```
cd /usr/lib/hue
source ./build/env/bin/activate
hue syncdb
hue migrate
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

- **MySQL**

1. Copy the hue.ini settings from your old hue.ini configuration file to new hue.ini configuration file.

2. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

3.23. Configure and Start Apache Knox

When working with the Apache Knox Gateway in your Hadoop cluster, it is important you have the latest version of Knox installed so you can take advantage of new features and enhancements, in addition to ensuring your instance of Knox is in sync with other Hadoop components (e.g., Ranger, Spark, Hive, Hue, etc.) for stability and performance. For example, if you need to upgrade your Hadoop cluster from 2.2 to 2.5.x, you should also make sure that your individual Hadoop components are also upgraded to the latest version.



Note

In this document, whenever you see a `{ }` with a value inside, this denotes a value you must define.

3.23.1. Upgrade the Knox Gateway

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Knox manually to the latest version. Because “rolling upgrades” are now supported in HDP 2.3.0, it is not important which version of Knox you are currently running, only that you have an instance of the Knox Gateway running.



Note

If you have not already installed Knox, refer to the “Install the Knox RPMs on the Knox Server” section of the *Non-Ambari Cluster Installation Guide* for instructions on how to install and configure the Knox Gateway.

Before upgrading the Knox Gateway, there are a several steps you must follow to ensure your configuration files, settings, and topology files can be copied to the new Knox Gateway instance when the upgrade is complete, which are described below.



Note

The `su` commands in this section use “knox” to represent the Knox Service user. If you are using another name for your Knox Service user, you need to substitute your Knox Service user name for “knox” in each of the `su` commands.

1. Back up your existing `conf` directory if you have not already done so.
2. Stop each Knox server if you have not already done so.

```
su -l knox /usr/hdp/{the current Knox version}/knox-server/bin/gateway.sh  
stop
```

3. Select the HDP server version you are upgrading to after you have stopped each Knox server if you have not already done so.

```
hdp-select set knox-server {the HDP server version}
```

4. Start the ldap service.

```
/usr/hdp/current/knox-server/bin/ldap.sh start
```

5. For HDP 2.5.0, the default paths for Knox change. Upgrade Knox in order to update these paths.

- a. Restore the backed up security directory. This places the master secret and keystores back in place for the new deployment.

- b. Start the Gateway:

```
su -l knox -c "/usr/hdp/current/knox-server/bin/gateway.sh start"
```

- c. Unzip your previously saved configuration directory (the `conf` directory you backed up in step 1) into the new `/var/log/knox/gateway.conf` directory to import these files.

- d. Restart the Knox server to complete the upgrade.

```
su -l knox -c "/usr/hdp/{the new HDP server version}/knox-server/bin/gateway.sh restart"
```

3.23.2. Verify the Knox Upgrade

To verify the upgrade was successful, follow the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful.
2. Verify you have cluster access using the `LISTSTATUS` WebHDFS API call.

```
curl -i -v -k -u {user}:{password} https://{knox host}:8443 /gateway/webhdfs/v1/?op=LISTSTATUS
```

3. Verify the Knox version using the Knox Admin service and Version API.

```
curl -i -v -k -u admin:admin-password -X GET "https://localhost:8443/gateway/admin/api/v1/version"
curl -i -v -k -u admin:admin-password -X GET "https://localhost:8443/gateway/sandbox/webhdfs/v1?op=LISTSTATUS"
```



Note

The Admin API requires you to be a member of an Admin group, as specified in the `admin.xml` authorization provider.

When you have verified the Knox upgrade was successful, you can begin using Knox.

3.24. Configure and Validate Apache Falcon



Note

In HDP 2.5.0, if authorization is enabled (for example, in the properties file with `*.falcon.security.authorization.enabled=true`) then Access Control List (ACL) is mandated for all entities.

Upgrade Apache Falcon after you have upgraded HDFS, Hive, Oozie, and Pig. Stop Oozie jobs while upgrading Falcon.

1. Replace your configuration after upgrading. Copy `/etc/falcon/conf` from the template to the conf directory in falcon hosts.
2. Check your Falcon entities. There should be no changes, but in some cases you may need to update your entities post-upgrade.
3. In HDP 2.5.0 for Falcon, TLS is enabled by default. When TLS is enabled, Falcon starts on `https://<falcon_host>.15443/`. You can disable TLS by adding the following line to the `startup.properties` file.

```
"*.falcon.enableTLS=false
```

4. If Transport Layer Security (TLS) is disabled, check the `client.properties` file to make sure the property `"falcon.uri"` is set as follows:

```
falcon.uri=http://<falcon_host>:15000/
```

3.25. Configure and Start Apache Storm

Apache Storm is fairly independent of changes to the HDP cluster, but you must upgrade Storm for rolling upgrade support in HDP 2.5.0 and be on the latest version of Storm.

1. After upgrading Storm, replace your configuration. In `/etc/storm/conf/` `storm.yaml` change references `backtype.storm` to `org.apache.storm`. This is needed because namespace migration in apache storm.
2. Replace your ZooKeeper configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
3. Ensure ZooKeeper is running. On the storm node, run the following command:

```
su - zookeeper -c "source /etc/zookeeper/conf/zookeeper-env.sh; export
ZOO_CFG_DIR=/etc/zookeeper/conf; /usr/hdp/current/zookeeper-server/bin/
zkServer.sh start >> $ZOO_LOG_DIR/zoo.out"
```

where

- `$ZOO_LOG_DIR` is the directory where ZooKeeper server logs are stored. For example, `/var/log/zookeeper`.
4. Start Storm processes, using a process controller (like `supervisord`) or manually. Start `nimbus`, then `supervisor/ui/drpc/logviewer`:

```
/usr/hdp/current/storm-nimbus/bin/storm nimbus.
```

You can use the same command syntax to start Storm using `nimbus/ui` and `logviewer`.

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm
supervisor
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm ui
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm logviewer
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm drpc
```

3.26. Configure and Start Apache Ranger

This section describes how to upgrade the Apache Ranger service.

3.26.1. Preparing Your Cluster to Upgrade Ranger

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Ranger manually to the latest version. This section describes the steps you need to follow to prepare your cluster for the Ranger upgrade.

1. Back up the following Ranger configuration directories:

- Ranger Policy Administration Service

```
/etc/ranger/admin/conf
```

- Ranger UserSync

```
/etc/ranger/usersync/conf
```

- Ranger Plugins:

- Hadoop

```
/etc/hadoop/conf
```

- Hive

```
/etc/hive/conf
```

- HBase

```
/etc/hbase/conf
```

- Knox

```
/etc/knox/conf
```

- Storm

```
/etc/storm/conf
```

2. Backup the Ranger Policy and Audit databases. Make sure to take note of the following details in the `install.properties` file:

- db_host
- db_name
- db_user
- db_password

- policy manager configuration
- LDAP directory configuration
- LDAP settings
- LDAP AD domain
- LDAP URL

```
mysqldump -u root -p root ranger > dest_dir/filename.sql  
mysqldump -u root -p root ranger_audit > dest_dir/audit_filename.sql
```

3.26.2. Stop the Ranger Services

Now that you have prepared your cluster for the Ranger upgrade, you need to stop the Ranger Admin and Ranger UserSync services. To stop the Ranger services, perform the steps described below.

1. Stop the Ranger Policy Admin service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-admin stop
```

2. Stop the Ranger UserSync service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service ranger-usersync stop
```

3. Stop the applicable services using the Ranger plugin (HDFS, HBase, Hive, Knox, Storm).

See [Stopping HDP Services](#) for more information.

3.26.3. Preparing the Cluster for Upgrade

Before you begin the upgrade process, you need to perform a series of steps to prepare the cluster for upgrade. These steps are described in the "*Getting Ready To Upgrade*" section of this guide, which you need to follow before continuing to upgrade Ranger. Some of these steps include:

- Backing up HDP directories
- Stopping all long-running applications and services.
- Backing up the Hive and Oozie metastore databases.
- Backing up Hue
- Backing up specific directories and configurations

3.26.4. Registering the HDP 2.4 Repo

After you have prepared your cluster for the upgrade, you need to register the HDP 2.5.0 repo. This requires you to perform the following steps:

1. (Optional) The Ranger components should already have installed in the at the beginning of the HDP upgrade process, but you can use the following commands to confirm that the Ranger packages have been installed:

```
hdp-select status ranger-admin
hdp-select status ranger-usersync
```

If the packages have not been installed, you can use the install commands specific to your OS. For example, for RHEL/CentOS you would use the following commands to install the packages.

```
yum install ranger_2_3_*-admin
yum install ranger_2_3_*-usersync
```

2. Select the Ranger Admin and Ranger UserSync versions you want to use.

```
hdp-select set ranger-admin <HDP_server_version>
hdp-select set ranger-usersync <HDP_server_version>
```

3. Change ownership of /etc/ranger/admin/conf/ and /etc/ranger/usersync/conf/ to ranger:ranger:

```
chown -R ranger:ranger /etc/ranger/admin/conf/
```

4. Solr must be installed and configured before installing RangerAdmin or any of the Ranger component plugins.

For information regarding installation and configuration of Solr, see [Using Apache Solr for Ranger Audits](#).

5. Update the `install.properties` file to migrate the database credentials properties and `POLICYMGR_EXTERNAL-URL` property from HDP 2.2 to HDP 2.5.0.

Table 3.4. Ranger_Admin install.properties names and values

Property Name	Property Value
DB_FLAVOR	MySQL(ORACLE POSTGRES MSSQL SQLA)
db_root_user	root
db_root_password	Password of db (eg: vagrant)
db_host	Hostname : where your db does exist
polycmgr_external_url	http://<hostname>:6080
polycmgr_http_enabled	true
authentication_method	UNIX(LDAP ACTIVE_DIRECTORY UNIX NONE)
audit_store	solr
audit_solr_urls	http://<solr_host>:6083/solr/ranger_audits



Note

When you migrate to new version, you have to remove `/user/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

6. Install the Ranger Admin component. Be sure to set the JAVA_HOME environment variable if it is not already set.

```
cd /usr/hdp/current/ranger-admin/

cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-site.xml ews/webapp/WEB-INF/classes/conf/

cp ews/webapp/WEB-INF/classes/conf.dist/ranger-admin-default-site.xml ews/webapp/WEB-INF/classes/conf/

cp ews/webapp/WEB-INF/classes/conf.dist/security-applicationContext.xml ews/webapp/WEB-INF/classes/conf/

./setup.sh
```

This should successfully install the Ranger Admin component.

7. Start the Ranger Admin component.

```
service ranger-admin start
```

8. Configure and setup the Ranger UserSync component by migrating the properties from the HDP 2.2 `install.properties` file (POLICY_MGR_URL, SYNC_SOURCE and LDAP/AD properties).

Table 3.5. Ranger_Usersync install.properties names and values

Property Name	Property Value
POLICY_MGR_URL	http://<hostname>:6080
SYNC_SOURCE	unix
SYNC_INTERVAL	5



Note

When you migrate to new version, you have to remove `/user/bin/ranger-admin`, which points to the older version of the ranger start file, `/usr/hdp/<version>/ranger-admin/ews/ranger-admin-services.sh`. After you remove this file, you have to run setup again.

9. Install the Ranger UserSync component. Be sure to set the JAVA_HOME component if it is not already set.

```
cd /usr/hdp/current/ranger-usersync/

./setup.sh
```

- 10 Start the Ranger UserSync component.

```
service ranger-usersync start
```

3.26.5. Install the Ranger Components

Next, you need to re-install each Ranger component again to ensure you have the latest version. Because you have already upgraded your HDP stack, you only need to follow the

instructions in the *Non-Ambari Cluster Installation Guide* to install each Ranger component. You must install the following Ranger components:

- Ranger Policy Admin
- Ranger UserSync
- Ranger Plugins:
 - HDFS
 - HBase
 - Hive
 - Knox
 - Storm



Note

When installing each Ranger component, you also need to make sure you upgrade each individual component to version 2.5.0 before restarting each service.

With this release, Ranger has also added support for the following components:

- Solr
- Kafka
- YARN

3.26.6. Restart the Ranger Services

Once you have re-installed each Ranger component, you then need to restart these components to ensure the new configurations are loaded in your cluster. This includes restarting the Policy Admin and UserSync components, NameNode, and each Ranger plugin.



Note

Before restarting the NameNode, make sure to remove the `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable the NameNode after finishing the upgrade.

The *Non-Ambari Cluster Installation Guide* describes how you can start the following Ranger services:

- Ranger Policy Admin service

```
service ranger-admin start
```

- Ranger UserSync service

```
service ranger-usersync start
```

3.26.7. Enable Ranger Plugins

The final step in the Ranger upgrade process requires you to re-enable the Ranger plugins. Although you are only required to enable HDFS in your cluster, you should re-enable all of the Ranger plugins because class names have changed for the 2.5.0 release, and to ensure smooth operation of Ranger services in your cluster.



Note

When you enable each Ranger plugin, be sure to remove all 2.2 class name values.



Note

Re-enabling a Ranger plugin does not affect policies you have already created. As long as you use the same database as the Policy store, all of your data remains intact.

To re-enable the Ranger plugins, use the links listed below to access instructions in the *Non-Ambari Cluster Installation* guide that describe editing the `install.properties` file and enabling the Ranger plugins:



Important

Before enabling the HDFS plugin, remove `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable this plugin after the upgrade is complete.

- [HDFS Plugin](#)
- [YARN Plugin](#)
- [Kafka Plugin](#)
- [HBase Plugin](#)
- [Hive Plugin](#)
- [Knox Plugin](#)
- [Storm Plugin](#)

3.26.8. Enable KMS Configuration

Configure and setup the Ranger KMS configuration by editing the KMS `install.properties` file:

Table 3.6. KMS `install.properties` names and values

Property Name	Property Name
DB_FLAVOR	MYSQL
db_root_user	root
db_root_password	<db password>
db_host	<db hostname>

Property Name	Property Name
db_name	rangerkms (default)
db_user	rangerkms (default)
db_password	<kms db password>
hadoop_conf	/etc/hadoop/conf
POLICY_MGR_URL	http://<hostname>:6080
REPOSITORY_NAME	kmsdev
XAAUDIT.SOLR.ENABLE	false (default), change to true if desired
XAAUDIT.SOLR.URL	http://<solr_host>:6083/solr/ranger_audits

3.26.9. Configure and Start Apache Ranger on a Kerberized Cluster

Beginning with HDP 2.5.0, kerberos authentication is supported for Ranger and its plugins. Use the this section if you have an HDP 2.2 cluster in a kerberized environment with Ranger in simple authentication mode that you want to upgrade.

For additional information regarding Kerberos, refer to [Kerberos Overview](#) in the *Hadoop Security Guide*.

3.26.9.1. Create Keytabs and Principals

1. Follow these steps to create keytabs and principals:

For Ranger Admin:

- a. Create rangeradmin/<FQDN of Ranger Admin>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangeradmin/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangeradmin/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangeradmin.keytab rangeradmin/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

For Ranger Lookup:

- a. Create rangerlookup/<FQDN of Ranger Admin>@<REALM>.

```
b.
> kadmin.local
> addprinc -randkey rangerlookup/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerlookup/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/<FQDN of
Ranger Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerlookup.keytab rangerlookup/ranger-
upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Usersync:

a. Create rangerusersync/<FQDN>@<REALM>.

b.

```
> kadmin.local
> addprinc -randkey rangerusersync/<FQDN of Ranger usersync>
Eg: addprinc -randkey rangerusersync/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerusersync.keytab rangerusersync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger Tagsync

a. Create rangertagsync/<FQDN>@<REALM>.

b.

```
> kadmin.local
> addprinc -randkey rangertagsync/<FQDN of Ranger tagsync>
Eg: addprinc -randkey rangertagsync/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
<FQDN>@<REALM>
Eg: xst -k /etc/security/keytabs/rangertagsync.keytab rangertagsync/
ranger-upgrade-0707-2.openstacklocal@EXAMPLE.COM
> exit
```

For Ranger KMS:

a. Create rangerkms/<FQDN of Ranger Admin>@<REALM>

b.

```
> kadmin.local
> addprinc -randkey rangerkms/<FQDN of Ranger Admin>
Eg: addprinc -randkey rangerkms/ranger-upgrade-0707-2.openstacklocal
> xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/<FQDN of Ranger
Admin>@<REALM>
Eg: xst -k /etc/security/keytabs/rangerkms.keytab rangerkms/ranger-
upgrade-0707-2.openstacklocal/EXAMPLE.COM
> exit
```

2. If the Kerberos server and admin are on different hosts, copy the keytab on the admin host, assign permission to user ranger, and change the permissions.

```
scp the <rangeradmin_keytab_file> to <new_path>
chown ranger <rangeradmin_keytab_path>
chmod 400 <rangeradmin_keytab_path>
```

3. Use **kdestroy** to delete the Kerberos credentials cache file.
4. Set the following properties and values in the `/etc/hadoop/conf/core-site.xml` file:

Table 3.7. Properties for the `/etc/hadoop/conf/core-site.xml` file

Property Name	Property Value
fs.defaultFS	hdfs://ranger-upgrade-0707-2.openstacklocal:8020

Property Name	Property Value
hadoop.security.authentication	kerberos
hadoop.security.authorization	true
hadoop.security.auth_to_local	RULE:[1:\$1@\$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/\$1@\$0](.*@EXAMPLE.COM)s/@.*/ // RULE:[2:\$1@\$0](dn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0](nn@EXAMPLE.COM)s/.*/hdfs/ RULE:[2:\$1@\$0](rangeradmin@EXAMPLE.COM)s/.*/ranger/ RULE:[2:\$1@\$0](rangerusersync@EXAMPLE.COM)s/.*/rangerusersync/ DEFAULT

The following is an example of a `core-site.xml` file with the properties set for Kerberos:

```
<configuration>

  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://ranger-upgrade-0707-2.openstacklocal:8020</value>
    <final>true</final>
  </property>

  <property>
    <name>hadoop.security.authentication</name>
    <value>kerberos</value>
  </property>

  <property>
    <name>hadoop.security.authorization</name>
    <value>true</value>
  </property>

  <property>
    <name>hadoop.security.auth_to_local</name>
    <value>RULE:[1:$1@$0](ambari-qa-cluster1@EXAMPLE.COM)s/.*/ambari-qa/
RULE:[1:$1@$0](.*@EXAMPLE.COM)s/@.*/ //
RULE:[2:$1@$0](dn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](nn@EXAMPLE.COM)s/.*/hdfs/
RULE:[2:$1@$0](rangeradmin@EXAMPLE.COM)s/.*/ranger/
RULE:[2:$1@$0](rangerusersync@EXAMPLE.COM)s/.*/rangerusersync/
DEFAULT</value>
  </property>

</configuration>
```

3.26.9.2. Run Setup Again for Ranger Admin

1. Stop the ranger-admin service.
2. Add the following properties and values to the ranger-admin `install.properties` file:

Table 3.8. Ranger-admin `install.properties`

Property Name	Property Value
spnego_principal	HTTP/<FQDN_OF_Ranger_Admin_Cluster>@<REALM>

Property Name	Property Value
spnego_keytab	<HTTP keytab path>
token_valid	30
cookie_domain	<FQDN_OF_Ranger_Admin_Cluster>
cookie_path	/
admin_principal	rangeradmin/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
admin_keytab	<rangeradmin keytab path>
lookup_principal	rangerlookup/ <FQDN_OF_Ranger_Admin_Cluster>@<REALM>
lookup_keytab	<rangerlookup keytab path>
hadoop_conf	/etc/hadoop/conf

3. Execute the `setup.sh` script.
4. Start the ranger-admin service.
5. Stop the ranger-usersync service.
6. Add the following properties and values to the ranger-sync `install.properties` file:

Table 3.9.

Property Name	Property Value
usersync_principal	rangerusersync/<FQDN>@<REALM>
usersync_keytab	<rangerusersync keytab path>
hadoop_conf	/etc/hadoop/conf

7. Execute the `setup.sh` script.
8. Start the ranger-sync service.
9. Stop the ranger-tagsync service.
10. Add the following properties and values to the ranger-tagsync `install.properties` file:

Table 3.10. Ranger-tagsync install.properties and values

Property Name	Property Value
tagsync_principal	rangertagsync/<FQDN>@<REALM>
tagsync_keytab	<rangertagsync keytab path>
hadoop_conf	/etc/hadoop/conf

11. Execute the `setup.sh` script.
12. Start the ranger-tagsync service.
13. Stop the ranger-kms service.
14. Add the following properties and values to the ranger-kms `install.properties` file:

Table 3.11. Ranger-kms install.properties and values

Property Name	Property Value
kms_principal	rangerkms/<FQDN>@<REALM>
kms_keytab	<rangerkms keytab path>
hadoop_conf	/etc/hadoop/conf

3.26.9.3. Install and Enable the Ranger HDFS Plugin

This section documents how to install and enable the Ranger HDFS plugin. You might want to consider making similar changes for the other Ranger plugins that you are using.

The Ranger HDFS plugin is located at `/usr/hdp/<version>/ranger-hdfs-plugin`.

Follow these steps to install and enable the Ranger HDFS Plugin:

1. Edit the relevant lines in the `install.properties` file:

```
POLICY_MGR_URL=http://<FQDN of ranger admin host>:6080
REPOSITORY_NAME=hadoopdev
Audit info (Solr/HDFS options available)
```

2. Enter the following commands to enable the HDFS plugin:

```
export JAVA_HOME=<JAVA Path>
./enable-hdfs-plugin.sh
```

3. Start and stop the namenode:

```
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh stop
namenode"
su hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start
namenode"
```

4. In the custom repo configuration file, add the component user, for example, `hdfs`, as a value for each of the following properties:
 - `policy.download.auth.users` or `policy.grantrevoke.auth.users`
 - `tag.download.auth.users`
5. Verify that the plugin communicates with Ranger admin using the **Audit # plugins** tab.
6. Set the following properties in the `hdfs-site.xml` file:

Table 3.12. hdfs-site.xml Property Names and Values

Property Name	Property Value
<code>dfs.permissions.enabled</code>	<code>true</code>
<code>dfs.permissions.supergroup</code>	<code>hdfs</code>
<code>dfs.block.access.token.enable</code>	<code>true</code>
<code>dfs.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.secondary.namenode.kerberos.principal</code>	<code>nn/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.principal</code>	<code>HTTP/_HOST@EXAMPLE.COM</code>
<code>dfs.web.authentication.kerberos.keytab</code>	<code>/etc/security/keytabs/spnego.service.keytab</code>

Property Name	Property Value
dfs.datanode.kerberos.principal	dn/_HOST@EXAMPLE.COM
dfs.namenode.keytab.file	/etc/security/keytabs/nn.service.keytab
dfs.secondary.namenode.keytab.file	/etc/security/keytabs/nn.service.keytab
dfs.datanode.keytab.file	/etc/security/keytabs/dn.service.keytab
dfs.https.port	50470
dfs.namenode.https-address	Example:ip-10-111-59-170.ec2.internal:50470
dfs.datanode.data.dir.perm	750
dfs.cluster administrators	hdfs
dfs.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}
dfs.secondary.namenode.kerberos.internal.spnego.principal	\${dfs.web.authentication.kerberos.principal}

The following is an example of a `hdfs-site.xml` file with the properties set for Kerberos:

```
<property>
  <name>dfs.permissions</name>
  <value>true</value>
  <description> If "true", enable permission checking in
HDFS. If "false", permission checking is turned
off, but all other behavior is
unchanged. Switching from one parameter value to the other does
not change the mode, owner or group of files or
directories. </description>
</property>

<property>
  <name>dfs.permissions.supergroup</name>
  <value>hdfs</value>
  <description>The name of the group of
super-users.</description>
</property>

<property>
  <name>dfs.namenode.handler.count</name>
  <value>100</value>
  <description>Added to grow Queue size so that more
client connections are allowed</description>
</property>

<property>
  <name>ipc.server.max.response.size</name>
  <value>5242880</value>
</property>

<property>
  <name>dfs.block.access.token.enable</name>
  <value>true</value>
  <description> If "true", access tokens are used as capabilities
for accessing datanodes. If "false", no access tokens are checked on
accessing datanodes. </description>
</property>

<property>
```

```
<name>dfs.namenode.kerberos.principal</name>
<value>nn/_HOST@EXAMPLE.COM</value>
<description> Kerberos principal name for the
NameNode </description>
</property>

<property>
  <name>dfs.secondary.namenode.kerberos.principal</name>
  <value>nn/_HOST@EXAMPLE.COM</value>
  <description>Kerberos principal name for the secondary NameNode.
  </description>
</property>

<property>
  <!--cluster variant -->
  <name>dfs.secondary.http.address</name>
  <value>ip-10-72-235-178.ec2.internal:50090</value>
  <description>Address of secondary namenode web server</description>
</property>

<property>
  <name>dfs.secondary.https.port</name>
  <value>50490</value>
  <description>The https port where secondary-namenode
  binds</description>
</property>

<property>
  <name>dfs.web.authentication.kerberos.principal</name>
  <value>HTTP/_HOST@EXAMPLE.COM</value>
  <description> The HTTP Kerberos principal used by Hadoop-Auth in the
  HTTP endpoint.
  The HTTP Kerberos principal MUST start with 'HTTP/' per Kerberos HTTP
  SPNEGO specification.
  </description>
</property>

<property>
  <name>dfs.web.authentication.kerberos.keytab</name>
  <value>/etc/security/keytabs/spnego.service.keytab</value>
  <description>The Kerberos keytab file with the credentials for the HTTP
  Kerberos principal used by Hadoop-Auth in the HTTP endpoint.
  </description>
</property>

<property>
  <name>dfs.datanode.kerberos.principal</name>
  <value>dn/_HOST@EXAMPLE.COM</value>
  <description>
  The Kerberos principal that the DataNode runs as. "_HOST" is replaced
  by the real
  host name.
  </description>
</property>

<property>
  <name>dfs.namenode.keytab.file</name>
  <value>/etc/security/keytabs/nn.service.keytab</value>
  <description>
  Combined keytab file containing the namenode service and host
```

```
    principals.  
    </description>  
</property>  
  
<property>  
  <name>dfs.secondary.namenode.keytab.file</name>  
  <value>/etc/security/keytabs/nn.service.keytab</value>  
  <description>  
    Combined keytab file containing the namenode service and host  
    principals.  
  </description>  
</property>  
  
<property>  
  <name>dfs.datanode.keytab.file</name>  
  <value>/etc/security/keytabs/dn.service.keytab</value>  
  <description>  
    The filename of the keytab file for the DataNode.  
  </description>  
</property>  
  
<property>  
  <name>dfs.https.port</name>  
  <value>50470</value>  
  <description>The https port where namenode  
    binds</description>  
</property>  
  
<property>  
  <name>dfs.https.address</name>  
  <value>ip-10-111-59-170.ec2.internal:50470</value>  
  <description>The https address where namenode binds</description>  
</property>  
  
<property>  
  <name>dfs.datanode.data.dir.perm</name>  
  <value>750</value>  
  <description>The permissions that should be there on  
    dfs.data.dir directories. The datanode will not come up if the  
    permissions are different on existing dfs.data.dir directories. If  
    the directories don't exist, they will be created with this  
    permission.</description>  
</property>  
  
<property>  
  <name>dfs.access.time.precision</name>  
  <value>0</value>  
  <description>The access time for HDFS file is precise upto this  
    value.The default value is 1 hour. Setting a value of 0  
    disables access times for HDFS.  
  </description>  
</property>  
  
<property>  
  <name>dfs.cluster.administrators</name>  
  <value> hdfs</value>  
  <description>ACL for who all can view the default  
    servlets in the HDFS</description>  
</property>
```

```

<property>
  <name>ipc.server.read.threadpool.size</name>
  <value>5</value>
  <description></description>
</property>

<property>
  <name>dfs.namenode.kerberos.internal.spnego.principal</name>
  <value>${dfs.web.authentication.kerberos.principal}</value>
</property>

<property>
  <name>dfs.secondary.namenode.kerberos.internal.spnego.principal</name>
  <value>${dfs.web.authentication.kerberos.principal}</value>
</property>

```

7. For Download Policy to be successful, use the Ranger UI to update the service configuration with the following custom properties:

```

policy.download.auth.users=<Component service user>
tag.download.auth.users=<Component service user>(if tag download)

```

8. For Grant/Revoke for Hive and Hbase to be successful, use the Ranger UI to update the service configuration with the following custom property:

```

policy.grantrevoke.auth.users = <Component service user>

```

9. For Test Connection and Resource Lookup to be successful, use the Ranger UI to add lookup user in the permission list of the policies.

3.27. Configuring and Upgrading Apache Spark

Instructions in this section are specific to HDP-2.3.4 and later. For earlier versions for HDP, refer to the version-specific documentation.

1. Add the node where you want Apache Spark 1.5.2 History Server to run. Install the version corresponding to the HDP version you currently have installed.

- a. `su - root`
- b. `yum install spark_spark_2_3_4_0_3371-master -y`
- c. To use Python: `yum install spark_spark_2_3_4_0_3371-python`
- d. `conf-select create-conf-dir --package spark --stack-version spark_2_3_4_0_3371 --conf-version 0`
- e. `cp /etc/spark/spark_2_3_4_0_3371/0/* /etc/spark/spark_2_3_4_0_3371/0/`
- f. `conf-select set-conf-dir --package spark --stack-version spark_2_3_4_0_3371 --conf-version 0`
- g. `hdp-select set spark-client spark_2_3_4_0_3371`
- h. `hdp-select set spark-historyserver spark_2_3_4_0_3371`

2. In HDP 2.4.0, the Spark History Server runs on top of HDFS, not YARN ATS, as in previous versions. Modify Spark configuration files as follows:

- a. As the hdfs service user, create an HDFS directory called spark-history with user:spark, user group:hadoop, and permissions = 777:

```
hdfs dfs -mkdir /spark-history
hdfs dfs -chown -R spark:hadoop /spark-history
hdfs dfs -chmod -R 777 /spark-history
```

- b. Edit the spark-defaults.conf file.

- Add the following properties and values:

```
spark.eventLog.dir to hdfs:///spark-history
spark.eventLog.enabled to true
spark.history.fs.logDirectory to hdfs:///spark-history
spark.history.provider to org.apache.spark.deploy.history.
FsHistoryProvider
```

- Delete the spark.yarn.services property.

- c. Edit the spark-thrift-sparkconf.conf file

- Add the following properties and values:

```
spark.eventLog.dir to hdfs:///spark-history
spark.eventLog.enabled to true
spark.history.fs.logDirectory to hdfs:///spark-history
spark.history.provider to org.apache.spark.deploy.history.
FsHistoryProvider
```

3. Restart Spark on YARN in either yarn-cluster mode or yarn-client mode:

- **yarn-cluster mode:** `./usr/hdp/current/spark-client/bin/spark-submit --class path.to.your.Class --master yarn-cluster [options] [app options]`
- **yarn-client mode:** `./usr/hdp/current/spark-client/bin/spark-shell --master yarn-client`

4. Validate the Spark installation. As user spark, run SparkPI example:

- a. `sudo su spark`
 - b. `cd /usr/hdp/current/spark-client`
 - c. `./bin/run-example SparkPi 10`

3.28. Upgrade Apache Slider

To upgrade Apache Slider, simply upgrade the Slider client.

1. Upgrade Slider client:

```
hdp-select set slider-client 2.5.0.0-<version>
```

```
slider version
```

3.29. Upgrade Apache Kafka

To upgrade Apache Kafka 2.5.0, you must stop all of the Kafka brokers before you start the upgrade.

1. Shut down the current Kafka daemon, switch to the new version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"
hdp-select set kafka-broker 2.5.0.0-1100
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.
3. If the upgrade process fails, follow the steps in "Downgrading Kafka" to return to your previous version of Kafka.

3.29.1. Downgrading Kafka

Downgrade each Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the previous version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"
hdp-select set kafka-broker 2.2.9.0-3393
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.

3.30. Finalize the Upgrade



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.

1. Verify your file system health before finalizing the upgrade. After you finalize an upgrade, your backup is discarded!
2. As the `$HDFS_USER`, enter:

```
su - hdfs -c "hdfs dfsadmin -finalizeUpgrade"
```


3.31. Migrate the Audit Logs from DB to Solr

Beginning with HDP 2.5.0, audit log to DB support has been removed. If your logs were previously stored on DB, you can migrate the logs to Solr. Refer to [Migrating Audit Logs from DB to Solr](#).

3.32. Install New HDP 2.5.0 Services

1. Install new HDP 2.5.0 Services (see the [Non-Ambari Cluster Installation Guide](#)):
 - Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
 - SmartSense – a next generation subscription model that features upgrade and configuration recommendations.
2. As Ranger has upgraded to a new version, you can Tag-sync, using the following Tagsync > install.properties, if required.

Table 3.13. Ranger Tagsync_install.properties and values

Property Name	Property Value
TAG_DEST_RANGER_ENDPOINT	http://<hostname>:6080
If TAG_SOURCE_ATLAS_ENABLED=true	
TAG_SOURCE_ATLAS_KAFKA_BOOTSTRAP_SERVERS	<hostname>:6667
TAG_SOURCE_ATLAS_KAFKA_ZOOKEEPER_CONNECT	<hostname>:2181
TAG_SOURCE_ATLAS_KAFKA_ENTITIES_GROUP_ID	ranger_entities_consumer
TAG_SOURCE_ATLAS_KAFKA_SERVICE_NAME	kafka
TAG_SOURCE_ATLAS_KAFKA_SECURITY_PROTOCOL	PLAINTEXTSASL
If TAG_SOURCE_ATLASREST_ENABLED = true	
TAG_SOURCE_ATLASREST_ENDPOINT	http://<hostname>:21000
TAG_SOURCE_ATLASREST_DOWNLOAD_INTERVAL_IN_MINUTES	60000
TAG_SOURCE_ATLASREST_USERNAME	
TAG_SOURCE_ATLASREST_PASSWORD	
If TAG_SOURCE_FILE_ENABLED = true	
TAG_SOURCE_FILE_ENABLED	true
TAG_SOURCE_FILE_FILENAME	/etc/ranger/data/tags.json (where the tag-sync is installed)

4. Upgrade from HDP 2.1 to HDP 2.5.0 Manually

This chapter provides instructions on how to manually upgrade to HDP 2.5.0 from the HDP 2.1.0 release. It assumes the existing HDP 2.1 was also installed manually.



Important

If you installed and manage HDP 2.1 with Ambari, **you must use the [Ambari Upgrade Guide](#)** to perform the the HDP 2.3 to HDP 2.5.0 upgrade.

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

The HDP packages for a complete installation of HDP 2.5.0 consumes about 6.5 GB of disk space.

Starting with HDP-2.2, HDP supports side-by-side installation of HDP 2.2 and subsequent releases, which lets you perform rolling upgrades on your cluster and improve execution times for in-place upgrade. To support side-by-side installation, the HDP package version naming convention for both RPMs and Debs has changed to include the HDP product version. For example, `hadoop-hdfs` in HDP-2.2.4 is now `hadoop_2.2.4.0.hdfs`. HDP-2.2 marked the first release where HDP rpms, debs, and directories contained versions in the names to permit side-by-side installations of later HDP releases. To select from the releases you have installed side-by-side, Hortonworks provides `hdp-select`, a command that lets you select the active version of HDP from the versions you have installed.

However, because HDP-2.1 did not support side-by-side installation, you upgrade to HDP-2.3.0 in a way very similar to previous minor-version upgrades. Subsequent upgrades after 2.5.0 is easier, and if you choose to add Ambari to your cluster, you can able to use the Rolling Upgrade feature.



Important

You cannot perform a side-by-side rolling upgrade from HDP 2.1 to HDP-2.5.0; only upgrade in place is supported. Side-by-side and rolling upgrade support starts with releases HDP 2.2 and above.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.



Important

You are strongly encouraged to read completely through this entire document before starting the Manual Upgrade process, in order to understand the interdependencies and the order of the steps.

The following steps are required to upgrade from HDP-2.1 to the latest release of HDP-2.5.0:

1. Get Ready to Upgrade
2. Upgrade HDP 2.1 Components
3. Symlink Directories with hdp-select
4. Configure and Start Apache ZooKeeper
5. Configure and Start Hadoop
6. Start HDFS
7. Configure and Start Apache HBase
8. Configure and Start Apache Phoenix
9. Configure and Start Apache Accumulo
10. Configure and Start Apache Tez
11. Configure and Start Apache Hive and Apache HCatalog
12. Configure and Start Apache Oozie
13. Configure and Start Apache WebHCat (Templeton)
14. Configure and Start Apache Pig
15. Configure and Start Apache Sqoop
16. Configure and Start Apache Flume
17. Configure and Validate Apache Mahout
18. Configure and Start Hue
19. Configure and Start Apache Knox
20. Configure and Start Apache Falcon
21. Finalize the Upgrade
22. Install new HDP-2.5.0 services if desired, such as Ranger or Spark

4.1. Getting Ready to Upgrade

HDP Stack upgrade involves upgrading from HDP 2.1 to HDP-2.5.0 versions and adding the new HDP-2.5.0 services. These instructions change your configurations.

**Note**

You must use **kinit** before running the commands as any particular user.

Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP-2.5.0 consumes about 6.5 GB of disk space.

The first step is to ensure you keep a backup copy of your HDP 2.1 configurations.

**Note**

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Back up the HDP directories for any hadoop components you have installed.

The following is a list of all HDP directories:

- `/etc/hadoop/conf`
- `/etc/hbase/conf`
- `/etc/hive-hcatalog/conf`
- `/etc/hive-webhcat/conf`
- `/etc/accumulo/conf`
- `/etc/phoenix/conf`
- `/etc/hive/conf`
- `/etc/pig/conf`
- `/etc/sqoop/conf`
- `/etc/flume/conf`
- `/etc/mahout/conf`
- `/etc/oozie/conf`
- `/etc/hue/conf`
- `/etc/spark/conf`
- `/etc/storm/conf`

- `/etc/storm-slider/conf`

- /etc/zookeeper/conf
 - /etc/tez/conf
 - /etc/storm/conf
 - **Optional** - Back up your userlogs directories, `${mapred.local.dir}/userlogs`.
2. Navigate to the `$HIVE_HOME/lib` directory. Backup the JDBC jar file for the type of Hive metastore you are using (Postgre, MySQL etc).
 3. Run the `fsck` command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.)

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-old-fsck-1.log"
```

4. Use the following instructions to compare status before and after the upgrade.

The following commands must be executed by the user running the HDFS service (by default, the user is hdfs).

- a. Capture the complete namespace of the file system. (The following command does a recursive listing of the root file system.)



Important

Make sure the NameNode is started.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-old-lsr-1.log"
```



Note

In secure mode you must have Kerberos credentials for the hdfs user.

- b. Run the report command to create a list of DataNodes in the cluster.

```
su - "hdfs dfsadmin -c "-report > dfs-old-report-1.log"
```
 - c. **Optional:** You can copy all or unrecoverable only data storelibext-customer directory in HDFS to a local file system or to a backup instance of HDFS.
 - d. **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.
5. Save the namespace by executing the following commands:

```
su - hdfs  
  
hdfs dfsadmin -safemode enter  
  
hdfs dfsadmin -saveNamespace
```

6. Backup your NameNode metadata.

- a. Copy the following checkpoint files into a backup directory:

The NameNode metadata is stored in a directory specified in the `hdfs-site.xml` configuration file under the configuration value `"dfs.namenode.dir"`.

For example, if the configuration value is:

```
<property>
  <name>dfs.namenode.name.dir</name>
  <value>/hadoop/hdfs/namenode</value>
</property>
```

Then, the NameNode metadata files are all housed inside the directory `/hadoop.hdfs/namenode`.

- b. Store the layoutVersion of the namenode.

```
${dfs.namenode.name.dir}/current/VERSION
```

7. Finalize any prior HDFS upgrade, if you have not done so already.

```
su - hdfs -c "hdfs dfsadmin -finalizeUpgrade"
```

8. If you have the Hive component installed, back up the Hive Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, see your database documentation.

Table 4.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hive < /tmp/mydir/backup_hive.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hive < /tmp/mydir/backup_hive.sql</pre>
Oracle	<p>Export the database:</p> <pre>exp username/password@database full=yes file=output_file.dmp</pre>	<p>Import the database:</p> <pre>imp username/password@database file=input_file.dmp</pre>

9. If you have the Oozie component installed, back up the Oozie metastore database.

These instructions are provided for your convenience. Please check your database documentation for the latest backup instructions.

Table 4.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sql</pre> <p>For example:</p> <pre>mysqldump oozie > /tmp/mydir/backup_oozie.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sql</pre> <p>For example:</p> <pre>mysql oozie < /tmp/mydir/backup_oozie.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump oozie > /tmp/mydir/backup_oozie.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres psql oozie < /tmp/mydir/backup_oozie.sql</pre>
Oracle	<p>Export the database:</p> <pre>exp username/password@database full=yes file=output_file.dmp</pre>	<p>Import the database:</p> <pre>imp username/password@database file=input_file.dmp</pre>

10.Optional: Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 4.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hue > /tmp/mydir/backup_hue.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hue < /tmp/mydir/backup_hue.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hue > /tmp/mydir/backup_hue.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hue < /tmp/mydir/backup_hue.sql</pre>
Oracle	<p>Connect to the Oracle database using sqlplus. Export the database.</p> <p>For example:</p> <pre>exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr</pre>	<p>Import the database:</p> <p>For example:</p> <pre>imp username/password@database file=input_file.dmp</pre>
SQLite	<pre>/etc/init.d/hue stop</pre> <pre>su \$HUE_USER</pre> <pre>mkdir ~/hue_backup</pre>	<pre>/etc/init.d/hue stop</pre> <pre>cd /var/lib/hue</pre> <pre>mv desktop.db desktop.db.old</pre>

Database Type	Backup	Restore
	sqlite3 desktop.db .dump > ~/hue_backup/desktop.bak /etc/init.d/hue start	sqlite3 desktop.db < ~/hue_backup/desktop.bak /etc/init.d/hue start

11 Stop all services (including MapReduce) and client applications deployed on HDFS:

Component	Command
Knox	cd /usr/lib/knox/ su knox -c "bin/gateway.sh stop"
Oozie	su \$OOZIE_USER /usr/lib/oozie/bin/oozied.sh stop
WebHCat	su - hcat -c "/usr/lib/hive-hcatalog/sbin/webhcat_server.sh stop"
Hive	Run this command on the Hive Metastore and Hive Server2 host machine: ps aux awk '{print \$1,\$2}' grep hive awk '{print \$2}' xargs kill >/dev/null 2>&1
HBase RegionServers	su - hbase -c "/usr/lib/hbase/bin/hbase-daemon.sh --config /etc/hbase/conf stop regionserver"
HBase Master host machine	su - hbase -c "/usr/lib/hbase/bin/hbase-daemon.sh --config /etc/hbase/conf stop master"
YARN	Run this command on all NodeManagers: su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-yarn/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop nodemanager" Run this command on the History Server host machine: su - mapred -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-mapreduce/sbin/mr-jobhistory-daemon.sh --config /etc/hadoop/conf stop historyserver" Run this command on the ResourceManager host machine(s): su - yarn -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-yarn/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop resourcemanager" Run this command on the YARN Timeline Server node: su -l yarn -c "export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec && /usr/lib/hadoop-yarn/sbin/yarn-daemon.sh --config /etc/hadoop/conf stop timelineserver"
HDFS	On all DataNodes:

Component	Command
	<pre>su - hdfs -c "/usr/lib/hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode"</pre> <p>If you are not running a highly available HDFS cluster, stop the Secondary NameNode by executing this command on the Secondary NameNode host machine:</p> <pre>su - hdfs -c "/usr/lib/hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop secondarynamenode"</pre> <p>On the NameNode host machine(s):</p> <pre>su - hdfs -c "/usr/lib/hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop namenode"</pre> <p>If you are running NameNode HA, stop the ZooKeeper Failover Controllers (ZKFC) by executing this command on the NameNode host machine:</p> <pre>su - hdfs -c "/usr/lib/hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop zkfc"</pre> <p>If you are running NameNode HA, stop the JournalNodes by executing these commands on the JournalNode host machines:</p> <pre>su - hdfs -c "/usr/lib/hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf stop journalnode"</pre>
ZooKeeper Host machines	<pre>su - zookeeper -c "export ZOOCFGDIR=/etc/zookeeper/conf ; export ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-env.sh ; /usr/lib/zookeeper/bin/zkServer.sh stop"</pre>
Ranger (XA Secure)	<pre>service xapolicymgr stop</pre> <pre>service uxugsync stop</pre>

12. Verify that edit logs in `${dfs.namenode.name.dir}/current/edits*` are empty.

- Run: `hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out`
- Verify `edits.out` file. It should only have `OP_START_LOG_SEGMENT` transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

- If `edits.out` has transactions other than `OP_START_LOG_SEGMENT`, run the following steps and then verify edit logs are empty.

- Start the existing version NameNode.
- Ensure there is a new FS image file.
- Shut the NameNode down:

```
hdfs dfsadmin - saveNamespace
```

13. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path
component in this version of HDFS.

Please rollback and delete or rename this path, or upgrade with the
-renameReserved key-value pairs option to automatically rename these
paths during upgrade.
```

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup.

For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, specify `-upgrade -renameReserved .snapshot=.my-snapshot, .reserved=.my-reserved`.

If no key-value pairs are specified with `-renameReserved`, the NameNode then suffixes reserved paths with:

```
.<LAYOUT-VERSION>.UPGRADE_RENAMED
```

For example: `.snapshot.-51.UPGRADE_RENAMED`.



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

14. Upgrade the JDK on all nodes to JDK 7 or JDK 8 before upgrading HDP.

4.2. Upgrade HDP 2.1 Components

The upgrade process to HDP-2.5.0 involves the following steps. See the HDP-2.5.0 [Release Notes](#) for repo information.

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.1 components. This command uninstalls the HDP 2.1 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
yum list installed | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP-2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<version>/2.x/
updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hue" "hive"
"tez" "storm" "kafka" "flume" "phoenix" "accumulo" "mahout" "knox"
"hdp_mon_nagios_addons"
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

RHEL/CentOS/Oracle 5

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.1 components. This command uninstalls the HDP 2.1 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
yum list installed | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

5. Install the HDP-2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<versions>/2.x/
updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo namestatus
HDP-2.5.0 Hortonworks Data Platform Version - HDP-2.5.0
```

6. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "hadoop-lzo" "openssl" "hive-webhcat" "hive-
hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase"
"hue" "hive" "tez" "storm" "falcon" "flume" "phoenix" "accumulo" "mahout"
"knox" "kafka" "spark" "ranger" "slider" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.1 components. This command uninstalls the HDP 2.1 components. It leaves the user data, and metadata, but removes your configurations:

```
zypper rm "hadoop*" "webhcat*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*"
"sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*"
"phoenix*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell" "knox*"
"hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
yum list installed | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

5. Download the HDP-2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.5.
0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog"
"oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive"
"hue" "tez" "storm" "flume" "phoenix" "accumulo" "mahout" "knox"
"kafka" "spark" "slider*" "hdp_mon_nagios_addons" "hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5
```

```
zypper install oozie-client
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep
hcatalog
```

No component files names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.1 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*"
"pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*"
"falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common"
"hue-shell" "knox*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
zypper search -i | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.list
```

5. Download the HDP-2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/suselisp3/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

6. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
zypper install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn"
"hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hue" "hive"
"tez" "storm" "kafka" "flume" "phoenix" "accumulo" "mahout" "knox"
"hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5
```

```
zypper install oozie-client
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

7. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep hcatalog
```

No component files names should appear in the returned list.

Ubuntu 12

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.1 components. This command uninstalls the HDP 2.1 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez.*"
"storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue.*"
"knox*" "hdp_mon_nagios_addons" --purge
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
yum list installed | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP-2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2.5.0.0/hdp.list -O /etc/apt/sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm" "kafka" "flume" "phoenix" "accumulo" "mahout" "knox" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

Debian

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.1 components. This command uninstalls the HDP 2.1 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*" "tez*" "storm*" "falcon*" "flume*" "phoenix*" "accumulo*" "mahout*" "hue*" "knox*" "hdp_mon_nagios_addons"
```

3. Validate that all HDP 2.1 component binaries are uninstalled:

```
yum list installed | grep @HDP2.1
```

4. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

5. Download the HDP-2.5.0 hdp.repo file:


```
wget -nv http://public-repo-1.hortonworks.com/HDP/
debian<version>/2.x/updates/2.5.0.0/hdp.list - O /etc/apt/
sources.list.d/hdp.list
```

6. Run an update:

```
apt-get update
```

7. Install the HDP-2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:



Important

For Falcon, we recommend upgrading from HDP 2.1 to HDP 2.2 and then upgrading HDP 2.2 to HDP 2.5. See [Configure and Validate Falcon](#) for more information.

```
apt-get install "hadoop" "hadoop-hdfs" "libhdfs0" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "tez" "storm" "flume"
"phoenix" "accumulo" "mahout" "knox" "hdp_mon_nagios_addons"
```



Note

If you installed Apache Argus, it is now Apache Ranger. See [Upgrade Ranger](#) for more information on the upgrade path.

Install the Compression Libraries. See [Install Compression Libraries](#).

4.3. Symlink Directories with hdp-select



Warning

HDP-2.5.0 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

1. Before you run hdp-select, remove one link:

```
rm /usr/bin/oozie
```

2. Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5-2800
```

4.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available HDFS cluster, configure and restart ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.

1. Replace your configuration after upgrading on all the Apache ZooKeeper nodes. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "export ZOOCFGDIR=/usr/hdp/current/zookeeper-server/
conf ;
export ZOOCFG=zoo.cfg; source /usr/hdp/current/zookeeper-server/conf/
zookeeper-env.sh ;
/usr/hdp/current/zookeeper-server/bin/zkServer.sh start"
```

4.5. Configure Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed in HDP-2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`.
2. Paths have changed in HDP-2.5.0. Make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

Ubuntu/Debian

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP-2.5.0. make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

4.6. Start Hadoop Core



Warning

Before you start HDFS on a highly available HDFS cluster, you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

To start HDFS, run commands as the `$HDFS_USER`.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading on all the HDFS nodes. Replace the HDFS template configuration in `/etc/hdfs/conf`.
2. If you are upgrading from a highly available HDFS cluster configuration, start all JournalNodes. On each JournalNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-  
daemon.sh start journalnode"
```



Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

3. If you are running HDFS on a highly available namenode, you must first start the Zookeeper service.



Note

Perform this step only if you are on a highly available HDFS cluster.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-  
daemon.sh start zkfc"
```

4. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually.

On the active NameNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the `-upgrade` option only once. After you have completed this step, you can bring up the NameNode using this command without including the `-upgrade` option.

To check if the Upgrade is in progress, check that the `"\previous"` directory has been created in the `\NameNode` and `\JournalNode` directories. The `"\previous"` directory contains a snapshot of the data before upgrade.

In a highly available HDFS cluster configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode.

To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the `'-bootstrapStandby'` flag. Do NOT start this standby NameNode with the `'-upgrade'` flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The `bootstrapStandby` command downloads the most recent fsimage from the active NameNode into the `$dfs.name.dir` directory of the standby NameNode. You can enter that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

5. Verify that the NameNode is up and running:

```
ps -ef | grep -i NameNode
```

6. If you do not have a highly available HDFS cluster configuration (non_HA namenode), start the Secondary NameNode.



Note

Do not perform this step if you have a highly available HDFS cluster configuration.

On the Secondary NameNode host machine, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

7. Verify that the Secondary NameNode is up and running.



Note

Do not perform this step if you have a highly available HDFS cluster environment.

```
ps -ef|grep SecondaryNameNode
```

8. Start DataNodes.

On each of the DataNodes, enter the following command. Note: If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start datanode"
```

9. Verify that the DataNode process is up and running:

```
ps -ef|grep DataNode
```

10. Verify that NameNode can go out of safe mode.

```
>su - hdfs -c "hdfs dfsadmin -safemode wait"
```

You should see the following result: Safe mode is OFF

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

4.7. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

1. Run the `fsck` command on namenode as \$HDFS_USER:

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-new-fsck-1.log"
```

You should see feedback that the filesystem under path / is HEALTHY.

2. Run `hdfs namespace` and report.

a. List directories.

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-new-lsr-1.log"
```

b. Open the `dfs-new-lsr-1.log` and confirm that you can see the file and directory listing in the namespace.

c. Run report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-new-report-1.log"
```

d. Open the `dfs-new-report` file and validate the admin report.

3. Compare the namespace report before the upgrade and after the upgrade. Verify that user files exist after upgrade.

The file names are listed below:

```
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log
```

```
dfs-old-report-1.log < -- > dfs-new-report-1.log
```



Note

You must do this comparison manually to catch all errors.

4. From the NameNode WebUI, determine if all DataNodes are up and running.

```
http://<namenode>:<namenodeport>
```

5. If you are on a highly available HDFS cluster, go to the StandbyNameNode web UI to see if all DataNodes are up and running:

```
http://<standbynamenode>:<namenodeport>
```

6. If you are **not** on a highly available HDFS cluster, go to the SecondaryNameNode web UI to see if the secondary node is up and running:

```
http://<secondarynamenode>:<secondarynamenodeport>
```

7. Verify that read and write to `hdfs` works successfully.

```
hdfs dfs -put [input file] [output file]
```

```
hdfs dfs -cat [output file]
```

4.8. Configure YARN and MapReduce

After you upgrade Hadoop, complete the following steps to update your configs.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

1. Ensure that all HDFS directories configured in `yarn-site.xml` configuration files (for example, `yarn.timeline-service.entity-group-fs-store.active-dir`, `yarn.timeline-service.entity-group-fs-store.done-dir`) in HDFS.

2. Upload the MapReduce tarball to HDFS. As the HDFS user, for example 'hdfs':

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/mapreduce.tar.gz /hdp/apps/2.5.0.0-<$version>/mapreduce/"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

```
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/mapreduce"
```

```
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/mapreduce/mapreduce.tar.gz"
```

3. Make the following changes to `/etc/hadoop/conf/mapred-site.xml`:

- Add:

```
<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}
    /mapreduce/mapreduce.tar.gz#mr-framework
  </value>
</property>

<property>
  <name>yarn.app.mapreduce.am.admin-command-opts</name>
  <value>Dhdp.version=${hdp.version}</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
```

```

<value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}
/hadoop/lib/native:/usr/hdp/${hdp.version}/hadoop/
lib/native/Linux-amd64-64
</value>
</property>

<property>
<name>mapreduce.admin.map.child.java.opts</name>
<value>-server -Djava.net.preferIPv4Stack=true
-Dhdp.version=${hdp.version}
</value>
<final>true</final>
</property>

<property>
<name>mapreduce.admin.reduce.child.java.opts</name>
<value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
<final>true</final>
</property>

<property>
<name>mapreduce.application.classpath</name>
<value>${PWD}/mr-framework/hadoop/share/hadoop/mapreduce/*,
${PWD}/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
${PWD}/mr-framework/hadoop/share/hadoop/common/*,
${PWD}/mr-framework/hadoop/share/hadoop/common/lib/*,
${PWD}/mr-framework/hadoop/share/hadoop/yarn/*,
${PWD}/mr-framework/hadoop/share/hadoop/yarn/lib/*,
${PWD}/mr-framework/hadoop/share/hadoop/hdfs/*,
${PWD}/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
/usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
/etc/hadoop/conf/secure</value>
</property>

```



Note

You do not need to modify `${hdp.version}`.

- Remove the following properties from `/etc/hadoop/conf/mapred-site.xml`: `mapreduce.task.tmp.dir`, `mapreduce.job.speculative.slownodethreshold` (deprecated), and `mapreduce.job.speculative.speculativecap` (deprecated).

4. Add the following properties to `/etc/hadoop/conf/yarn-site.xml`:

```

<property>
<name>yarn.application.classpath</name>
<value>${HADOOP_CONF_DIR},/usr/hdp/${hdp.version}/hadoop-client/*,
/usr/hdp/${hdp.version}/hadoop-client/lib/*,
/usr/hdp/${hdp.version}/hadoop-hdfs-client/*,
/usr/hdp/${hdp.version}/hadoop-hdfs-client/lib/*,
/usr/hdp/${hdp.version}/hadoop-yarn-client/*,
/usr/hdp/${hdp.version}/hadoop-yarn-client/lib/*</value>
</property>

```

5. On secure clusters only, add the following properties to `/etc/hadoop/conf/yarn-site.xml`:


```

<property>
  <name>yarn.timeline-service.recovery.enabled:</name>
  <value>TRUE</value>
</property>

  <property>
    <name>yarn.timeline-service.state-store.class: org.apache.hadoop.yarn.
server.timeline.recovery:</name>
    <value>LevelDbTimelineStateStore</value>
  </property>

  <property>
    <name>yarn.timeline-service.leveldb-state-store.path:</name>
    <value><the same as the default of "yarn.timeline-service-leveldb-
timeline-store.path</value>
  </property>

```

6. Modify the following property to `/etc/hadoop/conf/yarn-site.xml`:

```

<property>
  <name>mapreduce.application.classpath</name>
  <value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
    $PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/common/*,
    $PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/yarn/*,
    $PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
    $PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/share/hadoop/tools/lib/*,
    /usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
    /etc/hadoop/conf/secure</value>
</property>

```

7. Make the following change to the `/etc/hadoop/conf/yarn-env.sh`:

Change `export HADOOP_YARN_HOME=/usr/lib/hadoop-yarn`

to

```
export HADOOP_YARN_HOME=/usr/hdp/current/hadoop-yarn-
nodemanager/
```

8. Make the following change to the `/etc/hadoop/conf/yarn-env.sh`:

Change

```
export HADOOP_LIBEXEC_DIR=/usr/lib/hadoop/libexec
```

to

```
HADOOP_LIBEXEC_DIR=/usr/hdp/current/hadoop-client/libexec/
```

9. For secure clusters, you must create and configure the `container-executor.cfg` configuration file:

- Create the `container-executor.cfg` file in `/etc/hadoop/conf/`.

- Insert the following properties:

```
yarn.nodemanager.linux-container-executor.group=hadoop
banned.users=hdfs,yarn,mapred
min.user.id=1000
```

- `yarn.nodemanager.linux-container-executor.group` - Configured value of `yarn.nodemanager.linux-container-executor.group`. This must match the value of `yarn.nodemanager.linux-container-executor.group` in `yarn-site.xml`.
- `banned.users` - Comma-separated list of users who can not run container-executor.
- `min.user.id` - Minimum value of user id. This prevents system users from running container-executor.
- `allowed.system.users` - Comma-separated list of allowed system users.
- Set the file `/etc/hadoop/conf/container-executor.cfg` file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/bin/container-executor
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/bin/container-executor
```

4.9. Start YARN/MapReduce Services

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.



Note

The `su` commands in this section use "yarn" to represent the YARN Service user and `mapreduce` to represent the MAPREDUCE Service user. If you are using another name for these Service users, you need to substitute your Service user name for "yarn" or "mapreduce" in each of the `su` commands.

1. If you have a highly available HDFS cluster configuration, manually clear the ResourceManager state store.

```
su - yarn -c "yarn resourcemanager -format-state-store"
```

2. Start the ResourceManager on all your ResourceManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh start resourcemanager"
ps -ef | grep -i resourcemanager
```

3. Start the TimelineServer on your TimelineServer host.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-timeline-servers/sbin/yarn-daemon.  
sh start timelineserver"  
ps -ef | grep -i timelineserver
```

4. Start the NodeManager on all your NodeManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh  
start nodemanager"  
ps -ef | grep -i nodemanager
```

5. To start MapReduce, run the following commands:

```
su - mapreduce -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-  
jobhistory-daemon.sh start historyserver"  
ps -ef | grep -i jobhistoryserver
```

4.10. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user.

The job uses MapReduce to write 100MB of data into HDFS with RandomWriter

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.  
jar  
randomwriter -Dtest.randomwrite.total_bytes=10000000 test-after-  
upgrade.
```

You should see messages similar to:

```
map 0% reduce 0%  
...map 100% reduce 100%  
Job ... completed successfully
```

MapReduce upgraded successfully. You can now upgrade your other components.

Basic Troubleshooting

To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

`http://<resource manager host>:8088/cluster/nodes`

The number of active nodes should be equal to the number of nodemanagers.

Accessing error messages:

1. Access the ApplicationMaster WebUI to view the container logs.
2. At your console logs for MapReduce job, look for a line with this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track the job: http://<resource  
manager host>:8088/proxy/application_1380673658357_0007/
```
3. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

4.11. Configure and Start Apache HBase



Note

The `su` commands in this section use "hbase" to represent the HBASE Service user. If you are using another name for your HBASE Service user, you need to substitute your HBASE Service user name for "hbase" in each of the `su` commands.

The `hbase.bucketcache.percentage.in.combinedcache` is removed in HDP-2.5.0. This simplifies the configuration of block cache. BucketCache configurations from HDP 2.1 need to be recalculated to attain identical memory allotments in HDP-2.5.0. The L1 LruBlockCache be whatever `hfile.block.cache.size` is set to and the L2 BucketCache is whatever `hbase.bucketcache.size` is set to.

1. Upgrade Apache HBase.

```
hbase upgrade -execute
```

2. Replace your configuration after upgrading. Replace the HBase template configuration in `/etc/hbase/conf`.

3. Start services. From root, assuming that `$HBASE_USER=hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

4. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

4.12. Configure Apache Phoenix

To configure Apache Phoenix, complete the following steps:

1. Add the following property to the `/etc/hbase/hbase-site.xml` file on all HBase nodes, the MasterServer, and all RegionServers to prevent deadlocks from occurring during maintenance on global indexes:

```
<property>
  <name>hbase.regionserver.wal.codec</name>
  <value>org.apache.hadoop.hbase.regionserver.wal.IndexedWALEditCodec</
value>
</property>
```

2. To enable user-defined functions, configure the following property in `/etc/hbase/conf` on all Hbase nodes.

```
<property>
  <name>phoenix.functions.allowUserDefinedFunctions</name>
  <value>true</value>
  <description>enable UDF functions</description>
</property>
```

3. Ensure the client side hbase-site.xml matches the server side configuration.
4. **(Optional)** To use local indexing, set the following three properties. These properties ensure collocation of data table and local index regions:



Note

The local indexing feature is a technical preview and considered under development. Do not use this feature in your production systems. If you have questions regarding this feature, contact Support by logging a case on our Hortonworks Support Portal at <http://hortonworks.com/support/>.

Add the following two properties, if they don't already exist, to the master side configurations:

```
<property>
  <name>hbase.master.loadbalancer.class</name>
  <value>org.apache.phoenix.hbase.index.balancer.IndexLoadBalancer</value>
</property>

<property>
  <name>hbase.coprocessor.master.classes</name>
  <value>org.apache.phoenix.hbase.index.master.IndexMasterObserver</value>
</property>
```

Add the following, if it doesn't already exist, to the Region server side configurations:

```
<property>
  <name>hbase.coprocessor.regionserver.classes</name>
  <value>org.apache.hadoop.hbase.regionserver.LocalIndexMerger</value>
</property>
```

5. If the folder specified in hbase.tmp.dir property on hbase-site.xml does not exist, create that directory with adequate permissions.
6. Set the following property in the hbase-site.xml file for all RegionServers, but not on the client side:

```
<property>
  <name>hbase.rpc.controllerfactory.class</name>
  <value>org.apache.hadoop.hbase.ipc.controller.ServerRpcControllerFactory</value>
</property>
```

7. Restart the HBase Master and Region Servers.

Configuring Phoenix to Run in a Secure Cluster

Perform the following additional steps to configure Phoenix to run in a secure Hadoop cluster:

1. To link the HBase configuration file with the Phoenix libraries:

```
ln -sf HBASE_CONFIG_DIR/hbase-site.xml PHOENIX_HOME/bin/hbase-site.xml
```

2. To link the Hadoop configuration file with the Phoenix libraries:

```
ln -sf HADOOP_CONFIG_DIR/core-site.xml PHOENIX_HOME/bin/core-site.xml
```



Note

When running the `psql.py` and `sqlline.py` Phoenix scripts in secure mode, you can safely ignore the following warnings.

```
14/04/19 00:56:24 WARN util.NativeCodeLoader:
Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable

14/04/19 00:56:24 WARN util.DynamicClassLoader: Failed to identify the fs of
dir hdfs://<HOSTNAME>:8020/apps/hbase/data/lib, ignored java.io.IOException:
No FileSystem for scheme: hdfs
```

4.13. Configure and Start Apache Accumulo



Note

The `su` commands in this section use "accumulo" to represent the Accumulo Service user. If you are using another name for your Apache Accumulo Service user, you need to substitute your Accumulo Service user name for "accumulo" in each of the `su` commands.

Upon upgrade from HDP 2.1 to HDP-2.5.0, Accumulo automatically changes the HDFS and Zookeeper data stored on Accumulo. This change is not backward compatible, and HDP 2.1 does not run on this updated data.

After upgrading from HDP 2.1 to HDP-2.5.0, Accumulo automatically upgrades the internal metadata, notably the data in ZooKeeper, when the Accumulo Master for HDP-2.5.0 first starts. This change is not backward compatible and HDP 2.1 no longer runs against the data.

Prior to performing the following steps, you need to initialize Accumulo. See [Initialization](#).

1. You must replace your configuration after upgrading. Copy `/etc/accumulo/conf` from the template to the `conf` directory in Accumulo hosts.
2. In HDP-2.5.0, the `instance.dfs.dir` and `instance.dfs.uri` properties are deprecated with the `instance.volumes` property. If it does not already exist, add the `instance.volumes` property to the `accumulo-site.xml` file. Do **not** remove the `instance.dfs.dir` and `instance.dfs.uri` properties. You can extrapolate the value for the `instance.volumes` property from the `instance.dfs.dir` and `instance.dfs.uri` properties.

For example:

```
<property>
```

```

    <name>instance.dfs.dir</name>
    <value>/accumulo</value>
  </property>
</property>
    <name>instance.dfs.uri</name>
    <value>hdfs://my_namenode:8020</value>
  </property>
</property>
    <name>instance.volumes</name>
    <value>hdfs://my_namenode:8020/accumulo</value>
  </property>

```

3. Start the services:

```

su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh
`hostname` master"

su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh
`hostname` tserver"

su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh
`hostname` gc"

su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh
`hostname` tracer"

su - accumulo -c "/usr/hdp/current/accumulo-master/bin/start-server.sh
`hostname` monitor"

```

4. Check that the processes are running

```
ps -ef | grep accumulo
```

or visit <http://<hostname>:50095> in your browser

4.14. Configure and Start Apache Tez



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.



Important

If you have a secure server, you need Kerberos credentials for hdfs user access.

To upgrade Apache Tez:

1. Copy your previously backed-up copy of `tez-site.xml` into the `/etc/tez/conf` directory.
2. Check to see if the HDP-2.5.0 Tez tarball libraries are already in the `/hdp/apps/<hdp 2.5>/tez/ version>` directory - if so, then skip this step. If not, put the HDP-2.5.0 Tez tarball libraries in the `/hdp/apps` directory in HDFS, so that submitted Tez applications to this cluster have access to these shared Tez tarball libraries. Execute this step on any host that has the Tez client installed.

```

su - hdfs
hdfs dfs -mkdir -p /hdp/apps/<hdp_version>/tez/
hdfs dfs -put /usr/hdp/<hdp_version>/tez/lib/tez.tar.gz /hdp/apps/
<hdp_version>/tez/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/<hdp_version>/tez
hdfs dfs -chmod -R 444 /hdp/apps/<hdp_version>/tez/tez.tar.gz

```

Where <hdp_version> is the current HDP version, for example 2.5.0-2800.

3. Edit the `tez.lib.uris` property in the `tez-site.xml` file to point to `/hdp/apps/<hdp_version>/tez/tez.tar.gz`

```

...
<property>
  <name>tez.lib.uris</name>
  <value>/hdp/apps/<hdp_version>/tez/tez.tar.gz</value>
</property>
...

```

Where <hdp_version> is the current HDP version, for example 2.5.0-2800.

4. **Optional** Earlier releases of Tez did not have access control. In the current version of Tez, the default behavior restricts the ability to view the Tez history to only the owner of the job. To retain unrestricted access for non-secure clusters, set `tez.am.view-acls` set to `"*"`.
5. Update the following `tez-site.xml` properties to their new names.

Old Property Name	New Property Name
<code>tez.am.java.opts</code>	<code>tez.am.launch.cmd-opts</code>
<code>tez.am.env</code>	<code>tez.am.launch.env</code>
<code>tez.am.shuffle-vertex-manager.min-src-fraction</code>	<code>tez.shuffle-vertex-manager.min-src-fraction</code>
<code>tez.am.shuffle-vertex-manager.max-src-fraction</code>	<code>tez.shuffle-vertex-manager.max-src-fraction</code>
<code>tez.am.shuffle-vertex-manager.enable.auto-parallel</code>	<code>tez.shuffle-vertex-manager.enable.auto-parallel</code>
<code>tez.am.shuffle-vertex-manager.desired-task-input-size</code>	<code>tez.shuffle-vertex-manager.desired-task-input-size</code>
<code>tez.am.shuffle-vertex-manager.min-task-parallelism</code>	<code>tez.shuffle-vertex-manager.min-task-parallelism</code>
<code>tez.am.grouping.split-count</code>	<code>tez.grouping.split-count</code>
<code>tez.am.grouping.by-length</code>	<code>tez.grouping.by-length</code>
<code>tez.am.grouping.by-count</code>	<code>tez.grouping.by-count</code>
<code>tez.am.grouping.max-size</code>	<code>tez.grouping.max-size</code>
<code>tez.am.grouping.min-size</code>	<code>tez.grouping.min-size</code>
<code>tez.am.grouping.rack-split-reduction</code>	<code>tez.grouping.rack-split-reduction</code>
<code>tez.am.am.complete.cancel.delegation.tokens</code>	<code>tez.cancel.delegation.tokens.on.completion</code>
<code>tez.am.max.task.attempts</code>	<code>tez.am.task.max.failed.attempts</code>
<code>tez.generate.dag.viz</code>	<code>tez.generate.debug.artifacts</code>
<code>tez.runtime.intermediate-output.key.comparator.class</code>	<code>tez.runtime.key.comparator.class</code>
<code>tez.runtime.intermediate-output.key.class</code>	<code>tez.runtime.key.class</code>
<code>tez.runtime.intermediate-output.value.class</code>	<code>tez.runtime.value.class</code>
<code>tez.runtime.intermediate-output.should-compress</code>	<code>tez.runtime.compress</code>
<code>tez.runtime.intermediate-output.compress.codec</code>	<code>tez.runtime.compress.codec</code>

Old Property Name	New Property Name
tez.runtime.intermediate-input.key.secondary.comparator.class	tez.runtime.key.secondary.comparator.class
tez.runtime.broadcast.data-via-events.enabled	tez.runtime.transfer.data-via-events.enabled
tez.runtime.broadcast.data-via-events.max-size	tez.runtime.transfer.data-via-events.max-size
tez.runtime.shuffle.input.buffer.percent	tez.runtime.shuffle.fetch.buffer.percent
tez.runtime.task.input.buffer.percent	tez.runtime.task.input.post-merge.buffer.percent
tez.runtime.job.counters.max	tez.am.counters.max.keys
tez.runtime.job.counters.group.name.max	tez.am.counters.group-name.max.keys
tez.runtime.job.counters.counter.name.max	tez.am.counters.name.max.keys
tez.runtime.job.counters.groups.max	tez.am.counters.groups.max.keys
tez.task.merge.progress.records	tez.runtime.merge.progress.records
tez.runtime.metrics.session.id	tez.runtime.framework.metrics.session.id
tez.task.scale.memory.additional.reservation.fraction.per-io	tez.task.scale.memory.additional-reservation.fraction.per-io
tez.task.scale.memory.additional.reservation.fraction.max	tez.task.scale.memory.additional-reservation.fraction.max
tez.task.initial.memory.scale.ratios	tez.task.scale.memory.ratios
tez.resource.calculator.process-tree.class	tez.task.resource.calculator.process-tree.class

For more information on setting Tez configuration parameters in HDP-2.5.0, see [Installing and Configuring Apache Tez](#) in the Non-Ambari Cluster Installation Guide.

6. Change the value of the `tez.tez-ui.history-url.base` property to the url for the upgraded Tez View. For more details, see [Using the Tez View](#) in the HDP Ambari Views Guide.

4.15. Configure and Start Apache Hive and Apache HCatalog



Note

The `su` commands in this section use "hive" to represent the Apache Hive Service user. If you are using another name for your Hive Service user, you need to substitute your Hive Service user name for "hive" in each of the `su` commands.

1. Prior to starting the upgrade process, set the following in your hive configuration file:

```
datanucleus.autoCreateSchema=false
```

2. Copy the jdbc connector jar from `OLD_HIVE_HOME/lib` to `CURRENT_HIVE_HOME/lib`.
3. Restore the JDBC jar files backed up into the `$HIVE_HOME/lib` directory. Make sure to restore all Metastore-related properties, (such as, `ConnectionURL`, `user` etc), from your older hive installation.
4. Upgrade the Hive Metastore database schema. Restart the Hive Metastore database and run:

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -upgradeSchema -dbType <$databaseType>"
```

The value for `$databaseType` can be `derby`, `mysql`, `oracle`, or `postgres`.



Important

When you use MySQL as your Hive metastore, you must use `mysql-connector-java-5.1.35.zip` or later JDBC driver.



Note

If you are using PostgreSQL, you should reset the Hive Metastore database owner to `<HIVE_USER>`:

```
sudo <POSTGRES_USER>
```

Start the PostgreSQL CLU using the `psql` command.

Execute: `ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER TO <HIVE_USER>`



Note

If you are using Oracle 11, you might see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
optimize.mapjoin.mapreduce does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
heapspace does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
server2.enable.impersonation does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
semantic.analyzer.factory.impl does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.
convert.sortmerge.join.noconditionaltask does not exist
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.
internal:1521/XE
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver
Metastore connection User: hiveuser
Starting upgrade metastore schema from version 0.13.0 to 0.14.0
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql
Error: ORA-00955: name is already used by an existing object
(state=42000,code=955)
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.
oracle.sql: Schema script failed, errorcode 2
Completed upgrade-0.13.0-to-0.14.0.oracle.sql
schemaTool completed
```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.



Note

Copy only the necessary configuration files. Do not copy the `env.sh` files, for example, `hadoop-env.sh`, `hive-env.sh`, and so forth. Additionally, all `env.sh` files must be properly configured.

5. Edit the hive-site.xml file and modify the properties based on your environment.

a. Edit the following properties in the hive-site.xml file:

```
<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache
  </description>
</property>
```

b. **Optional:** To enable the Hive builtin authorization mode, make the following changes. If you want to use the advanced authorization provided by Ranger, refer to the [Ranger](#) instructions.

Set the following Hive authorization parameters in the hive-site.xml file:

```
<property>
  <name>hive.server2.enable.doAs</name>
  <value>>false</value>
</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
    StorageBasedAuthorizationProvider,org.apache.hadoop.hive.ql.security.
    authorization.MetaStoreAuthzAPIAuthorizeEmbedOnly</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.plugin.sqlstd.
    SQLStdConfOnlyAuthorizeFactory</value>
</property>
```

Also set hive.users.in.admin.role to the list of comma-separated users who need to be added to admin role. A user who belongs to the admin role needs to run the "set role" command before getting the privileges of the admin role, as this role is not in the current roles by default.

Set the following in the hiveserver2-site.xml file.

```
<property>
  <name>hive.security.authenticator.manager</name>
  <value>org.apache.hadoop.hive.ql.security.
    SessionStateUserAuthenticator</value>
</property>

<property>
  <name>hive.security..authorization.enabled</name>
  <value>true</value>
</property>
```

```
<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive ql.security.authorization.plugin.sqlstd.
SQLStdHiveAuthorizeFactory</value>
</property>
```

- c. For a remote Hive metastore database, set the IP address (or fully-qualified domain name) and port of the metastore host using the following hive-site.xml property value.

```
<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server.
    To enable HiveServer2, leave the property value empty.
  </description>
</property>
```

You can further fine-tune your configuration settings based on node hardware specifications, using the HDP utility script.

6. Start Hive Metastore.

On the Hive Metastore host machine, run the following command:

```
su - hive -c "nohup /usr/hdp/current/hive-metastore/bin/hive
--service metastore -hiveconf hive.log.file=hivemetastore.log
>/var/log/hive/hivemetastore.out 2>/var/log/hive/
hivemetastoreerr.log &"
```

7. Start Hive Server2.

On the Hive Server2 host machine, run the following command:

```
su - hive

nohup /usr/hdp/current/hive-server2/bin/hiveserver2 -hiveconf
hive.metastore.uris=" " -hiveconf hive.log.file=hiveserver2.log
>/var/log/hive/hiveserver2.out 2> /var/log/hive/
hiveserver2err.log &
```

4.16. Configure and Start Apache Oozie



Note

The duration of the Oozie upgrade is dependent on the amount of job history stored in ooziedb. This history must be backed up and restored during the upgrade process. Best practice when planning for upgrade is to backup ooziedb from your production oozie server and restore it to a test or development oozie server. This can help you estimate the time that is required to upgrade Oozie during your production upgrade.



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and "oozie" to represent the Oozie Service user. If you are using another name for your HDFS Service user or your Oozie Service user, you need to substitute your Service user names for "hdfs" or "oozie" in each of the `su` commands.

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore `oozie-site.xml` from your backup to the `conf` directory on each oozie server and client.

2. Copy the JDBC jar to `libext-customer`:

- a. Create the `/usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer` directory.

```
cd /usr/hdp/2.5.0.0-<version>/oozie-server
```

```
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the `libext-customer` directory.

```
chmod -R 777 /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer
```

3. Copy these files to the `libext-customer` directory

```
cp /usr/hdp/2.5.0.0-<version>/hadoop-client/lib/hadoop*lzo*.jar /usr/hdp/current/oozie-server/libext-customer
```

```
cp /usr/share/HDP-oozie/ext.zip /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer/
```

Also, copy Oozie db jar in `libext-customer`.

4. If Falcon was also installed and configured before upgrade in HDP 2.2.x, then after upgrade you might also need to do the following:

```
cp /usr/hdp/current/falcon-server/oozie/ext/falcon-oozie-el-extension-*.jar /usr/hdp/current/oozie-server/libext-customer
```

5. Extract share-lib.

```
/usr/hdp/current/oozie-server/bin/oozie-setup.sh sharelib create -fs hdfs://<namenode>:8020
```

To verify that the sharelibs extracted correctly, run the following command:

```
oozie admin -oozie http://<oozie server host address>:11000/oozie -shareliblist
```

There should be:

- Available ShareLib
- oozie
- hive
- distcp
- hcatalog
- sqoop
- mapreduce-streaming
- pig

Change the ownership and permissions of the oozie directory:

```
su -l hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie"
```

```
su -l hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

6. If a previous version of Oozie was created using auto schema creation, you must run an SQL query.

Use the `oozie-site.xml` properties:

- `oozie.service.JPAService.jdbc.username`
- `oozie.service.JPAService.jdbc.username`
- `oozie.service.JPAService.jdbc.url`

to obtain the password, username and db to run the query.

Run the SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

7. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh  
upgrade -run"
```

8. As root, prepare the Oozie WAR file.

```
chown oozie:oozie /usr/hdp/current/oozie-server/oozie-server/  
conf/server.xml
```

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-setup.sh  
prepare-war -d /usr/hdp/current/oozie-server/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar
New Oozie WAR file with added 'JARs' at /var/lib/oozie/oozie-server/webapps/
oozie.war
```

9. Make sure that following property is added in `oozie-log4j.properties`:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where `${oozie.instance.id}` is determined by oozie, automatically.

10. If you have custom Oozie actions, you must define them in `oozie-site.xml`. Edit the `/etc/oozie/conf/oozie-site.xml` file and add the following property:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>[Comma seperated list of custom actions]</value>
</property>
```

For example, if you have added Spark Action, enter the following:

```
<property>
  <name>oozie.service.SchemaService.wf.ext.schemas</name>
  <value>spark-action-0.1.xsd</value>
</property>
```

11. Configure HTTPS for the Oozie server.

- a. Create a self signed certificate or get certificate from a trusted CA for the Oozie Server
- b. Import the certificate to the client JDK trust store on all client nodes.
- c. In the Ambari Oozie configuration, set the following environment variables in `oozie-env.sh`, adding them if it does not exist:

```
export OOZIE_HTTPS_PORT=11443
export OOZIE_HTTPS_KEYSTORE_FILE=/home/oozie/.keystore
export OOZIE_HTTPS_KEYSTORE_PASS=password
```

- d. Change `OOZIE_HTTP_PORT={{oozie_server_port}}` to `OOZIE_HTTP_PORT=11000`.
- e. Set the `oozie.base.url` to the HTTPS address.
- f. Save the configuration, and restart the Oozie components.

12. Start Oozie as the Oozie user:

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-start.sh"
```

13. Check processes.

```
ps -ef | grep -i oozie
```

4.17. Configure and Start Apache WebHCat



Note

The `su` commands in this section use "hdfs" to represent the HDFS Service user and webhcat to represent the Apache WebHCat Service user. If you are using another name for these Service users, you need to substitute your Service user name for "hdfs" or "webhcat" in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/hive-webhcat/conf` from the template to the `conf` directory in webhcat hosts.
2. Modify the WebHCat configuration files.
 - a. Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User` (in this example, `hdfs`):

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/"
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/"
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /
hdp/apps/2.5.0.0-<$version>/pig/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.
gz /hdp/apps/2.5.0.0-<$version>/hive/"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.
gz /hdp/apps/2.5.0.0-<$version>/sqoop/"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/pig.
tar.gz"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/
hive.tar.gz"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"
su - hdfs -c "hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/
sqoop.tar.gz"
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp"
```

- b. Update the following properties in the `webhcat-site.xml` configuration file, as their values have changed:

```
<property>
  <name>templeton.pig.archive</name>
```



```

<value>hdfs:///hdp/apps/${hdp.version}/pig/pig.tar.gz</value>
</property>

<property>
  <name>templeton.hive.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/hive/hive.tar.gz</value>
</property>

<property>
  <name>templeton.streaming.jar</name>
  <value>hdfs:///hdp/apps/${hdp.version}/mapreduce/
    hadoop-streaming.jar</value>
  <description>The hdfs path to the Hadoop streaming jar file.</
description>
</property>

<property>
  <name>templeton.sqoop.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/sqoop/sqoop.tar.gz</value>
  <description>The path to the Sqoop archive.</description>
</property>

<property>
  <name>templeton.sqoop.path</name>
  <value>sqoop.tar.gz/sqoop/bin/sqoop</value>
  <description>The path to the Sqoop executable.</description>
</property>

<property>
  <name>templeton.sqoop.home</name>
  <value>sqoop.tar.gz/sqoop</value>
  <description>The path to the Sqoop home in the exploded archive.
    </description>
</property>

```



Note

You do not need to modify `${hdp.version}`.

- c. Remove the following obsolete properties from `webhcat-site.xml`:

```

<property>
  <name>templeton.controller.map.mem</name>
  <value>1600</value>
  <description>Total virtual memory available to map tasks.</description>
</property>

<property>
  <name>hive.metastore.warehouse.dir</name>
  <value>/path/to/warehouse/dir</value>
</property>

```

- d. Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHcat to impersonate your additional HDP-2.5.0 groups and hosts:

```

<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>

```

```
</property>
<property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>
</property>
```

Where:

`hadoop.proxyuser.hcat.group`

Is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

`hadoop.proxyuser.hcat.hosts`

A comma-separated list of the hosts which are allowed to submit requests by 'hcat'.

3. Start WebHCat:

```
su - hcat -c "/usr/hdp/current/hive-webhcat/sbin/
webhcat_server.sh start"
```

4. Smoke test WebHCat.

a. At the WebHCat host machine, run the following command:

```
http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

b. If you are using a secure cluster, run the following command:

```
curl --negotiate -u: http://cluster.$PRINCIPAL.$REALM:50111/
templeton/v1/status {"status":"ok","version":"v1"}
[machine@acme]$
```

4.18. Configure Apache Pig

1. Replace your configuration after upgrading. Copy `/etc/pig/conf` from the template to the conf directory in pig hosts.

2. To validate the Apache Pig upgrade, complete the following steps:

a. On the host machine where Pig is installed, run the following commands:

```
su - $HDFS_USER -c "/usr/hdp/current/hadoop-client/bin/hadoop fs -
copyFromLocal /etc/passwd
passwd"
```

b. Create a Pig script file named `/tmp/id.pig` that contains the following Pig Latin commands:

```
A = load 'passwd' using PigStorage(':'); B = foreach A generate $0 as id;
store B into '/tmp/id.out';
```

c. Run the Pig script:

```
su - $HDFS_USER
pig -l /tmp/pig.log /tmp/id.pig
```

4.19. Configure and Start Apache Sqoop



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Copy `/etc/sqoop/conf` from the template to the conf directory in sqoop hosts.
2. As the HDFS Service user, upload the Apache Sqoop tarball to HDFS.

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
su - hdfs -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
```

3. If you are using the MySQL database as a source or target, then the MySQL connector jar must be updated to 5.1.29 or later.

Refer to the MySQL web site for information on updating the MySQL connector jar.

4. Because Sqoop is a client tool with no server component, you need to run your own jobs to validate the upgrade.

4.20. Configure, Start, and Validate Apache Flume

1. If you have not already done so, upgrade Apache Flume. On the Flume host machine, run the following command:

- For **RHEL/CentOS/Oracle Linux**:

```
yum upgrade flume
```

- For **SLES** :

```
zypper update flume
```

```
zypper remove flume
```

```
zypper se -s flume
```

You should see Flume in the output.

Install Flume:

```
zypper install flume
```

- For **Ubuntu/Debian**:

```
apt-get install flume
```

2. To confirm that Flume is working correctly, create an example configuration file. The following snippet is a sample configuration that can be set using the properties file. For more detailed information, see the “Flume User Guide.”

```
agent.sources = pstream
agent.channels = memoryChannel
agent.channels.memoryChannel.type = memory

agent.sources.pstream.channels = memoryChannel
agent.sources.pstream.type = exec
agent.sources.pstream.command = tail -f /etc/passwd

agent.sinks = hdfsSink
agent.sinks.hdfsSink.type = hdfs
agent.sinks.hdfsSink.channel = memoryChannel
agent.sinks.hdfsSink.hdfs.path = hdfs://tmp/flumetest
agent.sinks.hdfsSink.hdfs.fileType = SequenceFile
agent.sinks.hdfsSink.hdfs.writeFormat = Text
```

The source here is defined as an exec source. The agent runs a given command on startup, which streams data to stdout, where the source gets it. The channel is defined as an in-memory channel and the sink is an HDFS sink.

3. Given this configuration, you can start Flume by navigating to FLUME_HOME and executing the following command:

```
$ bin/flume-ng agent --conf ./conf --conf-file example.conf --name a1 -
Dflume.root.logger=INFO,console
```



Note

The directory specified for `--conf` argument would include a shell script `flume-env.sh` and potentially a `log4j` properties file. In this example, we pass a Java option to force Flume to log to the console and we go without a custom environment script.

4. After validating data in `hdfs://tmp/flumetest`, stop Flume and restore any backup files. Copy `/etc/flume/conf` to the `conf` directory in Flume hosts.

4.21. Configure and Validate Apache Mahout



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Replace your configuration after upgrading. Copy `/etc/mahout/conf` from the backup if it existed to the `conf` directory in mahout hosts.

To validate Apache Mahout:

1. Create a test user named "testuser" in the Linux cluster and in HDFS, and log in as that user.
2. Export the required environment variables for Mahout:

```
export JAVA_HOME="your_jdk_home_install_location_here"
export HADOOP_HOME=/usr/hdp/current/hadoop-client
export MAHOUT_HOME=/usr/hdp.current/mahout-client
export PATH="$PATH":$HADOOP_HOME/bin:$MAHOUT_HOME/bin
export CLASSPATH="$CLASSPATH":$MAHOUT_HOME
```

3. Upload a few megabytes of natural-language plain text to the Linux server as `/tmp/sample-test.txt`.
4. Transfer the `sample-test.txt` file to a subdirectory of the testusers's HDFS home directory.

```
hdfs dfs -mkdir /user/testuser/testdata
hdfs dfs -put /tmp/sample-test.txt /user/testuser/testdata
```

5. Enter the mahout command to convert the plain text file `sample-test.txt` into a sequence file stored in the output directory `mahouttest`:

```
mahout seqdirectory --input /user/testuser/testdata --output /user/testuser/
mahouttest -ow --charset utf-8
```

4.22. Configure and Start Hue



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

For HDP-2.5.0, use the Hue version shipped with HDP-2.5.0. If you have a previous version of Hue, use the following steps to upgrade Hue.

1. Migrate `hue.ini` setting from your old `hue.ini` configuration file to new `hue.ini` configuration file.

2. If you are using the embedded SQLite database, remember to restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database. For example:

```
su - hdfs
cd /var/lib/hue
mv desktop.db desktop.db.old
sqlite3 desktop.db < ~/hue_backup/desktop.bak
exit
```

3. Synchronize Database

```
cd /usr/lib/hue
source ./build/env/bin/activate
hue syncdb
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

4.23. Configure and Start Apache Knox

When working with the Apache Knox Gateway in your Hadoop cluster, it is important you have the latest version of Knox installed so you can take advantage of new features and enhancements, in addition to ensuring your instance of Knox is in sync with other Hadoop components (e.g. Ranger, Spark, Hive, Hue, etc.) for stability and performance. For example, if you need to upgrade your Hadoop cluster from 2.1 to 2.2.x, you should also make sure that your individual Hadoop components are also upgraded to the latest version.

HDP enables you to perform a rolling upgrade in 2.2.x. A rolling upgrade means that you can upgrade a component, or the entire Hadoop stack, without losing service, and your users can continue to use the cluster and run jobs with no application or server downtime. The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple instances of the gateway and a load balancer serving each gateway instance from a single URL.

The main distinction between a rolling upgrade and a traditional upgrade implies the use of a Knox cluster for high availability capabilities. This means that you can require multiple instances of the gateway and a load balancer serving each gateway instance from a single URL. Once the upgrade process is completed, you are up and running with the latest version of Knox on each server you have designated as a Knox server.



Note

In this document, whenever you see a { } with a value inside, this denotes a value you must define.

4.23.1. Upgrade the Knox Gateway

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Knox manually to the latest version. Because “rolling upgrades” are now supported in HDP-2.5.0, it is not important which version of Knox you are currently running, only that you have an instance of the Knox Gateway running.



Note

If you have not already installed Knox, refer to the "Install the Knox RPMs on the Knox Server" section of the Non-Ambari Cluster Installation Guide for instructions on how to install and configure the Knox Gateway.

Before upgrading the Knox Gateway, there are a several steps you must follow to ensure your configuration files, settings, and topology files can be copied to the new Knox Gateway instance when the upgrade is complete, which are described below.

1. Back up your existing `conf` directory.
2. Stop each Knox server.

```
su -l Knox /usr/hdp/{the current Knox version}/knox/bin/gateway.sh stop
```

3. Select the HDP server version you are upgrading to after you have stopped each Knox server.

```
hdp-select set Knox-server {the HDP server version}
```

4. Unzip your previously saved configuration directory (the `conf` directory you backed up in step 1) into the new `/var/log/knox/gateway.conf` directory to import these files.

5. Create the Master Secret:

```
su - Knox -c "/usr/hdp/current/knox/bin/knoxcli.sh create-master"
```

6. Start the Gateway:

```
su - Knox -c "/usr/hdp/current/knox-server/bin/gateway.sh start"
```

7. Restart the Knox server to complete the upgrade.

```
su -l Knox /usr/hdp/{the new HDP server version}/knox/bin/gateway.sh start
```

4.23.2. Verify the Knox Upgrade

To verify the upgrade was successful, follow the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful.
2. Verify you have cluster access using the `LISTSTATUS` WebHDFS API call.

```
curl -ivk -u {user}:{password} https://{knox host}:8443 /gateway/webhdfs/v1/  
tmp?op=LISTSTATUS
```

3. Verify the Knox version using the Knox Admin service and Version API.

```
curl -ivk -u {adminuser}:{adminpassword} https://{knox host}:8443 /gateway/  
admin/v1/version
```



Note

The Admin API requires you to be a member of an Admin group, as specified in the `admin.xml` authorization provider.

When you have verified the Knox upgrade was successful, you can begin using Knox. If, however, the upgrade was unsuccessful, you need to downgrade the Knox Gateway to the previous version. The steps to downgrade the Knox Gateway are described in the next section.

4.23.3. Downgrade the Knox Gateway to the Previous Version

If the Knox Gateway upgrade was unsuccessful, you need to downgrade Knox to the previous version to ensure you have a working Knox Gateway for your cluster. To downgrade Knox, follow the steps listed below.

1. For each server running Knox, stop the server.

```
su -l knox /usr/hdp/{current HDP server version}/knox/bin/gateway.sh stop
```

2. Select the HDP server version you want to use to downgrade your Knox Gateway.

```
hdp-select set knox-server {previous HDP server version}
```

3. Restart the server.

```
su -l knox /usr/hdp/{previous HDP server version}/knox/bin/gateway.sh start
```

4.23.4. Verify the Knox Downgrade Was Successful

When the restart is complete, verify you are running an older version of Knox by following the steps listed below.

1. Navigate to the `/var/log/knox/gateway` directory and check the `gateway.log` file for errors and an acknowledgement that the server restart was successful
2. Verify you have cluster access using the `LISTSTATUSWebHDFS` API call.
3. Check the Knox version using the Knox Admin service and Version API using the following command:

```
curl -ivk -u {adminuser}:{adminpassword} https://{knox host}:8443 /gateway/  
admin/v1/version
```


4.24. Configure and Validate Apache Falcon



Note

In HDP-2.5, if authorization is enabled (for example, in the startup.properties file with *.falcon.security.authorization.enabled=true") then Access Control List (ACL) is mandated for all entities.

Upgrade Falcon after you have upgraded HDFS, Hive, Oozie, and Pig. Stop Oozie jobs while running Falcon.



Important

We recommend upgrading Falcon from HDP 2.1 to HDP 2.2 and then upgrading from HDP 2.2 to HDP 2.5.

Upgrade Apache Falcon from HDP 2.1 to HDP 2.2

1. Upgrade Falcon from HDP 2.1 to HDP 2.2.

- **RHEL/CentOS/Oracle 6**

- a. Download the HDP 2.2.9 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<version>/2.x/updates/2.2.9.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- b. Install the HDP 2.2.9 version of Falcon:

```
yum install "falcon"
```

- **SLES**

- a. Download the HDP 2.2.9 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.2.9.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

- b. Install the HDP 2.2.9 version of Falcon:

```
zypper install "falcon"
```

- **SLES**

- a. Download the HDP 2.2.0 hdp.repo file:

```
http://public-repo-1.hortonworks.com/HDP/susellsp3/2.x/updates/2.2.9.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

- b. Install the HDP 2.2.9 version of Falcon:

```
zypper install "falcon"
```

- **Ubuntu 12**

- a. Download the HDP 2.2 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2.2.9.0/hdp.list -O /etc/apt/sources.list.d/hdp.list
```

- b. Install the HDP 2.2.9 version of Falcon:

```
apt-get install "falcon"
```

- **Debian**

- a. Download the HDP 2.2 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/debian<version>/2.x/updates/2.2.9.0/hdp.list -O /etc/apt/sources.list.d/hdp.list
```

- b. Install the HDP 2.2.9 version of Falcon:

```
apt-get install "falcon"
```

2. Verify start.properties. In Falcon 2.2, TLS is enabled by default. When TLS is enabled, Falcon starts on `https://<falcon_host>:15443/`. You can disable TLS by adding the following line in startup.properties:

```
*.falcon.enableTLS=false
```

3. Verify the client.properties file. If TLS is disabled, make sure the property `falcon.url` is set as follows:

```
falcon.url=http://localhost:15000/
```

4. Install the hotfix provided by Hortonworks customer support to migrate entities in bulk.
5. Restart Falcon.

```
su - falcon -c "cd /usr/hdp/2.2.9.0-$BUILD/falcon/bin/falcon-start"
```

6. Run Falcon CLI to update Falcon entities from HDP 2.1 to HDP 2.2. This script also updates the ownership and permissions of staging and working directories of cluster entities. Run this script as user falcon.

```
falcon upgrade -owner hrt_qa -group users -configDir <falcon_config_store> -backupDir/tmp/
```

Where:

- Owner and group are used to set ACL for Falcon entities.
- configDir refers to the location where Falcon stores its entities. This directory should be the same location as the value of property `*.config.store.uri` in the file `conf/startup.properties`.

- `backupDir` is the location where Falcon 2.1 entitles are copied as backup. User can restore entitis from this directory.

7. Restart Falcon.

```
su - falcon -c "cd /usr/hdp/2.2.9.0-$BUILD/falcon/bin/falcon-start"
```

8. Replace your configuration after upgrading. Copy `/etc/falcon/conf` from the template to the conf directory in falcon hosts.

9. Check your Falcon entities. There should be no changes, but in some cases you may need to update your entities post-upgrade.

Upgrade Falcon from HDP 2.2 to HDP 2.5

1. RHEL/CentOS/Oracle 6

- On all Falcon hosts, clean the yum repository.

```
yum clean all
```

- Remove the HDP 2.2 Falcon component. This command uninstalls the HDP 2.2 component. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "falcon*"
```

- Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

- Remove the HDP 2.2 hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

- Install the HDP 2.5.0 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<version>/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

SLES

- On all Falcon hosts, clean the yum repository.

```
zypper clean -all
```

- Remove the Falcon HDP 2.2 component. This command uninstalls the HDP 2.2 components. It leaves the user data, and metadata, but removes your configurations:

```
zypper rm "falcon*"
```

- Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

- d. Remove the HDP 2.2 hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

- e. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/sles11sp1/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

- f. Install the HDP 2.5.0 version of the Falcon component.

```
zypper install "falcon"
```

SLES

- a. On all Falcon hosts, clean the zypper repository.

```
zypper clean -all
```

- b. Remove HDP 2.2 Falcon component.

```
zypper rm "falcon*"
```

- c. Validate that all HDP 2.2 component binaries are uninstalled:

```
zypper search --installed-only --repo HDP-2.2.6.0
```

- d. Remove the HDP 2.2 hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

- e. Download the HDP 2.5.0 hdp.repo file:

```
http://public-repo-1.hortonworks.com/HDP/susellsp3/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

- f. Install the HDP 2.5.0 version of the Falcon component. F

```
zypper install "falcon"
```

Ubuntu 12

- a. On all Falcon hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

- b. Remove the HDP 2.2 Falcon component. This command uninstalls the HDP 2.2 component. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "falcon*"
```

- c. Validate that all HDP 2.2 component binaries are uninstalled:

```
yum list installed | grep @HDP2.2
```

- d. Remove the HDP 2.2 hdp.repo file

```
rm /etc/apt/sources.list.d/hdp.list
```

- e. Download the HDP 2.5.0 hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2.5.0.0/hdp.list - O /etc/apt/sources.list.d/hdp.list
```

- f. Run an update:

```
apt-get update
```

- g. Install the HDP 2.5.0 versions of the Falcon component.

```
apt-get install "falcon"
```

Debian

- a. On all Falcon hosts, clean the apt-get repository.

```
apt-get clean
```

- b. Remove the HDP 2.2 Falcon component. This command uninstalls the HDP 2.2 component. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "falcon*"
```

- c. Validate that all HDP 2.2 component binaries are uninstalled:

```
dpkg -l | grep "^ii" | grep hadoop
```

- d. Remove the HDP 2.2 hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

- e. Download the HDP 2.5.0 hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/debian<version>/2.x/updates/2.5.0.0/hdp.list - O /etc/apt/sources.list.d/hdp.list
```

- f. Run an update:

```
apt-get update
```

- g. Install the HDP 2.5.0 versions of the Falco component.

```
apt-get install "falcon"
```

2. Verify start.properties. In Falcon 2.5, TLS is enabled by default. When TLS is enabled, Falcon starts on `https://<falcon_host>:15443/`. You can disable TLS by adding the following line in `startup.properties`:

```
*.falcon.enableTLS=false
```

3. Verify the `client.properties` file. If TLS is disabled, make sure the property `falcon.url` is set as follows:

```
falcon.url=http://localhost:15000/
```

4. Restart Falcon.

```
su - falcon -c "cd /usr/hdp/2.5.0.0-$BUILD/falcon/bin/falcon-start"
```

5. Run Falcon CLI to update Falcon entities from HDP 2.1 to HDP 2.2. This script also updates the ownership and permissions of staging and working directories of cluster entities. Run this script as user falcon.

```
falcon upgrade -owner hrt_qa -group users -configDir  
<falcon_config_store> -backupDir/tmp/
```

Where:

- Owner and group are used to set ACL for Falcon entities.
- `configDir` refers to the location where Falcon stores its entities. This directory should be the same location as the value of property `*.config.store.uri` in the file `conf/startup.properties`.
- `backupDir` is the location where Falcon 2.1 entities are copied as backup. User can restore entities from this directory.

6. Restart Falcon.

```
su - falcon -c "cd /usr/hdp/2.5.0.0-$BUILD/falcon/bin/falcon-start"
```

7. Replace your configuration after upgrading. Copy `/etc/falcon/conf` from the template to the `conf` directory in falcon hosts.
8. Check your Falcon entities. There should be no changes, but in some cases you may need to update your entities post-upgrade.

4.25. Configure and Start Apache Storm



Note

The `su` commands in this section use "zookeeper" to represent the Zookeeper Service user. If you are using another name for your Zookeeper Service user, you need to substitute your Zookeeper Service user name for "zookeeper" in each of the `su` commands.

Apache Storm is fairly independent of changes to the HDP cluster, but you must upgrade Storm for rolling upgrade support in HDP-2.5.0 and be on the latest version of Storm.

1. Deactivate all running topologies.
2. Delete all states under zookeeper:

```
/usr/hdp/current/zookeeper-client/bin/zkCli.sh (optionally in  
secure environment specify -server zk.server:port)
```

```
rmr /storm
```

3. Delete all states under the storm-local directory:

```
rm -rf <value of storm.local.dir>
```

4. Stop Storm services on the storm node.

5. Stop ZooKeeper services on the storm node.

```
su - zookeeper -c "export ZOOCFGDIR=/etc/zookeeper/conf ; export  
ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-env.sh ; /  
usr/lib/zookeeper/bin/zkServer.sh stop"
```

6. Remove Storm and Zookeeper from the storm node and install the HDP-2.5.0 version:

- For **RHEL/CentOS/Oracle Linux**:

```
yum erase storm
```

```
yum erase zookeeper
```

```
yum install storm
```

```
yum install zookeeper
```

- For **SLES**:

```
zypper rm storm
```

```
zypper rm zookeeper
```

```
zypper install storm
```

```
zypper install zookeeper
```

- For **Ubuntu/Debian**:

```
apt-get remove storm --purge
```

```
apt-get remove zookeeper --purge
```

```
apt-get install storm
```

```
apt-get install zookeeper
```

7. Replace your configuration after upgrading. Copy /etc/storm/conf from the template to the conf directory .

8. Replace your ZooKeeper configuration after upgrading. Replace the ZooKeeper template configuration in /etc/zookeeper/conf.

9. Start ZooKeeper. On the storm node, run the following command:

```
su - zookeeper -c "source /etc/zookeeper/conf/zookeeper-env.sh; export
ZOO_CFG_DIR=/etc/zookeeper/conf; /usr/hdp/current/zookeeper-server/bin/
zkServer.sh start >> $ZOO_LOG_DIR/zoo.out\""
```

where

- `$ZOO_LOG_DIR` is the directory where ZooKeeper server logs are stored. For example, `/var/log/zookeeper`.

10 Start nimbus, then supervisor/ui/drpc/logviewer:

```
/usr/hdp/current/storm-nimbus/bin/storm nimbus.
```

11 Start Storm, using a process controller, such as supervisor:

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm
supervisor
```

You can use the same command syntax to start Storm using nimbus/ui and logviewer.

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm nimbus
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm ui
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm logviewer
```

```
su - storm /usr/hdp/current/storm-supervisor/bin/storm drpc
```

4.26. Configure and Start Apache Ranger

This section describes how to upgrade the Apache Ranger service.



Note

XA Secure was an add-on component in HDP-2.1. Ranger is the new name for XA Secure. In HDP-2.2 and subsequent releases, Ranger is installed with HDP.

4.26.1. Preparing Your Cluster to Upgrade Ranger

If you are not currently using Ambari to manage your Hadoop cluster, you need to upgrade Ranger manually to the latest version. This section describes the steps you need to follow to prepare your cluster for the Ranger upgrade.

1. Back up the following Ranger configuration directories:

- Ranger Policy Administration Service

```
/usr/lib/xapolicymgr
```

- Ranger UserSync

```
/usr/lib/uxugsync or /etc/uxugsync (Depending on your installation)
```

- Ranger Plugins:

- Hadoop

```
/etc/hadoop/conf
```

- Hive

```
/etc/hive/conf
```

- HBase

```
/etc/hbase/conf
```

- Knox

```
/etc/knox/conf
```

- Storm

```
/etc/storm/conf
```

2. Backup the Ranger Policy and Audit databases. Make sure to take note of the following details in the `install.properties` file:

- db_host
- db_name
- db_user
- db_password
- policy manager configuration
- LDAP directory configuration
- LDAP settings
- LDAP AD domain
- LDAP URL

4.26.2. Stop the Ranger Services

Now that you have prepared your cluster for the Ranger upgrade, you need to stop the Ranger Admin and Ranger UserSync services. To stop the Ranger services, perform the steps described below.

1. Stop the Ranger Policy Admin service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service xapolicymgr stop
```

2. Stop the Ranger UserSync service. When the service is stopped, you receive an acknowledgement from the server that the service has been stopped.

```
service uxugsync stop
```

3. Stop each individual Ranger plugin (HDFS, HBase, Knox, Storm). You receive an acknowledgement from the server that the plugin has been stopped.

```
service <plugin name> stop
```

4.26.3. Install the Ranger Components

Next, you need to re-install each Ranger component again to ensure you have the latest version. Because you have already upgraded your HDP stack, you only need to follow the instructions in the [Non-Ambari Cluster Installation Guide](#) to install each Ranger component. The following components must be installed:

- Ranger Policy Admin service
- Ranger UserSync service
- Ranger Plugins:
 - HDFS
 - HBase
 - Hive
 - Knox
 - Storm

With this release, Ranger has also added support for the following components:

- Solr
- Kafka
- YARN

4.26.4. Restart the Ranger Services

Once you have re-installed each Ranger component, you then need to restart these components to ensure the new configurations are loaded in your cluster. The Non-Ambari Cluster Installation Guide describes how you can start the following Ranger services:

- Ranger Policy Admin service

```
service ranger-admin start
```

- Ranger UserSync service

```
service ranger-usersync start
```

4.26.5. Remove Existing Startup Files and Symbolic Links

In order to ensure that your Ranger components are upgraded correctly, and there are no conflicts between versions, you should remove any existing startup files and symbolic links

from the previous Ranger 2.2 version. The steps you need to follow to remove these files and links are described below.

1. Remove the Policy Manager startup files.

```
rm -f /etc/init.d/xapolicymgr
```

2. Remove the Policy Manager symbolic links.

```
rm -rf /etc/rc*.d/*xapolicymgr
```

3. Remove the UserSync startup files.

```
rm -f /etc/rc*.d/*uxugsync
```

4. Remove the UserSync symbolic links.

```
rm -rf /etc/rc*.d/uxugsync
```

5. Remove the Policy Manager library files.

```
rm -f /usr/lib/xapolicymgr
```

6. Remove the UserSync library files.

```
rm -f /usr/lib/uxugsync
```

4.26.6. Enable Ranger Plugins

The final step in the Ranger upgrade process requires you to re-enable the Ranger plugins. Although you are only required to enable HDFS in your cluster, you should re-enable all of the Ranger plugins because class names have changed for the 2.5.0 release.



Note

When you enable each Ranger plugin, make sure you remove all 2.1 class name values.

To re-enable the Ranger plugins, use the links listed below to access instructions in the *Non-Ambari Cluster Installation* guide that describe editing the `install.properties` file and enabling the Ranger plugins:



Important

Before enabling the HDFS plugin, remove `set-hdfs-plugin-env.sh` from `/etc/hadoop/conf`. You need to re-enable this plugin after the upgrade is complete.

- [HDFS Plugin](#)
- [YARN Plugin](#)
- [Kafka Plugin](#)
- [HBase Plugin](#)

- [Hive Plugin](#)
- [Knox Plugin](#)
- [Storm Plugin](#)

4.27. Configuring and Upgrading Apache Spark

Instructions in this section are specific to HDP-2.3.4 and later. For earlier versions for HDP, refer to the version-specific documentation.

1. Stop your current version of Apache Spark:

```
su - spark -c "/usr/hdp/current/spark-client/sbin/ stop-history-server.sh".
```

2. Remove your current version of Spark: `yum erase "spark*"`.

3. When you install Spark, two directories are created:

- `/usr/hdp/current/spark-client` for submitting Spark jobs
- `/usr/hdp/current/spark-history` for launching Spark master processes, such as the Spark history server

Search for Spark in the HDP repo:

- For RHEL or CentOS:

```
yum search spark
```

- For SLES:

```
zypper install spark
```

- For Ubuntu and Debian:

```
apt-cache spark
```

This shows all the versions of Spark available. For example,

```
spark_2_3_4_0_3371-master.noarch : Server for Spark master
spark_2_3_4_0_3371-python.noarch : Python client for Spark
spark_2_3_4_0_3371-worker.noarch : Server for Spark worker
```

4. Add the node where you want Spark 1.5.2 History Server to run. Install the version corresponding to the HDP version you currently have installed.

a. `su - root`

b. `yum install spark_spark_2_3_4_0_2950-master -y`

c. To use Python: `yum install spark_spark_2_3_4_0_2950-python`

d. `conf-select create-conf-dir --package spark --stack-version spark_2_3_4_0_2950 --conf-version 0`

- e. `cp /etc/spark/spark_2_3_4_0_2950/0/* /etc/spark/spark_2_3_4_0_2950/0/`
 - f. `conf-select set-conf-dir --package spark --stack-version spark_2_3_4_0_2950 --conf-version 0`
 - g. `hdp-select set spark-client spark_2_3_4_0_2950`
 - h. `hdp-select set spark-historyserver spark_2_3_4_0_2950`
5. Validate the Spark installation. As user `spark`, run SparkPI example:
- a. `su - spark -c "cd /usr/hdp/current/spark-client"`
 - b. `./bin/spark-submit --class org.apache.spark.examples.SparkPi --master yarn-client --num-executors 3 --driver-memory 512m --executor-memory 512m --executor-cores 1 lib/spark-examples*.jar 10`
6. Restart Spark on YARN in either `yarn-cluster` mode or `yarn-client` mode:
- **yarn-cluster mode:** `./usr/hdp/current/spark-client/bin/spark-submit --class path.to.your.Class --master yarn-cluster [options] [app options]`
 - **yarn-client mode:** `./usr/hdp/current/spark-client/bin/spark-shell --master yarn-client`

4.28. Upgrade Apache Slider

To upgrade Apache Slider, simply upgrade the Slider client.

1. Upgrade Slider client:

```
hdp-select set slider-client 2.5.0.0-<version>
slider version
```

4.29. Upgrade Apache Kafka

To upgrade Apache Kafka 2.5.0, you must stop all of the Kafka brokers before you start the upgrade.

1. Shut down the current Kafka daemon, switch to the new version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"
hdp-select set kafka-broker 2.2.4.0-2633
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.
3. If the upgrade process fails, follow the steps in "Downgrading Kafka" to return to your previous version of Kafka.

4.29.1. Downgrading Kafka

Downgrade each Kafka node one at a time. You can stop each Kafka broker and upgrade the component without downtime if you have enough replication for your topic.

1. Shut down the current Kafka daemon, switch to the previous version, and start the daemon:

```
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka stop"
hdp-select set kafka-broker 2.2.2.0-2041
su - kafka -c "/usr/hdp/current/kafka-broker/bin/kafka start"
```

2. To verify that the Kafka daemon joined the cluster, create a topic and submit it to Kafka. Send a test message for that topic, and then validate that it was received by a consumer.

4.30. Finalize the Upgrade



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.

1. Verify your filesystem health before finalizing the upgrade. After you finalize an upgrade, your backup is discarded!
2. As the `$HDFS_USER`, enter:

```
su - hdfs -c "dfsadmin -finalizeUpgrade"
```

4.31. Install New HDP 2.5 Services

Install new HDP-2.5.0 Services (see the [Non-Ambari Cluster Installation Guide](#)):

- Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
- SmartSense – a next generation subscription model that features upgrade and configuration recommendations.

5. Upgrade from HDP 2.0 to HDP 2.5.0 Manually

This chapter provides instructions on how to manually upgrade to HDP 2.5.0 from the HDP 2.0 release. It assumes the existing HDP 2.0 was also installed manually.



Important

If you installed and manage HDP 2.0 with Ambari, **you must use the [Ambari Upgrade Guide](#)** to perform the the HDP 2.3 to HDP 2.5.0 upgrade.

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

Starting with HDP-2.2, HDP supports side-by-side installation of HDP 2.2 and subsequent releases, which lets you perform rolling upgrades on your cluster and improve execution times for in-place upgrade. To support side-by-side installation, the HDP package version naming convention for both RPMs and Debs has changed to include the HDP product version. For example, `hadoop-hdfs` in HDP-2.2.4 is now `hadoop_2.2.4.0.hdfs`. HDP-2.2 marked the first release where HDP rpms, debs, and directories contained versions in the names to permit side-by-side installations of later HDP releases. To select from the releases you have installed side-by-side, Hortonworks provides `hdp-select`, a command that lets you select the active version of HDP from the versions you have installed.

However, because HDP-2.1 did not support side-by-side installation, you upgrade to HDP-2.3.0 in a way very similar to previous minor-version upgrades. Subsequent upgrades after 2.5.0 is easier, and if you choose to add Ambari to your cluster, you can able to use the Rolling Upgrade feature.



Important

You cannot perform a side-by-side rolling upgrade from HDP 2.1 to HDP-2.5.0; only upgrade in place is supported. Side-by-side and rolling upgrade support starts with releases HDP 2.2 and above.

HDP packages for a complete installation of HDP 2.5 consumes about 6.5 GB of disk space.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.

This document provides instructions on how to upgrade to HDP 2.5.0 from the HDP 2.0 release.



Note

These instructions cover the upgrade between two minor releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

1. Download HDP 2.5
2. Get Ready to Upgrade
3. Configure and Start Hadoop
4. Start HDFS
5. Upgrade Apache ZooKeeper
6. Upgrade Apache HBase
7. Upgrade Apache Hive and Apache HCatalog
8. Upgrade Apache Oozie
9. Upgrade Apache WebHCat (Templeton)
10. Upgrade Apache Pig
11. Upgrade Apache Sqoop
12. Upgrade Apache Flume
13. Upgrade Apache Mahout
14. Upgrade Hue
15. Finalize the Upgrade
16. Install additional HDP 2.5 services

5.1. Getting Ready to Upgrade

HDP Stack upgrade involves upgrading from HDP 2.0 to HDP 2.5 versions and adding the new HDP 2.5 services.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

The first step is to ensure you keep a backup copy of your HDP 2.0 configurations.

**Note**

You must use **kinit** before running the commands as any particular user.

1. Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP 2.5 consumes about 6.5 GB of disk space.

2. Back up the following HDP directories:

- /etc/hadoop/conf
- /etc/hbase/conf
- /etc/hive/conf
- /etc/pig/conf
- /etc/sqoop/conf
- /etc/flume/conf
- /etc/mahout/conf
- /etc/oozie/conf
- /etc/zookeeper/conf
- **Optional:** Back up your userlogs directories, \${mapred.local.dir}/userlogs.

3. Run the fsck command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.)

```
su - hdfs -c "hdfs fsck / -files -blocks -locations > dfs-old-fsck-1.log"
```

4. Use the following instructions to compare status before and after the upgrade:

The following commands must be executed by the user running the HDFS service (by default, the user is hdfs).

a. Capture the complete namespace of the file system. (The second command does a recursive listing of the root file system.)

```
su - hdfs -c "hdfs dfs -ls -R / > dfs-old-lsr-1.log"
```

**Note**

In secure mode you must have kerberos credentials for the hdfs user.

b. Run the report command to create a list of DataNodes in the cluster.

```
su - hdfs -c "hdfs dfsadmin -report > dfs-old-report-1.log"
```

- c. **Optional:** You can copy all or unrecoverable only data storelibext-customer directory in HDFS to a local file system or to a backup instance of HDFS.
 - d. **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.
5. As the HDFS user, save the namespace by executing the following command:

```
su - hdfs

hdfs dfsadmin -safemode enter

hdfs dfsadmin -saveNamespace
```

6. Back up your NameNode metadata.

- Copy the following checkpoint files into a backup directory. If these directories do not exist, create them.

```
dfs.namenode.dir/edits

dfs.namenode.name.dir/image/fsimage

dfs.namenode.name.dir/current/fsimage
```

- Store the layoutVersion of the namenode.

```
${dfs.namenode.name.dir}/current/VERSION
```

7. Finalize any **PRIOR** HDFS upgrade, if you have not done so already.

```
su - hdfs -c "hdfs dfsadmin -finalizeUpgrade"
```

8. **Optional:** Back up the Hive Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation.

Table 5.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hive < /tmp/mydir/backup_hive.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hive > /tmp/mydir/backup_hive.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hive < /tmp/mydir/backup_hive.sql</pre>

Database Type	Backup	Restore
Oracle	Connect to the Oracle database using sqlplus export the database: exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr For example: mysql hive < /tmp/mydir/backup_hive.sql	Import the database: imp username/password@database file=input_file.dmp

9. Optional: Back up the Oozie metastore database.

These instructions are provided for your convenience. Please check your database documentation for the latest back up instructions.

Table 5.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	mysqldump \$dbname > \$outputfilename.sql For example: mysqldump oozie > /tmp/mydir/backup_oozie.sql	mysql \$dbname < \$inputfilename.sql For example: mysql oozie < /tmp/mydir/backup_oozie.sql
PostgreSQL	sudo -u \$username pg_dump \$databasename > \$outputfilename.sql For example: sudo -u postgres pg_dump oozie > /tmp/mydir/backup_oozie.sql	sudo -u \$username psql \$databasename < \$inputfilename.sql For example: sudo -u postgres psql oozie < /tmp/mydir/backup_oozie.sql

10.Optional: Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 5.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	mysqldump \$dbname > \$outputfilename.sqlsbr For example: mysqldump hue > /tmp/mydir/backup_hue.sql	mysql \$dbname < \$inputfilename.sqlsbr For example: mysql hue < /tmp/mydir/backup_hue.sql
PostgreSQL	sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr For example: sudo -u postgres pg_dump hue > /tmp/mydir/backup_hue.sql	sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr For example: sudo -u postgres psql hue < /tmp/mydir/backup_hue.sql

Database Type	Backup	Restore
Oracle	Connect to the Oracle database using sqlplus. Export the database. For example: exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr	Import the database: For example: imp username/password@database file=input_file.dmp
SQLite	/etc/init.d/hue stop su \$HUE_USER mkdir ~/hue_backup sqlite3 desktop.db .dump > ~/hue_backup/desktop.bak /etc/init.d/hue start	/etc/init.d/hue stop cd /var/lib/hue mv desktop.db desktop.db.old sqlite3 desktop.db < ~/hue_backup/desktop.bak /etc/init.d/hue start

11 Stop all services (including MapReduce) and client applications deployed on HDFS using the instructions provided in the [Stopping HDP Services](#).

12. Verify that edit logs in \${dfs.namenode.name.dir}/current/edits* are empty.

a. Run: `hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out`

b. Verify edits.out file. It should only have OP_START_LOG_SEGMENT transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

c. If edits.out has transactions other than OP_START_LOG_SEGMENT run the following steps and then verify edit logs are empty.

- Start the existing version NameNode.
- Ensure there is a new FS image file.
- Shut the NameNode down.

```
hdfs dfsadmin - saveNamespace
```

13. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path component in this version of HDFS.
```

Please rollback and delete or rename this path, or upgrade with the `-renameReserved` key-value pairs option to automatically rename these paths during upgrade.

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup.

For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, specify:

```
-upgrade -renameReserved .snapshot=.my-snapshot,.reserved=.my-reserved.
```

If no key-value pairs are specified with `-renameReserved`, the NameNode suffixes reserved paths with `.<LAYOUT-VERSION>.UPGRADE_RENAMED`. For example:

```
.snapshot.-51.UPGRADE_RENAMED.
```



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

5.2. Upgrade HDP 2.0 Components

The upgrade process to HDP 2.5 involves the following steps. Select your OS:

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.0 components. This command uninstalls the HDP 2.0 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*"
"flume*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old `hdp.repo` file:

```
rm /etc/yum.repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the `hdp.repo` file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
centos<version>/2.x/updates/2.5.0.0/hdp.repo -O /etc/
yum.repos.d/hdp.repo
```

- Confirm that the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0.0 Hortonworks Data Platform Version - HDP-2.5.0.0
```

5. Install the HDP 2.5.0 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.0 components:

```
yum install "hadoop" "hadoop-hdfs" "hadoop-libhdfs" "hadoop-yarn" "hadoop-
mapreduce" "hadoop-client" "openssl" "webhcat" "hcatalog" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hue" "hive" "storm" "flume"
"accumulo" "mahout" "kafka" "spark" "slider" "hdp_mon_nagios_addons"
```

6. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

Check to make sure no component file names from the previous stack appear in the returned list.

RHEL/CentOS/Oracle

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 2.0 components. This command uninstalls the HDP 2.0 components. It leaves the user data, and metadata, but removes your configurations:

```
yum erase "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*"
"flume*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
centos<version>/2.x/updates/2.5.0.0/hdp.repo -O /etc/
yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0.0 Hortonworks Data Platform Version - HDP-2.5.0.0
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.1 components:

```
yum install "hadoop" "hadooplzo" "webhcat" "hcatalog" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive"
"flume" "accumulo" "mahout" "hue" "hdp_mon_nagios_addons"

yum install install hive-webhcat webhcat-tar-hive webhcat-tar-
pig
```

6. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.0 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*"
"flume*" "phoenix*" "accumulo*" "mahout*" "hue" "hue-common"
"hue-shell" "hdp_mon_nagios_addons"
```

3. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
sles11sp1/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/
hdp.repo
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.0 components:

```
zypper install "hadoop" "hadooplzo" "webhcat" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "flume"
"accumulo" "mahout" "hue" "hdp_mon_nagios_addons"

zypper install hive-webhcat webhcat-tar-hive webhcat-tar-pig

zypper up -r HDP-2.5.0.0

zypper install oozie-client
```

6. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatalog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep
hcatalog
```

No component files names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 2.0 components.

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*" "hive*"
"flume*" "accumulo*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.list
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
http://public-repo-1.hortonworks.com/HDP/suse11sp3/2.x/
updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/hdp.repo
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.0 components:

```
zypper install "hadoop" "hadooplzo" "webhcat" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "flume"
"accumulo" "mahout" "hue" "knox" "hdp_mon_nagios_addons"
```



```
zypper install hive-webhcat webhcat-tar-hive webhcat-tar-pig
zypper up -r HDP-2.5.0.0
zypper install oozie-client
```

6. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatlog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep
hcatalog
```

No component files names should appear in the returned list.

Ubuntu 12

1. On all hosts, clean the apt-get repository.

```
apt-get clean -&-all
```

2. Remove your old HDP 2.0 components.

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*"
"collectd*" "gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*"
"hbase*" "hive*" "flume*" "accumulo*" "mahout*" "hue.*"
"hdp_mon_nagios_addons" --purge
```

3. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/
updates/2.5.0.0/hdp.list -O /etc/apt/sources.list.d/hdp.list
```

5. Run an update:

```
apt-get update
```

6. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.0 components:

```
apt-get install "hadoop" "hadooplzo" "webhcat" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive"
"flume" "accumulo" "mahout" "knox" "hdp_mon_nagios_addons"

apt-get install hive-webhcat webhcat-tar-hive webhcat-tar-pig
```

Debian

1. On all hosts, clean the apt-get repository.

```
apt-get clean all
```

2. Remove your old HDP 2.0 components. This command uninstalls the HDP 2.0 components. It leaves the user data, and metadata, but removes your configurations:

```
apt-get remove "hadoop*" "webhcat*" "hcatalog*" "oozie*"
"collectd*" "gccxml*" "pig" "hdfs*" "sqoop*" "zookeeper*"
"hbase*" "hive*" "flume*" "mahout*" "hue"
```

3. Remove your old hdp.repo file:

```
rm /etc/apt/sources.list.d/hdp.list
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget http://public-repo-1.hortonworks.com/HDP/
debian<version>/2.x/updates/2.5.0.0/hdp.list -O /etc/apt/
sources.list.d/hdp.list
```

5. Run an update:

```
apt-get update
```

6. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 2.0 components:

```
apt-get install "hadoop" "hadooplzo" "webhcat" "oozie"
"collectd" "gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive"
"flume" "accumulo" "mahout" "knox" "hdp_mon_nagios_addons"

apt-get install hive-webhcat webhcat-tar-hive webhcat-tar-pig
```

5.3. Symlink Directories with hdp-select



Warning

HDP 2.5 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

- Before you run hdp-select, remove one link:

```
rm /usr/bin/oozie
```

- Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5.0.0-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5.0.0-2
```

5.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available cluster, upgrade Apache ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "zookeeper" is used to represent the Zookeeper Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all the ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "export ZOOCFGDIR=/usr/hdp/current/zookeeper-server/conf ; export ZOOCFG=zoo.cfg; source /usr/hdp/current/zookeeper-server/conf/zookeeper-env.sh ; /usr/hdp/current/zookeeper-server/bin/zkServer.sh start"

/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc
```

5.5. Configure Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/CPU settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5. make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/CPU settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5. make sure you remove old path specifications from `hadoop-env.sh`, such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load, as this is not where lzo is installed.

Ubuntu/Debian

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/CPU settings in `yarn-site.xml` and `mapred-site.xml`
2. Paths have changed in HDP 2.5. make sure you remove lines from `hadoop-env.sh` such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your `hadoop-env.sh` file, the lzo compression code does not load as this is not where lzo is installed.

5.6. Set RPC Authentication

To successfully upgrade to HDP 2.5, you must change the value of `hadoop.rpc.protection` to authentication in `core-site.xml`:

```
<property>  
  <name>hadoop.rpc.protection</name>  
  <value>authentication</value>  
</property>
```

5.7. Start Hadoop Core



Warning

Before you start HDFS on an HA system you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

To start HDFS, run commands as the `$HDFS_USER`.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. If you are upgrading from an HA NameNode configuration, start all JournalNodes. On each JournalNode host, run the following command:

```
su - hdfs -c "/usr/hdp/current/hadoop-hdfs-journalnode/./  
hadoop/sbin/hadoop-daemon.sh start journalnode"
```



Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

2. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually. On the active NameNode host, run the following command:

```
su - hdfs -c "/usr/hdp/current/hadoop-hdfs-namenode/./hadoop/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the `-upgrade` option only once. After you have completed this step, you can bring up the NameNode using this command without including the `-upgrade` option.

To check if the Upgrade is in progress, check that the `\previous` directory has been created in `\NameNode` and `\JournalNode` directories. The `\previous` directory contains a snapshot of the data before upgrade.

In a NameNode HA configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode. To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the `'-bootstrapStandby'` flag. Do NOT start this standby NameNode with the `'-upgrade'` flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The `bootstrapStandby` command downloads the most recent fsimage from the active NameNode into the `$dfs.name.dir` directory of the standby NameNode. You can enter that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

3. Verify that the NameNode is up and running:

```
ps -ef | grep -i NameNode
```

4. Start the Secondary NameNode. On the Secondary NameNode host machine, run the following command:

```
su -l hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

5. Verify that the Secondary NameNode is up and running:

```
ps -ef|grep SecondaryNameNode
```

6. Start DataNodes.

On each of the DataNodes, enter the following command. If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-hdfs-datanode/./hadoop/sbin/hadoop-daemon.sh --config /etc/hadoop/conf start datanode"
```

7. Verify that the DataNode process is up and running:

```
ps -ef|grep DataNode
```

8. Verify that NameNode can go out of safe mode.

```
hdfs dfsadmin -safemode wait
```

```
Safemode is OFF
```

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

5.8. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.

1. Run the fsck command on namenode as \$HDFS_USER:

```
hdfs fsck / -files -blocks -locations > dfs-new-fsck-1.log
```

2. Run hdfs namespace and report.

- a. List directories.

```
hdfs dfs -ls -R / > dfs-new-lsr-1.log
```

- b. Run report command to create a list of DataNodes in the cluster.

```
hdfs dfsadmin -report > dfs-new-report-1.log
```

3. Compare the namespace report from before the upgrade with the report after the upgrade. Verify that user files exist after upgrade:

```
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log dfs-old-lsr-1.log < -- > dfs-new-lsr-1.log
```



Note

You must do this comparison manually to catch all errors.

4. From the NameNode WebUI, see if all DataNodes are up and running.

`http://<namenode>:50070`

5.9. Configure YARN and MapReduce

After you upgrade hadoop, complete the following steps to update your configs.

1. Upload the MapReduce tarball to HDFS. As the HDFS user, for example 'hdfs':

```
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/mapreduce/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/mapreduce.tar.gz /hdp/apps/
2.5.0.0-<$version>/mapreduce/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/mapreduce
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/mapreduce/mapreduce.tar.
gz
```

2. Make the following changes to `mapred-site.xml`:

- Add:

```
<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>
```

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
  <value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}/hadoop/lib/native:/usr/
hdp/${hdp.version}/hadoop/lib/native/Linux-amd64-64</value>
</property>

<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.admin.reduce.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}/mapreduce/mapreduce.tar.gz#mr-framework</
value>
</property>

<property>
  <name>mapreduce.application.classpath</name>
```

```
<value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
$PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/common/*,
$PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/yarn/*,
$PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
$PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
$PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
/usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
/etc/hadoop/conf/secure</value>
</property>
```

3. Add the following property to yarn-site.xml:

```
<property>
  <name>yarn.application.classpath</name>
  <value>$HADOOP_CONF_DIR,/usr/hdp/${hdp.version}/hadoop-client/*,/usr/hdp/
${hdp.version}/hadoop-
client/lib/*,/usr/hdp/${hdp.version}/hadoop-hdfs-client/*,/usr/hdp/${hdp.
version}/hadoop-hdfs-
client/lib/*,/usr/hdp/${hdp.version}/hadoop-yarn-client/*,/usr/hdp/${hdp.
version}/hadoop-yarn-client/lib/*<
  /value>
</property>
```

4. For secure clusters, you must create and configure the container-executor.cfg configuration file:

- Create the container-executor.cfg file in /etc/hadoop/conf/.
- Insert the following properties:

```
yarn.nodemanager.linux-container-executor.group=hadoop
banned.users=hdfs,yarn,mapred
min.user.id=1000
```

- yarn.nodemanager.linux-container-executor.group - Configured value of yarn.nodemanager.linux-container-executor.group. This must match the value of yarn.nodemanager.linux-container-executor.group in yarn-site.xml.
- banned.users - Comma-separated list of users who can not run container-executor.
- min.user.id - Minimum value of user id. This prevents system users from running container-executor.
- allowed.system.users - Comma-separated list of allowed system users.
- Set the file /etc/hadoop/conf/container-executor.cfg file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/
bin/container-executor
```



```
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn-server-nodemanager/bin/
container-executor
```

5.10. Start YARN/MapReduce Services



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Application Timeline Server has changed. If you installed ATS with HDP 2.0, you need to add the following properties to `yarn-site.xml` to continue using ATS:

```
yarn.timeline-service.leveldb-timeline-store.path=/var/log/hadoop-yarn/
timeline
yarn.timeline-service.leveldb-timeline-store.ttl-interval-ms=300000

** If you are upgrading to HDP 2.1.3 or higher, use the following setting: **
yarn.timeline-service.store-class=org.apache.hadoop.yarn.server.timeline.
LevelDbTimelineStore

** If you are upgrading to HDP 2.1.2, use the following setting: **
yarn.timeline-service.store-class=org.apache.hadoop.yarn.server.
applicationhistoryservice.timeline.LevelDbTimelineStore

yarn.timeline-service.ttl-enable=true
yarn.timeline-service.ttl-ms=2678400000
yarn.timeline-service.generic-application-history.store-class=org.apache.
hadoop.yarn.server.applicationhistoryservice.NullApplicationHistoryStore
yarn.timeline-service.webapp.address={PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:8188
yarn.timeline-service.webapp.https.address=
{PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:8190
yarn.timeline-service.address={PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:10200
```

Add the following properties to `hive-site.xml`:

```
hive.execution.engine=mr
hive.exec.failure.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.exec.post.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.exec.pre.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.tez.container.size={map-container-size}

** If mapreduce.map.memory.mb > 2GB then set it equal to mapreduce.map.memory.
Otherwise, set it equal to mapreduce.reduce.memory.mb **

hive.tez.java.opts=-server -Xmx800m -Djava.net.preferIPv4Stack=true -
XX:NewRatio=8
-XX:+UseNUMA -XX:+UseParallelGC
```

Use configuration values appropriate for your environment. For example, the value "800" in the preceding example is an example, not a requirement.

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.

1. If you are using an HA enabled system, you must upgrade the ZooKeeper service and it must be running.
2. Start the ResourceManager on all your ResourceManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/sbin/yarn-daemon.sh start resourcemanager"
```

```
ps -ef | grep -i resourcemanager
```

3. Start the NodeManager on all your NodeManager hosts.

```
su - yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh start nodemanager"
```

```
ps -ef | grep -i nodemanager
```

4. To start MapReduce, run the following commands:

```
su - yarn -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-jobhistory-daemon.sh
```

```
start historyserver"
```

```
ps -ef | grep -i jobhistoryserver
```

5.11. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user.

The job uses MapReduce to write 100MB of data into HDFS with RandomWriter.

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.jar randomwriter -Dtest.randomwrite.total_bytes=10000000 test-after-upgrade.
```

You should see messages similar to:

```
map 0% reduce 0%
...map 100% reduce 100%
Job ... completed successfully
```

MapReduce works successfully. You can now upgrade your other components.

Basic Troubleshooting

To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

```
http://<resource manager host>:8088/cluster/nodes
```

The number of active nodes should be equal to the number of nodemanagers.

Accessing error messages:

1. Access the ApplicationMaster WebUI to view the container logs.

2. In your console log for the MapReduce job, look for a line in this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track
the job: http://<resource manager host>:8088/proxy/
application_1380673658357_0007/
```

3. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

5.12. Configure and Start Apache HBase



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hbase" is used to represent the Apache HBase Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Start the ZooKeeper Server and NameNode services.

2. As the HBase user, run an upgrade:

```
su - hbase -c "hbase upgrade -execute"
```

You should see a completed Znode upgrade with no errors.

3. Replace your configuration after upgrading. Replace the HBase template configuration in `/etc/hbase/conf`.

4. Start services. Run as root and suppose `$HBASE_USER = hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

5. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

5.13. Configure and Start Apache Hive and Apache HCatalog



Warning

In HDP 2.1.3 (Apache Hive 0.13.0) the Decimal data type is now treated as the type `Decimal(10,0)`: 10 digits of precision and 0 scale. This is a change from the variable precision and scale that was available in Hive 0.11.0 and Hive 0.12.0, which allowed up to 38 digits of precision and unlimited scale.

To avoid unintended "rounding" of decimal data, sites that were previously running Hive 0.11.0 and Hive 0.12.0 may need to migrate tables with Decimal columns after upgrading to Hive 0.13.0. For details, see the [Apache Hive wiki](#). For assistance with upgrades that involve Decimal data, please contact Hortonworks Support.



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hive" is used to represent the Hive Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Prior to starting the upgrade process, set the following in your hive configuration file:

```
datanucleus.autoCreateSchema=false
```

2. Ensure that the required driver .jar file (`oracle.jdbc.driver.OracleDriver`) is available in the `/hive/metastore/lib` folder.
3. Upgrade the Hive Metastore database schema. Restart the Hive Metastore database and run:

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -
upgradeSchema -dbType <$databaseType>"
```

The value for `$databaseType` can be `derby`, `mysql`, `oracle` or `postgres`.



Note

If you are using PostgreSQL, you should reset the Hive Metastore database owner to `<HIVE_USER>`, run the following commands:

```
su - <POSTGRES_USER>
```

```
ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER
TO<HIVE_USER>
```



Note

If you are using Oracle 11, you may see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
optimize.mapjoin.mapreduce does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
heapsize does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
server2.enable.impersonation does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.
semantic.analyzer.factory.impl does not exist
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.
convert.sortmerge.join.noconditionaltask does not exist
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.
internal:1521/XE
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver
```

```

Metastore connection User: hiveuser
Starting upgrade metastore schema from version 0.13.0 to 0.14.0
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql
Error: ORA-00955: name is already used by an existing object
      (state=42000,code=955)
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.
oracle.sql: Schema script failed, errorcode 2
Completed upgrade-0.13.0-to-0.14.0.oracle.sql
schemaTool completed

```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.

4. Download and extract the HDP companion files.

Copy the hive-site.xml file in the configuration_files/hive directory of the extracted companion files to the etc/hive/conf directory on your Hive host machine. This new version of hive-site.xml contains new properties for HDP 2.5.0 features.

5. Edit hive-site.xml and modify the properties based on your environment. Search for TODO in the file for the properties to replace.

- Edit the connection properties for your Hive metastore database in hive-site.xml:

```

<property>
  <name>javax.jdo.option.ConnectionURL</name>
  <value>jdbc:mysql://TODO-HIVE-METASTORE-DB-SERVER:TODO-HIVE-METASTORE-DB-
PORT/
  TODO-HIVE-METASTORE-DB-NAME?createDatabaseIfNotExist=true</value>
  <description>Enter your Hive Metastore Connection URL, for example if
MySQL: jdbc:mysql://localhost:3306/mysql?createDatabaseIfNotExist=true</
description>
</property>

<property>
  <name>javax.jdo.option.ConnectionUserName</name>
  <value>TODO-HIVE-METASTORE-DB-USER-NAME</value>
  <description>Enter your Hive Metastore database user name.</description>
</property>

<property>
  <name>javax.jdo.option.ConnectionPassword</name>
  <value>TODO-HIVE-METASTORE-DB-PASSWORD</value>
  <description>Enter your Hive Metastore database password.</description>
</property>

<property>
  <name>javax.jdo.option.ConnectionDriverName</name>
  <value>TODO-HIVE-METASTORE-DB-CONNECTION-DRIVER-NAME</value>
  <description>Enter your Hive Metastore Connection Driver Name, for
example if MySQL:
com.mysql.jdbc.Driver</description>
</property>

```

- Edit the following properties in the hive-site.xml file:

```

<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>

```

```

    <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache</
description>
</property>

```

- **Optional:** If you want Hive Authorization, set the following Hive authorization parameters in the hive-site.xml file:

```

<property>
  <name>hive.security.authorization.enabled</name>
  <value>>true</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
StorageBasedAuthorizationProvider</value>
</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
StorageBasedAuthorizationProvider</value>
</property>

<property>
  <name>hive.security.authenticator.manager</name>
  <value>org.apache.hadoop.hive.ql.security.ProxyUserAuthenticator</value>
</property>

```

- For a remote Hive metastore database, set the IP address (or fully-qualified domain name) and port of the metastore host using the following hive-site.xml property value.

(To enable HiveServer2, leave this property value empty.)

```

<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server. sbrTo enable
HiveServer2,
  leave the property value empty. </description>
</property>

```

You can further fine-tune your configuration settings based on node hardware specifications, using the HDP utility script.

- Disable autocreation of schemas to match HDP 2.1+ configurations. Edit hive-site.xml to set the value of datanucleus.autoCreateSchema to false.

```

<property>
  <name>datanucleus.autoCreateSchema</name>
  <value>>false</value>

```

```
<description>Creates necessary schema on a startup if one doesn't exist.
</description>
</property>
```

6. Start Hive Metastore.

On the Hive Metastore host machine, run the following command:

```
su - hive -c "nohup /usr/hdp/current/hive-metastore/bin/hive
--service metastore -hiveconf hive.log.file=hivemetastore.log
>/var/log/hive/hivemetastore.out 2>/var/log/hive/
hivemetastoreerr.log &"
```

7. Start Hive Server2.

On the Hive Server2 host machine, run the following commands:

```
su - hive -c "/usr/hdp/current/hive-server2/bin/hiveserver2 >/
var/log/hive/hiveserver2.out 2> /var/log/hive/hiveserver2.log &"
```

5.14. Configure and Start Apache Oozie



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore `oozie-site.xml` from your backup to the `conf` directory on each oozie server and client.
2. Copy the JDBC jar to `libext-customer`.

- a. Create the `/usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer` directory.

```
cd /usr/hdp/2.5.0.0-<version>/oozie-server
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the `libext-customer` directory.

```
chmod -R 777 /usr/hdp/2.5.0.0-<version>/oozie-server/libext-
customer
```

3. Copy these files to the `libext-customer` directory

```
cp /usr/hdp/2.5.0.0-<version>/hadoop-client/lib/
hadoop*lzo*.jar /usr/hdp/current/oozie-server/libext-customer
```

```
cp /usr/share/HDP-oozie/ext.zip /usr/hdp/2.5.0.0-<version>/
oozie-server/libext-customer/
```

Also, copy Oozie db jar in libext-customer.

4. Extract share-lib.

```
cd /usr/hdp/2.5.0.0-<version>/oozie
tar xzvf /usr/hdp/2.5.0.0-<version>/oozie/oozie-sharelib.tar.gz
su - hdfs -c "hdfs dfs -mkdir -p /user/oozie"
su - hdfs -c "hdfs dfs -copyFromLocal /usr/hdp/current/oozie/
share /user/oozie/."
```

You can expect to see complaints that some files already exist. Delete any existing / oozie/share and replace it with the newly-extracted files.

```
su - hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie" su -l
hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

5. if a previous version of Oozie was created using auto schema creation, run the following SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

6. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh
upgrade -run"
```

7. As root, prepare the Oozie WAR file.

```
su - oozie -c "/usr/hdp/current/oozie/bin/oozie-setup.sh
prepare-war -d /usr/hdp/current/oozie/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar
New Oozie WAR file with added 'JARS' at /var/lib/oozie/oozie-server/webapps/
oozie.war
```

8. Add the following property to oozie-log4j.properties:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where \${oozie.instance.id} is determined by oozie, automatically.

9. Start Oozie as the Oozie user:

```
su - oozie -c "cd /grid/0/var/log/oozie; /usr/hdp/current/oozie/
oozie-server/bin/catalina.sh /usr/hdp/current/oozie/oozie-
```



```
server/bin/setclasspath.sh /usr/hdp/current/oozie-server/bin/
oozied.sh start"
```

10. Check processes.

```
ps -ef | grep -i oozie
```

5.15. Configure and Start Apache WebHCat (Templeton)



Note

The `su` commands in this section use keywords to represent the Service user. For example, "webhcat" is used to represent the Apache WebHCat Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. You must replace your configuration after upgrading. Copy `/etc/webhcat/conf` from the template to the `conf` directory in webhcat hosts.
2. Modify the WebHCat configuration files.

- Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User`. In this example, `hdfs`:

```
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /hdp/apps/2.5.0.0-<$version>/pig/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.gz /hdp/apps/2.5.0.0-<$version>/hive/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/pig.tar.gz
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/hive.tar.gz
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz
hdfs dfs -chown -R hdfs:hadoop /hdp
```

- Update the following properties in the `webhcat-site.xml` configuration file, as their values have changed:

```
<property>
  <name>templeton.pig.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/pig/pig.tar.gz</value>
</property>

<property>
  <name>templeton.hive.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/hive/hive.tar.gz</value>
</property>

<property>
  <name>templeton.streaming.jar</name>
```

```

<value>hdfs:///hdp/apps/${hdp.version}/mapreduce/hadoop-streaming.jar</value>
<description>The hdfs path to the Hadoop streaming jar file.</description>
</property>

<property>
  <name>templeton.sqoop.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/sqoop/sqoop.tar.gz</value>
  <description>The path to the Sqoop archive.</description>
</property>

<property>
  <name>templeton.sqoop.path</name>
  <value>sqoop.tar.gz/sqoop/bin/sqoop</value>
  <description>The path to the Sqoop executable.</description>
</property>

<property>
  <name>templeton.sqoop.home</name>
  <value>sqoop.tar.gz/sqoop</value>
  <description>The path to the Sqoop home in the exploded archive.</description>
</property>

```



Note

You do not need to modify `${hdp.version}`.

- Remove the following obsolete properties from `webhcat-site.xml`:

```

<property>
  <name>templeton.controller.map.mem</name>
  <value>1600</value>
  <description>Total virtual memory available to map tasks.</description>
</property>

<property>
  <name>hive.metastore.warehouse.dir</name>
  <value>/path/to/warehouse/dir</value>
</property>

```

- Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHcat to impersonate your additional HDP 2.5 groups and hosts:

```

<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>
</property>

<property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>
</property>

```

Where:

`hadoop.proxyuser.hcat.group`

is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

```
hadoop.proxyuser.hcat.hosts
```

A comma-separated list of the hosts that are allowed to submit requests by 'hcat'.

3. Start WebHCat:

```
su - webhcat -c "/usr/hdp/current/hive-hcatalog/sbin/
webhcat_server.sh start"
```

4. Smoke test WebHCat.

- At the WebHCat host machine, run the following command:

```
http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

- If you are using a secure cluster, run the following command:

```
curl --negotiate -u: http://cluster.$PRINCIPAL.$REALM:50111/
templeton/v1/status {"status":"ok","version":"v1"}
[machine@acme]$
```

5.16. Configure and Start Apache Pig

1. Replace your configuration after upgrading. Copy `/etc/pig/conf` from the template to the `conf` directory in pig hosts.
2. To validate the Apache Pig upgrade, complete the following steps:
 - a. On the host machine where Pig is installed, run the following commands:

```
su - $HDFS_USER/usr/hdp/current/hadoop/bin/hadoop
fs -copyFromLocal /etc/passwd passwd
```

- b. Create a Pig script file named `/tmp/id.pig` that contains the following Pig Latin commands:

```
A = load 'passwd' using PigStorage(':');B = foreach A generate $0 as id;
store B into '/tmp/id.out';
```

- c. Run the Pig script:

```
su - $HDFS_USER
pig -l /tmp/pig.log /tmp/id.pig
```

5.17. Configure and Start Apache Sqoop



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hdfs" is used to represent the HDFS Service user. If you are using

another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

1. Replace your configuration after upgrading. Copy `/etc/sqoop/conf` from the template to the conf directory in sqoop hosts.
2. Upload the Apache Sqoop tarball to HDFS. As the `<HDFS_User>`, for example 'hdfs':

```
su - hdfs -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"
```

```
su - hdfs -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"
```

```
su - hdfs -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"
```

```
su - hdfs -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"
```

```
su - hdfs -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz" hdfs
```

3. To validate the Sqoop upgrade, run the following command. You should see the Sqoop version information displayed.

```
sqoop version | grep 'Sqoop [0-9].*'
```

Because Sqoop is a client tool with no server component, you need to run your own jobs to further validate the upgrade.

5.18. Configure, Start, and Validate Apache Flume

1. If you have not already done so, upgrade Apache Flume. On the Flume host machine, run the following command:

- For **RHEL/CentOS/Oracle Linux**:

```
yum upgrade flume
```

- For **SLES**:

```
zypper update flume
```

```
zypper remove flume
```

```
zypper se -s flume
```

You should see Flume in the output.

Install Flume:

```
zypper install flume
```

- For **Ubuntu/Debian**:

```
apt-get install flume
```

2. To confirm that Flume is working correctly, create an example configuration file. The following snippet is a sample configuration that can be set using the properties file. For more detailed information, see the "Flume User Guide."

```
agent.sources = pstream
agent.channels = memoryChannel
agent.channels.memoryChannel.type = memory

agent.sources.pstream.channels = memoryChannel
agent.sources.pstream.type = exec
agent.sources.pstream.command = tail -f /etc/passwd

agent.sinks = hdfsSink
agent.sinks.hdfsSink.type = hdfs
agent.sinks.hdfsSink.channel = memoryChannel
agent.sinks.hdfsSink.hdfs.path = hdfs://tmp/flumetest
agent.sinks.hdfsSink.hdfs.fileType = SequenceFile
agent.sinks.hdfsSink.hdfs.writeFormat = Text
```

The source here is defined as an exec source. The agent runs a given command on startup, which streams data to stdout, where the source gets it. The channel is defined as an in-memory channel and the sink is an HDFS sink.

3. Given this configuration, you can start Flume as follows:

```
$ bin/flume-ng agent --conf ./conf --conf-file example.conf --name a1 -
Dflume.root.logger=INFO,console
```



Note

The directory specified for `--conf` argument would include a shell script `flume-env.sh` and potentially a `log4j` properties file. In this example, we pass a Java option to force Flume to log to the console and we go without a custom environment script.

4. After validating data in `hdfs://tmp/flumetest`, stop Flume and restore any backup files. Copy `/etc/flume/conf` to the `conf` directory in Flume hosts.

5.19. Configure, Start, and Validate Apache Mahout

Replace your configuration after upgrading. Copy `/etc/mahout/conf` from the template to the `conf` directory in mahout hosts.

To validate Apache Mahout:

1. Create a test user named "testuser" in the Linux cluster and in HDFS, and log in as that user.
2. Export the required environment variables for Mahout:

```
export JAVA_HOME="your_jdk_home_install_location_here"
export HADOOP_HOME=/usr/hdp/current/hadoop-client
export MAHOUT_HOME=/usr/hdp/current/mahout-client
export PATH="$PATH":$HADOOP_HOME/bin:$MAHOUT_HOME/bin
export CLASSPATH="$CLASSPATH":$MAHOUT_HOME
```

3. Upload a few megabytes of natural-language plain text to the Linux server as `/tmp/sample-test.txt`.
4. Transfer the `sample-test.txt` file to a subdirectory of the testusers's HDFS home directory.

```
hdfs dfs -mkdir /user/testuser/testdata
hdfs dfs -put /tmp/sample-test.txt /user/testuser/testdata
```

5. Enter the mahout command to convert the plain text file `sample-test.txt` into a sequence file stored in the output directory `mahouttest`:

```
mahout seqdirectory --input /user/testuser/testdata --output /user/testuser/
mahouttest -ow --charset utf-8
```

5.20. Configure and Start Hue



Note

The `su` commands in this section use keywords to represent the Service user. For example, "hue" is used to represent the Hue Service user. If you are using another name for your Service users, you need to substitute your Service user name in each of the `su` commands.

For HDP 2.5, use the Hue version shipped with HDP 2.5. If you have a previous version of Hue, use the following steps to upgrade Hue.

1. Migrate the `hue.ini` setting from your old `hue.ini` configuration file to new `hue.ini` configuration file.
2. If you are using the embedded SQLite database, remember to restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database.

For example:

```
su - hue

cd /var/lib/hue

mv desktop.db desktop.db.old

sqlite3 desktop.db < ~/hue_backup/desktop.bak

exit
```

3. Synchronize the database.

```
cd /usr/lib/hue  
  
source ./build/env/bin/activate  
  
hue syncdb  
  
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

5.21. Finalize the Upgrade

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.

1. Verify your filesystem health before finalizing the upgrade and removing the old config.
2. As the \$HDFS_USER, enter:

```
hdfs dfsadmin -finalizeUpgrade
```

5.22. Install New HDP 2.5 Services

Install new HDP 2.5.0 Services (see the [Non-Ambari Cluster Installation Guide](#)):

- Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
- SmartSense – a next generation subscription model that features upgrade and configuration recommendations.

6. Upgrade from HDP 1.3 to HDP 2.5.0 Manually

This chapter provides instruction for how to manually upgrade to HDP 2.5.0 from the HDP 1.3 release.

These instructions cover the upgrade between two major releases. If you need to upgrade between two maintenance releases, follow the upgrade instructions in the HDP Release Notes.

HDP 2.2 supports side-by-side installation of HDP 2.2 and above releases, which lets you perform rolling upgrades on your cluster and improve execution times for in-place upgrade. To support side-by-side installation, the HDP package version naming convention for both RPMs and Debs has changed to include the HDP 2.2 product version. For example, `hadoop-hdfs` is now `hadoop-2.3.0.0-hdfs`. HDP 2.2 marks the first release where HDP rpms, debs, and directories contain versions in the names to permit side-by-side installations of later HDP releases.

To select from the releases you have installed side-by-side, Hortonworks provides `hdp-select`, a command that lets you select the active version of HDP from the versions you have selected.

The HDP packages for a complete installation of HDP 2.5 consumes about 6.5 GB of disk space.



Warning

Until the upgrade is finalized, no HDFS data is deleted from the cluster. Be sure to review your capacity and ensure that you have extra space available during the upgrade window.

Use the following instructions to upgrade to the latest release of HDP from HDP 1.3:

1. Download HDP 2.5
2. Getting Ready to Upgrade
3. Upgrade Hadoop
4. Migrate the HDP Configurations
5. Create Local Directories
6. Start HDFS
7. Upgrade Apache ZooKeeper
8. Upgrade Apache HBase
9. Upgrade Apache Hive and Apache HCatalog
10. Upgrade Apache Oozie
11. Upgrade Apache WebHCat (Templeton)

- 12.Upgrade Apache Pig
- 13.Upgrade Apache Sqoop
- 14.Upgrade Apache Flume
- 15.Upgrade Apache Mahout
- 16.Upgrade Hue
- 17.Finalize the Upgrade
- 18.Install new HDP 2.5 Services

6.1. Getting Ready to Upgrade

HDP Stack upgrade involves removing HDP 1.x MapReduce and replacing it with HDP 2.x YARN and MapReduce2. Before you begin, review the upgrade process and complete the backup steps.



Note

You must use **kinit** before running the commands as any particular user.

1. Hardware recommendations

Although there is no single hardware requirement for installing HDP, there are some basic guidelines. The HDP packages for a complete installation of HDP 2.5 consumes about 6.5 GB of disk space.

2. Back up the following HDP 1.x directories:

- `/etc/hadoop/conf`
- `/etc/hbase/conf`
- `/etc/hcatalog/conf` (Note: With HDP 2.5, `/etc/hcatalog/conf` is divided into `/etc/hive-hcatalog/conf` and `/etc/hive-webhcat/conf`. You cannot use `/etc/hcatalog/conf` in HDP 2.5.)
- `/etc/hive/conf`
- `/etc/pig/conf`
- `/etc/sqoop/conf`
- `/etc/flume/conf`
- `/etc/mahout/conf`
- `/etc/oozie/conf`
- `/etc/hue/conf`
- `/etc/zookeeper/conf`

3. **Optional:** Back up your userlogs directories, `${mapred.local.dir}/userlogs`.

Run the `fsck` command as the HDFS Service user and fix any errors. (The resulting file contains a complete block map of the file system.) For example, in a situation where clusters are unsecure and Kerberos credentials are not required for the HDFS user:

```
su -l <HDFS_USER>
```

```
hadoop fsck / -files -blocks -locations > dfs-old-fsck-1.log
```

where `$HDFS_USER` is the HDFS Service user. For example, `hdfs`.

4. As the user running the HDFS service (by default, the user is `hdfs`), run the following commands:

- Capture the complete namespace of the file system. (The following command does a recursive listing of the root file system.)

```
su -l <HDFS_USER>
hadoop dfs -lsr / > dfs-old-lsr-1.log
```

where `$HDFS_USER` is the HDFS Service user. For example, `hdfs`.

- Run the report command to create a list of DataNodes in the cluster.

```
su -l <HDFS_USER>
hadoop dfsadmin -report > dfs-old-report-1.log
```

where `$HDFS_USER` is the HDFS Service user. For example, `hdfs`

- **Optional:** You can copy all or unrecoverable only data stored in HDFS to a local file system or to a backup instance of HDFS.
- **Optional:** You can also repeat the steps 3 (a) through 3 (c) and compare the results with the previous run to ensure the state of the file system remained unchanged.

5. HBase 0.96.0 and subsequent releases discontinue support for the HFileV1 file format, a common format prior to HBase 0.94. Before you upgrade, check for V1-format files as follows:

- [Download](#) the Apache 0.94.24+HBase tarball in a machine. Run the binaries.
- On the machine running the HBase 0.94 binaries, point the `hbase-site.xml` configuration file to a 0.94 cluster.
- Check for HFiles in V1 format as follows:

```
./bin/hbase org.apache.hadoop.hbase.util.HFileV1Detector -p
<hbase root data path>
```

When you run the upgrade check, if "Count of HFileV1" returns any files, start the HBase shell to use major compaction for regions that have HFileV1 format. For example, the following sample output indicates that you need to compact two regions, `fa02dac1f38d03577bd0f7e666f12812` and `ecdd3eaee2d2fcf8184ac025555bb2af`:

```
Tables Processed:
```

```

hdfs://localhost:41020/myHBase/.META.
hdfs://localhost:41020/myHBase/usertable
hdfs://localhost:41020/myHBase/TestTable
hdfs://localhost:41020/myHBase/tCount of HFileV1: 2 HFileV1:
hdfs://localhost:41020/myHBase/usertable/fa02dac1f38d03577bd0f7e666f12812/
family/249450144068442524
hdfs://localhost:41020/myHBase/usertable/ecdd3eae2d2fcf8184ac025555bb2af/
family/249450144068442512

Count of corrupted files: 1
Corrupted Files:
hdfs://localhost:41020/myHBase/usertable/fa02dac1f38d03577bd0f7e666f12812/
family/1
Count of Regions with HFileV1: 2
Regions to Major Compact:
hdfs://localhost:41020/myHBase/usertable/fa02dac1f38d03577bd0f7e666f12812
hdfs://localhost:41020/myHBase/usertable/ecdd3eae2d2fcf8184ac025555bb2af

```

6. **Optional:** If you are upgrading HBase on a secure cluster, flush the ACL table by running the following HBase shell command as the \$HBase_User.

```
flush '_acl_'
```

7. Stop all HDP 1.3 services (including MapReduce) except HDFS:

- Stop Nagios. On the Nagios host machine, run the following command:

```
service nagios stop
```

- Stop Ganglia.

- Run this command on the Ganglia server host machine:

```
/etc/init.d/hdp-gmetad stop
```

- Run this command on all the nodes in your Hadoop cluster:

```
/etc/init.d/hdp-gmond stop
```

- Stop Oozie. On the Oozie server host machine, run the following command:

```
sudo su -l $OOZIE_USER -c "cd $OOZIE_LOG_DIR; /usr/lib/oozie/
bin/oozie-stop.sh"
```

where:

\$OOZIE_USER is the Oozie Service user. For example, oozie

\$OOZIE_LOG_DIR is the directory where Oozie log files are stored (for example: /var/log/oozie).

- Stop WebHCat. On the WebHCat host machine, run the following command:

```
su -l $WEBHCAT_USER -c "/usr/lib/hcatalog/sbin/
webhcat_server.sh stop"
```

where \$WEBHCAT_USER is the WebHCat Service user. For example, hcat.

- Stop Hive. On the Hive Metastore host machine and Hive Server2 host machine, run the following command:

```
ps aux | awk '{print $1,$2}' | grep hive | awk '{print $2}' |
xargs kill >/dev/null 2>&1
```

This stops the Hive Metastore and HCatalog services.

- Stop HBase.
 - Run these commands on all RegionServers:

```
su -l $HBASE_USER -c "/usr/lib/hbase/bin/hbase-daemon.sh --
config /etc/hbase/conf stop regionserver"
```

- Run these commands on the HBase Master host machine:

```
su -l $HBASE_USER -c "/usr/lib/hbase/bin/hbase-daemon.sh --
config /etc/hbase/conf stop master"
```

where `$HBASE_USER` is the HBase Service user. For example, `hbase`.

- Stop ZooKeeper. On the ZooKeeper host machine, run the following command:

```
su - $ZOOKEEPER_USER -c "export ZOOCFGDIR=/etc/zookeeper/conf ;
export ZOOCFG=zoo.cfg ;source /etc/zookeeper/conf/zookeeper-
env.sh ; /usr/lib/zookeeper-server/bin/zkServer.sh stop"
```

where `$ZOOKEEPER_USER` is the ZooKeeper Service user. For example, `zookeeper`.

- Stop MapReduce

- Run these commands on all TaskTrackers slaves:

```
su -l $MAPRED_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh
--config /etc/hadoop/conf stop tasktracker"
```

- Run these commands on the HistoryServer host machine:

```
su -l $MAPRED_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh
--config /etc/hadoop/conf stop historyserver"
```

- Run these commands on the node running the JobTracker host machine:

```
su -l $MAPRED_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh
--config /etc/hadoop/conf stop jobtracker"
```

where `$MAPRED_USER` is the MapReduce Service user. For example, `mapred`.

8. As the HDFS user, save the namespace by executing the following command:

```
su -l <HDFS_USER>
```

```
hadoop dfsadmin -safemode enter
```

```
hadoop dfsadmin -saveNamespace
```

9. Backup your NameNode metadata.

- Copy the following checkpoint files into a backup directory. If these directories do not exist, create them.
 - dfs.name.dir/edits
 - dfs.name.dir/image/fsimage
 - dfs.name.dir/current/fsimage
- Store the layoutVersion of the namenode.

```
${dfs.name.dir}/current/VERSION
```

10.If you have a prior HDFS upgrade in progress, finalize it if you have not done so already.

```
su -l <HDFS_USER>
```

```
hadoop dfsadmin -finalizeUpgrade
```

11.Optional: Back up the Hive Metastore database.

These instructions are provided for your convenience. Please check your database documentation for the latest backup instructions.

The following instructions are provided for your convenience. For the latest backup instructions, please check your database documentation.

Table 6.1. Hive Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sql</pre> <p>For example:</p> <pre>mysqldump hive > /tmp/mydir/ backup_hive.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sql</pre> <p>For example:</p> <pre>mysql <dbname> </tmp/mydir/ backup_hive.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hive > / tmp/mydir/ backup_hive.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres psql hive </tmp/ mydir/ backup_hive.sql</pre>
Oracle	<p>Connect to the Oracle database using sqlplus.</p> <p>Export the database:</p> <pre>exp username/ password@database full=yes file=output_file.dmp</pre>	<p>Import the database: imp</p> <pre>username/password@database ile=input_file.dmp</pre>

12.Optional: Back up the Oozie Metastore database.

The following instructions are provided for your convenience. For the latest backup instructions, please check your database documentation.

Table 6.2. Oozie Metastore Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sql</pre> <p>For example:</p> <pre>mysqldump oozie > /tmp/ mydir/backup_oozie.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sql</pre> <p>For example:</p> <pre>mysql oozie </tmp/mydir/ backup_oozie.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump oozie > /tmp/mydir/ backup_oozie.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sql</pre> <p>For example:</p> <pre>sudo -u postgres psql oozie </tmp/mydir/ backup_oozie.sql</pre>

13.Optional: Back up the Hue database.

The following instructions are provided for your convenience. For the latest backup instructions, please see your database documentation. For database types that are not listed below, follow your vendor-specific instructions.

Table 6.3. Hue Database Backup and Restore

Database Type	Backup	Restore
MySQL	<pre>mysqldump \$dbname > \$outputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysqldump hue > /tmp/mydir/ backup_hue.sql</pre>	<pre>mysql \$dbname < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>mysql hue < /tmp/mydir/ backup_hue.sql</pre>
PostgreSQL	<pre>sudo -u \$username pg_dump \$databasename > \$outputfilename.sql sbr</pre> <p>For example:</p> <pre>sudo -u postgres pg_dump hue > / tmp/mydir/backup_hue.sql</pre>	<pre>sudo -u \$username psql \$databasename < \$inputfilename.sqlsbr</pre> <p>For example:</p> <pre>sudo -u postgres psql hue < /tmp/ mydir/backup_hue.sql</pre>
Oracle	<p>Connect to the Oracle database using sqlplus. Export the database.</p> <p>For example:</p> <pre>exp username/password@database full=yes file=output_file.dmp mysql \$dbname < \$inputfilename.sqlsbr</pre>	<p>Import the database:</p> <p>For example:</p> <pre>imp username/password@database file=input_file.dmp</pre>
SQLite	<pre>/etc/init.d/hue stop</pre> <pre>su \$HUE_USER</pre>	<pre>/etc/init.d/hue stop</pre> <pre>cd /var/lib/hue</pre>

Database Type	Backup	Restore
	<pre>mkdir ~/hue_backup sqlite3 desktop.db .dump > ~/hue_backup/desktop.bak /etc/init.d/hue start</pre>	<pre>mv desktop.db desktop.db.old sqlite3 desktop.db < ~/hue_backup/desktop.bak /etc/init.d/hue start</pre>

14. Stop HDFS.

- a. Run these commands on all DataNodes:

```
su -l $HDFS_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh --config /etc/hadoop/conf stop datanode"
```

- b. Run these commands on the Secondary NameNode host machine:

```
su -l $HDFS_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh --config /etc/hadoop/conf stop secondarynamenode"
```

- c. Run these commands on the NameNode host machine:

```
su -l $HDFS_USER -c "/usr/lib/hadoop/bin/hadoop-daemon.sh --config /etc/hadoop/conf stop namenode"
```

where \$HDFS_USER is the HDFS Service user. For example, hdfs.

15. Verify that edit logs in \${dfs.namenode.name.dir}/current/edits* are empty.

- a. Run the following command:

```
hdfs oev -i ${dfs.namenode.name.dir}/current/edits_inprogress_* -o edits.out
```

- b. Verify edits.out file. It should only have OP_START_LOG_SEGMENT transaction. For example:

```
<?xml version="1.0" encoding="UTF-8"?>
<EDITS>
<EDITS_VERSION>-56</EDITS_VERSION>
<RECORD>
<OPCODE>OP_START_LOG_SEGMENT</OPCODE>
<DATA>
<TXID>5749</TXID>
</DATA>
</RECORD>
```

- c. If edits.out has transactions other than OP_START_LOG_SEGMENT, run the following steps and then verify edit logs are empty.

- Start the existing version NameNode.
- Ensure there is a new FS image file.
- Shut the NameNode down.

16. Rename or delete any paths that are reserved in the new version of HDFS.

When upgrading to a new version of HDFS, it is necessary to rename or delete any paths that are reserved in the new version of HDFS. If the NameNode encounters a reserved path during upgrade, it prints an error such as the following:

```
/.reserved is a reserved path and .snapshot is a reserved path
component in this version of HDFS. Please rollback and delete or
rename this path, or upgrade with the -renameReserved key-value
pairs option to automatically rename these paths during upgrade.
```

Specifying `-upgrade -renameReserved` optional key-value pairs causes the NameNode to automatically rename any reserved paths found during startup. For example, to rename all paths named `.snapshot` to `.my-snapshot` and change paths named `.reserved` to `.my-reserved`, a user would specify:

```
-upgrade -renameReserved .snapshot=.my-snapshot,.reserved=.my-
reserved.
```

If no key-value pairs are specified with `-renameReserved`, the NameNode suffixes reserved paths with `.<LAYOUT-VERSION>.UPGRADE_RENAMED`, for example:

```
.snapshot.-51.UPGRADE_RENAMED.
```



Note

We recommend that you perform a `-saveNamespace` before renaming paths (running `-saveNamespace` appears in a previous step in this procedure). This is because a data inconsistency can result if an edit log operation refers to the destination of an automatically renamed file.

Also note that running `-renameReserved` renames all applicable existing files in the cluster. This may impact cluster applications.

6.2. Upgrade HDP 1.3 Components

The upgrade process to HDP 2.5 involves the following steps. Select your OS:

RHEL/CentOS/Oracle 6

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 1.3 components. This command uninstalls the HDP 1.3 components. It leaves the user data, metadata and modified configurations in place and does not delete them:

```
yum erase "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*"
"hive*" "flume*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old `hdp.repo` file:


```
rm /etc/yum.repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<version>/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0.0 Hortonworks Data Platform Version - HDP-2.5.0.0
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 1.3 components:

```
yum install "hadoop" "hadooplzo" "webhcat" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "flume"
"mahout" "hue" "knox" "hdp_mon_nagios_addons"
```

6. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

RHEL/CentOS/Oracle

1. On all hosts, clean the yum repository.

```
yum clean all
```

2. Remove your old HDP 1.3 components. This command uninstalls the HDP 1.3 components. It leaves the user data and metadata but deletes your modified configurations:

```
yum erase "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*"
"hive*" "flume*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old hdp.repo file:

```
rm /etc/yum.repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/centos<versions>/2.x/updates/2.5.0.0/hdp.repo -O /etc/yum.repos.d/hdp.repo
```

- Confirm the HDP repository is configured.

```
yum repolist
```

You should see something like this. Verify that you have the HDP-2.5.0.0 directory:

```
Loaded plugins: fastestmirror, security
Loading mirror speeds from cached hostfile
* base: mirrors.cat.pdx.edu
* extras: linux.mirrors.es.net
* updates: mirrors.usc.edu
repo id repo name status
HDP-2.5.0.0 Hortonworks Data Platform Version - HDP-2.5.0.0
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 1.3 components:

```
yum install "hadoop" "hadooplzo" "webhcat" "oozie" "collectd"
"gccxml" "pig" "sqoop" "zookeeper" "hbase" "hive" "flume"
"mahout" "hue" "knox" "hdp_mon_nagios_addons"
```

6. Verify that the components were upgraded.

```
yum list installed | grep HDP-<old.stack.version.number>
```

No component file names should appear in the returned list.

SLES

1. On all hosts, clean the yum repository.

```
zypper clean -all
```

2. Remove your old HDP 1.3 components. This command uninstalls the HDP 1.3 components. It leaves the user data and metadata but deletes your modified configurations:

```
zypper rm "hadoop*" "webhcat*" "hcatalog*" "oozie*" "collectd*"
"gccxml*" "pig*" "hdfs*" "sqoop*" "zookeeper*" "hbase*"
"hive*" "flume*" "mahout*" "hue" "hue-common" "hue-shell"
"hdp_mon_nagios_addons"
```

3. Remove your old hdp.repo file:

```
rm /etc/zypp/repos.d/hdp.repo
```

4. Install the HDP 2.5 repo:

- Download the hdp.repo file:

```
wget -nv http://public-repo-1.hortonworks.com/HDP/
sles11sp1/2.x/updates/2.5.0.0/hdp.repo -O /etc/zypp/repos.d/
hdp.repo
```

5. Install the HDP 2.5 versions of the components that you want to upgrade. For example, if you installed and want to upgrade all HDP 1.3 components:

```
zypper install "hadoop" "webhcat" "oozie" "collectd" "gccxml"
"pig" "hdfs" "sqoop"
```

```
"zookeeper" "hbase" "hive" "flume" "mahout" "hue"
"hdp_mon_nagios_addons"
```

```
zypper install webhcat-tar-hive webhcat-tar-pig
```

```
zypper up -r HDP-2.5.0.0
```

```
zypper install oozie-client
```

6. Verify that the components were upgraded. For example, to verify hdfs, hive, and hcatlog:

```
rpm -qa | grep hdfs, && rpm -qa | grep hive && rpm -qa | grep
hcatalog
```

No component files names should appear in the returned list.

6.3. Symlink Directories with hdp-select



Warning

HDP 2.5 installs hdp-select automatically with the installation or upgrade of the first HDP component. If you have not already upgraded ZooKeeper, hdp-select has not been installed.

To prevent version-specific directory issues for your scripts and updates, Hortonworks provides hdp-select, a script that symlinks directories to hdp-current and modifies paths for configuration directories.

- Before you run hdp-select, you must remove one link:

```
rm /usr/bin/oozie
```

- Run hdp-select set all on your NameNode and all your DataNodes:

```
hdp-select set all 2.5.0.0-<$version>
```

For example:

```
/usr/bin/hdp-select set all 2.5.0.0-2
```

6.4. Configure and Start Apache ZooKeeper



Tip

If you are running a highly available cluster, upgrade Apache ZooKeeper **before** you upgrade HDFS. This best practice lets the upgraded ZKFC work with your primary NameNode and your Standby NameNode.

RHEL/CentOS/Oracle Linux

1. Replace your configuration after upgrading ZooKeeper. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all the ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "export ZOOCFGDIR=/usr/hdp/current/zookeeper-server/conf ; export ZOOCFG=zoo.cfg; source /usr/hdp/current/zookeeper-server/conf/zookeeper-env.sh ; /usr/hdp/current/zookeeper-server/bin/zkServer.sh start"

/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc
```

SLES

1. Replace your configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all the ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "export ZOOCFGDIR=/usr/hdp/current/zookeeper-server/conf ; export ZOOCFG=zoo.cfg; source /usr/hdp/current/zookeeper-server/conf/zookeeper-env.sh ; /usr/hdp/current/zookeeper-server/bin/zkServer.sh start"

/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc
```

Ubuntu/Debian

1. Replace your configuration after upgrading. Replace the ZooKeeper template configuration in `/etc/zookeeper/conf`.
2. Start ZooKeeper.

On all the ZooKeeper server host machines, run the following command to start ZooKeeper and the ZKFC:

```
su - zookeeper -c "export ZOOCFGDIR=/usr/hdp/current/zookeeper-server/conf ; export ZOOCFG=zoo.cfg; source /usr/hdp/current/
```

```
zookeeper-server/conf/zookeeper-env.sh ; /usr/hdp/current/  
zookeeper-server/bin/zkServer.sh start "  
  
/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc
```

6.5. Configure and Start Hadoop

RHEL/CentOS/Oracle Linux

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in yarn-site.xml and mapred-site.xml
2. Paths have changed in . make sure you remove lines from hadoop-env.sh such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your hadoop-env.sh file, the lzo compression code does not load as this is not where lzo is installed.

SLES

1. Use the HDP Utility script to calculate memory configuration settings. You must update the memory/cpu settings in yarn-site.xml and mapred-site.xml
2. Paths have changed in HDP 2.5. make sure you remove lines from hadoop-env.sh such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-  
amd64-64
```

If you leave these paths in your hadoop-env.sh file, the lzo compression code does not load as this is not where lzo is installed.

6.6. Migrate the HDP Configurations

Configurations and configuration file names have changed between HDP 1.3.2 (Hadoop 1.2.x) and HDP 2.5. To upgrade to HDP 2.5, back up your current configuration files, download the new HDP 2.5 files, and compare. The following tables provide mapping information to make the comparison between releases easier.

To migrate the HDP Configurations

1. Edit `/etc/hadoop/conf/core-site.xml` and set `hadoop.rpc.protection` from none to authentication.



Note

Hadoop lets cluster administrators control the quality of protection in the configuration parameter “hadoop.rpc.protection” in core-site.xml. It is an optional parameter in HDP 2.5. If not present, the default QOP setting of “auth” is used, which implies “authentication only”.

Valid values for this parameter are: "authentication" : Corresponds to "auth"
 "integrity" : Corresponds to "auth-int" "privacy" : Corresponds to "auth-conf"

The default setting is authentication-only because integrity checks and encryption are a performance cost.

2. Copy your /etc/hcatalog/conf configurations to /etc/hive-hcatalog/conf and /etc/hive-webhcat as appropriate.
3. Copy log4j.properties from the hadoop config directory of the companion files to /etc/hadoop/conf. The file should have owners and permissions similar to other files in /etc/hadoop/conf.
4. Download the your HDP 2.5.x companion files from Download Companion Files and migrate your HDP 1.x configuration.
5. Copy these configurations to all nodes in your clusters.
 - /etc/hadoop/conf
 - /etc/hbase/conf
 - /etc/hcatalog/conf
 - /etc/hive/conf
 - /etc/pig/conf
 - /etc/sqoop/conf
 - /etc/flume/conf
 - /etc/mahout/conf
 - /etc/oozie/conf
 - /etc/zookeeper/conf



Note

Upgrading the repo using `yum` or `zypper` resets all configurations. Prepare to replace these configuration directories each time you perform a `yum` or `zypper` `rmgrade`.

6. Review the following HDP 1.3.2 Hadoop Core configurations and the new configurations or locations in HDP 2.5.x.

Table 6.4. HDP 1.3.2 Hadoop Core Site (core-site.xml)

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
<code>fs.default.name</code>	<code>core-site.xml</code>	<code>fs.defaultFS</code>	<code>core-site.xml</code>
<code>fs.checkpoint.dir</code>	<code>core-site.xml</code>	<code>dfs.namenode.checkpoint.dir</code>	<code>hdfs-site.xml</code>

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
fs.checkpoint.edits.dir	core-site.xml	dfs.namenode.checkpoint.edits.dir	hdfs-site.xml
fs.checkpoint.period	core-site.xml	dfs.namenode.checkpoint.period	hdfs-site.xml
io.bytes.per.checksum	core-site.xml	dfs.bytes-per-checksum	hdfs-site.xml
dfs.df.interval	hdfs-site	fs.df.interval	core-site.xml
hadoop.native.lib	core-site.xml	io.native.lib.available	core-site.xml
hadoop.configured.node.mapping	–	net.topology.configured.node.mapping	core-site.xml
topology.node.switch.mapping.impl	core-site.xml	net.topology.node.switch.mapping.impl	core-site.xml
topology-script.file.name	core-site.xml	net.topology.script.file.name	core-site.xml
topology.script.number.args	core-site.xml	net.topology.script.number.args	core-site.xml



Note

The `hadoop.rpc.protection` configuration property in `core-site.xml` needs to specify authentication, integrity and/or privacy. No value defaults to authentication, but an invalid value such as "none" causes an error.

- Review the following 1.3.2 HDFS site configurations and their new configurations and files in HDP 2.x.

Table 6.5. HDP 1.3.2 Hadoop Core Site (hdfs-site.xml)

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
dfs.block.size	hdfs-site.xml	dfs.blocksize	hdfs-site.xml
dfs.write.packet.size	hdfs-site.xml	dfs.client-write-packet-size	hdfs-site.xml
dfs.https.client.keystore.resource	hdfs-site.xml	dfs.client.https.keystore.resource	hdfs-site.xml
dfs.https.need.client.auth	hdfs-site.xml	dfs.client.https.need-auth	hdfs-site.xml
dfs.read.prefetch.size	hdfs-site.xml	dfs.bytes-per-checksum	hdfs-site.xml
dfs.socket.timeout	hdfs-site.xml	dfs.client.socket-timeout	hdfs-site.xml
dfs.balance.bandwidthPerSec	hdfs-site.xml	dfs.datanode.balance.bandwidthPerSec	hdfs-site.xml
dfs.data.dir	hdfs-site.xml	dfs.datanode.data.dir	hdfs-site.xml
dfs.datanode.max.xcievers	hdfs-site.xml	dfs.datanode.max.transfer.threads	hdfs-site.xml
session.id	hdfs-site.xml	dfs.metrics.session-id	hdfs-site.xml
dfs.access.time.precision	hdfs-site.xml	dfs.namenode.accesstime.precision	hdfs-site.xml

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
dfs.backup.address	hdfs-site.xml	dfs.namenode.backup.address	hdfs-site.xml
dfs.backup.http.address	hdfs-site.xml	dfs.namenode.backup.http-address	hdfs-site.xml
fs.checkpoint.dir	hdfs-site.xml	dfs.namenode.checkpoint.dir	hdfs-site.xml
fs.checkpoint.edits.dir	hdfs-site.xml	dfs.namenode.checkpoint.edits.dir	hdfs-site.xml
fs.checkpoint.period	hdfs-site.xml	dfs.namenode.checkpoint.period	hdfs-site.xml
dfs.name.edits.dir	hdfs-site.xml	dfs.namenode.edits.dir	hdfs-site.xml
heartbeat.recheck.interval	hdfs-site.xml	dfs.namenode.heartbeat.recheck-interval	hdfs-site.xml
dfs.http.address	hdfs-site.xml	dfs.namenode.http-address	hdfs-site.xml
dfs.https.address	hdfs-site.xml	dfs.namenode.https-address	hdfs-site.xml
dfs.max.objects	hdfs-site.xml	dfs.namenode.max.objects	hdfs-site.xml
dfs.name.dir	hdfs-site.xml	dfs.namenode.name.dir	hdfs-site.xml
dfs.name.dir.restore	hdfs-site.xml	dfs.namenode.name.dir.restore	hdfs-site.xml
dfs.replication.considerLoad	hdfs-site.xml	dfs.namenode.replication.considerLoad	hdfs-site.xml
dfs.replication.interval	hdfs-site.xml	dfs.namenode.replication.interval	hdfs-site.xml
dfs.max-repl-streams	hdfs-site.xml	dfs.namenode.replication.max-streams	hdfs-site.xml
dfs.replication.min	hdfs-site.xml	dfs.namenode.replication.min	hdfs-site.xml
dfs.replication.pending.timeout.sec	hdfs-site.xml	dfs.namenode.replication.pending.timeout-sec	hdfs-site.xml
dfs.safemode.extension	hdfs-site.xml	dfs.namenode.safemode.extension	hdfs-site.xml
dfs.safemode.threshold.pct	hdfs-site.xml	dfs.namenode.secondary.threshold-pct	
dfs.secondary.http.address	hdfs-site.xml	dfs.namenode.secondary.http-address	hdfs-site.xml
dfs.permissions	hdfs-site.xml	dfs.permissions.enabled	hdfs-site.xml
dfs.permissions.supergroup	hdfs-site.xml	dfs.permissions.superusergroup	hdfs-site.xml
dfs.df.interval	hdfs-site.xml	fs.df.interval	core-site.xml
dfs.umaskmode	hdfs-site.xml	fs.permissions.umask-mode	hdfs-site.xml

8. Review the following HDP 1.3.2 MapReduce Configs and their new HDP 2.x mappings.

Table 6.6. HDP 1.3.2 Configs now in Capacity Scheduler for HDP 2.x (mapred-site.xml)

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
mapred.map.child.java.opts	mapred-site.xml	mapreduce.map.java.opts	mapred-site.xml
mapred.job.map.memory.mb	mapred-site.xml	mapred.job.map.memory.mb	mapred-site.xml
mapred.reduce.child.java.opts	mapred-site.xml	mapreduce.reduce.java.opts	mapred-site.xml
mapreduce.job.reduce.memory.mb	mapred-site.xml	mapreduce.reduce.memory.mb	mapred-site.xml
security.task.umbilical.protocol.acl	mapred-site.xml	security.job.task.protocol.acl	mapred-site.xml

9. Review the following HDP 1.3.2 Configs and their new HDP 2.x Capacity Scheduler mappings.

Table 6.7. HDP 1.3.2 Configs now in capacity scheduler for HDP 2.x (capacity-scheduler.xml)

HDP 1.3.2 config	HDP 1.3.2 config file	HDP 2.5 config	HDP 2.5 config file
mapred.queue.names	mapred-site.xml	yarn.scheduler.capacity.root.queues	capacity-scheduler.xml
mapred.queue.default.acl-submit.job	mapred-queue-acls.xml	yarn.scheduler.capacity.root.default.acl_submit_jobs	capacity-scheduler.xml
mapred.queue.default.acl.administer-jobs	mapred-queue-acls.xml	yarn.scheduler.capacity.root.default.acl_administer_jobs	capacity-scheduler.xml
mapred.capacity-scheduler.queue.default.capacity	capacity-scheduler.xml	yarn.scheduler.capacity.root.default.capacity	capacity-scheduler.xml
mapred.capacity-scheduler.queue.default.user-limit-factor	capacity-scheduler.xml	yarn.scheduler.capacity.root.default.user-limit-factor	capacity-scheduler.xml
mapred.capacity-scheduler.queue.default.maximum-capacity	capacity-scheduler.xml	yarn.scheduler.capacity.root.default.maximum-capacity	capacity-scheduler.xml
mapred.queue.default.state	capacity-scheduler.xml	yarn.scheduler.capacity.root.default.state	capacity-scheduler.xml

10. Compare the following HDP 1.3.2 configs in `hadoop-env.sh` with the new configs in HDP 2.x.

Paths have changed in HDP 2.5 to `/usr/hdp/current`. You must remove lines such as:

```
export JAVA_LIBRARY_PATH=/usr/lib/hadoop/lib/native/Linux-
amd64-64
```

Table 6.8. HDP 1.3.2 Configs and HDP 2.x for hadoop-env.sh

HDP 1.3.2 config	HDP 2.5 config	Description
JAVA_HOME	JAVA_HOME	Java implementation to use
HADOOP_HOME_WARN_SUPPRESS	HADOOP_HOME_WARN_SUPPRESS	–
HADOOP_CONF_DIR	HADOOP_CONF_DIR	Hadoop configuration directory
not in hadoop-env.sh.	HADOOP_HOME	–
not in hadoop-env.sh.	HADOOP_LIBEXEC_DIR	–
HADOOP_NAMENODE_INIT_HEAPSIZE	HADOOP_NAMENODE_INIT_HEAPSIZE	–
HADOOP_OPTS	HADOOP_OPTS	Extra Java runtime options; empty by default
HADOOP_NAMENODE_OPTS	HADOOP_NAMENODE_OPTS	Command-specific options appended to HADOOP-OPTS
HADOOP_JOBTRACKER_OPTS	not in hadoop-env.sh.	Command-specific options appended to HADOOP-OPTS
HADOOP_TASKTRACKER_OPTS	not in hadoop-env.sh.	Command-specific options appended to HADOOP-OPTS
HADOOP_DATANODE_OPTS	HADOOP_DATANODE_OPTS	Command-specific options appended to HADOOP-OPTS
HADOOP_BALANCER_OPTS	HADOOP_BALANCER_OPTS	Command-specific options appended to HADOOP-OPTS
HADOOP_SECONDARYNAMENODE_OPTS	HADOOP_SECONDARYNAMENODE_OPTS	Command-specific options appended to HADOOP-OPTS
HADOOP_CLIENT_OPTS	HADOOP_CLIENT_OPTS	Applies to multiple commands (fs, dfs, fsck, distcp, etc.)
HADOOP_SECURE_DN_USER	not in hadoop-env.sh.	Secure datanodes, user to run the datanode as
HADOOP_SSH_OPTS	HADOOP_SSH_OPTS	Extra ssh options.
HADOOP_LOG_DIR	HADOOP_LOG_DIR	Directory where log files are stored in the secure data environment.
HADOOP_SECURE_DN_LOG_DIR	HADOOP_SECURE_DN_LOG_DIR	Directory where pid files are stored; /tmp by default.
HADOOP_PID_DIR	HADOOP_PID_DIR	Directory where pid files are stored, /tmp by default.
HADOOP_SECURE_DN_PID_DIR	HADOOP_SECURE_DN_PID_DIR	Directory where pid files are stored, /tmp by default.
HADOOP_IDENT_STRING	HADOOP_IDENT_STRING	String representing this instance of hadoop. \$USER by default
not in hadoop-env.sh.	HADOOP_MAPRED_LOG_DIR	–
not in hadoop-env.sh.	HADOOP_MAPRED_PID_DIR	–
not in hadoop-env.sh.	JAVA_LIBRARY_PATH	–
not in hadoop-env.sh.	JSVC_HOME	For starting the datanode on a secure cluster



Note

Some of the configuration settings refer to the variable `HADOOP_HOME`. The value of `HADOOP_HOME` is automatically inferred from the location of the startup scripts. `HADOOP_HOME` is the parent directory of the `bin` directory that holds the Hadoop scripts. In many instances this is `$HADOOP_INSTALL/hadoop`.

11 Add the following properties to the `yarn-site.xml` file:

```
<property>
  <name>yarn.resourcemanager.scheduler.class</name>
  <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.
    CapacityScheduler</value>
</property>

<property>
  <name>yarn.resourcemanager.resource-tracker.address</name>
  <value>$resourcemanager.full.hostname:8025</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.scheduler.address</name>
  <value>$resourcemanager.full.hostname:8030</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.address</name>
  <value>$resourcemanager.full.hostname:8050</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.admin.address</name>
  <value>$resourcemanager.full.hostname:8141</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.nodemanager.local-dirs</name>
  <value>/grid/hadoop/yarn/local,/grid1/hadoop/yarn/local</value>
  <description>Comma-separated list of paths. Use the list of directories
    from $YARN_LOCAL_DIR. For example, /grid/hadoop/yarn/local,/grid1/hadoop/
    yarn/local.</description>
</property>

<property>
  <name>yarn.nodemanager.log-dirs</name>
  <value>/grid/hadoop/yarn/log</value>
  <description>Use the list of directories from $YARN_LOCAL_LOG_DIR. For
    example, /grid/hadoop/yarn/log,/grid1/hadoop/yarn/log,/grid2/hadoop/yarn/
    log</description>
</property>

<property>
  <name>yarn.log.server.url</name>
```

```

<value>http://$jobhistoryserver.full.hostname:19888/jobhistory/logs/</value>
<description>URL for job history server</description>
</property>

<property>
  <name>yarn.resourcemanager.webapp.address</name>
  <value>$resourcemanager.full.hostname:8088</value>
  <description>URL for job history server</description>
</property>

<property>
  <name>yarn.nodemanager.admin-env</name>
  <value>MALLOC_ARENA_MAX=$MALLOC_ARENA_MAX</value>
  <description>Restrict the number of memory arenas to prevent
    excessive VMEM use by the glib arena allocator.
    For example, MALLOC_ARENA_MAX=4</description>
</property>

```

12 Add the following properties to the yarn-site.xml file:

```

<property>
  <name>yarn.resourcemanager.scheduler.class</name>
  <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.
    CapacityScheduler</value>
</property>

<property>
  <name>yarn.resourcemanager.resource-tracker.address</name>
  <value>$resourcemanager.full.hostname:8025</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.scheduler.address</name>
  <value>$resourcemanager.full.hostname:8030</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.address</name>
  <value>$resourcemanager.full.hostname:8050
  </value><description>Enter your ResourceManager hostname.
  </description></property>

<property>
  <name>yarn.resourcemanager.admin.address</name>
  <value>$resourcemanager.full.hostname:8141</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.nodemanager.local-dirs</name>
  <value>/grid/hadoop/yarn/local,/grid1/hadoop/yarn/local</value>
  <description>Comma separated list of paths. Use the list of directories
    from $YARN_LOCAL_DIR. For example,
    /grid/hadoop/yarn/local,/grid1/hadoop/yarn/local.
  </description>
</property>

```

```

<property>
  <name>yarn.nodemanager.log-dirs</name>
  <value>/grid/hadoop/yarn/log</value>
  <description>Use the list of directories from $YARN_LOCAL_LOG_DIR.
    For example, /grid/hadoop/yarn/log,/grid1/hadoop/yarn/log,/
    grid2/hadoop/yarn/log
</description>
</property>

<property>
  <name>yarn.log.server.url</name>
  <value>http://$jobhistoryserver.full.hostname:19888/jobhistory/logs/</
  value>
  <description>URL for job history server</description>
</property>

<property>
  <name>yarn.resourcemanager.webapp.address</name>
  <value>$resourcemanager.full.hostname:8088</value>
  <description>URL for job history server</description>
</property>

<property>
  <name>yarn.nodemanager.admin-env</name>
  <value>MALLOC_ARENA_MAX=$MALLOC_ARENA_MAX</value>
  <description>Restrict the number of memory arenas to prevent excessive VMEM
    use by
    the glib arena allocator.For example, MALLOC_ARENA_MAX=4</description>
</property>

```

13 Add the following properties to the yarn-site.xml file:

```

<property>
  <name>yarn.resourcemanager.scheduler.class</name>
  <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.capacity.
    CapacityScheduler</value>
</property>

<property>
  <name>yarn.resourcemanager.resource-tracker.address</name>
  <value>$resourcemanager.full.hostname:8025</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.scheduler.address</name>
  <value>$resourcemanager.full.hostname:8030</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.address</name>
  <value>$resourcemanager.full.hostname:8050</value>
  <description>Enter your ResourceManager hostname.</description>
</property>

<property>
  <name>yarn.resourcemanager.admin.address</name>
  <value>$resourcemanager.full.hostname:8141</value>
  <description>Enter your ResourceManager hostname.</description>

```

```

</property>

<property>
  <name>yarn.nodemanager.local-dirs</name>
  <value>/grid/hadoop/yarn/local,/grid1/hadoop/yarn/local</value>
  <description>Comma separated list of paths. Use the list of directories
    from $YARN_LOCAL_DIR. For example,
    /grid/hadoop/yarn/local,/grid1/hadoop/yarn/local.
  </description>
</property>

<property>
  <name>yarn.nodemanager.log-dirs</name>
  <value>/grid/hadoop/yarn/log</value>
  <description>Use the list of directories from $YARN_LOCAL_LOG_DIR.
    For example, /grid/hadoop/yarn/log,
    /grid1/hadoop/yarn/log,/grid2/hadoop/yarn/log
  </description>
</property>

<property>
  <name>yarn.log.server.url</name>
  <value>http://$jobhistoryserver.full.hostname:19888/jobhistory/logs/
  </value>
  <description>URL for job history server</description>
</property>

<property>
  <name>yarn.resourcemanager.webapp.address</name>
  <value>$resourcemanager.full.hostname:8088</value>
  <description>URL for job history server</description>
</property>

<property>
  <name>yarn.nodemanager.admin-env</name>
  <value>MALLOC_ARENA_MAX=$MALLOC_ARENA_MAX</value>
  <description>Restrict the number of memory arenas to prevent excessive
    VMEM use by the glib arena allocator. For example,
    MALLOC_ARENA_MAX=4</description>
</property>

```

14 Adding the following properties to the `mapred-site.xml` file:

```

<property>
  <name>mapreduce.jobhistory.address</name>
  <value>$jobhistoryserver.full.hostname:10020</value>
  <description>Enter your JobHistoryServer hostname.</description>
</property>

<property>
  <name>mapreduce.jobhistory.webapp.address</name>
  <value>$jobhistoryserver.full.hostname:19888</value>
  <description>Enter your JobHistoryServer hostname.</description>
</property>

<property>
  <name>mapreduce.shuffle.port</name>
  <value>13562</value>
</property>

```

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
```

15 For a secure cluster, add the following properties to `mapred-site.xml`:

```
<property>
  <name>mapreduce.jobhistory.principal</name>
  <value>jhs/_PRINCIPAL@$REALM.ACME.COM</value>
  <description>Kerberos principal name for the MapReduce JobHistory Server.
</description>
</property>

</property>
<name>mapreduce.jobhistory.keytab</name>
<value>/etc/security/keytabs/jhs.service.keytab</value>
<description>Kerberos keytab file for the MapReduce JobHistory Server.</
description>
</property>
```

16 For a secure cluster, you must also update `hadoop.security.auth_to_local` in `core-site.xml` to include a rule regarding the `mapreduce.jobhistory.principal` value you set in the previous step:

RULE:[2:\$1@\$0](PRINCIPAL@\$REALM.ACME.COM)s/./mapred/

where `PRINCIPAL` and `REALM` are the kerberos principal and realm you specified in `mapreduce.jobhistory.principal`.

17 Delete any remaining HDP1 properties in the `mapred-site.xml` file.

18 Replace the default memory configuration settings in `yarn-site.xml` and `mapred-site.xml` with the YARN and MapReduce memory configuration settings you calculated previously.

6.7. Create Local Directories

You must create local directories for YARN on each NodeManager host in your cluster (in HDP-2, the NodeManager replaces the TaskTracker) and set the appropriate permissions for the YARN log directories.

1. Set the permissions in the `yarn.nodemanager.local-dirs` directories. Run these commands on all DataNodes in your cluster.

```
chown -R yarn:hadoop ${yarn.nodemanager.local-dirs}
```

```
chmod 755 ${yarn.nodemanager.local-dirs}
```

where `${yarn.nodemanager.local-dirs}` is your local directory.

2. Change the permissions of the directories in `yarn.nodemanager.log-dirs`. If these directories do not exist, you can create them using the instructions in the [Create Directories](#) section of the Non-Ambari Cluster Installation Guide. Run these commands on all DataNodes in your cluster.

```
chown -R yarn:hadoop ${yarn.nodemanager.log-dirs}
```

```
chmod 755 ${yarn.nodemanager.log-dirs}
```

where `${yarn.nodemanager.log-dirs}` is your log directory.

3. Create directories for YARN_LOG_DIR and YARN_PID_DIR.

- Open `/etc/hadoop/conf/yarn-env.sh`
- Write down your values for YARN_LOG_DIR and YARN_PID_DIR, as the following instructions require values for the `${YARN_LOG_DIR}` and `${YARN_PID_DIR}`.

For example, in `yarn-env.sh`:

```
YARN_LOG_DIR=/grid/0/var/log/hadoop/yarn
```

```
YARN_PID_DIR=/grid/0/var/run/hadoop/yarn
```

4. Make directories for `${YARN_LOG_DIR}` and `${YARN_PID_DIR}` and set the appropriate permissions for them.

```
mkdir ${YARN_LOG_DIR}
```

```
chown yarn:hadoop ${YARN_LOG_DIR}
```

```
chmod 755 ${YARN_LOG_DIR}
```

```
mkdir ${YARN_PID_DIR}
```

```
chown yarn:hadoop ${YARN_PID_DIR}
```

```
chmod 755 ${YARN_PID_DIR}
```

6.8. Start Hadoop Core



Warning

Before you start HDFS on an HA cluster you must start the ZooKeeper service. If you do not start the ZKFC, there can be failures.

Start HDFS, executing commands as `$HDFS_USER`.

1. Replace your configuration after upgrading on all the HDFS nodes. Replace the HDFS template configuration in `/etc/hdfs/conf`.
2. If you are upgrading from a highly available HDFS cluster configuration, start all JournalNodes. On each JournalNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start journalnode"
```




Important

All JournalNodes must be running when performing the upgrade, rollback, or finalization operations. If any JournalNodes are down when running any such operation, the operation fails.

3. If you are running HDFS on a highly available namenode, you must first start the Zookeeper service.



Note

Perform this step only if you are on a highly available HDFS cluster.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start zkfc"
```

4. Start the NameNode.

Because the file system version has now changed you must start the NameNode manually.

On the active NameNode host, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start namenode -upgrade"
```

On a large system, this can take a long time to complete.



Note

Run this command with the `-upgrade` option only once. After you have completed this step, you can bring up the NameNode using this command without including the `-upgrade` option.

To check if the Upgrade is in progress, check that the `"\previous"` directory has been created in the `\NameNode` and `\JournalNode` directories. The `"\previous"` directory contains a snapshot of the data before upgrade.

In a highly available HDFS cluster configuration, this NameNode does not enter the standby state as usual. Rather, this NameNode immediately enters the active state, perform an upgrade of its local storage directories, and also perform an upgrade of the shared edit log. At this point, the standby NameNode in the HA pair is still down. It is out of sync with the upgraded active NameNode.

To synchronize the active and standby NameNode, re-establishing HA, re-bootstrap the standby NameNode by running the NameNode with the `'-bootstrapStandby'` flag. Do NOT start this standby NameNode with the `'-upgrade'` flag.

```
su - hdfs -c "hdfs namenode -bootstrapStandby -force"
```

The `bootstrapStandby` command downloads the most recent fsimage from the active NameNode into the `$dfs.name.dir` directory of the standby NameNode. You can enter

that directory to make sure the fsimage has been successfully downloaded. After verifying, start the ZKFailoverController, then start the standby NameNode. You can check the status of both NameNodes using the Web UI.

5. Verify that the NameNode is up and running:

```
ps -ef|grep -i NameNode
```

6. If you do not have a highly available HDFS cluster configuration (non_HA namenode), start the Secondary NameNode.



Note

Do not perform this step if you have a highly available HDFS cluster configuration.

On the Secondary NameNode host machine, run the following commands:

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start secondarynamenode"
```

7. Verify that the Secondary NameNode is up and running.



Note

Do not perform this step if you have a highly available HDFS cluster environment.

```
ps -ef|grep SecondaryNameNode
```

8. Start DataNodes.

On each of the DataNodes, enter the following command. Note: If you are working on a non-secure DataNode, use \$HDFS_USER. For a secure DataNode, use root.

```
su - hdfs -c "/usr/hdp/current/hadoop-client/sbin/hadoop-daemon.sh start datanode"
```

9. Verify that the DataNode process is up and running:

```
ps -ef|grep DataNode
```

10. Verify that NameNode can go out of safe mode.

```
>su - hdfs -c "hdfs dfsadmin -safemode wait"
```

You should see the following result: Safe mode is OFF

In general, it takes 5-10 minutes to get out of safemode. For thousands of nodes with millions of data blocks, getting out of safemode can take up to 45 minutes.

6.9. Verify HDFS Filesystem Health

Analyze if the filesystem is healthy.

1. Run the fsck command on namenode as \$HDFS_USER:

```
su - $HDFS_USER
hdfs fsck / -files -blocks -locations > dfs-new-fsck-1.log
```

2. Run hdfs namespace and report.

- a. List directories.

```
su -l <HDFS_USER>
hdfs dfs -ls -R / > dfs-new-lsr-1.log
```

- b. Run report command to create a list of DataNodes in the cluster.

```
su -l <HDFS_USER>
hdfs dfsadmin -report > dfs-new-report-1.log
```

- c. Compare the namespace report before the upgrade and after the upgrade. Verify that user files exist after upgrade.

```
su - $HDFS_USER
dfs-old-fsck-1.log < -- > dfs-new-fsck-1.log dfs-old-lsr-1.log
< -- > dfs-new-lsr-1.log
```



Note

You must do this comparison manually to catch all errors.

- d. From the Namenode WebUI, see if all DataNodes are up and running.

```
http://<namenode>:50070
```

6.10. Configure YARN and MapReduce

After you install Hadoop, modify your configs.

1. Upload the MapReduce tarball to HDFS. As the HDFS user, for example hdfs:

```
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/mapreduce/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hadoop/
mapreduce.tar.gz /hdp/apps/2.5.0.0-<$version>/mapreduce/
hdfs dfs -chown -R hdfs:hadoop /hdp
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/mapreduce
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/mapreduce/
mapreduce.tar.gz
```

2. Make the following changes to mapred-site.xml:

- Add the following property:

```
<property>
  <name>mapreduce.application.framework.path</name>
  <value>/hdp/apps/${hdp.version}
    /mapreduce/mapreduce.tar.gz#mr-framework
  </value>
</property>

<property>
  <name>yarn.app.mapreduce.am.admin-command-opts</name>
  <value>Dhdp.version=${hdp.version}</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

- Modify the following existing properties to include `${hdp.version}`:

```
<property>
  <name>mapreduce.admin.user.env</name>
  <value>LD_LIBRARY_PATH=/usr/hdp/${hdp.version}
    /hadoop/lib/native:/usr/hdp/${hdp.version}/hadoop/
    lib/native/Linux-amd64-64
  </value>
</property>

<property>
  <name>mapreduce.admin.map.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true
    -Dhdp.version=${hdp.version}
  </value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.admin.reduce.child.java.opts</name>
  <value>-server -Djava.net.preferIPv4Stack=true -Dhdp.version=${hdp.
version}</value>
  <final>true</final>
</property>

<property>
  <name>mapreduce.application.classpath</name>
  <value>$PWD/mr-framework/hadoop/share/hadoop/mapreduce/*,
    $PWD/mr-framework/hadoop/share/hadoop/mapreduce/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/common/*,
    $PWD/mr-framework/hadoop/share/hadoop/common/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/yarn/*,
    $PWD/mr-framework/hadoop/share/hadoop/yarn/lib/*,
    $PWD/mr-framework/hadoop/share/hadoop/hdfs/*,
    $PWD/mr-framework/hadoop/share/hadoop/hdfs/lib/*,
    /usr/hdp/${hdp.version}/hadoop/lib/hadoop-lzo-0.6.0.${hdp.version}.jar,
    /etc/hadoop/conf/secure</value>
</property>
```



Note

You do not need to modify `${hdp.version}`.

3. Add the following property to `yarn-site.xml`:

```
<property>
  <name>yarn.application.classpath</name>
  <value>$HADOOP_CONF_DIR,/usr/hdp/${hdp.version}/hadoop-client/*,
    /usr/hdp/${hdp.version}/hadoop-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/*,
    /usr/hdp/${hdp.version}/hadoop-hdfs-client/lib/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/*,
    /usr/hdp/${hdp.version}/hadoop-yarn-client/lib/*</value>
</property>
```

4. For secure clusters, you must create and configure the `container-executor.cfg` configuration file:

- Create the `container-executor.cfg` file in `/etc/hadoop/conf/`.
- Insert the following properties:

```
yarn.nodemanager.linux-container-executor.group=hadoop
banned.users=hdfs,yarn,mapred
min.user.id=1000
```

- `yarn.nodemanager.linux-container-executor.group` - Configured value of `yarn.nodemanager.linux-container-executor.group`. This must match the value of `yarn.nodemanager.linux-container-executor.group` in `yarn-site.xml`.
- `banned.users` - Comma-separated list of users who can not run container-executor.
- `min.user.id` - Minimum value of user id. This prevents system users from running container-executor.
- `allowed.system.users` - Comma-separated list of allowed system users.
- Set the file `/etc/hadoop/conf/container-executor.cfg` file permissions to only be readable by root:

```
chown root:hadoop /etc/hadoop/conf/container-executor.cfg
```

```
chmod 400 /etc/hadoop/conf/container-executor.cfg
```

- Set the container-executor program so that only root or hadoop group users can run it:

```
chown root:hadoop /usr/hdp/${hdp.version}/hadoop-yarn-server-
nodemanager/bin /container-executor
```

```
chmod 6050 /usr/hdp/${hdp.version}/hadoop-yarn-server-
nodemanager/bin /container-executor
```

6.11. Start YARN/MapReduce Services

Add the following properties to `yarn-site.xml` to configure Application Timeline Server (ATS):

```
yarn.timeline-service.leveldb-timeline-store.path=/var/log/hadoop-yarn/
timeline
yarn.timeline-service.leveldb-timeline-store.ttl-interval-ms=300000

** If you are upgrading to HDP 2.1.3 or higher, use the following setting:
yarn.timeline-service.store-class=org.apache.hadoop.yarn.server.timeline.
LevelDbTimelineStore**

** If you are upgrading to HDP 2.1.2, use the following setting:
yarn.timeline-service.store-class=org.apache.hadoop.yarn.server.
applicationhistoryservice.timeline.LevelDbTimelineStore **

yarn.timeline-service.ttl-enable=true
yarn.timeline-service.ttl-ms=2678400000
yarn.timeline-service.generic-application-history.store-class=org.apache.
hadoop.yarn.server.applicationhistoryservice.NullApplicationHistoryStore
yarn.timeline-service.webapp.address={PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:8188
yarn.timeline-service.webapp.https.address=
{PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:8190
yarn.timeline-service.address={PUT_THE_FQDN_OF_ATS_HOST_NAME_HERE}:10200

HIVE (hive-site.xml)
hive.execution.engine=mr
hive.exec.failure.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.exec.post.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.exec.pre.hooks=org.apache.hadoop.hive ql.hooks.ATSHook
hive.tez.container.size={map-container-size}

*If mapreduce.map.memory.mb > 2GB then set it equal to mapreduce.map.memory.
Otherwise, set it equal to mapreduce.reduce.memory.mb*

hive.tez.java.opts=-server -Xmx800m -Djava.net.preferIPv4Stack=true -
XX:NewRatio=8 -XX:+UseNUMA -XX:+UseParallelGC
```

Use configuration values appropriate for your environment. For example, the value "800" in the preceding example is an example, not a requirement.

To start YARN, run commands as a YARN user. To start MapReduce, run commands as a MapReduce user.

1. If you have a secure cluster, create the following principals and keytabs for YARN before you start the YARN service:
2. Start the ResourceManager on your previous JobTracker host.

```
su -l yarn -c "/usr/hdp/current/hadoop-yarn-resourcemanager/
sbin/yarn-daemon.sh start resourcemanager"

ps -ef | grep -i resourcemanager
```

3. Prepare the NodeManager hosts.

- a. Change permissions for `/usr/hdp/current/hadoop-yarn/bin/container-executor.cfg`:

```
chown yarn:hadoop /usr/hdp/current/hadoop-yarn/bin/container-executor

chmod -R 650 /usr/hdp/current/hadoop-yarn/bin/container-executor
```

- b. On all NodeManager hosts, add the yarn user to the hadoop group.

For example, if you are using CentOS 6:

```
usermod -a -G hadoop yarn
```

4. Start the NodeManager on your previous TaskTracker hosts.

```
su -l yarn -c "/usr/hdp/current/hadoop-yarn-nodemanager/sbin/yarn-daemon.sh start nodemanager"

ps -ef | grep -i nodemanager
```

5. To start MapReduce, run the following commands:

```
su -l yarn -c "/usr/hdp/current/hadoop-mapreduce-historyserver/sbin/mr-jobhistory-daemon.sh start historyserver"

ps -ef | grep -i jobhistoryserver
```

6.12. Run Hadoop Smoke Tests

To smoke test your Hadoop upgrade, you can run the following MapReduce job as a regular user. The job uses MapReduce to write 100MB of data into HDFS with RandomWriter.

```
hadoop jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-mapreduce-examples.jar randomwriter -
Dtest.randomwrite.total_bytes=100000000 test-after-upgrade
```

You should see messages similar to:

```
map 0% reduce 0%
...map 100% reduce 100%
Job ... completed successfully
```

You have successfully submitted your first MapReduce job in HDP 2.x. The next steps upgrade your other components to 2.x.

Basic troubleshooting:

1. To find the number of active nodes and NodeManagers, access the ResourceManager web UI:

```
http://<resource manager host>:8088/cluster/nodes
```

The number of active nodes should be equal to the number of nodemanagers.

2. Error messages. Access the ApplicationMaster WebUI to view the container logs.

a. Looking at your console logs for MapReduce job, there is a line in this format:

```
13/10/02 17:57:21 INFO mapreduce.Job: The url to track
the job: http://<resource manager host>:8088/proxy/
application_1380673658357_0007/
```

b. Go to the URL and select the job link.

c. Select the logs link under ApplicationMaster table. It redirects you to the container logs. Error messages display here.

6.13. Configure and Start Apache HBase

The `hbase.bucketcache.percentage.in.combinedcache` is removed in HDP-2.5.0. This simplifies the configuration of block cache. BucketCache configurations from HDP 2.1 need to be recalculated to attain identical memory allotments in HDP-2.5.0. The L1 LruBlockCache is whatever `hfile.block.cache.size` is set to and the L2 BucketCache is whatever `hbase.bucketcache.size` is set to.

1. Upgrade Apache HBase.

```
hbase upgrade -execute
```

2. Replace your configuration after upgrading. Replace the HBase template configuration in `/etc/hbase/conf`.

3. Start services. From root, assuming that `$HBASE_USER=hbase`:

```
su - hbase -c "/usr/hdp/current/hbase-master/bin/hbase-daemon.sh
start master; sleep 25"
```

```
su - hbase -c "/usr/hdp/current/hbase-regionserver/bin/hbase-
daemon.sh start regionserver"
```

4. Check processes.

```
ps -ef | grep -i hmaster
```

```
ps -ef | grep -i hregion
```

6.14. Configure and Start Apache Hive and Apache HCatalog



Warning

In HDP 2.1.3 and later, the Decimal data type is now treated as the type `Decimal(10,0)`: 10 digits of precision and 0 scale. This is a change from the variable precision and scale available in Apache Hive 0.11.0 and Hive 0.12.0, which allowed up to 38 digits of precision and unlimited scale.

To avoid unintended rounding of decimal data, sites that were previously running Hive 0.11.0 and Hive 0.12.0 may need to migrate tables with Decimal columns after upgrading to Hive 0.13.0. For details, see the [Apache Hive wiki](#). For assistance with upgrades that involve Decimal data, please contact Hortonworks Support.

1. Prior to starting the upgrade process, set the following in your hive configuration file:

```
datanucleus.autoCreateSchema=false
```

2. Stop the Hive Metastore, if you have not done so already.
3. Upgrade the Hive Metastore database schema by running the upgrade scripts included in HDP for your metastore database and then running the schematool to upgrade to Hive 14:

Database	Run as user...	This metastore upgrade script
MySQL	root	<pre>>cd /usr/hdp/current/hive-metastore/ scripts/metastore/upgrade/mysql > mysql hivemysql > source upgrade-0.11.0- to-0.12.0.mysql.sql > source upgrade-0.12.0- to-0.13.0.mysql.sql > source upgrade-0.13.0- to-0.14.0.mysql.sql</pre>
PostgreSQL	root	<pre>cd /usr/hdp/current/hive-metastore/ scripts/metastore/upgrade/postgres psql -d hive -a -f upgrade-0.11.0- to-0.12.0.postgres.sql psql -d hive -a -f upgrade-0.12.0- to-0.13.0.postgres.sql psql -d hive -a -f upgrade-0.13.0- to-0.14.0.postgres.sql</pre>
Oracle	root	<pre>cd /usr/hdp/2.5.0.0-<version>/ hive-metastore/scripts/metastore/ upgrade/oracle sqlplus SQL> @upgrade-0.11.0- to-0.12.0.oracle.sql SQL> @upgrade-0.12.0- to-0.13.0.oracle.sql SQL> @upgrade-0.13.0- to-0.14.0.oracle.sql</pre>

```
su - hive -c "/usr/hdp/current/hive-metastore/bin/schematool -
upgradeSchema -dbType <$databaseType>"
```



Note

If you are using PostgreSQL, reset the Hive Metastore database owner to the <HIVE_USER>:

```
sudo <POSTGRES_USER>
```

```
ALTER DATABASE <HIVE-METASTORE-DB-NAME> OWNER TO  
<HIVE_USER>
```



Note

If you are using Oracle 11, you may see the following error message:

```
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.  
optimize.mapjoin.mapreduce does not exist  
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.  
heapsize does not exist  
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.  
server2.enable.impersonation does not exist  
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.  
semantic.analyzer.factory.impl does not exist  
14/11/17 14:11:38 WARN conf.HiveConf: HiveConf of name hive.auto.  
convert.sortmerge.join.noconditionaltask does not exist  
Metastore connection URL: jdbc:oracle:thin:@//ip-172-31-42-1.ec2.  
internal:1521/XE  
Metastore Connection Driver : oracle.jdbc.driver.OracleDriver  
Metastore connection User: hiveuser  
Starting upgrade metastore schema from version 0.13.0 to 0.14.0  
Upgrade script upgrade-0.13.0-to-0.14.0.oracle.sql  
Error: ORA-00955: name is already used by an existing object  
(state=42000,code=955)  
Warning in pre-upgrade script pre-0-upgrade-0.13.0-to-0.14.0.  
oracle.sql: Schema script failed, errorcode 2  
Completed upgrade-0.13.0-to-0.14.0.oracle.sql  
schemaTool completed
```

You can safely ignore this message. The error is in the pre-upgrade script and can be ignored; the schematool succeeded.

4. Edit hive-site.xml to update the value of hive.metastore.export.location to the new, joint hive-hcatalog jar (previously hcatalog-core.jar):

```
<property>  
  <name>hive.metastore.export.location</name>  
  <value>export HIVE_AUX_JARS_PATH=/usr/hdp/2.5.0.0-<version>/hive-hcatalog/  
share/hcatalog/hive-hcatalog-core.jar</value>  
</property>
```

5. Edit hive-env.sh to point to the new hive-hcatalog.jar:

```
if [ "${HIVE_AUX_JARS_PATH}" != " " ]; then  
  export HIVE_AUX_JARS_PATH=/usr/hdp/current/hive-hcatalog/share/hcatalog/  
hive- hcatalog-core.jar:${HIVE_AUX_JARS_PATH}  
else  
  export HIVE_AUX_JARS_PATH=/usr/hdp/current/hive-hcatalog/share/hcatalog/  
hive- hcatalog-core.jar  
fi
```

6. Edit the hive-site.xml file and modify the properties based on your environment. Search for TODO in the file for the properties to replace.

- Edit the connection properties for your Hive metastore database in hive-site.xml:

```

<property>
  <name>javax.jdo.option.ConnectionURL</name>
  <value>jdbc:mysql://TODO-HIVE-METASTORE-DB-SERVER:TODO-HIVE-METASTORE-DB-
PORT/
TODO-HIVE-METASTORE-DB-NAME?createDatabaseIfNotExist=true</value>
  <description>Enter your Hive Metastore Connection URL,
    for example if MySQL:
      jdbc:mysql://localhost:3306/mysql?createDatabaseIfNotExist=true
  </description>
</property>

<property>
  <name>javax.jdo.option.ConnectionUserName</name>
  <value>TODO-HIVE-METASTORE-DB-USER-NAME</value>
  <description>Enter your Hive Metastore database user name.</description>
</property>

<property>
  <name>javax.jdo.option.ConnectionPassword</name>
  <value>TODO-HIVE-METASTORE-DB-PASSWORD</value>
  <description>Enter your Hive Metastore database password.</description>
</property>

<property>
  <name>javax.jdo.option.ConnectionDriverName</name>
  <value>TODO-HIVE-METASTORE-DB-CONNECTION-DRIVER-NAME</value>
  <description>Enter your Hive Metastore Connection Driver Name, for
    example if
    MySQL: com.mysql.jdbc.Driver</description>
</property>

```

- Edit the following properties in the hive-site.xml file:

```

<property>
  <name>fs.file.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.file.impl.disable.cache</
description>
</property>

<property>
  <name>fs.hdfs.impl.disable.cache</name>
  <value>>false</value>
  <description>Set to false or remove fs.hdfs.impl.disable.cache</
description>
</property>

```

- **Optional:** If you want Hive Authorization, set the following Hive authorization parameters in the hive-site.xml file:

```

<property>
  <name>hive.security.authorization.enabled</name>
  <value>true</value>
</property>

<property>
  <name>hive.security.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
StorageBasedAuthorizationProvider</value>

```

```

</property>

<property>
  <name>hive.security.metastore.authorization.manager</name>
  <value>org.apache.hadoop.hive.ql.security.authorization.
    StorageBasedAuthorizationProvider</value>
</property>

<property>
  <name>hive.security.authenticator.manager</name>
  <value>org.apache.hadoop.hive.ql.security.ProxyUserAuthenticator</ value>
</property>

```

- For a remote Hive metastore database, use the following hive-site.xml property value to set the IP address (or fully-qualified domain name) and port of the metastore host.

To enable HiveServer2, leave this property value empty.

```

<property>
  <name>hive.metastore.uris</name>
  <value>thrift://$metastore.server.full.hostname:9083</value>
  <description>URI for client to contact metastore server.
    To enable HiveServer2, leave the property value empty.
  </description>
</property>

```



Note

You can also use the HDP utility script to fine-tune your configuration settings based on node hardware specifications.

7. Disable autocreation of schemas to match HDP 2.1+ configurations. Edit hive-site.xml to set the value of `datanucleus.autoCreateSchema` to `false`.

```

<property>
  <name>datanucleus.autoCreateSchema</name>
  <value>true</value>
  <description>Creates necessary schema on a startup if one doesn't exist.
  </description>
</property>

```

8. Start Hive. On the Hive Metastore host machine, run the following commands:

```

su - hive

nohup /usr/hdp/current/hive-metastore/bin/hive --service
metastore>/var/log/hive/hive.out 2>/var/log/hive/hive.log &

```

9. Start Hive Server2.

On the Hive Server2 host machine, run the following commands:

```

su - hive

/usr/hdp/current/hive-server2/bin/hiveserver2 >/var/log/hive/
hiveserver2.out 2> /var/log/hive/hiveserver2.log&

```

where `$HIVE_USER` is the Hive Service user. For example, `hive`.

6.15. Configure and Start Apache Oozie

Upgrading Apache Oozie is a complex process. Although the instructions are straightforward, set aside a dedicated block of time to upgrade oozie clients and servers.

Perform the following preparation steps on each oozie server host:

1. You must restore `oozie-site.xml` from your backup to the `conf` directory on each oozie server and client.

2. Copy the JDBC jar to `libext-customer`:

- a. Create the `/usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer` directory.

```
cd /usr/hdp/2.5.0.0-<version>/oozie-server
```

```
mkdir libext-customer
```

- b. Grant read/write/execute access to all users for the `libext-customer` directory.

```
chmod -R 777 /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer
```

3. Copy these files to the `libext-customer` directory

```
cp /usr/hdp/2.5.0.0-<version>/hadoop-client/lib/hadoop*lzo*.jar /usr/hdp/current/oozie-server/libext-customer
```

```
cp /usr/share/HDP-oozie/ext.zip /usr/hdp/2.5.0.0-<version>/oozie-server/libext-customer/
```

Also, copy Oozie db jar in `libext-customer`.

4. If Falcon was also installed and configured before upgrade in HDP 2.5, then after upgrade you might also need to do the following:

```
cp /usr/hdp/current/falcon-server/oozie/ext/falcon-oozie-el-extension-jar /usr/hdp/current/oozie-server/libext-customer
```

5. Extract `share-lib`.

```
cd /usr/hdp/2.5.0.0-<version>/oozie-server
```

```
tar xzvf /usr/hdp/2.5.0.0-<version>/oozie-server/oozie-sharelib.tar.gz
```

```
su - hdfs -c "hdfs dfs -mkdir -p /user/oozie"
```

```
su - hdfs -c "hdfs dfs -copyFromLocal /usr/hdp/current/oozie-server/share /user/oozie/."
```

You can expect warnings that some files already exist. Delete any existing `/oozie/share` and replace it with the newly-extracted files.

```
su - hdfs -c "hdfs dfs -chown oozie:hadoop /user/oozie"
```

```
su - hdfs -c "hdfs dfs -chmod -R 755 /user/oozie"
```

6. If a previous version of Oozie was created using auto schema creation, run the following SQL query:

```
insert into oozie_sys (name, data) values ('db.version', '2.5');
```

7. As the Oozie user (not root), run the upgrade.

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/ooziedb.sh  
upgrade -run"
```

8. As root, prepare the Oozie WAR file.

```
chown oozie:oozie /usr/hdp/current/oozie-server/oozie-server/  
conf/server.xml
```

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-setup.sh  
prepare-war -d /usr/hdp/current/oozie-server/libext-customer"
```

Look for console output to indicate success. For example, if you are using MySQL you should see something similar to:

```
INFO: Adding extension: libext-customer/mysql-connector-java.jar  
New Oozie WAR file with added 'JARS' at /var/lib/oozie/oozie-server/webapps/  
oozie.war
```

9. Make sure that following property is added in oozie-log4j.properties:

```
log4j.appender.oozie.layout.ConversionPattern=%d{ISO8601} %5p  
%c{1}:%L - SERVER[${oozie.instance.id}] %m%n
```

where `${oozie.instance.id}` is determined by oozie, automatically.

10. Configure HTTPS for the Oozie server.

- Create a self signed certificate or get certificate from a trusted CA for the Oozie Server
- Import the certificate to the client JDK trust store on all client nodes.
- In the Ambari Oozie configuration, set the following environment variables in `oozie-env.sh`, adding them if it does not exist:

```
export OOOIE_HTTPS_PORT=11443  
export OOOIE_HTTPS_KEYSTORE_FILE=/home/oozie/.keystore  
export OOOIE_HTTPS_KEYSTORE_PASS=password
```

- Change `OOOIE_HTTP_PORT={{oozie_server_port}}` to `OOOIE_HTTP_PORT=11000`.

- Set the `oozie.base.url` to the HTTPS address.

- f. Save the configuration, and restart the Oozie components.

11 Start Oozie as the Oozie user:

```
su - oozie -c "/usr/hdp/current/oozie-server/bin/oozie-start.sh"
```

12 Check processes.

```
ps -ef | grep -i oozie
```

6.16. Configure and Start Apache WebHCat (Templeton)

RHEL/CentOS/Oracle Linux

1. Copy the appropriate configurations from `/etc/hcatalog/conf` to `/etc/hive-webhcat/conf/`.
2. Copy the new Pig, Hive and Hadoop-streaming jars to HDFS using the path you specified in `./etc/hive-webhcat/conf/` and change ownership to the hcat user with 755 permissions. For example:

```
hdfs dfs -copyFromLocal /usr/share/HDP-webhcat/hive.tar.gz /usr/
share/HDP-webhcat/pig.tar.gz /usr/hdp/version/hadoop-mapreduce/
hadoop-streaming.jar hdfs:///apps/webhcat/.
```

```
hdfs dfs -chmod -R 755 hdfs:///apps/webhcat/*
```

```
hdfs dfs -chown -R hcat hdfs:///apps/webhcat/*
```

3. Replace your WebHCat configuration after upgrading. Copy your modified `/etc/webhcat/conf` from the template to the configuration directory in all your WebHCat hosts.
4. Start WebHCat:

```
sudo su -l $WEBHCAT_USER -c "//hive-hcatalog/sbin/
webhcat_server.sh start"
```

5. Smoke test WebHCat.

On the WebHCat host machine, run the following command:

```
http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

If you are using a secure cluster, run the following command:

```
curl --negotiate -u:http://cluster.$PRINCIPAL.$REALM:50111/
templeton/v1/ status{"status":"ok","version":"v1"}
[machine@acme]$
```

6. Remove shared libraries from old Templeton installation.

On the WebHCat host machine, run the following command:

```
sudo su -l $HDFS_USER -c "hdfs dfs -rmr -skipTrash /apps/
templeton" rm -rf /usr/share/HDP-templeton
```

where

- `$WEBHCAT_USER` is the WebHCat Service user. For example, `hcat`.
- `$HDFS_USER` is the HDFS Service user. For example, `hdfs`.

SLES

1. Copy the appropriate configurations from `/etc/hcatalog/conf` to `/etc/hive-webhcat/conf/`.
2. Copy the new Pig, Hive and Hadoop-streaming jars to HDFS using the path you specified in `/etc/hive-webhcat/conf/` and change ownership to the `hcat` user with 755 permissions. For example:

```
hdfs dfs -copyFromLocal /usr/share/HDP-webhcat/hive.tar.gz /usr/
share/HDP-webhcat/pig.tar.gz/usr/hdp/version/hadoop-mapreduce/
hadoop-streaming.jar hdfs:///apps/webhcat/.
```

```
hdfs dfs -chmod -R 755 hdfs:///apps/webhcat/*
```

```
hdfs dfs -chown -R hcat hdfs:///apps/webhcat/*
```

3. Replace your WebHCat configuration after upgrading. Copy your modified `/etc/webhcat/conf` from the template to the configuration directory in all your WebHCat hosts.
4. Modify the WebHCat configuration files.

- Upload Pig, Hive and Sqoop tarballs to HDFS as the `$HDFS_User`. In this example, `hdfs`:

```
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/pig/
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/hive/
hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/pig/pig.tar.gz /hdp/apps/2.5.0.
0-<$version>/pig/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/hive/hive.tar.gz /hdp/apps/2.5.
0.0-<$version>/hive/
hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.
5.0.0-<$version>/sqoop/
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/pig
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/pig/pig.tar.gz
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/hive
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/hive/hive.tar.gz
hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop
hdfs dfs -chmod -R 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz
hdfs dfs -chown -R hdfs:hadoop /hdp
```

- Update the following properties in the `webhcat-site.xml` configuration file, as their values have changed:


```

<property>
  <name>templeton.pig.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/pig/pig.tar.gz</value>
</property>

</property>
  <name>templeton.hive.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/hive/hive.tar.gz</value>
</property>

</property>
  <name>templeton.streaming.jar</name>
  <value>hdfs:///hdp/apps/${hdp.version}/mapreduce/hadoop-streaming.jar</value>
  <description>The hdfs path to the Hadoop streaming jar file.</description>
</property>

</property>
  <name>templeton.sqoop.archive</name>
  <value>hdfs:///hdp/apps/${hdp.version}/sqoop/sqoop.tar.gz</value>
  <description>The path to the Sqoop archive.</description>
</property>

</property>
  <name>templeton.sqoop.path</name>
  <value>sqoop.tar.gz/sqoop/bin/sqoop</value>
  <description>The path to the Sqoop executable.</description>
</property>

</property>
  <name>templeton.sqoop.home</name>
  <value>sqoop.tar.gz/sqoop</value>
  <description>The path to the Sqoop home in the exploded archive.</description>
</property>

```



Note

You do not need to modify `${hdp.version}`.

- Remove the following obsolete properties from `webhcat-site.xml`:

```

<property>
  <name>templeton.controller.map.mem</name>
  <value>1600</value>
  <description>Total virtual memory available to map tasks.</description>
</property>

</property>
  <name>hive.metastore.warehouse.dir</name>
  <value>/path/to/warehouse/dir</value>
</property>

```

- Add new proxy users, if needed. In `core-site.xml`, make sure the following properties are also set to allow WebHcat to impersonate your additional HDP 2.5 groups and hosts:

```
<property>
  <name>hadoop.proxyuser.hcat.groups</name>
  <value>*</value>
</property>

</property>
  <name>hadoop.proxyuser.hcat.hosts</name>
  <value>*</value>
</property>
```

Where:

`hadoop.proxyuser.hcat.group`

Is a comma-separated list of the Unix groups whose users may be impersonated by 'hcat'.

`hadoop.proxyuser.hcat.hosts`

A comma-separated list of the hosts which are allowed to submit requests by 'hcat'.

5. Start WebHCat:

```
su -l hcat -c "/usr/hdp/current/hive-webhcat/sbin/
webhcat_server.sh start"
```

6. Smoke test WebHCat.

On the WebHCat host machine, run the following command:

```
http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

If you are using a secure cluster, run the following command:

```
curl --negotiate -u:http://cluster.$PRINCIPAL.$REALM:50111/
templeton/v1/status{"status":"ok","version":"v1"}
[machine@acme]$
```

7. Remove shared libraries from old Templeton installation.

On the WebHCat host machine, run the following command:

```
sudo su -l $HDFS_USER -c "hdfs dfs -rmr -skipTrash /apps/
templeton" rm -rf /usr/share/HDP-templeton
```

where

- `$WEBHCAT_USER` is the WebHCat Service user. For example, hcat.
- `$HDFS_USER` is the HDFS Service user. For example, hdfs.

Ubuntu/Debian

1. Copy the appropriate configurations from `/etc/hcatalog/conf` to `/etc/hive-webhcat/conf/`.

2. Copy the new Pig, Hive and Hadoop-streaming jars to HDFS using the path you specified in `./etc/hive-webhcat/conf/` and change ownership to the hcat user with 755 permissions. For example:

```
hdfs dfs -copyFromLocal /usr/share/HDP-webhcat/hive.tar.gz /usr/share/HDP-webhcat/pig.tar.gz/usr/hdp/version/hadoop-mapreduce/hadoop-streaming.jar hdfs:///apps/webhcat/.
```

```
hdfs dfs -chmod -R 755 hdfs:///apps/webhcat/*
```

```
hdfs dfs -chown -R hcat hdfs:///apps/webhcat/*
```

3. Replace your WebHCat configuration after upgrading. Copy your modified `/etc/webhcat/conf` from the template to the configuration directory in all your WebHCat hosts.

4. Start WebHCat:

```
sudo su -l $WEBHCAT_USER -c "/usr/lib/hive-hcatalog/sbin/webhcat_server.sh start"
```

5. Smoke test WebHCat.

On the WebHCat host machine, run the following command:

```
http://$WEBHCAT_HOST_MACHINE:50111/templeton/v1/status
```

If you are using a secure cluster, run the following command:

```
curl -&-negotiate -u:http://cluster.$PRINCIPAL.$REALM:50111/templeton/v1/status{"status":"ok","version":"v1"}  
[machine@acme]$
```

6. Remove shared libraries from old Templeton installation.

On the WebHCat host machine, run the following command:

```
sudo su -l $HDFS_USER -c "hdfs dfs -rmr -skipTrash /apps/templeton" rm -rf /usr/share/HDP-templeton
```

where

- `$WEBHCAT_USER` is the WebHCat Service user. For example, hcat.
- `$HDFS_USER` is the HDFS Service user. For example, hdfs.

6.17. Configure and Start Apache Pig

1. On all Apache Pig clients, run the following command:

- **For RHEL/CentOS/Oracle Linux:**

```
yum erase pig
```

```
yum install pig
```

- **For SLES:**

```
zypper rm pig
zypper install pig
```

- **For Ubuntu or Debian:**

```
apt-get remove pig --purge
apt-get install pig
```

2. Replace /etc/hive-webhcat/conf/ your configuration after upgrading.

Copy /etc/pig/conf from the template to the conf directory in pig hosts.

6.18. Configure and Start Apache Sqoop

1. Replace your configuration after upgrading.

Copy /etc/sqoop/conf from the template to the conf directory in sqoop hosts.

2. Upload the Apache Sqoop tarball to HDFS. As the <HDFS_User>, for example 'hdfs':

```
su -c "hdfs dfs -mkdir -p /hdp/apps/2.5.0.0-<$version>/sqoop"

hdfs su -c "hdfs dfs -chmod -R 555 /hdp/apps/2.5.0.0-<$version>/sqoop"

hdfs su -c "hdfs dfs -chown -R hdfs:hadoop /hdp/apps/2.5.0.0-<$version>/sqoop"

hdfs su -c "hdfs dfs -put /usr/hdp/2.5.0.0-<$version>/sqoop/sqoop.tar.gz /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz"

hdfs su -c "hdfs dfs -chmod 444 /hdp/apps/2.5.0.0-<$version>/sqoop/sqoop.tar.gz" hdfs
```

6.19. Configure, Start, and Validate Apache Flume

Upgrade Apache Flume. On the Flume host machine, run the following command:

- **For RHEL/CentOS/Oracle Linux:**

```
yum upgrade flume
```

- **For SLES:**

```
zypper update flume
zypper remove flume
zypper se -s flume
```

You should see Flume in the output.

Install Flume:

```
zypper install flume
```

- For **Ubuntu/Debian**:

```
apt-get install flume
```

Replace your configuration after upgrading. Copy `/etc/flume/conf` from the template to the `conf` directory in Flume hosts.

By default on installation Flume does not start running immediately. To validate, replace your default `conf/flume.conf` with the provided `flume.conf` file, and restart Flume. See if the data is flowing by examining the destination.

Use this `flume.conf` file:

```
1. Name the components on this agent
a1.sources = r1
a1.sinks = k1
a1.channels = c1

2. Describe/configure the source
a1.sources.r1.type = seq

3. Describe the sink
a1.sinks.k1.type = file_roll
a1.sinks.k1.channel = c1
a1.sinks.k1.sink.directory = /tmp/flume

4. Use a channel which buffers events in memory
a1.channels.c1.type = memory

5. Bind the source and sink to the channel
a1.sources.r1.channels = c1
a1.sinks.k1.channel = c1
```

After starting Flume, check `/tmp/flume` to see if there are any files there. The files should contain simple sequential numbers.

After validating, stop Flume and revert changes to `flume.conf` to prevent your disk from filling up.

6.20. Configure, Start, and Validate Apache Mahout

Upgrade Apache Mahout. On the Mahout client machines, run the following command:

- For **RHEL/CentOS/Oracle Linux**:

```
yum upgrade mahout
```

- For **SLES**:

- `zypper remove mahout`
- `zypper se -s mahout`

You should see Mahout in the output.

- **Install Mahout:**

```
zypper up mahout
```

- **For Ubuntu:**

```
apt-get install mahout
```

Replace your configuration after upgrading. Copy `/etc/mahout/conf` from the template to the conf directory in mahout hosts.

To validate mahout:

1. Create a test user:

```
hadoop dfs -put /tmp/sample-test.txt /user/testuser
```

2. Set up mahout to convert the plain text file `sample-test.txt` into a sequence file that is in the output directory `mahouttest`.

```
mahout seqdirectory --input /user/testuser/sample-test.txt --  
output /user/testuser/mahouttest --charset utf-8
```

6.21. Configure and Start Hue

For HDP 2.5, use the Hue version shipped with HDP 2.5. If you have a previous version of Hue, use the following steps to upgrade Hue.

1. Migrate the `hue.ini` setting from your old `hue.ini` configuration file to new `hue.ini` configuration file.
2. If you are using the embedded SQLite database, remember to restore your database after upgrade.

To restore the database from a backup, make sure the destination database is empty before copying (if necessary, rename or remove the current destination database), then copy your backup to the destination database.

For example:

```
su - $HUE_USER
```

```
cd /var/lib/hue
```

```
mv desktop.db desktop.db.old
```

```
sqlite3 desktop.db < ~/hue_backup/desktop.bak
```

```
exit
```

3. Synchronize the database:

```
cd /usr/lib/hue  
source ./build/env/bin/activate  
hue syncdb  
deactivate
```

4. Start Hue. As a root user, run the following command on the Hue Server:

```
/etc/init.d/hue start
```

6.22. Finalize the Upgrade

You can start HDFS without finalizing the upgrade. When you are ready to discard your backup, you can finalize the upgrade.



Warning

You must verify your filesystem health before finalizing the upgrade. After you finalize an upgrade, you cannot roll back.

To finalize the upgrade process, run the following command as the \$HDFS_USER:

```
hadoop dfsadmin -finalizeUpgrade
```

6.23. Install New HDP 2.5 Services

Install new HDP 2.5 Services:

- Atlas – a low-level service, similar to YARN, that provides metadata services to the HDP platform.
- SmartSense – a next generation subscription model that features upgrade and configuration recommendations.