

Hortonworks Data Platform

Configuring Kafka for Kerberos Over Ambari

(Sep 30, 2015)

Hortonworks Data Platform: Configuring Kafka for Kerberos Over Ambari

Copyright © 2012-2015 Hortonworks, Inc. Some rights reserved.

The Hortonworks Data Platform, powered by Apache Hadoop, is a massively scalable and 100% open source platform for storing, processing and analyzing large volumes of data. It is designed to deal with data from many sources and formats in a very quick, easy and cost-effective manner. The Hortonworks Data Platform consists of the essential set of Apache Hadoop projects including MapReduce, Hadoop Distributed File System (HDFS), HCatalog, Pig, Hive, HBase, ZooKeeper and Ambari. Hortonworks is the major contributor of code and patches to many of these projects. These projects have been integrated and tested as part of the Hortonworks Data Platform release process and installation and configuration tools have also been included.

Unlike other providers of platforms built using Apache Hadoop, Hortonworks contributes 100% of our code back to the Apache Software Foundation. The Hortonworks Data Platform is Apache-licensed and completely open source. We sell only expert technical support, [training](#) and partner-enablement services. All of our technology is, and will remain, free and open source.

Please visit the [Hortonworks Data Platform](#) page for more information on Hortonworks technology. For more information on Hortonworks services, please visit either the [Support](#) or [Training](#) page. Feel free to [contact us](#) directly to discuss your specific needs.



Except where otherwise noted, this document is licensed under
Creative Commons Attribution ShareAlike 3.0 License.
<http://creativecommons.org/licenses/by-sa/3.0/legalcode>

Table of Contents

1. Overview	1
2. Preparing the Cluster	2
3. Configuring the Kafka Broker for Kerberos	3
4. Creating Kafka Topics	4
5. Producing Events/Messages to Kafka on a Secured Cluster	5
6. Consuming Events/Messages from Kafka on a Secured Cluster	7
7. Authorizing Access when Kerberos is Enabled	9
7.1. Granting Permissions to a User	9
7.2. Granting Permissions to a Producer	10
7.3. Granting Permissions to a Consumer	10
7.4. Granting Permissions to a Consumer Group	11
7.5. Troubleshooting Authorizer Configuration	11
8. Appendix: Kafka Configuration Options	13
8.1. Server.properties key-value pairs	13
8.2. JAAS Configuration File for the Kafka Server	15
8.3. Configuration Setting for the Kafka Producer	15
8.4. JAAS Configuration File for the Kafka Client	15

1. Overview

This chapter describes how to configure Kafka for Kerberos security on an Ambari-managed cluster.

Kerberos security for Kafka is an optional feature. When security is enabled, features include:

- Authentication of client connections (consumer, producer) to brokers
- ACL-based authorization

2. Preparing the Cluster

Before you enable Kerberos, your cluster must meet the following prerequisites:

Prerequisite	References*
Ambari-managed cluster with Kafka installed. <ul style="list-style-type: none">• Ambari Version 2.1.0.0 or later• Stack version HDP 2.3.2 or later	Installing, Configuring, and Deploying a HDP Cluster in Automated Install with Ambari
Key Distribution Center (KDC) server installed and running	Installing and Configuring the KDC in the <i>Ambari Security Guide</i>
JCE installed on all hosts on the cluster (including the Ambari server)	Enabling Kerberos Security in the <i>Ambari Security Guide</i>

Links are for Ambari 2.1.2.0.

When all prerequisites are fulfilled, enable Kerberos security. (For more information see [Launching the Kerberos Wizard \(Automated Setup\)](#) in the *Ambari Security Guide*.)

3. Configuring the Kafka Broker for Kerberos

During the installation process, Ambari configures a series of Kafka settings and creates a JAAS configuration file for the Kafka server.

It is not necessary to modify these settings, but for more information see Appendix A, [Kafka Configuration Options](#).

4. Creating Kafka Topics

When you use a script, command, or API to create a topic, an entry is created under ZooKeeper. The only user with access to ZooKeeper is the service account running Kafka (by default, `kafka`). Therefore, the first step toward creating a Kafka topic on a secure cluster is to run `kinit`, specifying the Kafka service keytab. The second step is to create the topic.

1. Run `kinit`, specifying the Kafka service keytab. For example:

```
kinit -k -t /etc/security/keytabs/kafka.service.keytab kafka/  
c6401.ambari.apache.org@EXAMPLE.COM
```

2. Next, create the topic. Run the `kafka-topics.sh` command-line tool with the following options:

```
/bin/kafka-topics.sh --zookeeper <hostname>:<port> --create  
--topic <topic-name> --partitions <number-of-partitions> --  
replication-factor <number-of-replicating-servers>
```

For example:

```
/bin/kafka-topics.sh --zookeeper c6401.ambari.apache.org:2181 --create --  
topic test_topic --partitions 2 --replication-factor 2  
  
Created topic "test_topic".
```

For more information about `kafka-topics.sh` parameters, see [Basic Kafka Operations](#) on the Apache Kafka website.

Permissions

By default, permissions are set so that only the Kafka service user has access; no other user can read or write to the new topic. In other words, if your Kafka server is running with principal `$KAFKA-USER`, only that principal will be able to write to ZooKeeper.

For information about adding permissions, see [Authorizing Access when Kerberos is Enabled](#).

5. Producing Events/Messages to Kafka on a Secured Cluster

Prerequisite: Make sure that you have enabled access to the topic (via Ranger or native ACLs) for the user associated with the producer process. We recommend that you use Ranger to manage permissions. For more information, see the [Apache Ranger User Guide for Kafka](#).

During the installation process, Ambari configures a series of Kafka client and producer settings, and creates a JAAS configuration file for the Kafka client. It is not necessary to modify these settings, but for more information about them see Appendix A, [Kafka Configuration Options](#).

Note: Only the Kafka Java API is supported for Kerberos. Third-party clients are not supported.

To produce events/messages:

1. Specify the path to the JAAS configuration file as one of your JVM parameters:

```
-Djava.security.auth.login.config=/usr/hdp/current/kafka-broker/
config/kafka_client_jaas.conf
```

For more information about the `kafka_client_jaas` file, see "JAAS Configuration File for the Kafka Client" in [Kafka Configuration Options](#).

2. `kinit` with the principal's keytab.
3. Launch `kafka-console-producer.sh` with the following configuration options. (Note: these settings are the same as in previous versions, except for the addition of `--security-protocol PLAINTEXTSASL`.)

```
./bin/kafka-console-producer.sh --broker-list <hostname:port
[,hostname:port, ...]> --topic <topic-name> --security-protocol
PLAINTEXTSASL
```

For example:

```
./bin/kafka-console-producer.sh --broker-list
c6401.ambari.apache.org:6667,c6402.ambari.apache.org:6667 --
topic test_topic --security-protocol PLAINTEXTSASL
```

Troubleshooting

Issue: If you launch the producer without specifying the `security-protocol` option, you will see the following error:


```
2015-07-21 04:14:06,611] ERROR fetching topic metadata for topics
[Set(test_topic)] from broker
[ArrayBuffer(BrokerEndPoint(0,c6401.ambari.apache.org,6667),
BrokerEndPoint(1,c6402.ambari.apache.org,6667))] failed
(kafka.utils.CoreUtils$)
kafka.common.KafkaException: fetching topic metadata for topics
[Set(test_topic)] from broker
[ArrayBuffer(BrokerEndPoint(0,c6401.ambari.apache.org,6667),
BrokerEndPoint(1,c6402.ambari.apache.org,6667))] failed
    at kafka.client.ClientUtils$.fetchTopicMetadata(ClientUtils.scala:73)
Caused by: java.io.EOFException: Received -1 when reading from channel, socket
has likely been closed.
    at kafka.utils.CoreUtils$.read(CoreUtils.scala:193)
    at kafka.network.BoundedByteBufferReceive.
readFrom(BoundedByteBufferReceive.scala:54)
```

Solution: Add `--security-protocol PLAINTEXTSASL` to the `kafka-console-producer.sh` runtime options.

6. Consuming Events/Messages from Kafka on a Secured Cluster

Prerequisite: Make sure that you have enabled access to the topic (via Ranger or native ACLs) for the user associated with the consumer process. We recommend that you use Ranger to manage permissions. For more information, see the [Apache Ranger User Guide for Kafka](#).

During the installation process, Ambari configures a series of Kafka client and producer settings, and creates a JAAS configuration file for the Kafka client. It is not necessary to modify these values, but for more information see Appendix A, [Kafka Configuration Options](#).

Note: Only the Kafka Java API is supported for Kerberos. Third-party clients are not supported.

To consume events/messages:

1. Specify the path to the JAAS configuration file as one of your JVM parameters. For example:

```
-Djava.security.auth.login.config=/usr/hdp/current/kafka-broker/
config/kafka_client_jaas.conf
```

For more information about the `kafka_client_jaas` file, see "JAAS Configuration File for the Kafka Client" in [Kafka Configuration Options](#).

2. `kinit` with the principal's keytab.
3. Launch `kafka-console-consumer.sh` with the following configuration settings. (Note: these settings are the same as in previous versions, except for the addition of `--security-protocol PLAINTEXTSASL`.)

```
./bin/kafka-console-consumer.sh --zookeeper
c6401.ambari.apache.org:2181 --topic test_topic --from-beginning
--security-protocol PLAINTEXTSASL
```

Troubleshooting

Issue: If you launch the consumer without specifying the `security-protocol` option, you will see the following error:

```
2015-07-21 04:14:06,611] ERROR fetching topic metadata for topics
[Set(test_topic)] from broker
[ArrayBuffer(BrokerEndPoint(0,c6401.ambari.apache.org,6667),
BrokerEndPoint(1,c6402.ambari.apache.org,6667))] failed
(kafka.utils.CoreUtils$)
kafka.common.KafkaException: fetching topic metadata for topics
[Set(test_topic)] from broker
[ArrayBuffer(BrokerEndPoint(0,c6401.ambari.apache.org,6667),
BrokerEndPoint(1,c6402.ambari.apache.org,6667))] failed
    at kafka.client.ClientUtils$.fetchTopicMetadata(ClientUtils.scala:73)
Caused by: java.io.EOFException: Received -1 when reading from channel, socket
has likely been closed.
    at kafka.utils.CoreUtils$.read(CoreUtils.scala:193)
    at kafka.network.BoundedByteBufferReceive.
readFrom(BoundedByteBufferReceive.scala:54)
```

Solution: Add `--security-protocol PLAINTEXTSASL` to the `kafka-console-consumer.sh` runtime options.

7. Authorizing Access when Kerberos is Enabled

Kafka supports Access Control List (ACL) authorization when Kerberos is enabled. ACLs are stored in Zookeeper.

A Kafka ACL entry has the following general format, specified in the `kafka-acls.sh` commands in the remainder of this section:

```
[Allow/Deny] Operations <value1, value2, ...> on Resource R from  
Hosts H1,H2 for Principals P1,P2
```

where

- `Operations` can be one of: READ, WRITE, CREATE, DESCRIBE, or ALL. (The CLI lists other options like DELETE/ALTER; they are not currently supported.)
- `Resource` is either a topic name, a consumer group name, or the string "kafka-cluster" to indicate a cluster level resource (only used with a CREATE operation).
- `Hosts` is a comma-separated list of hosts, or * to indicate all hosts.
- `Principals` is a comma-separated list of principals, or * to indicate all principals.

For mappings between `Operations` values and Kafka protocol APIs, refer to the [Apache Kafka authorization documentation](#).

7.1. Granting Permissions to a User

To grant all permissions to a user – for all resources – add the user to the list of super users:

1. Add the user to the list of users specified in the `super.users` property.
2. Restart the cluster.
3. Restart all brokers.

Alternately, to grant full access to a specific topic, cluster, and consumer group, run the following commands. (This approach does not require you to restart the cluster.)

In the following example, substitute your own values for `<topic-name>` and `<user-name>`:

```
/usr/hdp/current/kafka-broker/bin/kafka-acls.sh --topic <topic-  
name> --add --allowprincipals user:<user-name> --operations ALL --  
config /usr/hdp/current/kafka-broker/config/server.properties
```

```
/usr/hdp/current/kafka-broker/bin/kafka-acls.sh --cluster --add --  
allowprincipals user:<user-name> --config /usr/hdp/current/kafka-  
broker/config/server.properties --operations ALL
```

```
/usr/hdp/current/kafka-broker/bin/kafka-acls.sh --consumer-group  
10 --add --allowprincipals user:<user-name> --operations ALL --  
config /usr/hdp/current/kafka-broker/config/server.properties
```

7.2. Granting Permissions to a Producer

To grant permissions to a producer:

- Grant WRITE permissions to the topic that the producer user will write to.
- Grant DESCRIBE permission on the cluster.

Note: Optionally, if you have set `auto.create.topics.enable` to `true` and the topic is not created before starting the producer, you must also grant CREATE permission on the cluster.

Example:

The following two commands grant principal `ambari-qa` access as a producer, to topic `test-topic`, from host `c6401.ambari.apache.org`. The commands grant WRITE permission to the `ambari-qa` user on that topic, and DESCRIBE permission on the cluster:

```
./bin/kafka-acls.sh --topic test-topic --add --allowhosts  
c6401.ambari.apache.org --allowprincipals user:ambari-qa --  
operations WRITE --config /usr/hdp/current/kafka-broker/config/  
server.properties
```

```
./bin/kafka-acls.sh --cluster --add --allowhosts  
c6401.ambari.apache.org --allowprincipals user:ambari-qa --  
config /usr/hdp/current/kafka-broker/config/server.properties --  
operations DESCRIBE
```

7.3. Granting Permissions to a Consumer

To grant permissions to a consumer:

- Grant READ permissions to the topic that the consumer is going to read from.
- Grant DESCRIBE permission on the cluster.

Example:

The following commands grant principal `ambari-qa` access as a consumer, to topic `test-topic` from host `c6401.ambari.apache.org`. This is done by granting READ permission to the `ambari-qa` user on that topic, and granting CREATE permission on the cluster:

```
./bin/kafka-acls.sh --topic yes --add --allowhosts  
c6401.ambari.apache.org --allowprincipals user:ambari-qa --
```

```
operations READ --config /usr/hdp/current/kafka-broker/config/
server.properties

./bin/kafka-acls.sh --cluster --add --allowhosts
c6401.ambari.apache.org --allowprincipals user:ambari-qa --
operations CREATE --config /usr/hdp/current/kafka-broker/config/
server.properties
```

7.4. Granting Permissions to a Consumer Group

If you are using a high-level consumer with a consumer group, you must also grant READ permission to the consumer group.

Example:

The following commands grant principal `ambari-qa` access to consumer group “10” from host `c6401.ambari.apache.org`. This is done by granting READ permission to the `ambari-qa` user on that group.

```
/usr/hdp/current/kafka-broker/bin/kafka-acls.sh --consumer-group
10 --add --allowhosts c6401.ambari.apache.org --allowprincipals
user:ambari-qa --operations READ --config /usr/hdp/current/kafka-
broker/config/server.properties
```

7.5. Troubleshooting Authorizer Configuration

Frequently-asked Questions:

When should I use Deny?

Never. By default, all principals that are not explicitly granted permissions get rejected so you should not have to use Deny ever. (Note: when defined, DENY takes precedence over ALLOW.)

Then why do we have deny?

Deny was introduced into Kafka for advanced use cases where negation was required. For example, if an admin wants to grant WRITE access to a topic from all hosts but `host1`, the admin can define two ACLs:

```
"grant WRITE to test-topic from hosts * for principal test-user"
"deny WRITE to test-topic from hosts host1 for principal test-
user"
```

Deny should only be used to negate a large allow, where listing all principals or hosts is cumbersome.

Can I define ACLs with principal as user@<realm>?

You can if you are not using `principal.to.local.class`, but if you have set this configuration property you must define your ACL with users without and REALM. This is a known issue in HDP 2.3.

I just gave a user CREATE Permissions on a cluster, but the user still can't create topics. Why?

Right now, Kafka create topic is not implemented as an API, but as a script that directly modifies ZooKeeper entries. In a secure environment only the Kafka broker user is allowed to write to ZooKeeper entries. Granting a user CREATE access does not allow that user to modify ZooKeeper entries.

However, if that user makes a producer request to the topic and has `auto.create.topics.enable` set to `true`, a topic will be created at the broker level.

8. Appendix: Kafka Configuration Options

8.1. Server.properties key-value pairs

Ambari configures the following Kafka values during the installation process. Settings are stored as key-value pairs stored in an underlying `server.properties` configuration file.

listeners

A comma-separated list of URIs that Kafka will listen on, and their protocols.

Required property with three parts:

`<protocol>:<hostname>:<port>`

Set `<protocol>` to `PLAINTEXTSASL`, to specify the protocol that server accepts connections. SASL authentication will be used over a plaintext channel. Once SASL authentication is established between client and server, the session will have the client's principal as an authenticated user. The broker can only accept SASL (Kerberos) connections, and there is no wire encryption applied. (Note: For a non-secure cluster, `<protocol>` should be set to `PLAINTEXT`.)

Set `hostname` to the hostname associated with the node you are installing. Kerberos uses this value and "principal" to construct the Kerberos service name. Specify `hostname 0.0.0.0` to bind to all interfaces. Leave `hostname` empty to bind to the default interface.

Set `port` to the Kafka service port. When Kafka is installed using Ambari, the default port number is 6667.

Examples of legal listener lists::

```
listeners=PLAINTEXTSASL://kafkal.witzend.com:6667
```

```
listeners=PLAINTEXT://myhost:9092, TRACE://:9091,  
PLAINTEXTSASL://0.0.0.0:9093
```

advertised.listeners

A list of listeners to publish to ZooKeeper for clients to use, if different than the listeners specified in the preceding section.

In IaaS environments, this value might need to be different from the interface to which the broker binds.

If `advertised.listeners` is not set, the value for `listeners` will be used.

Required value with three parts:

`<protocol>:<hostname>:<port>`

Set `protocol` to `PLAINTEXTSASL`, to specify the protocol that server accepts connections. SASL authentication will be used over a plaintext channel. Once SASL authentication is established between client and server, the session will have the client's principal as an authenticated user. The broker can only accept SASL (Kerberos) connections, and there is no wire encryption applied. (Note: For a non-secure cluster, `<protocol>` should be set to `PLAINTEXT`.)

Set `hostname` to the hostname associated with the node you are installing. Kerberos uses this and "principal" to construct the Kerberos service name.

Set `port` to the Kafka service port. When Kafka is installed using Ambari, the default port number is 6667.

For example:

```
advertised.listeners=PLAINTEXTSASL://kafka1.witzend.com:6667
```

security.inter.broker.protocol

Specifies the inter-broker communication protocol. In a Kerberized cluster, brokers are required to communicate over SASL. (This approach supports replication of topic data.) Set the value to `PLAINTEXTSASL`:

```
security.inter.broker.protocol=PLAINTEXTSASL
```

authorizer.class.name

Configures the authorizer class.

Set this value to `kafka.security.auth.SimpleAclAuthorizer`:

```
authorizer.class.name=kafka.security.auth.SimpleAclAuthorizer
```

For more information, see "Authorizing Access when Kerberos is Enabled."

principal.to.local.class

Transforms Kerberos principals to their local Unix usernames.

Set this value to `kafka.security.auth.KerberosPrincipalToLocal`:

```
principal.to.local.class=kafka.security.auth.KerberosPrincipalToLocal
```

super.users

Specifies a list of user accounts that will have all cluster permissions. Kafka is included in the list by default, because it is the principal associated with brokers – it needs the permissions.

Set this value to `user:<comma-separated-list-of-user-accounts>`; for example:

```
super.users=user:kafka
```

8.2. JAAS Configuration File for the Kafka Server

The Java Authentication and Authorization Service (JAAS) API supplies user authentication and authorization services for Java applications.

After enabling Kerberos, Ambari sets up a JAAS login configuration file for the Kafka server. This file is used to authenticate the Kafka broker against Kerberos. The file is stored at:

```
/usr/hdp/current/kafka-broker/config/kafka_server_jaas.conf
```

Ambari adds the following settings to the file. (Note: serviceName="kafka" is required for connections from other brokers.)

```
KafkaServer {
    com.sun.security.auth.module.Krb5LoginModule required
    useKeyTab=true
    keyTab="/etc/security/keytabs/kafka.service.keytab"
    storeKey=true
    useTicketCache=false
    serviceName="kafka"
    principal="kafka/c6401.ambari.apache.org@EXAMPLE.COM" ;
};

Client { // used for zookeeper connection
    com.sun.security.auth.module.Krb5LoginModule required
    useKeyTab=true
    keyTab="/etc/security/keytabs/kafka.service.keytab"
    storeKey=true
    useTicketCache=false
    serviceName="zookeeper"
    principal="kafka/c6401.ambari.apache.org@EXAMPLE.COM" ;
};
```

8.3. Configuration Setting for the Kafka Producer

After enabling Kerberos, Ambari sets the following key-value pair in the server.properties file:

```
security.protocol=PLAINTEXTSASL
```

8.4. JAAS Configuration File for the Kafka Client

After enabling Kerberos, Ambari sets up a JAAS login configuration file for the Kafka client. Settings in this file will be used for any client (consumer, producer) that connects to a Kerberos-enabled Kafka cluster. The file is stored at:

```
/usr/hdp/current/kafka-broker/config/kafka_client_jaas.conf
```

Ambari adds the following settings to the file. (Note: serviceName=kafka is required for connections from other brokers.)



Note

For command-line utilities like `kafka-console-producer` and `kafka-console-consumer`, use `kinit`. If you use a long-running process (for example, your own Producer), use `keytab`.

Kafka client configuration *with* keytab, for *producers*:

```
KafkaClient {  
  com.sun.security.auth.module.Krb5LoginModule required  
  useKeyTab=true  
  keyTab="/etc/security/keytabs/storm.service.keytab"  
  storeKey=true  
  useTicketCache=false  
  serviceName="kafka"  
  principal="storm@EXAMPLE.COM";  
};
```

Kafka client configuration *without* keytab, for *producers*:

```
KafkaClient {  
  com.sun.security.auth.module.Krb5LoginModule required  
  useTicketCache=true  
  renewTicket=true  
  serviceName="kafka";  
};
```

Kafka client configuration for *consumers*:

```
KafkaClient {  
  com.sun.security.auth.module.Krb5LoginModule required  
  useTicketCache=true  
  renewTicket=true  
  serviceName="kafka";  
};
```